
COMP 472: Artificial Intelligence

Machine Learning *part #2*


Decision Trees *video #4*

- Russell & Norvig: Sections 19.3


Today

1. Introduction to ML
2. Naive Bayes Classification
 - a. Application to Spam Filtering
3. **Decision Trees** 
4. (Evaluation
5. Unsupervised Learning)
6. Neural Networks
 - a. Perceptrons
 - b. Multi Layered Neural Networks

Guess Who?



Decision Trees




Ross
Quinlan

- Simplest, but most successful form of learning algorithm
- Very well-known algorithm is ID3 (Quinlan, 1987) and its successor C4.5

3-48 . . .

1. Rank features based on how good they are to indicate the result
2. Put the most discriminating feature as a node (as a question) of a tree
3. Split the examples so that those with different values for the chosen feature are in a different set
4. Repeat the same process with the next most discriminating feature



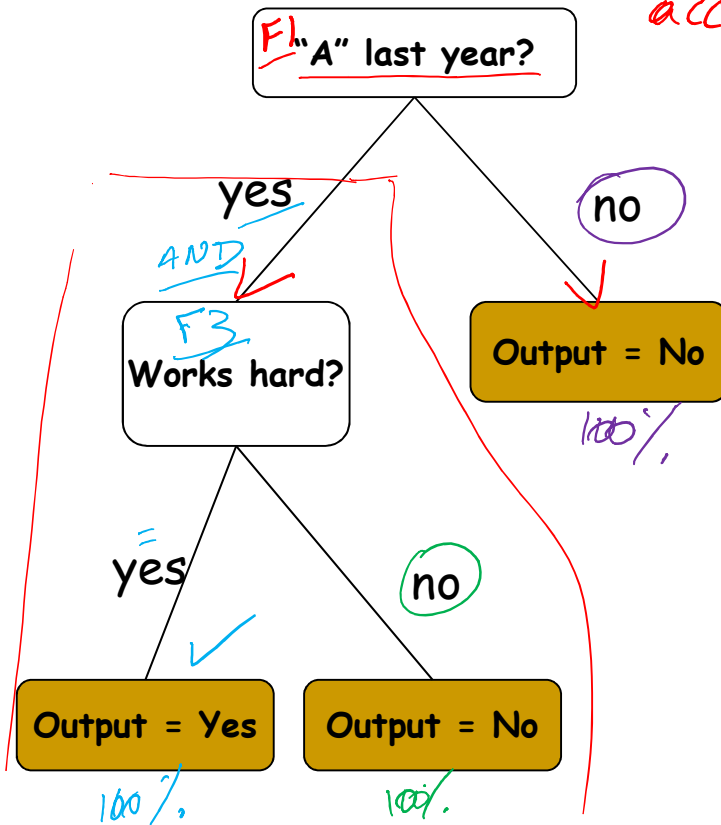
Example 1

Info on last year's students to determine if a student will get an 'A' this year

	Features (X)				Output f(X) <i>class</i>
Student	'A' last year? <i>F1</i>	Black hair? <i>F2</i>	Works hard? <i>F3</i>	Drinks? <i>F4</i>	'A' this year?
X1: <u>Richard</u>	<u>Yes</u>	<u>Yes</u>	<u>No</u>	<u>Yes</u>	<u>No</u>
X2: <u>Alan</u>	Yes	<u>Yes</u>	<u>Yes</u>	<u>No</u>	Yes
X3: Alison	No	<u>No</u>	<u>Yes</u>	No	No
X4: Jeff	No	Yes	No	Yes	No
X5: Gail	Yes	No	Yes	Yes	Yes
X6: Simon	No	Yes	Yes	Yes	No

Example 1

A random decision tree that fits the dataset *with 100% accuracy*



same data set as previous slide

	Features				Output f(X)
Student	'A' last year? F1	Black hair?	Works hard? F3	Drinks ?	'A' this year?
Richard	Yes	Yes	No	Yes	No
Alan	Yes	Yes	Yes	No	Yes
Alison	No	No	Yes	No	No
Jeff	No	Yes	No	Yes	No
Gail	Yes	No	Yes	Yes	Yes
Simon	No	Yes	Yes	Yes	No

- fits the training set with 100% accuracy
- but the features are chosen at random so there might be a shorter DT
 — // better

Example 2: The Restaurant

- Goal: learn whether one should wait for a table
- Attributes

1. Alternate: another suitable restaurant nearby
2. Bar: comfortable bar for waiting
3. Fri/Sat: true on Fridays and Saturdays
4. Hungry: whether one is hungry
5. Patrons: how many people are present (none, some, full)
6. Price: price range (\$, \$\$, \$\$\$)
7. Raining: raining outside
8. Reservation: reservation made
9. Type: kind of restaurant (French, Italian, Thai, Burger)
10. WaitEstimate: estimated wait by host (0-10 mins, 10-30, 30-60, >60)

10 features

Example 2: The Restaurant

■ Training data:

10 features.

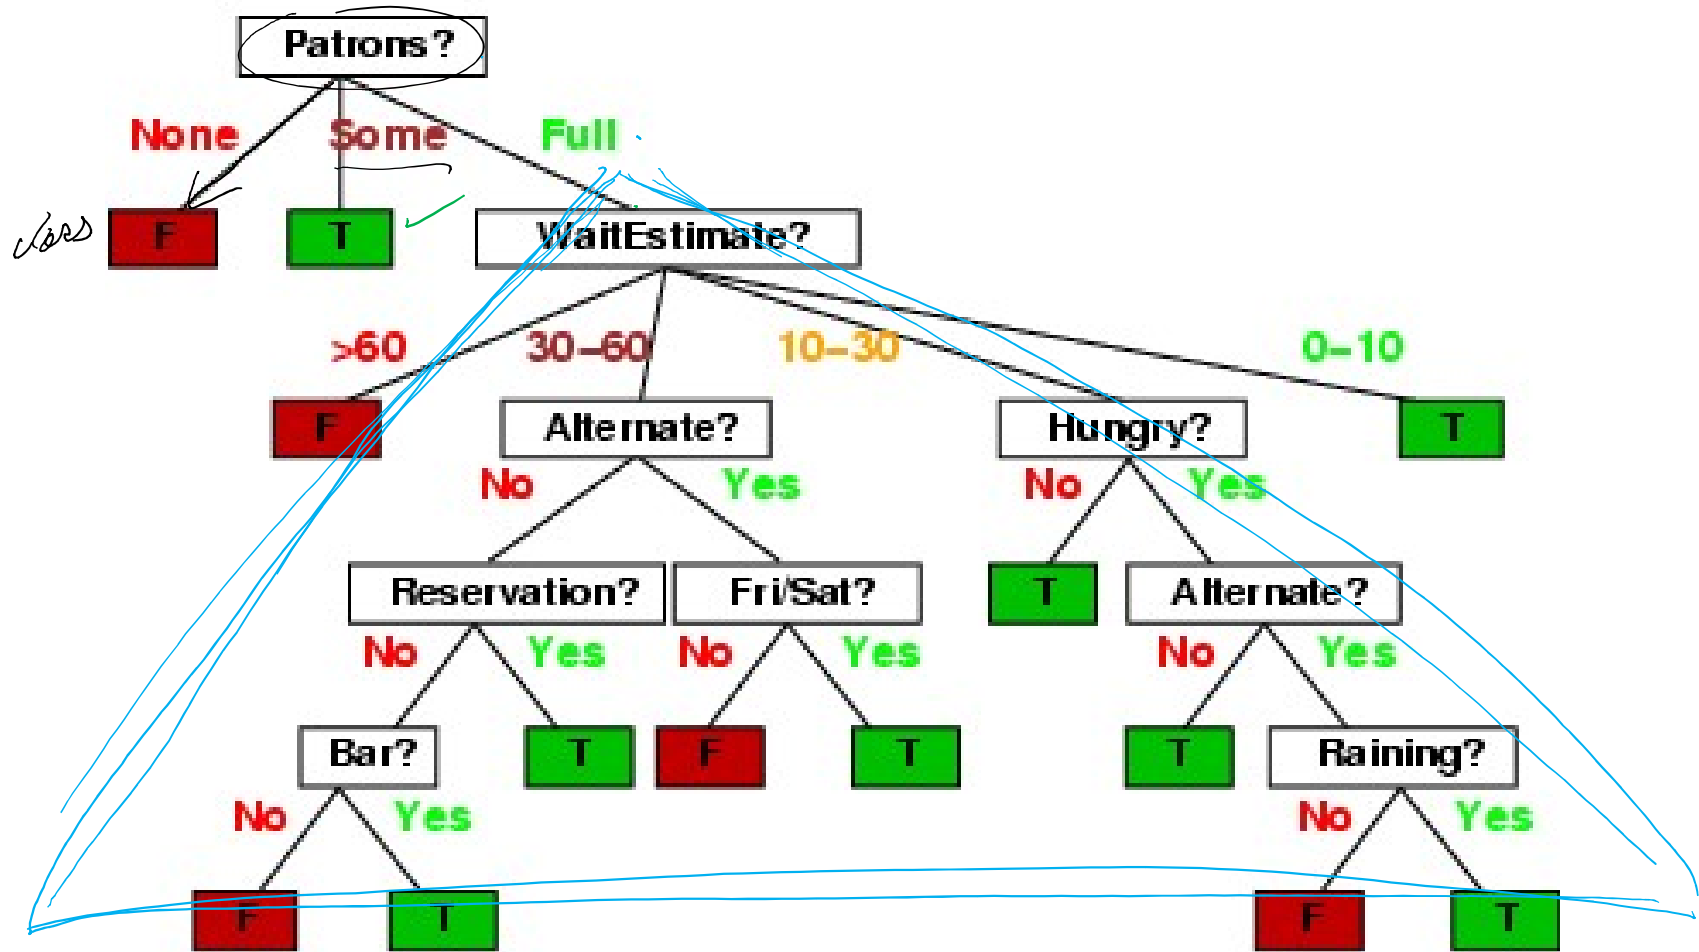
$f(x)$
class

12
instances

Example	Attributes										Target
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	Wait
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

source: Norvig (2003)

A First Decision Tree



- But is it the best decision tree we can build?

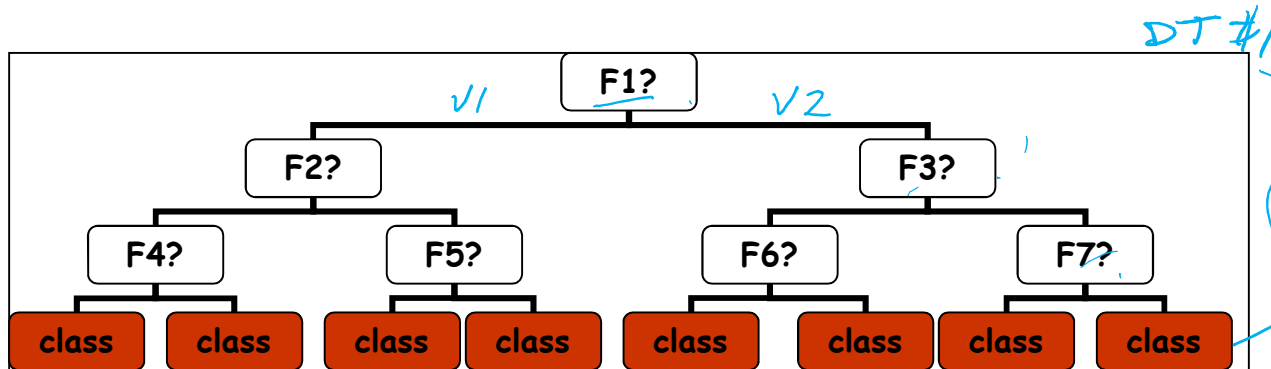
Ockham's Razor

It is vain to do more than can be done with less... Entities should not be multiplied beyond necessity.
[Ockham, 1324]

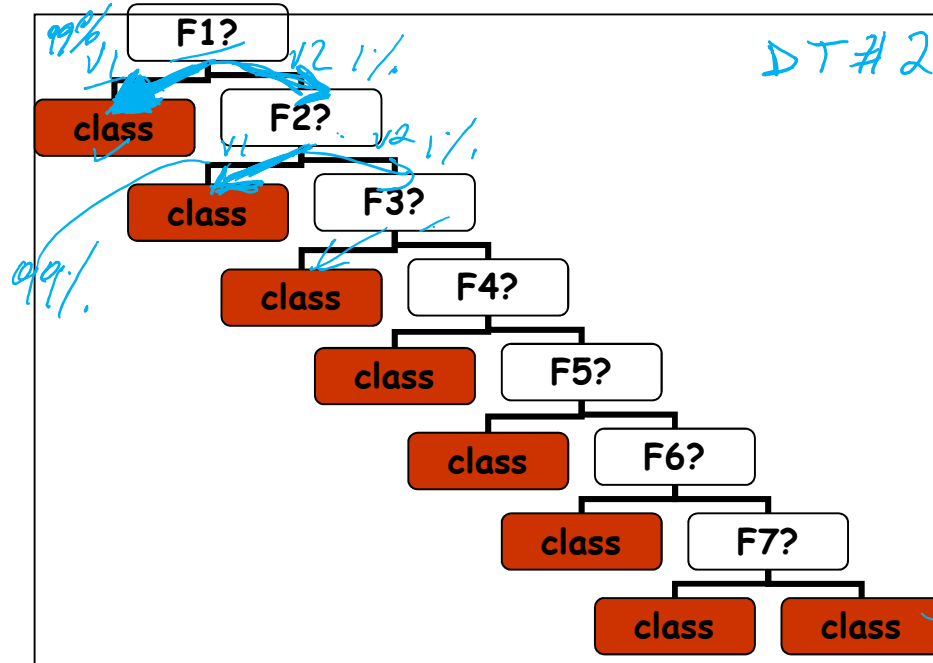


- In other words... always favor the simplest answer that correctly fits the training data
- i.e. the smallest tree on average
- This type of assumption is called inductive bias
 - inductive bias = making a choice beyond what the training instances contain

Which Tree is Best?




always asks
3 questions



may be the
shortest tree on
average depending
on the
probability of
each branch

A Better Decision Tree

- 4 tests instead of 9
- 11 branches instead of 21



Choosing the Next Feature

- The key problem is choosing which feature to split a given set of examples
- Different measures have been proposed
- ID3 uses Maximum Information-Gain
 - i.e. we choose the feature that has the largest information gain
 - we expect this feature to result in the smallest tree on average
 - based on information theory

Essential Information Theory

- Developed by ^{clau}Shannon in the 40s
- Shannon developed the notion of entropy (aka information content) of a random variable (RV)
- Entropy measures how "predictable" a RV is
 - If you already have a good idea about the answer (e.g. 90/10 split)
 - low entropy // *sure thing*
 - If you have no idea about the answer (e.g. 50/50 split)
 - high entropy // *total chaos*


Dartmouth Conference: The Founding Fathers of AI



John McCarthy



Marvin Minsky



Claude Shannon



Ray Solomonoff

Alan Newell



Herbert Simon



Arthur Samuel



And three others...
Oliver Selfridge
(Pandemonium theory)
Nathaniel Rochester
(IBM, designed 701)
Trenchard More
(Natural Deduction)

Entropy


- Let X be a discrete RV with i possible outcomes x_i
- Entropy (or information content) of X

$$H(X) = - \sum_{i=1}^n p(x_i) \log_2 p(x_i)$$

- measures:
 - the amount of information in a RV
 - average uncertainty of a RV
 - average length of a message needed to transmit an outcome x_i of that RV when encoded optimally over a binary channel
- measured in bits

base 2 always

x_1, x_2, \dots, x_n



Entropy of a Coin Toss

$$H(X) = - \sum_{x_i \in X} p(x_i) \log_2 p(x_i)$$

Entropy (or information content)

$$H(\text{fair coin toss}) = - \sum_{x_i \in X} p(x_i) \log_2 p(x_i) = H\left(\frac{1}{2}, \frac{1}{2}\right)$$

$$= - \left(\frac{1}{2} \log_2 \frac{1}{2} + \frac{1}{2} \log_2 \frac{1}{2} \right) = 1 \text{ bit}$$

$p(x_1) = \text{"head"}$
 $p(x_2) = \text{"tail"}$

entropy of a fair coin toss (the RV) with 2 possible outcomes, each with a probability of 1/2

a RV with only 2 outcomes x_1 and x_2 will have $1 \geq H(X) \geq 0$


// sure thing

// total chaos

(unpredictable)

Example: The Coin Toss

- Fair coin: $H(X) = -\sum_{i=1}^n p(x_i) \log_2 p(x_i) = -\left(\overset{\text{head}}{\frac{1}{2} \log_2 \frac{1}{2}} + \overset{\text{tail}}{\frac{1}{2} \log_2 \frac{1}{2}}\right) = \underline{1 \text{ bit}}$
- Rigged coin: $H(X) = -\sum_{i=1}^n p(x_i) \log_2 p(x_i) = -\left(\overset{\text{head}}{\frac{99}{100} \log_2 \frac{99}{100}} + \frac{1}{100} \log_2 \frac{1}{100}\right) = 0.08 \text{ bits}$



fair coin -> high entropy

rigged coin -> low entropy

P(head)

So what?

■ Training data:

for each attribute / feature I will compute how much knowing its values will reduce the entropy of the class

Example	Attributes										Target
	Alt	Bar	Fri	Hun	Pat	Price	Rain	Res	Type	Est	Wait
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

$$\text{entropy}(\text{Target Wait}) = 1$$

source: Norvig (2003)

Information Gain

■ information gain

- measure the entropy reduction of a RV, once a piece of information is known
- used to measure the "discriminating power" of an attribute A given a data set S
- Let $\text{Values}(A)$ = the set of values that attribute A can take
- Let S_v = the set of examples in the data set which have value v for attribute A (for each value v from $\text{Values}(A)$)

information gain (or
entropy reduction)

$$\begin{aligned}\text{gain}(S, A) &= H(S) - H(S|A) \\ &= H(S) - \sum_{v \in \text{values}(A)} \frac{|S_v|}{|S|} \times H(S_v)\end{aligned}$$

// see slide 15

Some Intuition

$$H(\text{output}) = 1 \quad // \text{ total chaos}$$

3 features

Size	Color	Shape	Output
Big	Red	Circle	+
Small	Red	Circle	+
Small	Red	Square	-
Big	Blue	Circle	+

- Size is the least discriminating attribute (i.e. smallest information gain)
- Shape and color are the most discriminating attributes (i.e. highest information gain)


shape.
 color. } more discriminating than size

A Small Example (1)

let's try $F2 = \text{color}$

Size	Color	Shape	Output
Big	Red	Circle	+
Small	Red	Circle	+
Small	Red	Square	-
Big	Blue	Circle	-

Values(Color) = {red, blue}



$$\text{gain}(S, \text{Color}) = H(S) - \sum_{v \in \text{values}(\text{Color})} \frac{|S_v|}{|S|} \times H(S_v)$$

entropy
of
output

$$H(S) = -\left(\frac{2}{4} \log_2 \frac{2}{4} + \frac{2}{4} \log_2 \frac{2}{4}\right) = 1$$

total chaos, not knowing anything

for each v of Values(Color)

$$H(S | \text{Color} = \text{red}) = H\left(\frac{2}{3}, \frac{1}{3}\right) = -\left(\frac{2}{3} \log_2 \frac{2}{3} + \frac{1}{3} \log_2 \frac{1}{3}\right) = 0.918$$

$$H(S | \text{Color} = \text{blue}) = H(1, 0) = -\left(\frac{1}{1} \log_2 \frac{1}{1}\right) = 0 \quad \text{// sure thing}$$

$$H(S | \text{Color}) = \frac{3}{4} (0.918) + \frac{1}{4} (0) = 0.6885 \quad \text{knowing the color}$$

Note: by definition,

- Log 0 = $-\infty$
- $0 \log 0$ is 0

1 time out of 4
we have a blue


$$\text{gain}(\text{Color}) = H(S) - H(S | \text{Color}) = 1 - 0.6885 = 0.3115$$

3 times out of 4 we have a red

A Small Example (2)

F3

Size	Color	Shape	Output
Big	Red	Circle	+
Small	Red	Circle	+
Small	Red	Square	-
Big	Blue	Circle	-



$$H(S) = -\left(\frac{2}{4}\log_2\frac{2}{4} + \frac{2}{4}\log_2\frac{2}{4}\right) = 1$$


$$H(S | \text{Shape}) = \frac{3}{4}(0.918) + \frac{1}{4}(0) = 0.6885$$

$$\text{gain}(\text{Shape}) = H(S) - H(S | \text{Shape}) = 1 - 0.6885 = 0.3115$$

A Small Example (3)

F1

Size	Color	Shape	Output
Big	Red	Circle	+
Small	Red	Circle	+
Small	Red	Square	-
Big	Blue	Circle	-



$$H(S) = -\left(\frac{2}{4}\log_2\frac{2}{4} + \frac{2}{4}\log_2\frac{2}{4}\right) = 1$$

$$H(S|Size) = \frac{2}{4}(1) + \frac{2}{4}(1)$$

$$\text{gain}(Size) = H(S) - H(S|Size) = 1 - 1 = 0$$

$H(\text{Output})$ without knowing any feature

$H(\text{Output}/Size)$ knowing the size

A Small Example (4)

Size	Color	Shape	Output
Big	Red	Circle	+
Small	Red	Circle	+
Small	Red	Square	-
Big	Blue	Circle	-

gain(Shape) = 0.3115 ✓

gain(Color) = 0.3115 ✓

gain(Size) = 0 ✓

- So first separate according to either color or shape (root of the tree)


A Small Example (4)

- Let's assume we pick Color for the root...

Size	Color	Shape	Output
Big	Red	Circle	+
Small	Red	Circle	+
Small	Red	Square	-
Big	Blue	Circle	-

S_2

S_2



$$H(S_2) = -\left(\frac{2}{3} \log_2 \frac{2}{3} + \frac{1}{3} \log_2 \frac{1}{3}\right)$$

for each v of Values(Size)

$$H(S_2 | \text{Size} = \text{big}) = H\left(\frac{1}{1}, \frac{0}{1}\right) = 0$$

$$H(S_2 | \text{Size} = \text{small}) = H\left(\frac{1}{2}, \frac{1}{2}\right) = 1$$

$$H(S_2 | \text{Size}) = \frac{1}{3}(0) + \frac{2}{3}(1)$$

$$\text{gain}(\text{Size}) = H(S_2) - H(S_2 | \text{Size})$$

only use the instances with color = red

for each v of Values(Shape)

$$H(S_2 | \text{Shape} = \text{circle}) = H\left(\frac{2}{2}, \frac{0}{2}\right) = 0$$

$$H(S_2 | \text{Shape} = \text{square}) = H\left(\frac{0}{1}, \frac{1}{1}\right) = 0$$

$$H(S_2 | \text{Shape})$$

$$\text{gain}(\text{Shape}) = H(S_2) - H(S_2 | \text{Shape})$$

only use the instances with color = blue

Highest gain is the root of the subtree

Back to the Restaurant

■ Training data:

pick max info gain

Example	Attributes										Target
	<i>Alt</i>	<i>Bar</i>	<i>Fri</i>	<i>Hun</i>	<i>Pat</i>	<i>Price</i>	<i>Rain</i>	<i>Res</i>	<i>Type</i>	<i>Est</i>	<i>Wait</i>
X_1	T	F	F	T	Some	\$\$\$	F	T	French	0-10	T
X_2	T	F	F	T	Full	\$	F	F	Thai	30-60	F
X_3	F	T	F	F	Some	\$	F	F	Burger	0-10	T
X_4	T	F	T	T	Full	\$	F	F	Thai	10-30	T
X_5	T	F	T	F	Full	\$\$\$	F	T	French	>60	F
X_6	F	T	F	T	Some	\$\$	T	T	Italian	0-10	T
X_7	F	T	F	F	None	\$	T	F	Burger	0-10	F
X_8	F	F	F	T	Some	\$\$	T	T	Thai	0-10	T
X_9	F	T	T	F	Full	\$	T	F	Burger	>60	F
X_{10}	T	T	T	T	Full	\$\$\$	F	T	Italian	10-30	F
X_{11}	F	F	F	F	None	\$	F	F	Thai	0-10	F
X_{12}	T	T	T	T	Full	\$	F	F	Burger	30-60	T

The Restaurant Example

gain(alt) = ... gain(bar) = ... gain(fri) = ... gain(hun) = ...

4 features

$$\begin{aligned} \text{gain}(\text{pat}) &= 1 - \left(\frac{2}{12} \times H\left(\frac{0}{2}, \frac{2}{2}\right) + \frac{4}{12} \times H\left(\frac{0}{4}, \frac{4}{4}\right) + \frac{6}{12} \times H\left(\frac{2}{6}, \frac{4}{6}\right) \right) \\ &= 1 - \left(\frac{2}{12} \times - \left(\frac{0}{2} \log_2 \frac{0}{2} + \frac{2}{2} \log_2 \frac{2}{2} \right) + \frac{4}{12} \times - \left(\frac{0}{4} \log_2 \frac{0}{4} + \frac{4}{4} \log_2 \frac{4}{4} \right) + \dots \right) \approx 0.541 \text{ bits} \end{aligned}$$

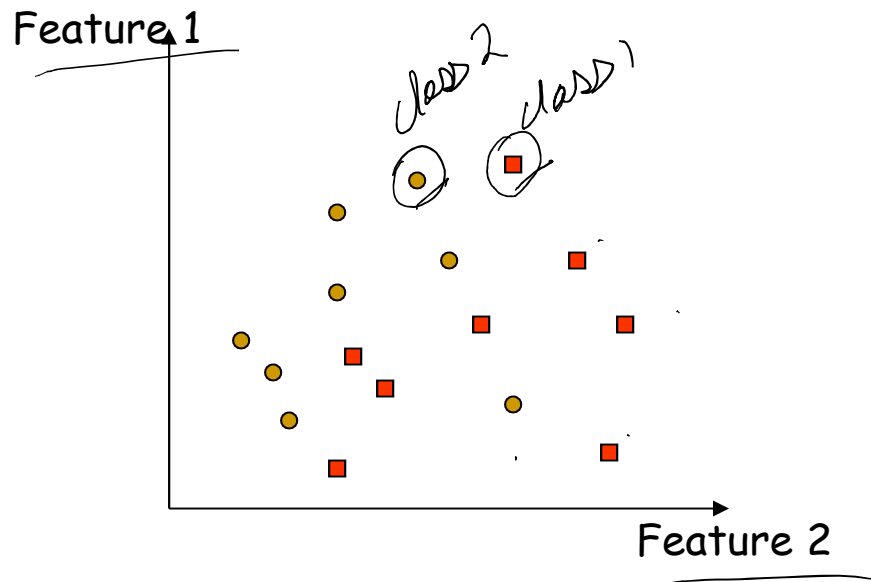
gain(price) = ... gain(rain) = ... gain(res) = ...

$$\text{gain}(\text{type}) = 1 - \left(\frac{2}{12} \times H\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{2}{12} \times H\left(\frac{1}{2}, \frac{1}{2}\right) + \frac{4}{12} \times H\left(\frac{2}{4}, \frac{2}{4}\right) + \frac{4}{12} \times H\left(\frac{2}{4}, \frac{2}{4}\right) \right) = 0 \text{ bits}$$



gain(est) = ...

- Attribute pat (Patron) has the highest gain, so root of the tree should be attribute *Patrons*
- do recursively for subtrees


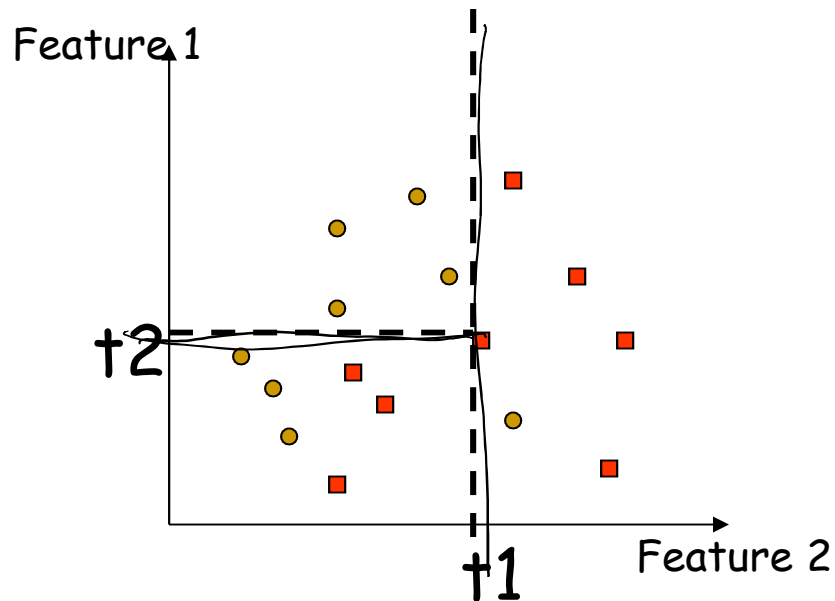
Decision Boundaries of Decision Trees



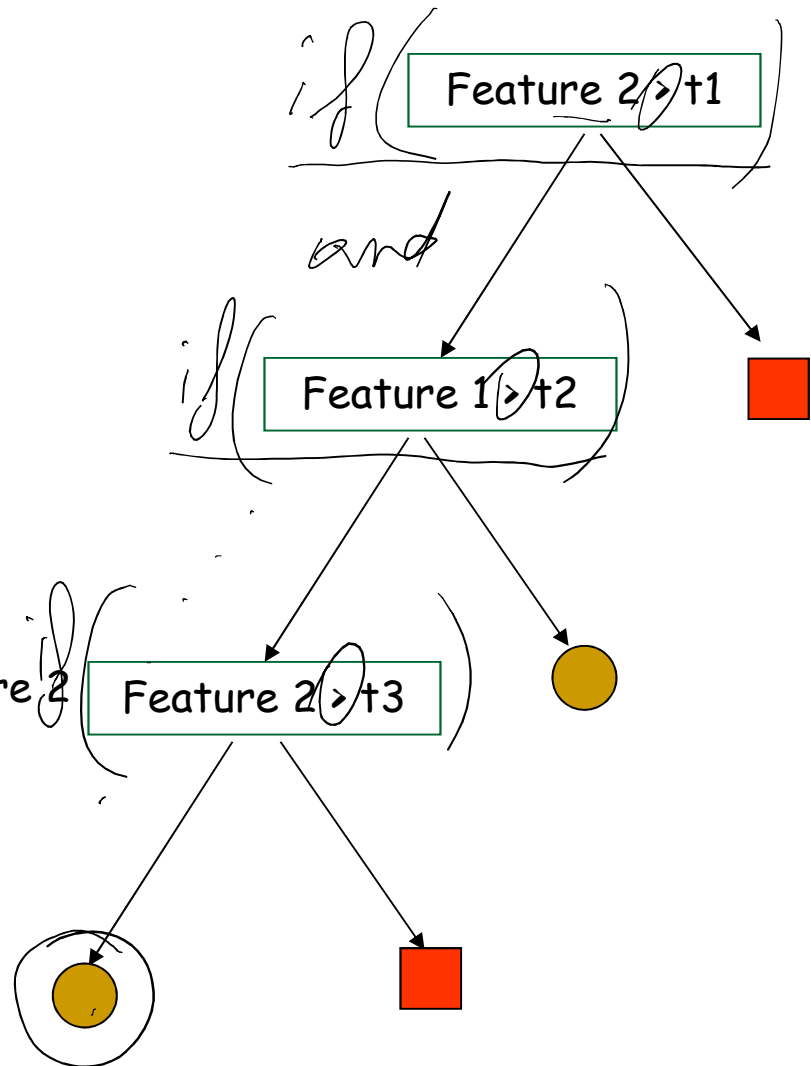

Decision Boundaries of Decision Trees



Decision Boundaries of Decision Trees







Decision Boundaries of Decision Trees



Applications of Decision Trees

- One of the most widely used learning methods in practice
 - Fast
 - Simple
 - Traceable (<-- very important!)

Today

1. Introduction to ML 
2. Naïve Bayes Classification 
 - a. Application to Spam Filtering 
3. Decision Trees  video #4
4. (Evaluation
5. Unsupervised Learning)
6. Neural Networks
 - a. Perceptrons
 - b. Multi Layered Neural Networks

Up Next

1. Introduction to ML
2. Naive Bayes Classification
 - a. Application to Spam Filtering
3. Decision Trees
4. (Evaluation)
5. Unsupervised Learning()
6. Neural Networks
 - a. Perceptrons
 - b. Multi Layered Neural Networks

supervised learning