


---

# COMP472: Artificial Intelligence Machine Learning Evaluation *video #5*

- Russell & Norvig: Sections 19.4

# Today

1. Introduction to ML
2. Naive Bayes Classification *Model #1*
  - a. Application to Spam Filtering
3. Decision Trees *Model #2*
4. Evaluation 
5. Unsupervised Learning )
6. Neural Networks
  - a. Perceptrons
  - b. Multi Layered Neural Networks



# Standard Methodology

1. Collect a large set of examples (all with correct classifications)

2. Divide collection into training, validation and test set

Loop:

3. Apply learning algorithm to the training set to learn the parameters

4. Measure performance with the validation set, and adjust hyper-parameters\* to improve performance

5. Measure performance with the test set

■ DO NOT LOOK AT THE TEST SET until step 5.

## Parameters:

basic values learned by the ML model. eg.

- for NB: prior & <sup>likelihood</sup> conditional probabilities
- for DTs: features to split
- for ANNs: weights

Hyper-parameters: parameters used to set up the ML model. eg.

- for NB: value of delta for smoothing,
- for DTs: pruning level
- for ANNs: nb of hidden layers, nb of nodes per layer...

ex:  
add -1 ✓  
0.5 ✓  
0.3 ✓


# Metrics

## 1 Accuracy

- % of instances of the test set the algorithm correctly classifies
- when all classes are equally important and represented  
 *$f(x)$  balanced dataset*

## 2 Recall, Precision & F-measure

- when one class is more important and the others



# Accuracy

- % of instances of the test set the algorithm correctly classifies
- when all classes are equally important and represented
- problem:
  - when one class (eg. sick) is more important and the others
  - eg. when data set is unbalanced

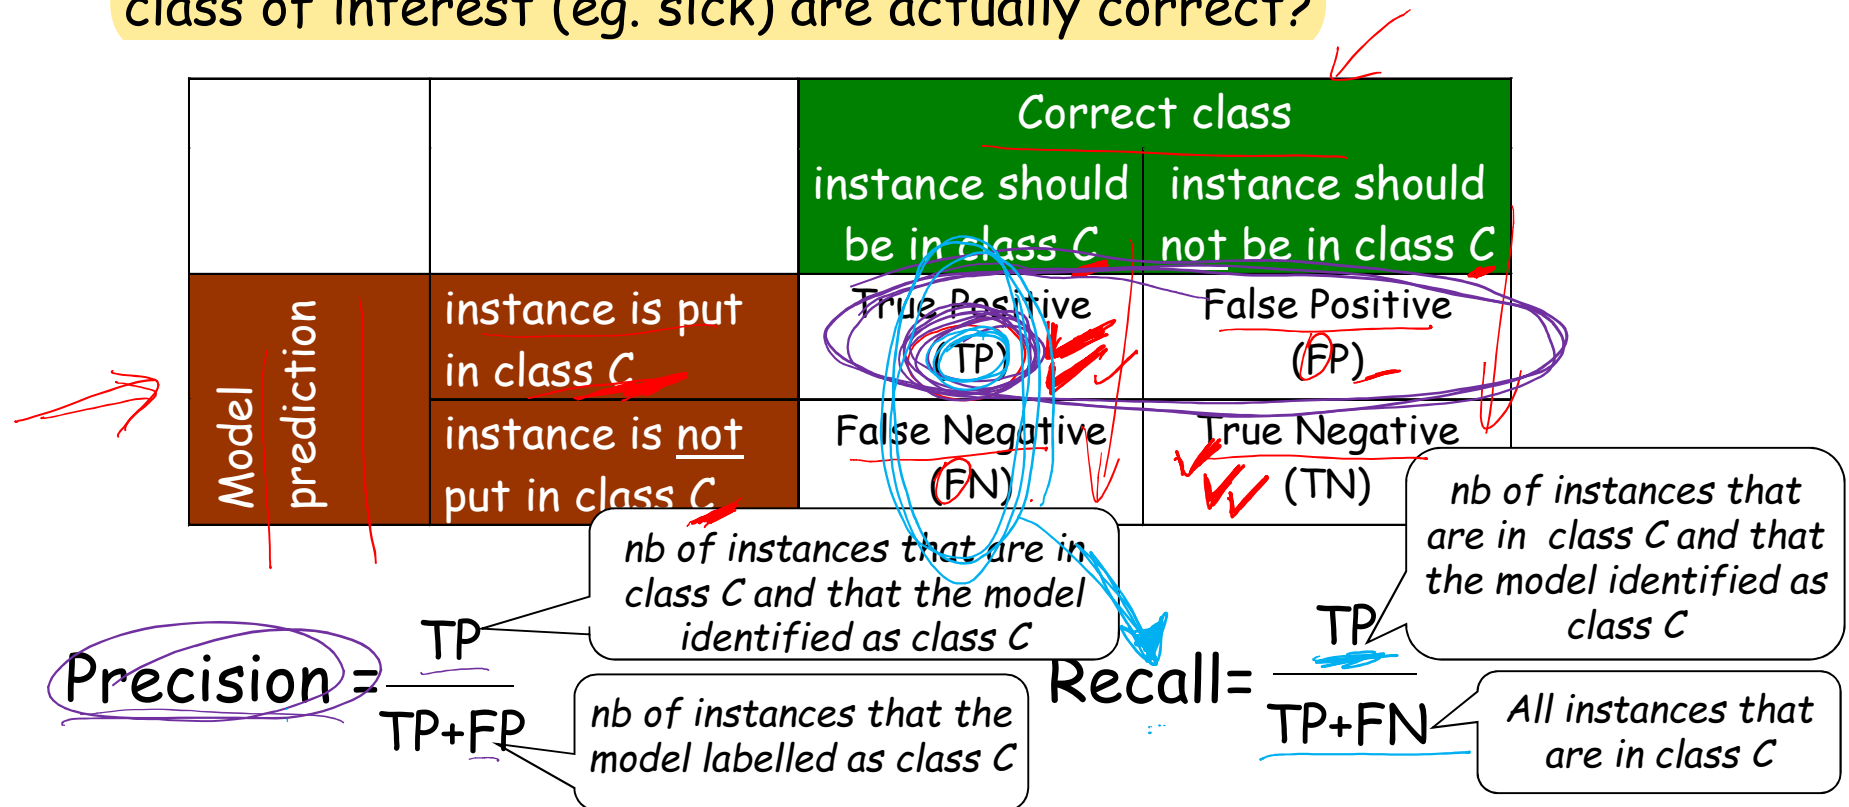
	Target	system 1
X1	sick	ok ✗
X2	sick	ok ✗
X3	sick	ok ✗
X4	sick	ok ✗
X5	sick	ok ✗
X6	ok	ok ✓
X7	ok	ok ✓
...	...	...
...	...	...
X500	ok	ok ✓
Accuracy		495/500 = 99% !

correct  
classification  
prediction

5 mistakes

# Recall, Precision

- Recall: What proportion of the instances in the class of interest (eg. sick) are labelled correctly?
- Precision: What proportion of instances labeled with the class of interest (eg. sick) are actually correct?



# Example

correct classification

class of interest = sick

	Target	system 1	system 2	system 3
X1	sick	sick ✓	sick ✓	ok ✗
X2	sick ✗	ok ✗	ok ✗	sick ✓
X3	sick	ok ✗	sick ✓	sick ✓
X4	sick	ok ✗	sick ✓	sick ✓
X5	sick ✗	ok ✗	ok ✗	sick ✓
X6	ok	ok ✓	ok ✓	sick ✗
X7	ok	ok ✓	ok ✓	sick ✗
..	ok	ok ✓	ok ✓	ok ✓
..	ok	ok ✓	ok ✓	ok ✓
X500	ok	ok ✓	ok ✓	ok ✓
Accuracy		496/500 = 99%	498/500 = 99.6%	496/500 = 99.2%
Precision		1/1 = 100%	3/3 = 100%	4/6 = 66.7%
Recall		1/5 = 20%	3/5 = 60%	4/5 = 80%

Which system is better?



# A Single Measure

## ■ cannot take mean of P&R

- if R = 50% P = 50% M = 50%
- if R = 100% P = 10% M = 55% (not fair)

## 1. take harmonic mean

- which penalizes extreme values

$$HM = \frac{2}{\frac{1}{R} + \frac{1}{P}}$$

HM is high only when both P&R are high

if R = 50% and P = 50% HM = 50%

if R = 100% and P = 10% HM = 18.2%

## 2. if P and R should not have the same importance in the problem domain, take weighted harmonic mean

$$WHM = \frac{1}{\frac{1}{2} \frac{1}{R} + \frac{1}{2} \frac{1}{P}} \quad // \text{ if weight } R = \text{weight } P = \frac{1}{2}$$

$$\frac{1}{2} + \frac{1}{2} = 1$$

$$WHM = \frac{1}{\frac{1}{a} \frac{1}{R} + \frac{1}{b} \frac{1}{P}} \quad // \text{ if weight } R = \frac{1}{a} \text{ weight } P = \frac{1}{b} \text{ and } \frac{1}{a} + \frac{1}{b} = 1$$

# Weighted Harmonic Mean of P&R

$$WHM = \frac{1}{\frac{1}{a} \frac{1}{R} + \frac{1}{b} \frac{1}{P}} \quad // \text{ if weight } R = \frac{1}{a} \quad \text{weight } P = \frac{1}{b} \quad \text{and } \frac{1}{a} + \frac{1}{b} = 1$$

1. let  $w_R = \frac{\delta}{\delta+1}$   $w_P = \frac{1}{\delta+1}$  // so that  $w_R + w_P = \frac{\delta+1}{\delta+1} = 1$

$$WHM = \frac{1}{\left(\frac{\delta}{\delta+1}\right)\frac{1}{R} + \left(\frac{1}{\delta+1}\right)\frac{1}{P}} = \frac{\delta+1}{\delta\frac{1}{R} + 1\frac{1}{P}} = \frac{(\delta+1)PR}{\delta P + 1R}$$

2. let  $\delta = \beta^2$

$$WHM = \frac{(\beta^2+1)PR}{\beta^2 P + 1R} \quad // \text{ called the F-measure}$$

# F-measure

- A weighted harmonic mean of precision and recall

$$F_{\beta} = \frac{(\beta^2 + 1)PR}{(\beta^2 P + R)}$$

- $\beta$  represents the relative importance of recall to precision
  - when  $\beta = 1$ 
    - F1 measure
    - precision & recall have same importance
  - when  $\beta > 1$ 
    - recall is given more weight  $\beta = 2$
    - e.g.  $F_2$  measure, recall is considered 2x more important than precision
  - when  $\beta < 1$ 
    - precision is given more weight  $\beta = 0.5$
    - e.g.  $F_{0.5}$  measure, precision is considered 2x more important than recall

# Example

	Target	system 1	system 2	system 3
X1	sick	sick ✓	sick ✓	ok ✗
X2	sick	ok ✗	ok ✗	sick ✓
X3	sick	ok ✗	sick ✓	sick ✓
X4	sick	ok ✗	sick ✓	sick ✓
X5	sick	ok ✗	ok ✗	sick ✓
X6	ok	ok ✓	ok ✓	sick ✗
X7	ok	ok ✓	ok ✓	sick ✗
..	ok	ok ✓	ok ✓	ok ✓
..	ok	ok ✓	ok ✓	ok ✓
X500	ok	ok ✓	ok ✓	ok ✓
Accuracy ✓		496/500 = 99%	498/500 = 99.6%	498/500 = 99.6%
Precision ✓		1/1 = 100%	3/3 = 100%	4/6 = 66.7%
Recall ✓		1/5 = 20%	3/5 = 60%	4/5 = 80%
F1-measure 2PR/(P+R)		$\frac{2 \cdot 100 \cdot 20}{100 + 20} = 33\%$	75%	72.9%

$\beta = 1$   $F_1 = \frac{2PR}{P+R}$

---

# P, R and F for Multiclass Classification

- previous P, R and F are ok when 1 particular class interests us (eg. sick)
- What if several classes interest us?
- then
  - compute *per-class* P, R, F
  - and to have a single measure for all classes: combine per-class F-measures via
    - macro F-measure, or
    - weighted-average F-measure

# Per-class Precision & Per-class Recall

*gold*

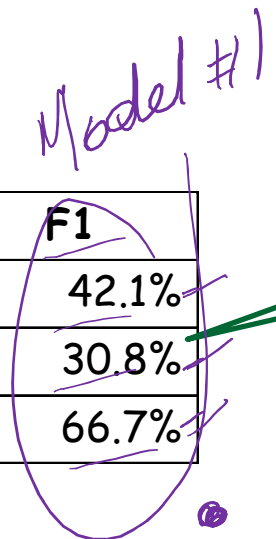
		Correct Class			
		Cat	Dog	Fish	Total
Predicted Class	Cat	4	6	3	13
	Dog	1	2	0	3
	Fish	1	2	6	9
	Total	6	10	9	25

*model prediction*

*images*

- precision of class Cat:  $4 / (4 + 6 + 3) = 30.8\%$
- precision of class Dog:  $2 / (1 + 2 + 0) = 66.7\%$
- precision of class Fish:  $6 / (1 + 2 + 6) = 66.7\%$
- recall of class Cat:  $4 / (4 + 1 + 1) = 66.7\%$
- recall of class Dog:  $2 / (2 + 6 + 2) = 20\%$
- recall of class Fish:  $6 / (3 + 0 + 6) = 66.7\%$

# Per-class F1-measure



	Precision	Recall	F1
<u>Cat</u>	30.8%	66.7%	42.1%
Dog	66.7%	20.0%	30.8%
Fish	66.7%	66.7%	66.7%


$$F1 = 2PR / (P + R)$$

$$F_{b=1}$$

- F1 of class Cat:  $(2 \times .308 \times .667) / (.308 + .667) = 0.421$
- F1 of class Dog:  $(2 \times .667 \times .200) / (.667 + .200) = 0.308$
- F1 of class Fish:  $(2 \times .667 \times .667) / (.667 + .667) = 0.667$

# Macro and Weighted-Average Measures

	Precision	Recall	F1
Cat	30.8%	66.7%	42.1%
Dog	66.7%	20.0%	30.8%
Fish	66.7%	66.7%	66.7%
<u>average</u>	$(30.8 + 66.7 + 66.7) / 3 = 54.7\%$	$(66.7 + 20.0 + 66.7) / 3 = 51.1\%$	$(42.1 + 30.8 + 66.7) / 3 = 46.5\%$
<u>weighted-average</u>	$(6 \times 30.8 // 6 \text{ cat} + 10 \times 66.7 // 10 \text{ dog} + 9 \times 66.7 // 9 \text{ fish}) / (6 + 10 + 9) // 25 \text{ samples} = 58.1\%$	$(6 \times 66.7 + 10 \times 20.0 + 6 \times 66.7) / 25 = 48.0\%$	$(6 \times 42.1 + 10 \times 30.8 + 6 \times 66.7) / 25 = 46.4\%$

■ To combine measures into a single one, we can:

□ take simple average

■ --> macro precision, macro recall, macro F1

□ take weighted average

■ ie. weight the average based on the nb of samples from each class

■ --> weighted averaged precision, weighted averaged recall, weighted averaged F1

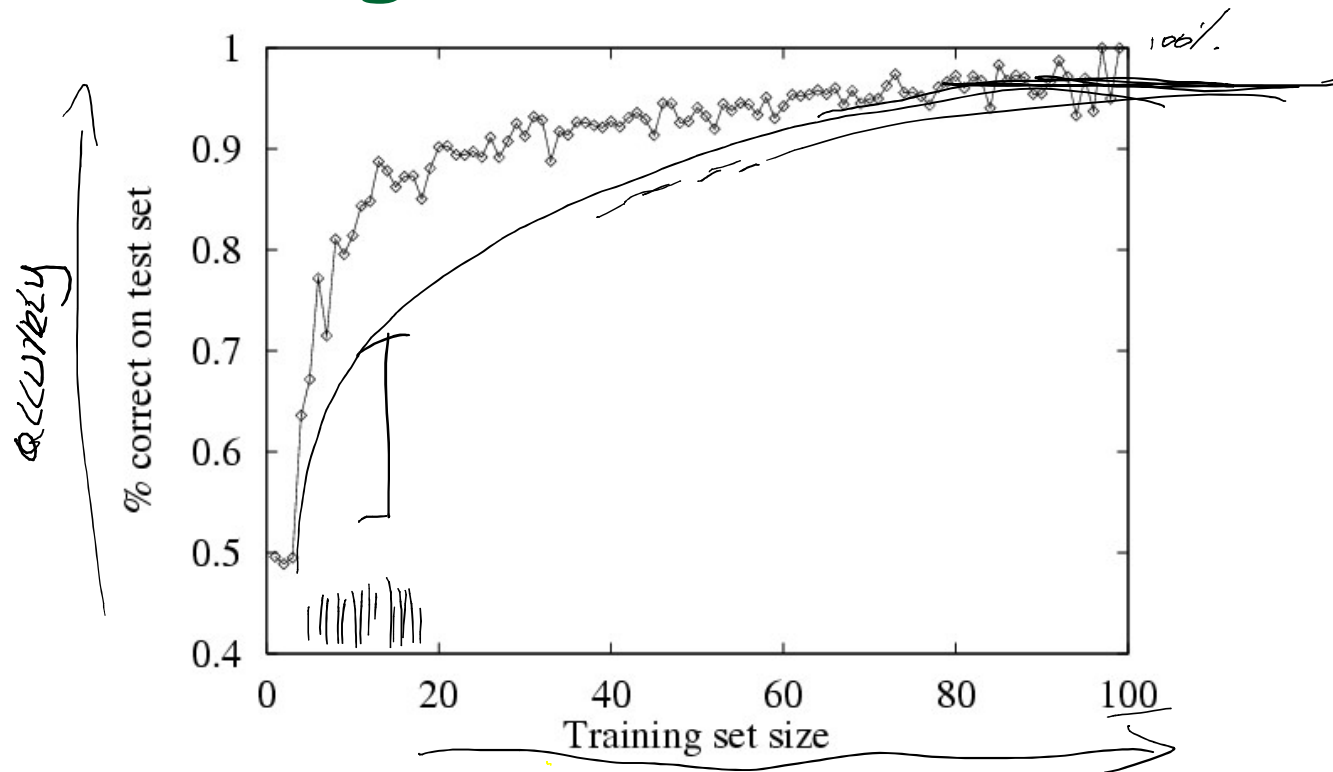


# Confusion Matrix

- to do an error analysis and find out where the model went wrong ?
- aka contingency table
- eg. 6 classes, 100 test instances

		Correct Class						
		C1	C2	C3	C4	C5	C6	Total
Predicted Class	C1	10✓	3×	0	0	3×	0	16
	C2	0	12✓	3×	4×	0	0	19
	C3	0	1×	9✓	2×	1×	2×	15
	C4	0	1×	3×	5✓	2×	0	11
	C5	0	0	3×	2×	10✓	3×	18
	C6	0	0	5×	0	5×	11✓	21
	Total	10	17	23	13	21	16	100

# A Learning Curve



- Size of training set
  - the more, the better
  - but after a while, not much improvement...

---

# Some Words on Training

- In all types of learning... watch out for:
  - Noisy input
  - Overfitting/underfitting the training data

# Noisy Input

- In all types of learning... watch out for:
  - **Noisy Input:**
    - Two examples have the same feature-value pairs, but different outputs

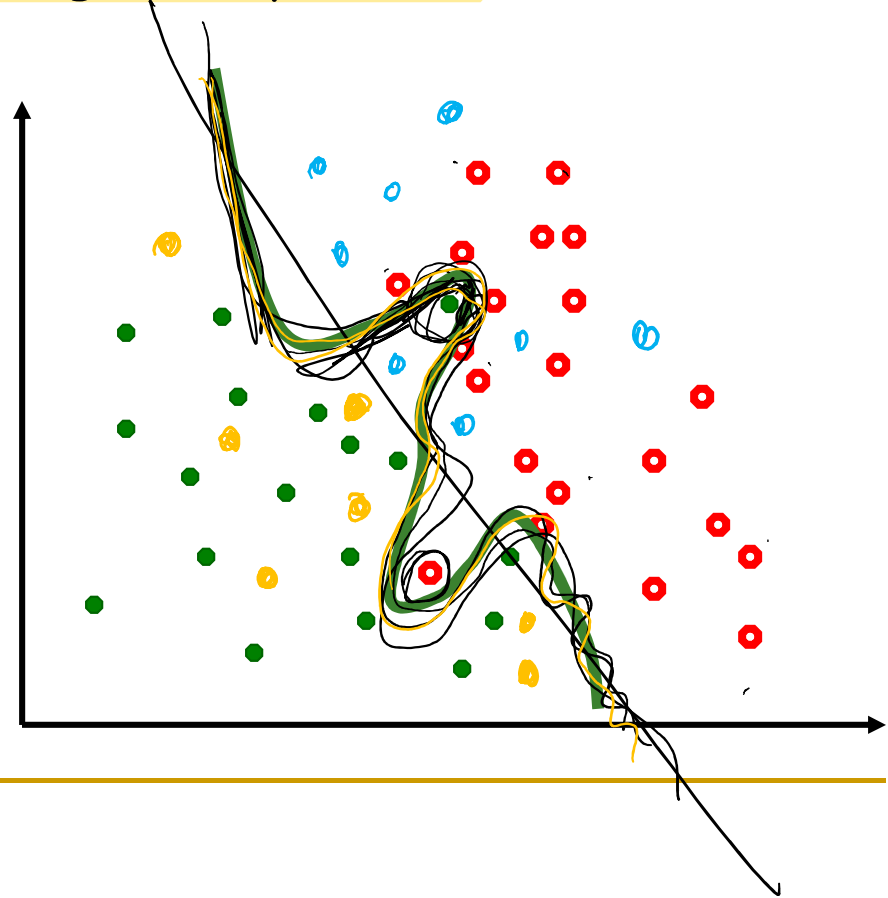
Size	Color	Shape	Output
Big	Red	Circle	+
Big	Red	Circle	+

- Some values of features are incorrect or missing (ex. errors in the data acquisition)
- Some relevant attributes are not taken into account in the data set

# Overfitting


- If a large number of irrelevant features are there, we may find meaningless regularities in the data that are particular to the training data but irrelevant to the general problem.
- Complicated boundaries *overfit* the data
- they are too tuned to the particular training data at hand
- They do not *generalize* well to the new data
- Extreme case: "rote learning"

- Training error is low
- Testing error is high



# Underfitting

- We can also underfit data, i.e. find a decision boundary that is too simple
- Model is not expressive enough (not enough features, or not enough capacity)
- eg. There is no way to fit a linear decision boundary so that the training examples are well separated



- Training error is high
- Testing error is high

# Cross-validation

## ■ K-fold cross-validation

- run k experiments, each time you test on 1/k of the data, and train on the rest
- then you average the results

## ■ ex: 10-fold cross validation

1. Collect a large set of examples (all with correct classifications)
2. Divide collection into two disjoint sets: training (90%) and test (10% = 1/k)
3. Apply learning algorithm to training set
4. Measure performance with the test set
5. Repeat steps 2-4, with the 10 different portions
6. Average the results of the 10 experiments

entire DS

exp 10

exp1:	train								test
exp2:	train								test
exp3:	train							test	train
...	...								

F-measure 8/60

F-measure 8/60

F - 83 99

F - 82 50

stable unstable

average

# Today

1. Introduction to ML ✓
2. Naïve Bayes Classification ✓
  - a. Application to Spam Filtering ✓
3. Decision Trees ✓
4. Evaluation ✓
5. Unsupervised Learning ✗
6. Neural Networks
  - a. Perceptrons
  - b. Multi Layered Neural Networks



---

# Up Next

1. Introduction to ML
2. Naive Bayes Classification
  - a. Application to Spam Filtering
3. Decision Trees
4. ( Evaluation
5. Unsupervised Learning )
6. Neural Networks
  - a. Perceptrons
  - b. Multi Layered Neural Networks