

COMP 472: Artificial Intelligence

Machine Learning

Neural Networks

part 2
video #8

- Russell & Norvig: Sections 19.6, 21.1

Today

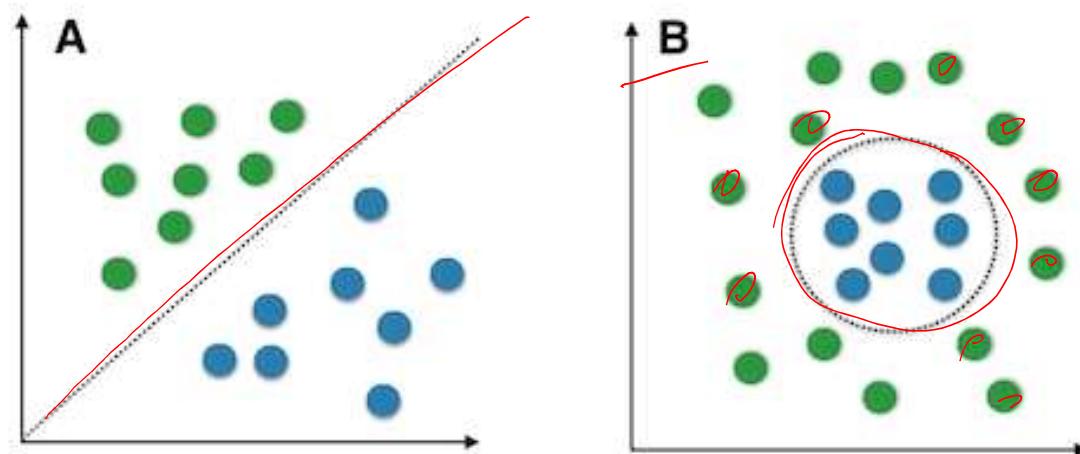
1. Introduction to ML
2. Naïve Bayes Classification
 - a. Application to Spam Filtering
3. Decision Trees
4. (Evaluation
5. Unsupervised Learning)
6. Neural Networks
 - a. Perceptrons
 - b. Multi Layered Neural Networks



Limits of the Perceptron

- can only model linear decision boundaries
- but real-world problems cannot always be represented by linearly-separable functions...

Linear vs. nonlinear problems



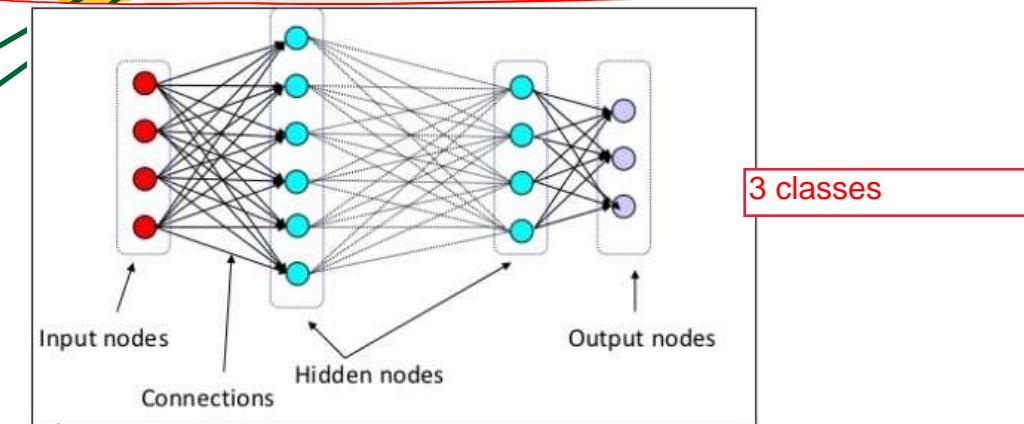
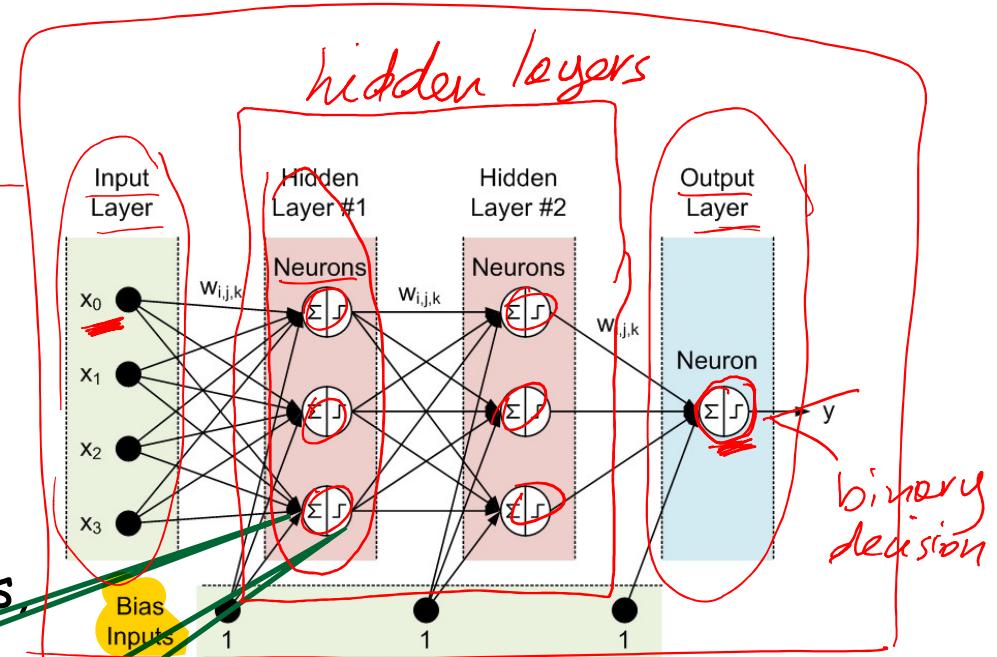
Multilayer Neural Networks

- Solution:

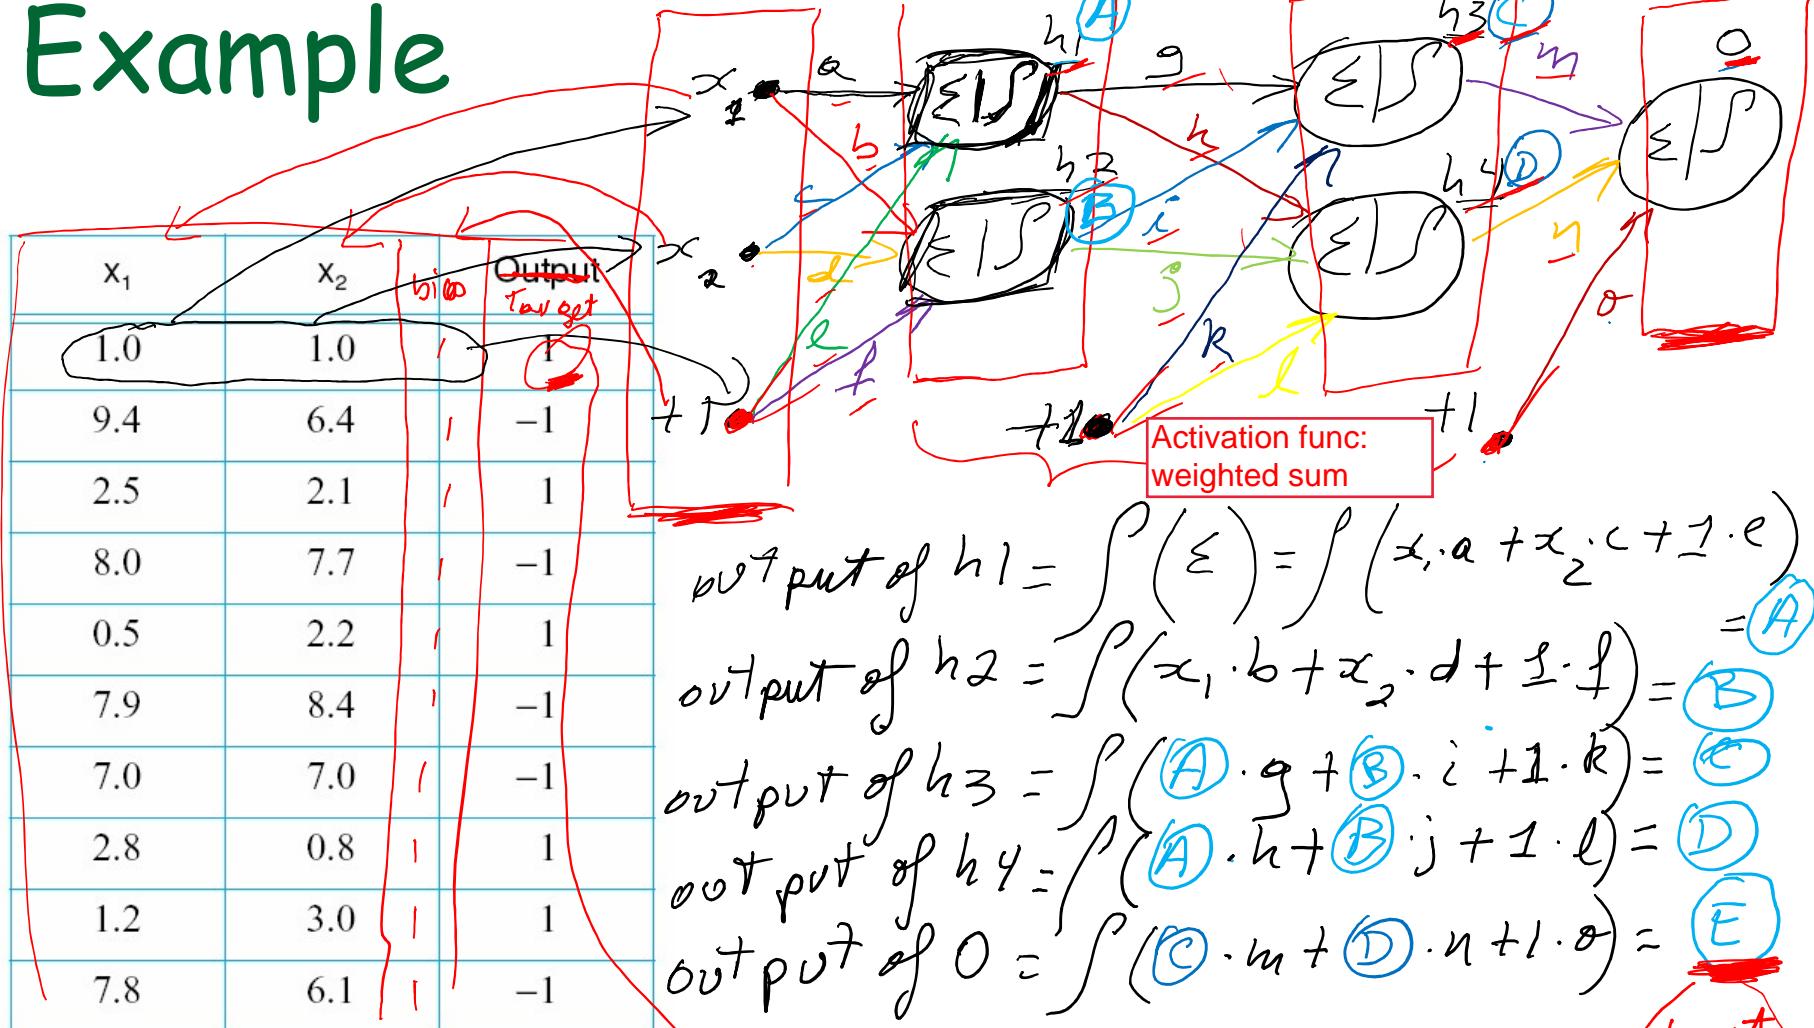
- use a non-linear activation function
- to learn more complex functions (more complex decision boundaries), have hidden nodes
- and for non-binary decisions, have multiple output nodes

usual transfer function

Non-linear, differentiable activation function



Example



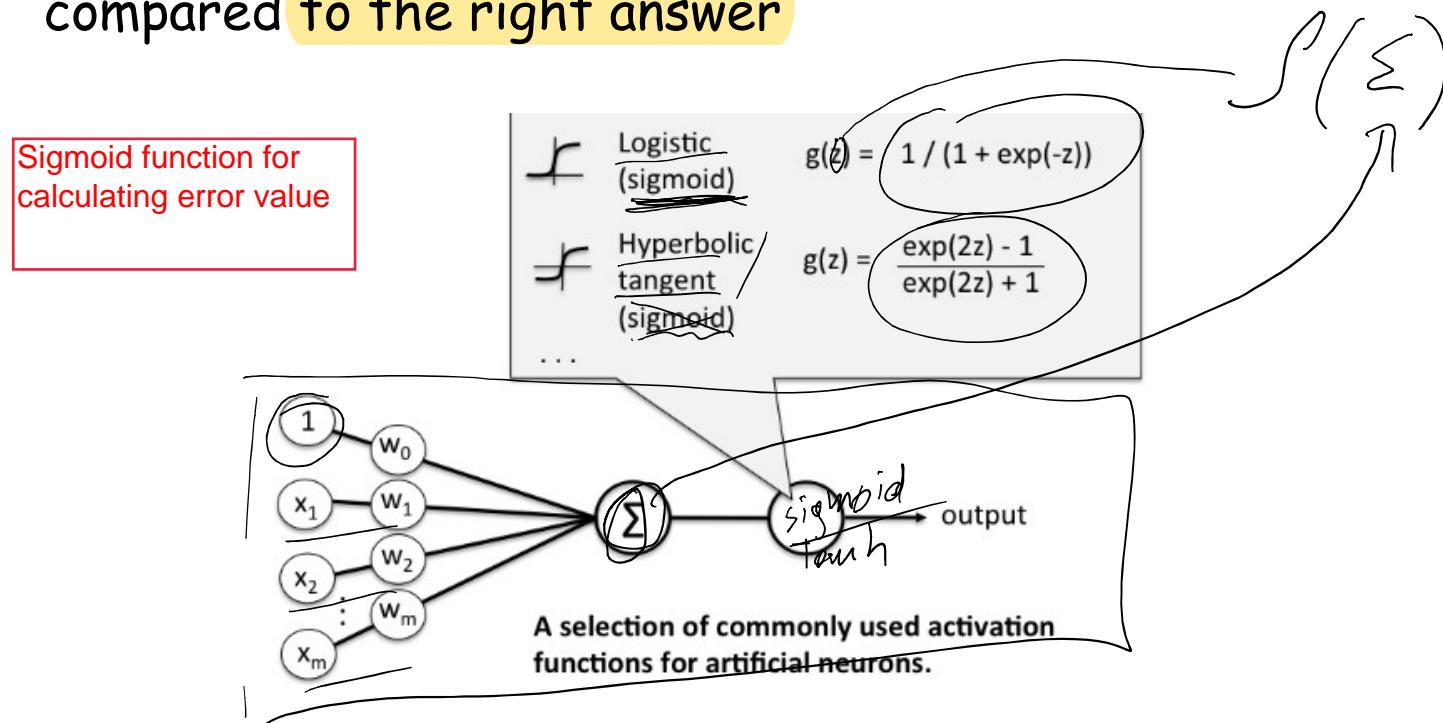
feed forward

output
of the
network

backprop

Activation Functions

- Backpropagation requires a differentiable activation function
- that returns continuous values within a range
 - e.g. a value between 0 and 1 (instead of 0 or 1 , like the perceptron)
- indicates how close/how far the output of the network is compared to the right answer



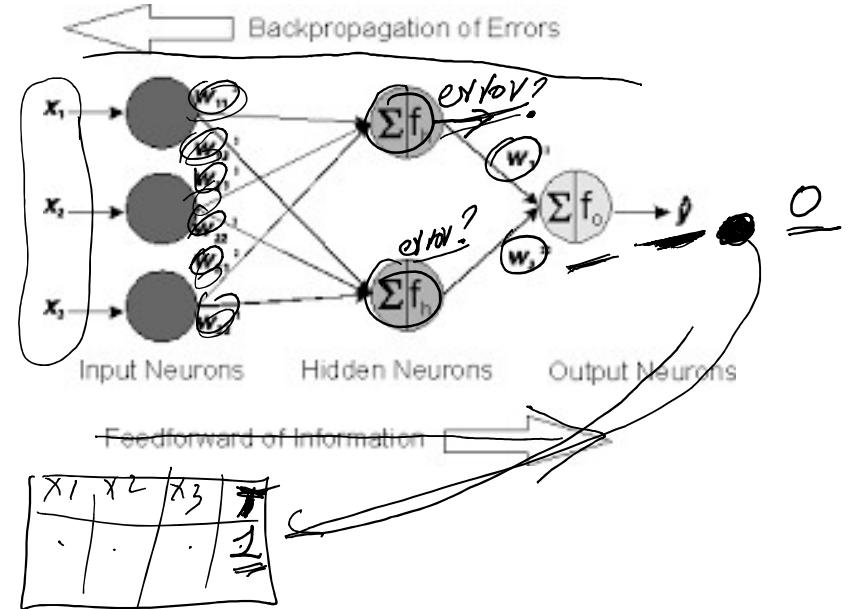
<https://medium.com/towards-data-science/activation-functions-and-its-types-which-is-better-a9a5310cc8f>

Learning in a Neural Network

- Learning is the same as in a perceptron:

1. Feed-forward:

- Input from the features is fed forward in the network from input layer towards the output layer



2. Backpropagation:

- Error rate flows backwards from the output layer to the input layer (to adjust the weights in order to minimize the error)

3. Iterate until error rate is minimized

- repeat the forward pass and back pass for the next data points until all data points are examined (1 epoch)
- repeat this entire exercise (several epochs) until the overall error is minimised

Typical Cost Functions

- Error of the network is computed via a cost function

1. Quadratic Cost:

- aka mean squared error (MSR)
- minimize the difference between output values and target values

$$C = \left(\frac{1}{n} \sum_{i=1}^n (T_i - O_i)^2 \right)$$

- used mostly for regression tasks

where n = nb of instances

$$\begin{aligned} i &= 1 & (T_1 - O_1)^2 \\ i &= 2 & (T_2 - O_2)^2 \\ i &= 3 & (T_3 - O_3)^2 \\ &\vdots & \sum \end{aligned}$$

2. Cross-entropy:

- used mostly for classification tasks
- where we don't care about the exact value of the network output, we only care about the final class

$$C = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K (T_{ik} \log(O_{ik}))$$

where:

n = nb of instances

K = nb of classes

\log = base 2, base 10, ln,

Example of Cross Entropy

- Assume 3 classes: red, blue and green

- i.e. $K = 3$

- for a specific instance i , the target is green

- i.e. $(\text{red}, \text{blue}, \text{green}) = (0, 0, 1)$ // real distribution

- $T^1 = 0$ $T^2 = 0$ $T^3 = 1$

- but the model predicts the probabilities as:

- $(\text{red}, \text{blue}, \text{green}) = (0.1, 0.4, 0.5)$ // predicted distribution

- $O^1 = 0.1$ $O^2 = 0.4$ $O^3 = 0.5$

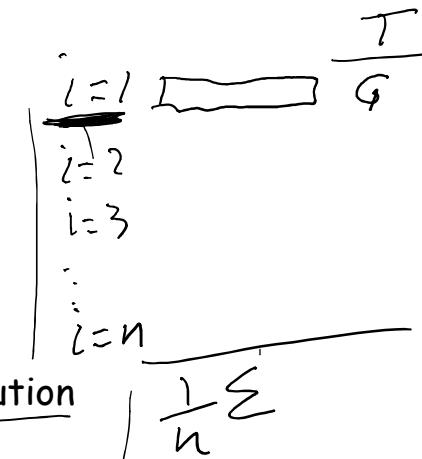
- the cross entropy of two distributions (real and predicted)

- $\sum_{k=1}^K (T^k \ln(O^k)) = -((0) \ln(0.1) + (0) \ln(0.4) + (1) \ln(0.5))$

- but we have several instances in the test set, so we will take the average of the cross-entropy across all instances i in the test set

$$C = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K (T_i^k \ln(O_i^k))$$

- see an example calculation computation in a few slides



Backpropagation

- In a multilayer network...

- Computing the error in the output layer is clear.
 - Computing the error in the hidden layers is not clear, because we don't know what output it should be

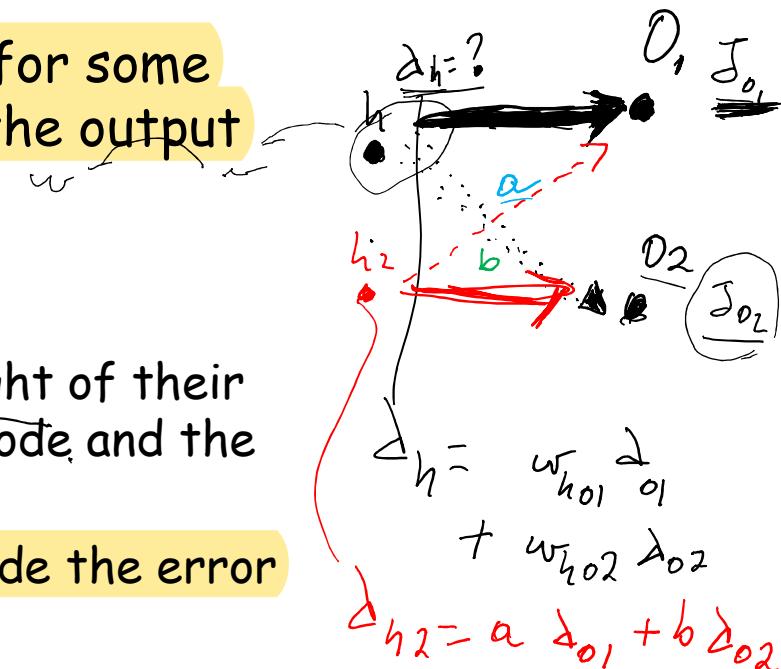
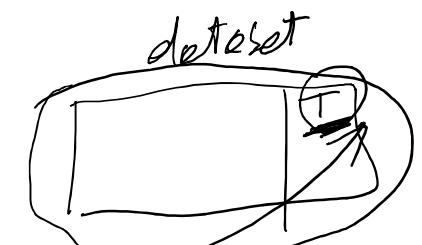
- Intuitively:

- A hidden node h is "responsible" for some fraction of the error in each of the output node to which it connects.

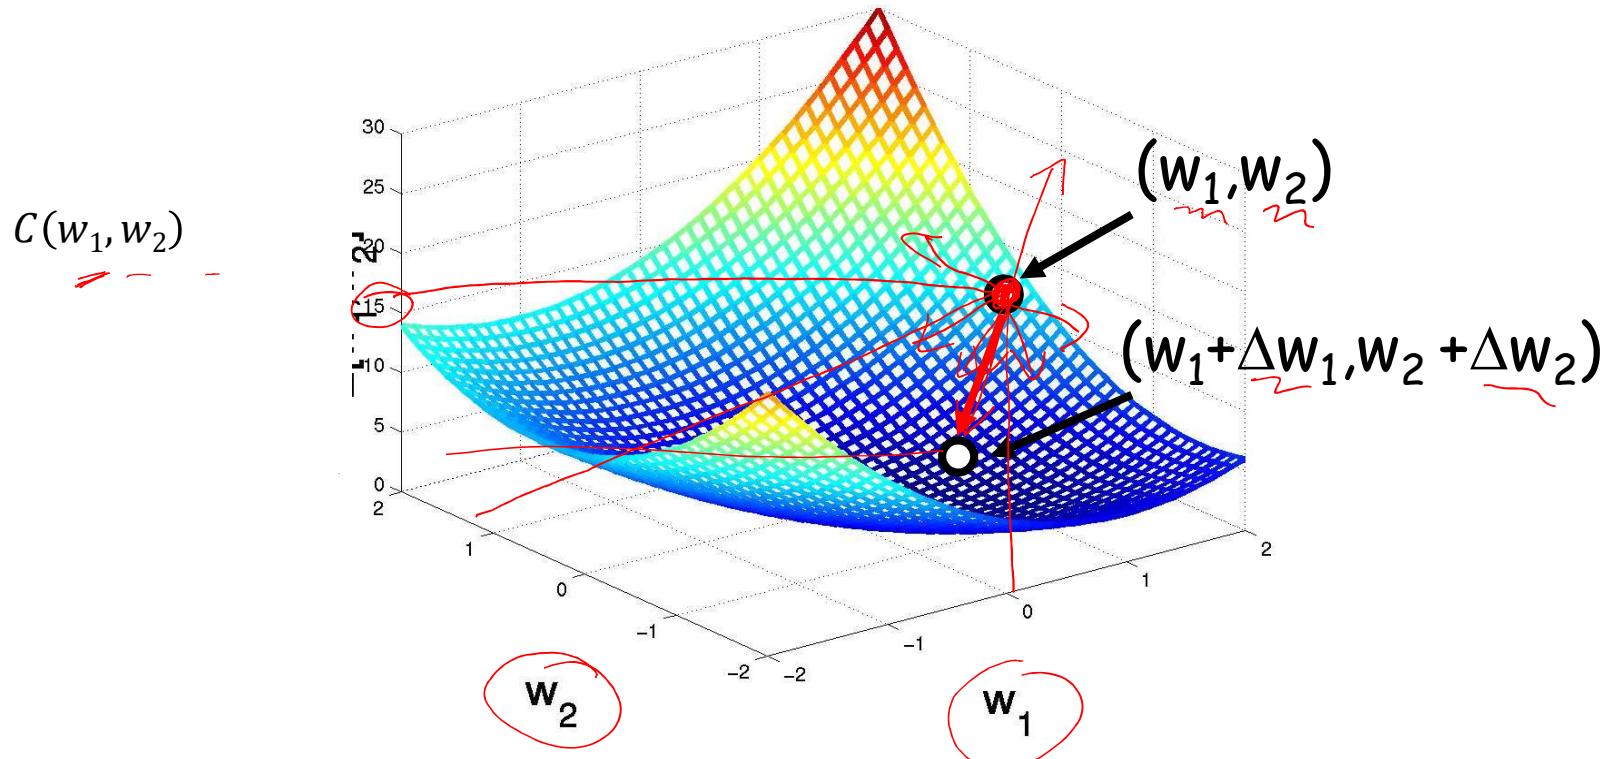
- So the error values (δ):

- are divided according to the weight of their connection between the hidden node and the output node

- and are propagated back to provide the error values (δ) for the hidden layer.



Gradient Descent - Adjusting the Weights

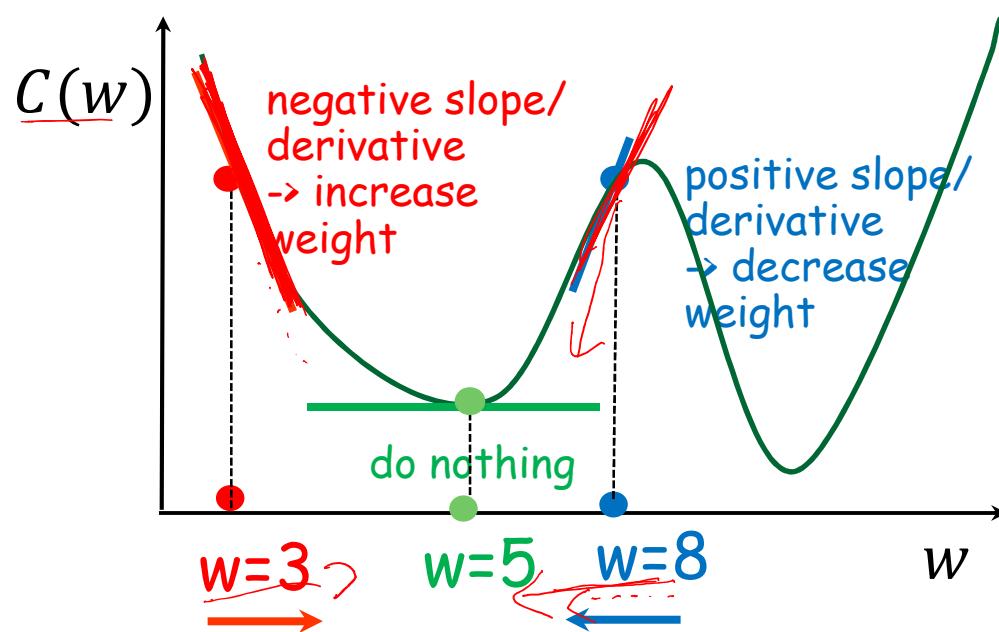


- Goal: minimize $C(w_1, w_2)$ by changing w_1 and w_2
- What is the best combination of change in w_1 and w_2 to minimize C faster?

Gradients

Gradient is just derivative in 1D

Eg: $C(w) = (w - 5)^2$ derivative is:
only 1 weight



$$\frac{\partial C}{\partial w} = 2(w - 5)$$

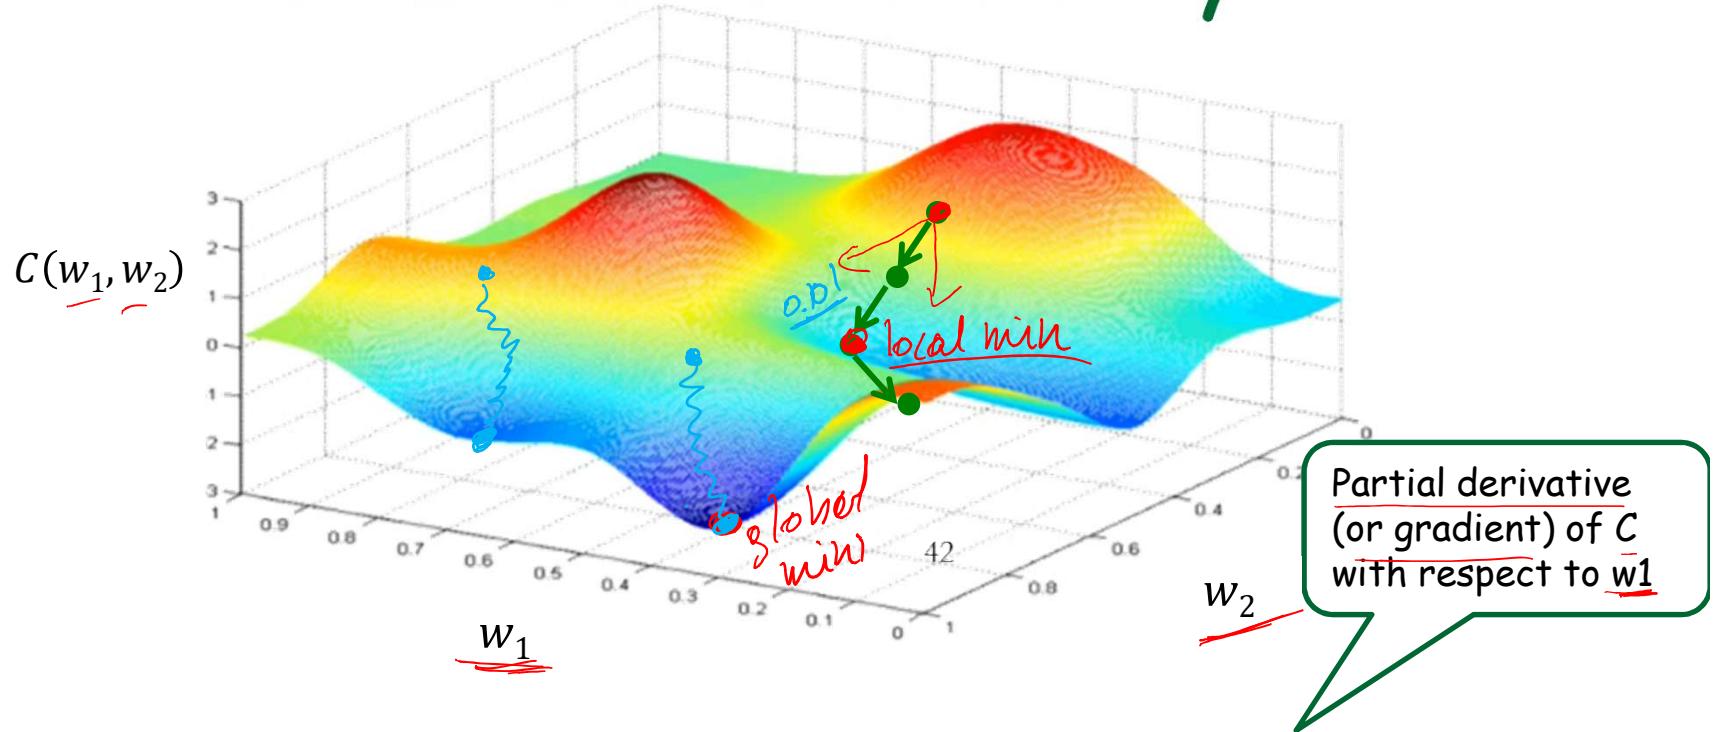
If $w=3$ $\frac{\partial C}{\partial w}(3) = 2(3 - 5) = -4$

derivative says increase w
(go in opposite direction
of derivative)

If $w=8$ $\frac{\partial C}{\partial w}(8) = 2(8 - 5) = 6$

derivative says decrease w
(go in opposite direction
of derivative)

Gradient Descent Visually



- need to know how much a change in w_1 will affect $C(w_1, w_2)$ i.e $\frac{\partial C}{\partial w_1}$
- need to know how much a change in w_2 will affect $C(w_1, w_2)$ i.e $\frac{\partial C}{\partial w_2}$
- Gradient ∇C points in the opposite direction of steepest decrease of $C(w_1, w_2)$
- i.e. hill-climbing approach...

Training the Network

After some calculus (see: <https://en.wikipedia.org/wiki/Backpropagation> we get...)

- Step 0: Initialise the weights of the network randomly

// feedforward

- Step 1: Do a forward pass through the network

$$O_i = g\left(\sum_j w_{ji} x_j\right) = \text{sigmoid}\left(\sum_j w_{ji} x_j\right) = \frac{1}{1 + e^{-\left(\sum_j w_{ji} x_j\right)}}$$

// propagate the errors backwards

- Step 2: For each output unit k , calculate its error term δ_k

$$\delta_k \leftarrow g'(x_k) \times Err_k = O_k(1 - O_k) \times (O_k - T_k)$$

- Step 3: For each hidden unit h , calculate its error term δ_h

$$\delta_h \leftarrow g'(x_h) \times Err_h = O_h(1 - O_h) \times \sum_{k \in \text{outputs}} w_{hk} \delta_k$$

- Step 4: Update each network weight w_{ij} :

$$w_{ij} \leftarrow w_{ij} \pm \Delta w_{ij} \text{ where } \Delta w_{ij} = -n_s \delta_j O_i$$

- Repeat steps 1 to 4 until the cost (C) is minimised

Note: To be consistent with Wikipedia, we'll use O-T instead of T-O, but we will subtract the error in the weight update

Derivative of sigmoid

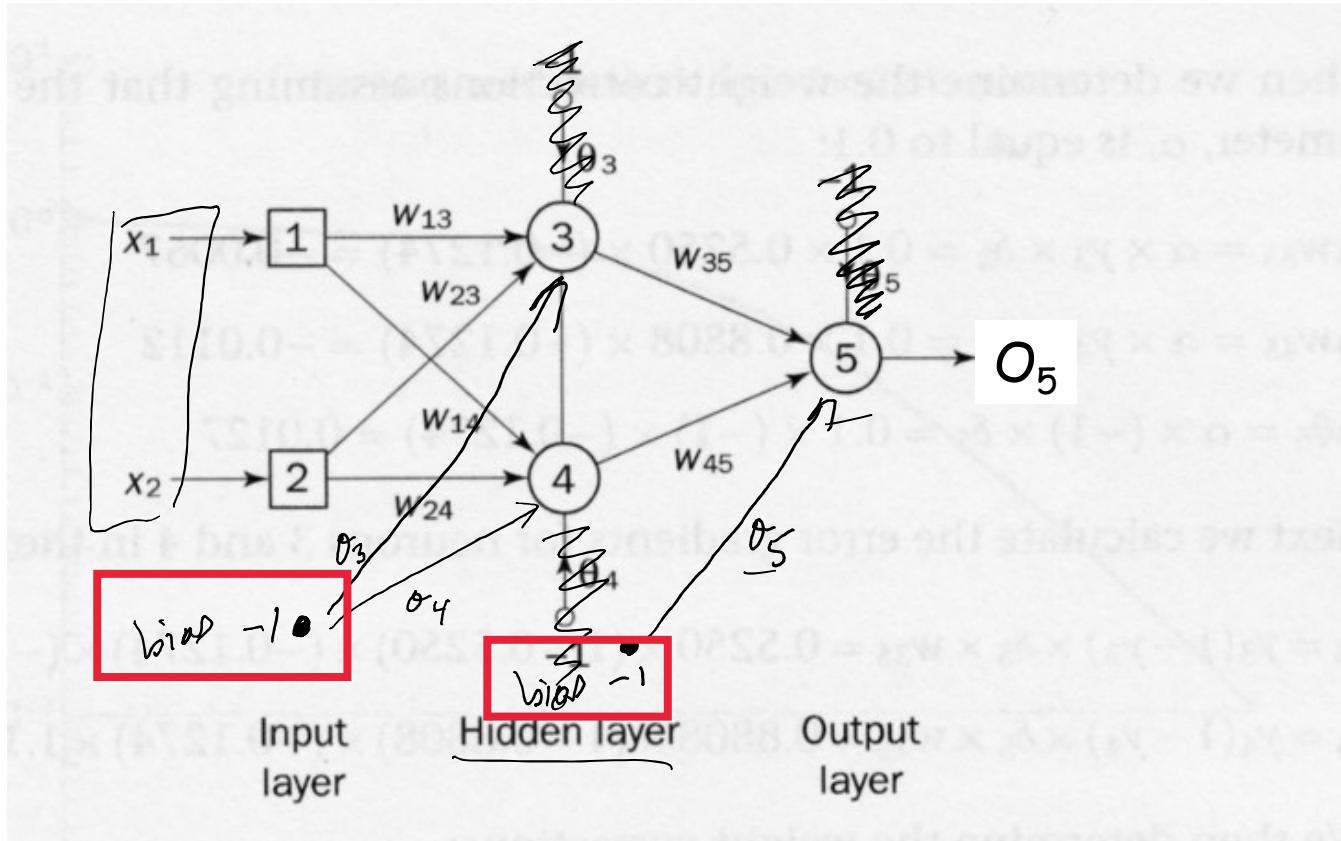
note, if we use $g = \text{sigmoid}$:
 $g'(x) = g(x)(1 - g(x))$

Sum of the weighted error term of the output nodes that h is connected to (ie. h contributed to the errors δ_k)

$\equiv x_i$ in slide #22

Example: XOR

$x_1 \text{ XOR } x_2$

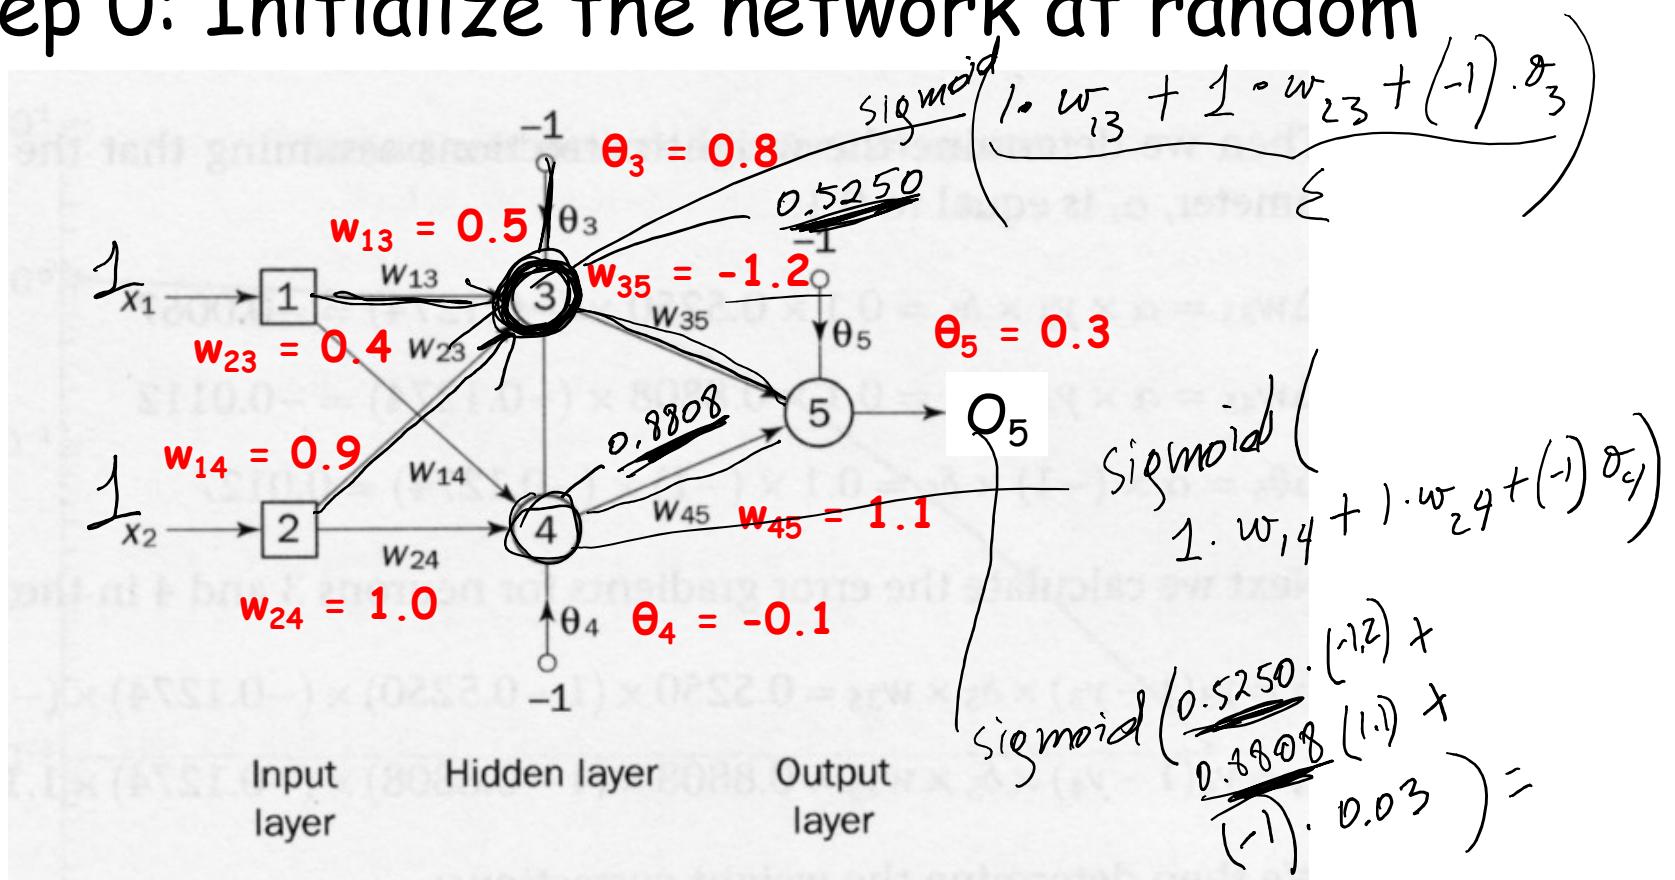


- 2 input nodes + 2 hidden nodes + 1 output node + 3 biases

source: Negnevitsky, Artificial Intelligence, p. 181

Example: Step 0 (initialization)

- Step 0: Initialize the network at random



Step 1: Feed Forward

- Step 1: Feed the inputs and calculate the output

$$O_i = \text{sigmoid} \left(\sum_j w_{ji} x_j \right) = \frac{1}{1 + e^{-\sum_j w_{ji} x_j}}$$

- With $(x_1=1, x_2=1)$ as input:

x_1	x_2	Target output T
1	1	0
0	0	0
1	0	1
0	1	1

- Output of the hidden node 3:

□ $O_3 = \text{sigmoid}(x_1 w_{13} + x_2 w_{23} - \theta_3) = 1 / (1 + e^{-(1 \times 5 + 1 \times 4 - 1 \times 8)}) = 0.5250 \checkmark$

- Output of the hidden node 4:

□ $O_4 = \text{sigmoid}(x_1 w_{14} + x_2 w_{24} - \theta_4) = 1 / (1 + e^{-(1 \times 9 + 1 \times 1.0 + 1 \times 0.1)}) = 0.8808$

- Output of neuron 5:

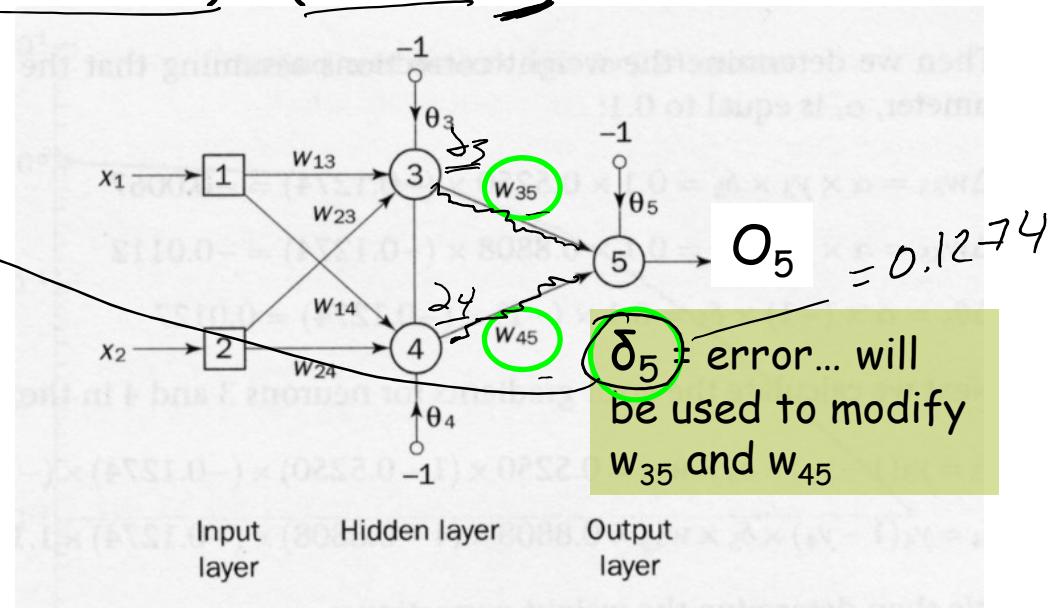
□ $O_5 = \text{sigmoid}(O_3 w_{35} + O_4 w_{45} - \theta_5) = 1 / (1 + e^{-(0.5250 \times 1.2 + 0.8808 \times 1.1 - 1 \times 0.3)}) = 0.5097$

Step 2: Calculate error term of output layer

$$\delta_k \leftarrow g'(x_k) \times \text{Err}_k = O_k(1 - O_k) \times (O_k - T_k)$$

- Error term of neuron 5 in the output layer:

□ $\delta_5 = O_5 (1 - O_5) (O_5 - T_5)$
= $(0.5097) \times (1 - 0.5097) \times (0.5097 - 0)$
= 0.1274



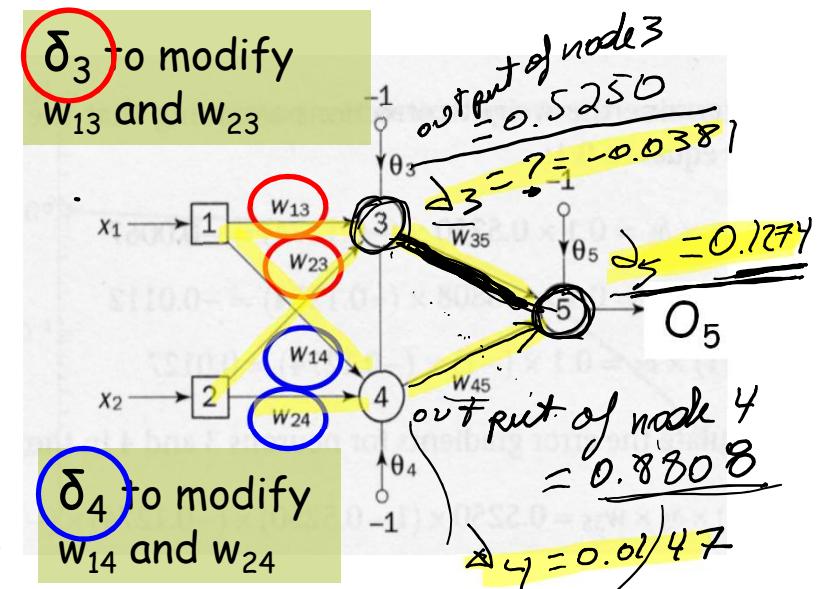
Step 3: Calculate error term of hidden layer

$$\delta_h \leftarrow g'(x_h) \times Err_h = O_k(1 - O_k) \times \sum_{k \in \text{outputs}} w_{kh} \delta_k$$

- Error term of neurons 3 & 4 in the hidden layer:

□ $\delta_3 = O_3(1-O_3) \delta_5 w_{35}$
 $= (0.5250) \times (1-0.5250) \times (0.1274) \times (-1.2)^{35}$
 $= -0.0381$

□ $\delta_4 = O_4(1-O_4) \delta_5 w_{45}$
 $= (0.8808) \times (1-0.8808) \times (0.1274) \times (1.1)$
 $= 0.0147$

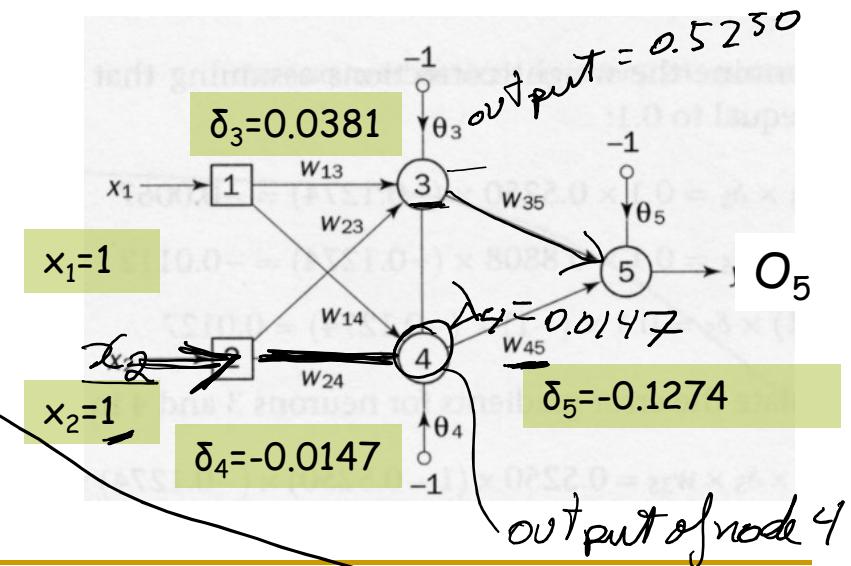


Step 4: Update Weights

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij} \text{ where } \Delta w_{ij} = -\eta \delta_j x_i$$

- Update all weights (assume a constant learning rate $\eta = 0.1$)

- $\Delta w_{13} = -\eta \delta_3 x_1 = -0.1 \times -0.0381 \times 1 = 0.0038$
- $\Delta w_{14} = -\eta \delta_4 x_1 = -0.1 \times 0.0147 \times 1 = -0.0015$
- $\Delta w_{23} = -\eta \delta_3 x_2 = -0.1 \times -0.0381 \times 1 = 0.0038$
- $\Delta w_{24} = -\eta \delta_4 x_2 = -0.1 \times 0.0147 \times 1 = -0.0015$
- $\Delta w_{35} = -\eta \delta_5 O_3 = -0.1 \times 0.1274 \times 0.5250 = -0.00669 // O_3 \text{ is seen as } x_5 \text{ (output of 3 is input to 5)}$
- $\Delta w_{45} = -\eta \delta_5 O_4 = -0.1 \times 0.1274 \times 0.8808 = -0.01122 // O_4 \text{ is seen as } x_5 \text{ (output of 4 is input to 5)}$
- $\Delta \theta_3 = -\eta \delta_3 (-1) = -0.1 \times -0.0381 \times -1 = -0.0038$
- $\Delta \theta_4 = -\eta \delta_4 (-1) = -0.1 \times 0.0147 \times -1 = -0.0015$
- $\Delta \theta_5 = -\eta \delta_5 (-1) = -0.1 \times 0.1274 \times -1 = -0.0127$

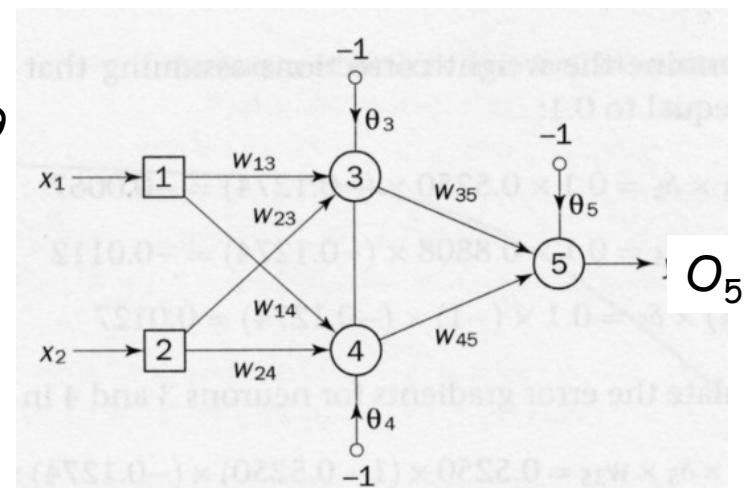


Step 4: Update Weights (con't)

$$w_{ij} \leftarrow w_{ij} + \Delta w_{ij} \text{ where } \Delta w_{ij} = -\eta \delta_j x_i$$

- Update all weights (assume a constant learning rate $\eta = 0.1$)

- $w_{13} = w_{13} + \Delta w_{13} = 0.5 + 0.0038 = 0.5038$
- $w_{14} = w_{14} + \Delta w_{14} = 0.9 - 0.0015 = 0.8985$
- $w_{23} = w_{23} + \Delta w_{23} = 0.4 + 0.0038 = 0.4038$
- $w_{24} = w_{24} + \Delta w_{24} = 1.0 - 0.0015 = 0.9985$
- $w_{35} = w_{35} + \Delta w_{35} = -1.2 - 0.00669 = -1.20669$
- $w_{45} = w_{45} + \Delta w_{45} = 1.1 - 0.01122 = 1.08878$
- $\theta_3 = \theta_3 + \Delta \theta_3 = 0.8 - 0.0038 = 0.7962$
- $\theta_4 = \theta_4 + \Delta \theta_4 = -0.1 + 0.0015 = -0.0985$
- $\theta_5 = \theta_5 + \Delta \theta_5 = 0.3 + 0.0127 = 0.3127$



Step 4: Iterate through data

- after adjusting all the weights, repeat the forward pass and back pass for the next data point until all data points are examined
- repeat this entire exercise until the cost function is minimised

□ Eg. $\underline{C} = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K (T^k_i \ln(O^k_i))$

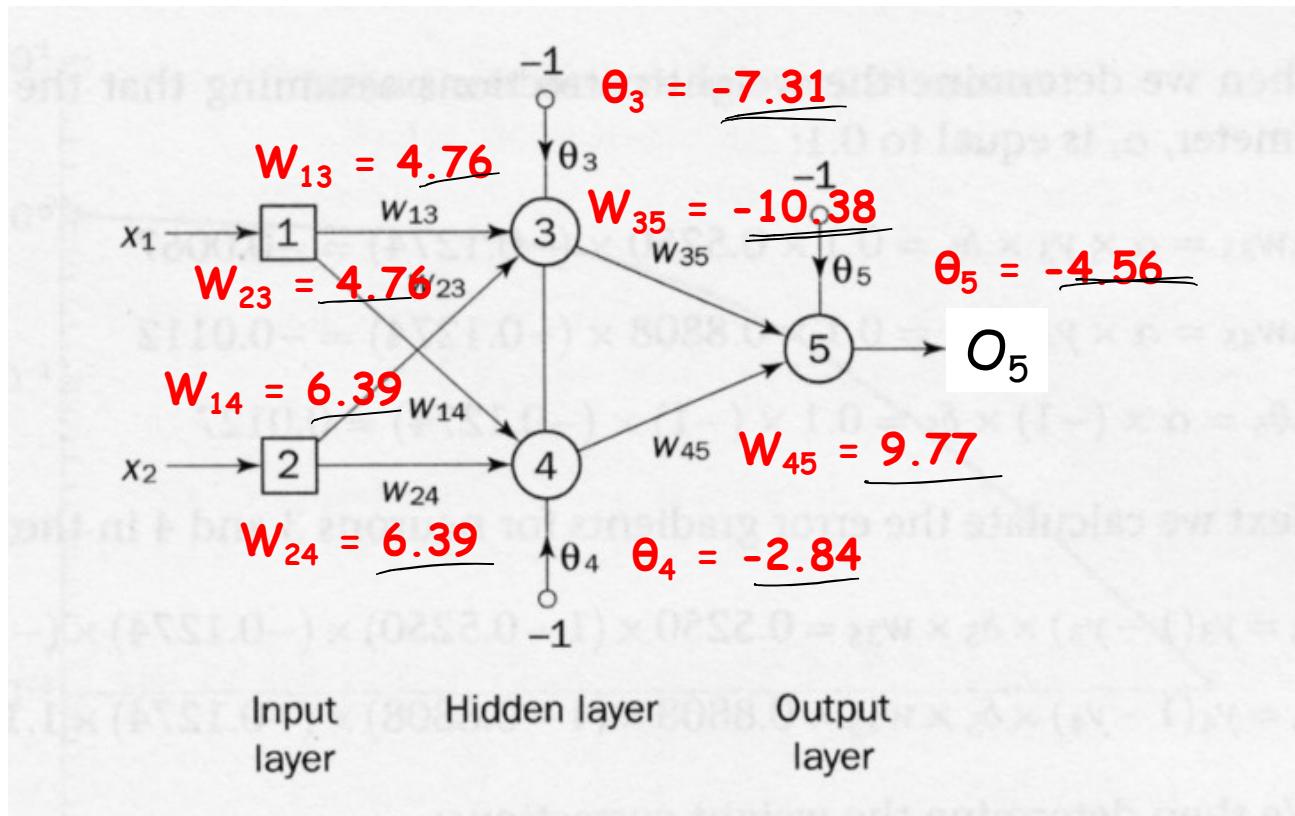
where

n = nb of training examples

K = nb of classes

The Result...

- After 224 epochs, we get:
 - (1 epoch = going through all data once)



Error is minimized

1-0.9845

Inputs		Target Output T	Actual Output O
x_1	x_2		
1	1	false (0)	0.0155
0	1	true (1)	0.9849
1	0	true (1)	0.9849
0	0	false (0)	0.0175

i	real distribution (false, true)	predicted distribution (false, true)
1	(1, 0)	0.9845, 0.0155
2	(0, 1)	0.0151, 0.9849
3	(0, 1)	(0.0151, 0.9849)
4	(1, 0)	(0.9825, 0.0175)

$$\begin{aligned}
 C = -\frac{1}{n} \sum_{i=1}^n \sum_{k=1}^K (T^k_i \ln(O^k_i)) &= -\frac{1}{n} \\
 &\quad (1) \ln(0.9845) + (0) \ln(0.0155) // \text{for } i=1 \\
 &\quad + (0) \ln(0.0151) + (1) \ln(0.9849) // \text{for } i=2 \\
 &\quad + (0) \ln(0.0151) + (1) \ln(0.9849) // \text{for } i=3 \\
 &\quad + (1) \ln(0.9825) + (0) \ln(0.0175) // \text{for } i=4 \\
 &= 0.01592 < \epsilon
 \end{aligned}$$

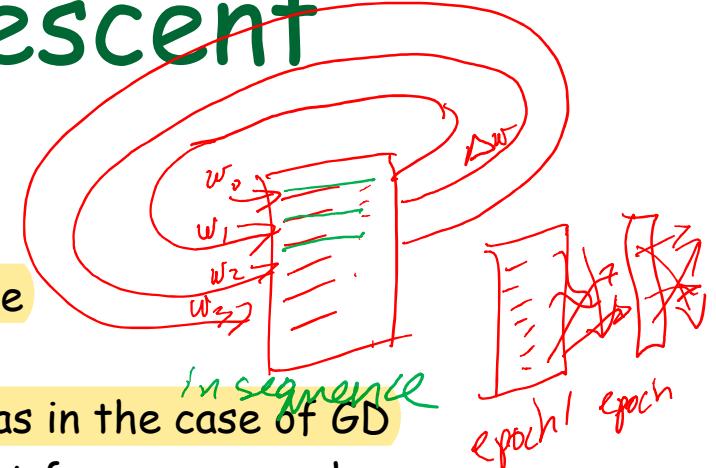


May be a local minimum...

Types of Gradient Descent

■ Stochastic Gradient Descent (SGD)

- updates the weights after each training example
- often converges faster compared to GD
- but the error function is not as well minimized as in the case of GD
- to obtain better results, shuffle the training set for every epoch



■ Batch Gradient Descent (GD)

- updates the weights after 1 epoch
- can be costly (time & memory) since we need to evaluate the whole training dataset before we take one step towards the minimum.



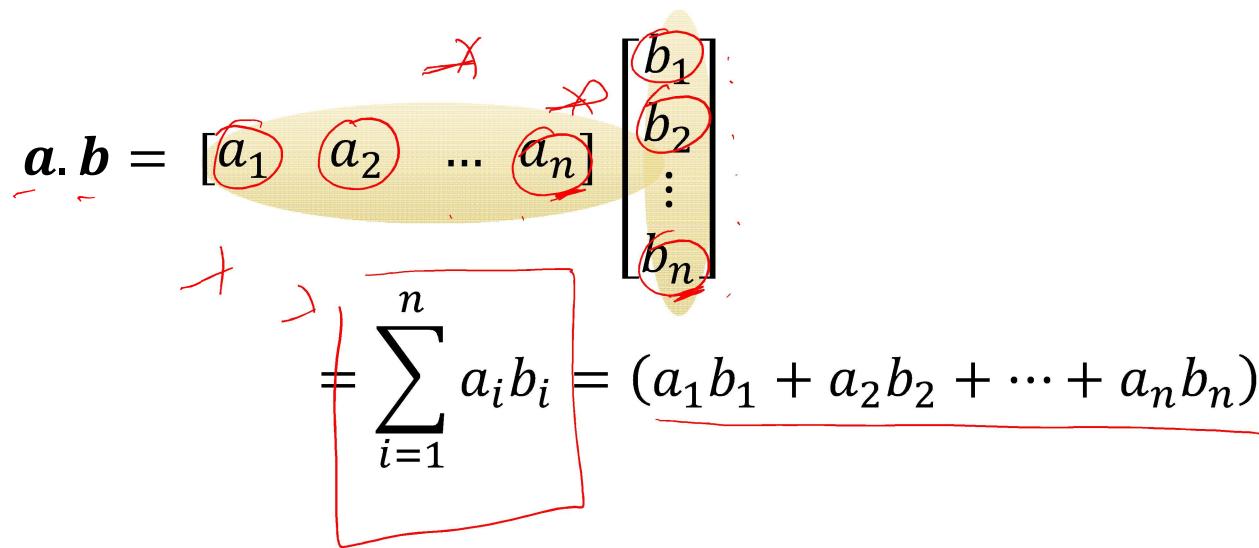
■ MiniBatch Gradient Descent:

- compromise between GD and SGD
- cut your dataset into sections, and update the weights after training on each section

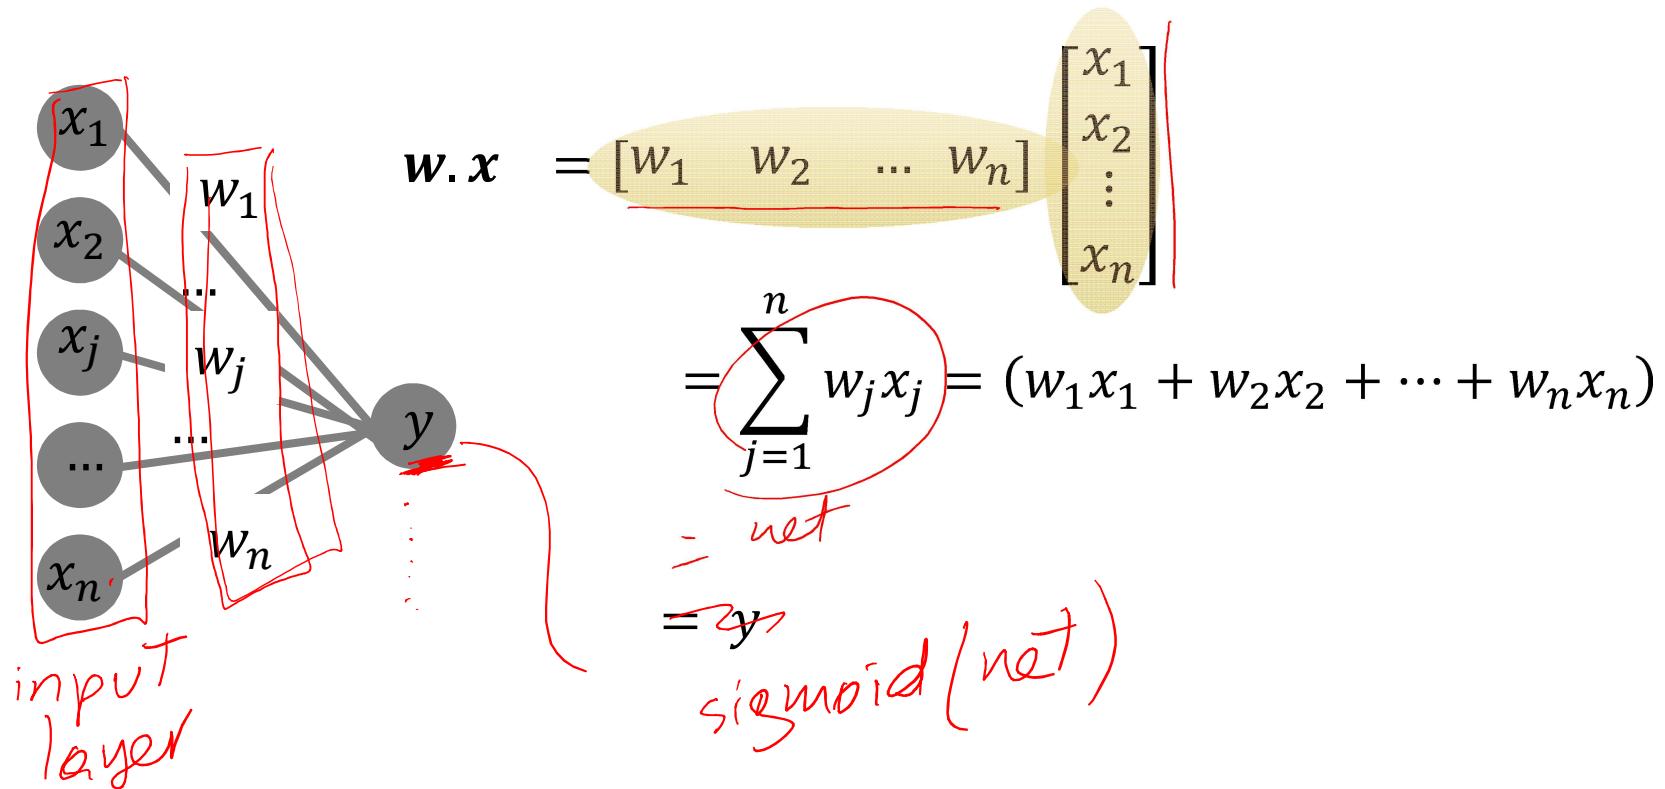


Remember your Linear Algebra

- Dot product (inner product) of 2 vectors

$$\begin{aligned} \mathbf{a} \cdot \mathbf{b} &= [a_1 \ a_2 \ \dots \ a_n] \begin{matrix} b_1 \\ b_2 \\ \vdots \\ b_n \end{matrix} \\ &= \sum_{i=1}^n a_i b_i = (a_1 b_1 + a_2 b_2 + \dots + a_n b_n) \end{aligned}$$


so what?



Remember your Linear Algebra

- matrix-vector product

$$\begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} \quad \begin{array}{c} m \times n \\ \curvearrowleft \\ \curvearrowleft \\ \curvearrowleft \\ \curvearrowleft \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array}$$
$$\begin{bmatrix} w_{11} & \dots & w_{1j} & \dots & w_{1n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w_{i1} & w_{i2} & w_{ij} & \dots & w_{in} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w_{m1} & w_{m2} & w_{mj} & \dots & w_{mn} \end{bmatrix} \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array} = \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_j \\ \vdots \\ x_n \end{bmatrix} \quad \begin{array}{c} n \times 1 \\ \curvearrowleft \\ \curvearrowleft \\ \curvearrowleft \\ \curvearrowleft \end{array} \quad \begin{array}{c} \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \\ \text{---} \end{array}$$
$$\begin{bmatrix} y_1 \\ \vdots \\ y_i \\ \vdots \\ y_m \end{bmatrix} \quad \begin{array}{c} m \times 1 \\ \curvearrowleft \\ \curvearrowleft \\ \curvearrowleft \\ \curvearrowleft \end{array}$$

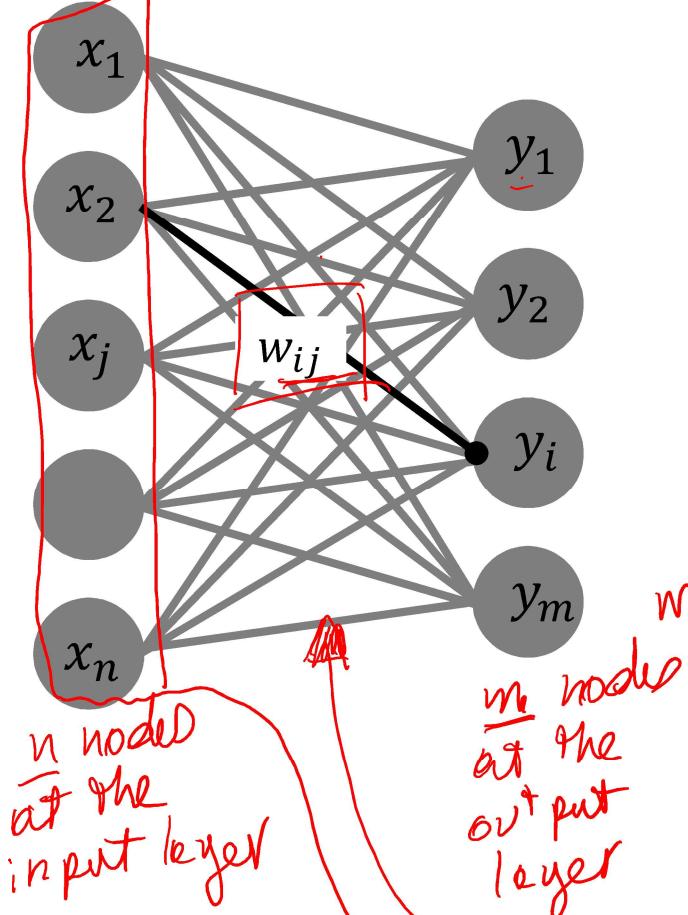
$j = 1 \rightarrow n$
 $i = 1 \rightarrow m$

where:

$y_i = \text{dot product of } i^{\text{th}} \text{ row of } W \text{ with } x$

$$y_i = \sum_{j=1}^n w_{ij} x_j$$

so what?



m hidden nodes

n input nodes

$$w_{ij} = \text{weight from node } x_j \text{ to } y_i$$

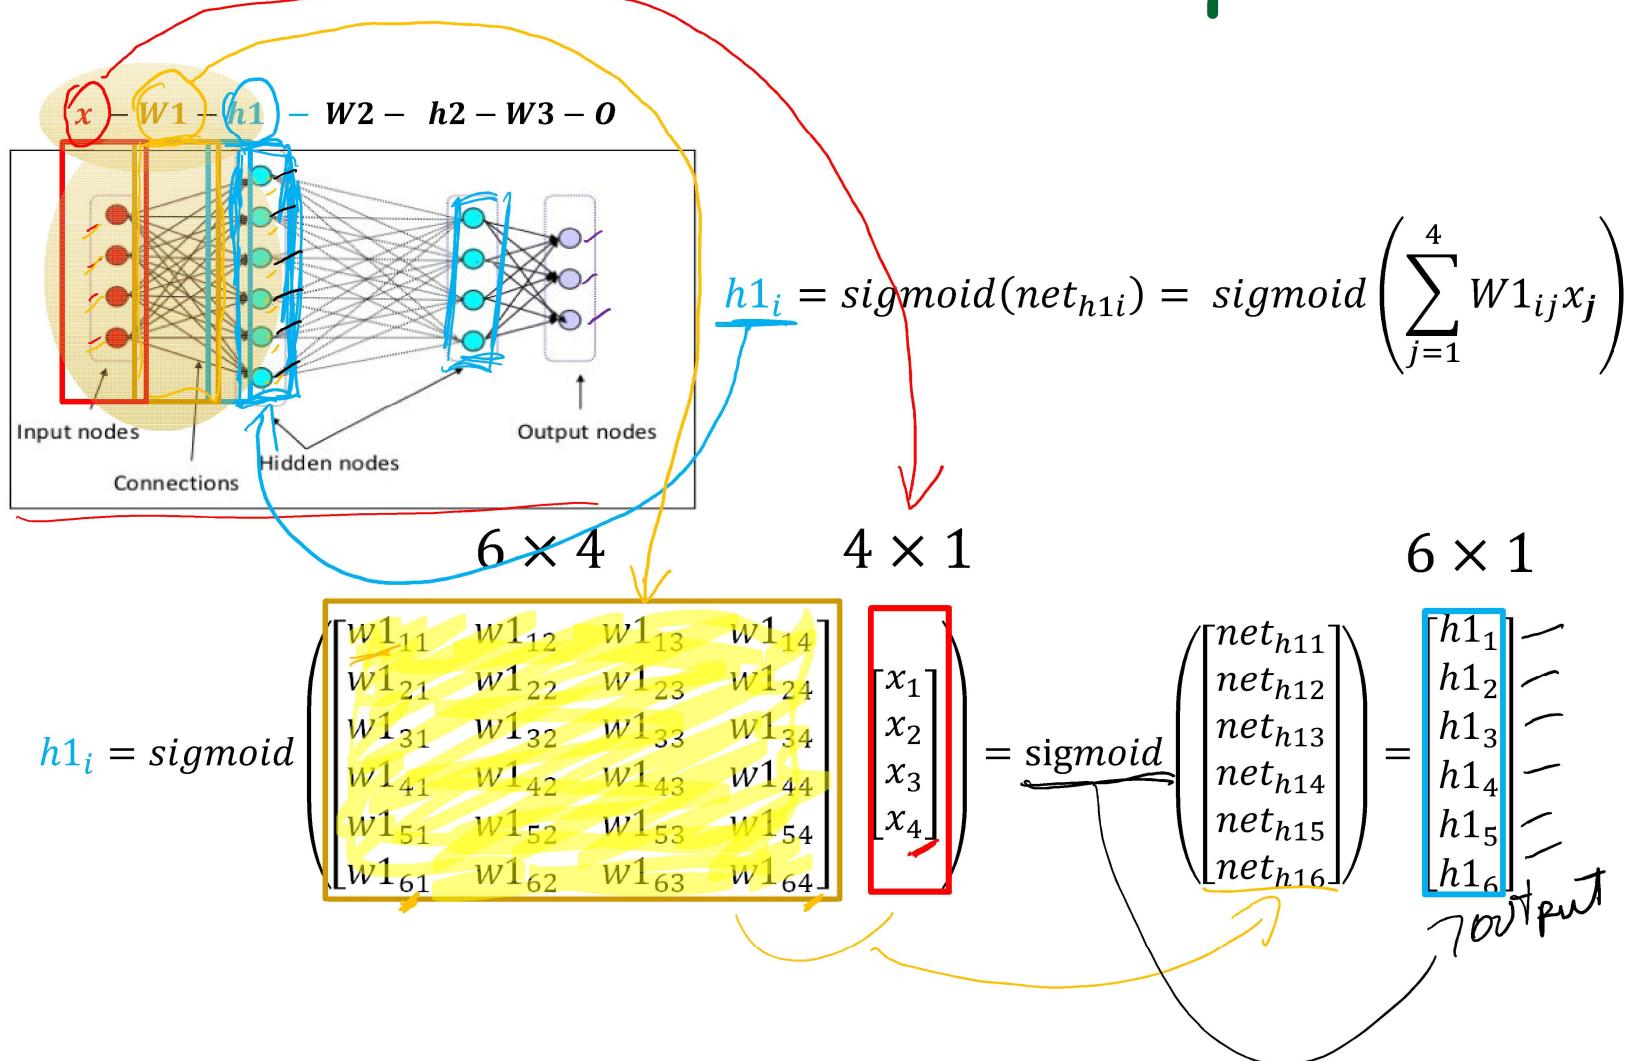
↑ put of y_i

$$\underline{y}_i = \text{sigmoid} \left(\sum_{j=1}^n w_{ij} x_j \right)$$

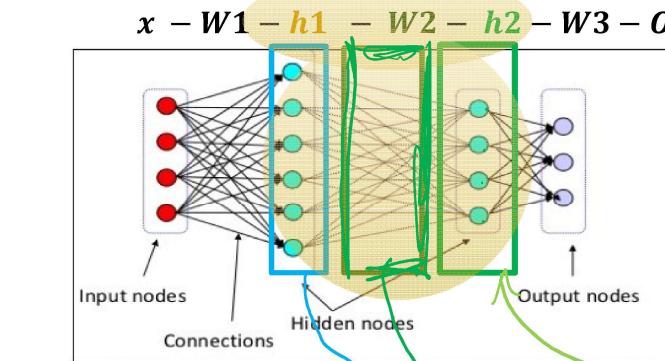
$$\begin{bmatrix} w_{11} & \dots & w_{1j} & \dots & w_{1n} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w_{i1} & w_{i2} & w_{ij} & \dots & w_{in} \\ \vdots & \vdots & \vdots & \ddots & \vdots \\ w_{m1} & w_{m2} & w_{mj} & \dots & w_{mn} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_j \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_i \\ \vdots \\ y_m \end{bmatrix}$$

$$y = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_m \end{bmatrix}$$

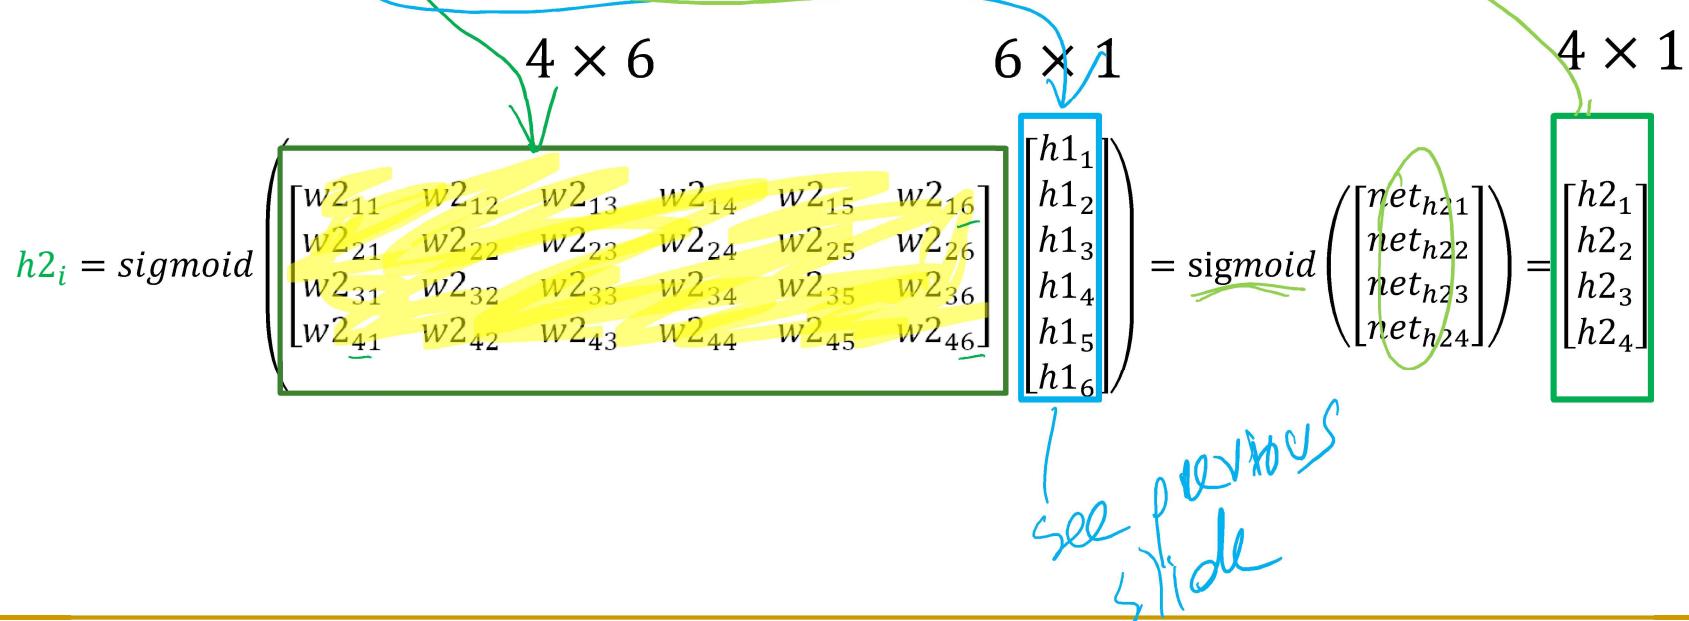
Matrix Notation - Example



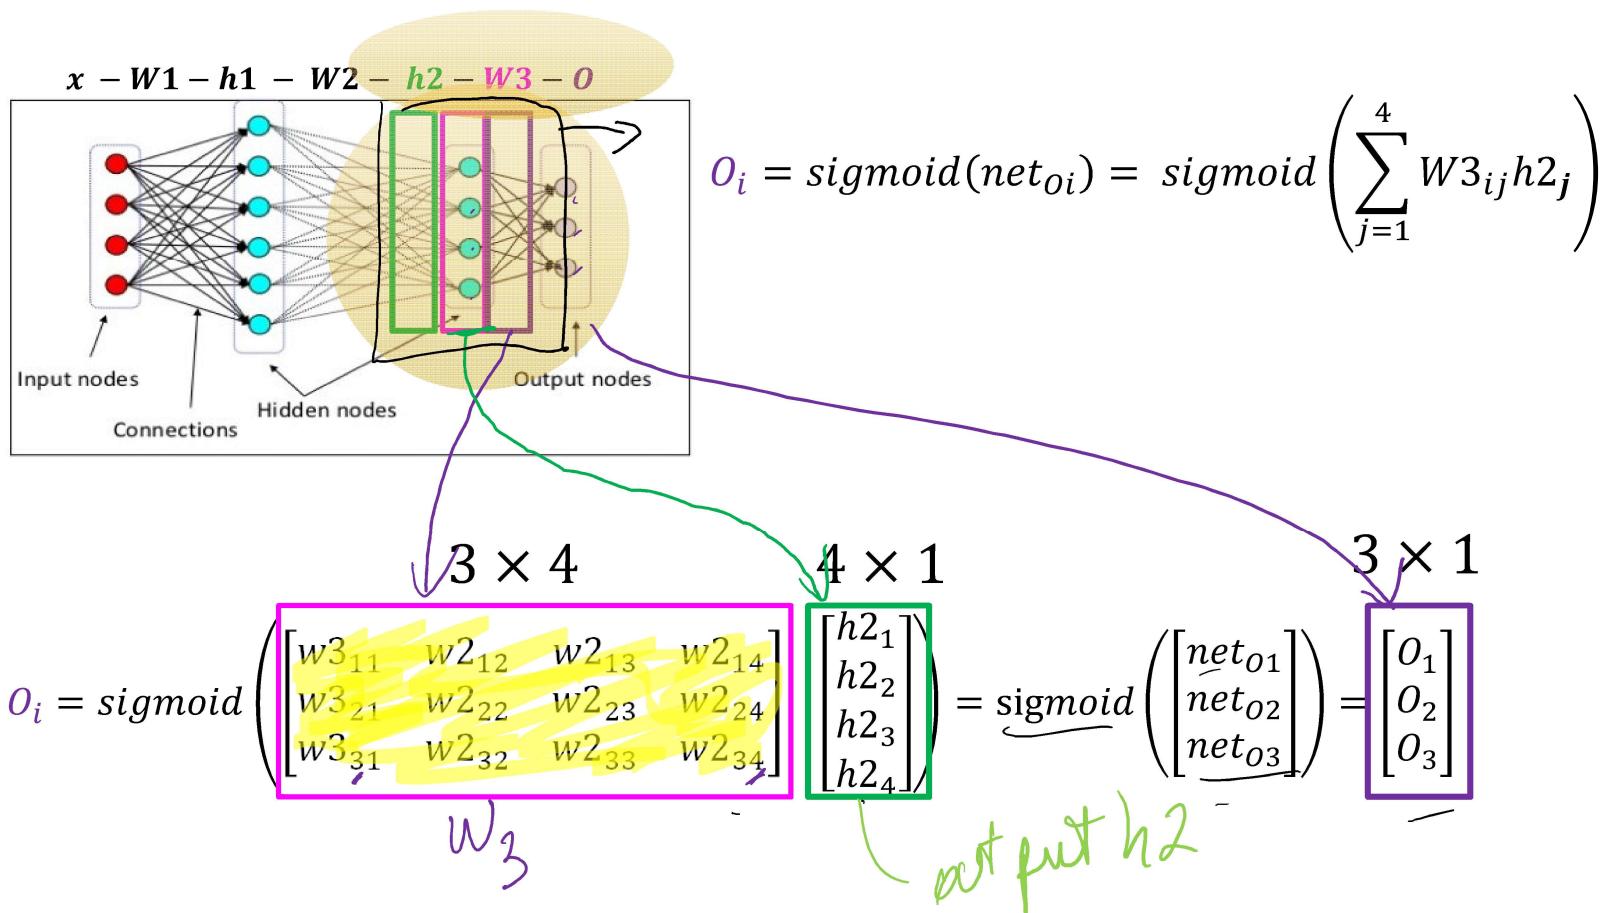
Repeat on next level



$$h_{2i} = \text{sigmoid}(\text{net}_{h2i}) = \text{sigmoid}\left(\sum_{j=1}^6 w_{2ij} h_{1j}\right)$$



Repeat on next level



Neural Networks

- Disadvantage:
 - result is not easy to understand by humans (set of weights compared to decision tree)... it is a black box
- Advantage:
 - robust to noise in the input (small changes in input do not normally cause a change in output) and graceful degradation

Today

1. Introduction to ML 
2. Naïve Bayes Classification

 - a. Application to Spam Filtering

3. Decision Trees
4. (Evaluation 
5. Unsupervised Learning) 
6. Neural Networks
 - a. Perceptrons 
 - b. Multi Layered Neural Networks 

Up Next

Part 4: Search

COMP 472: Artificial Intelligence

State Space Search part #3

State Space Representation video #1

- Russell & Norvig - Sections 3.1-3.3

Today

1. State Space Representation  video #1
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first and Depth-first
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Motivation



1970

- Many AI problems, can be expressed in terms of going from an **initial state** to a **goal state**
 - Ex: to solve a puzzle, to drive from home to Concordia...



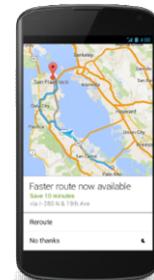
initial



goal

8-puzzle

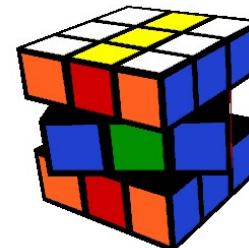
15 moves
200 moves



Google Maps

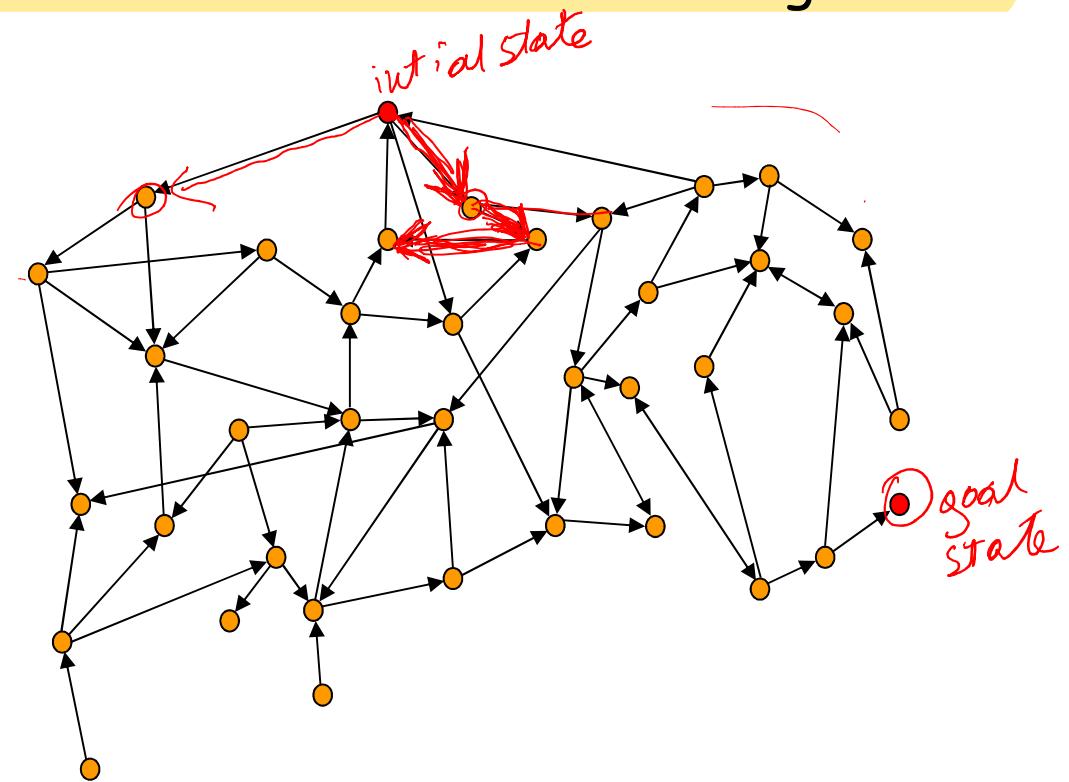


Rubik's cube



Motivation

- Often, there is no direct way to find a solution to go from the initial state to a goal state
 - but we can list the possibilities and search through them

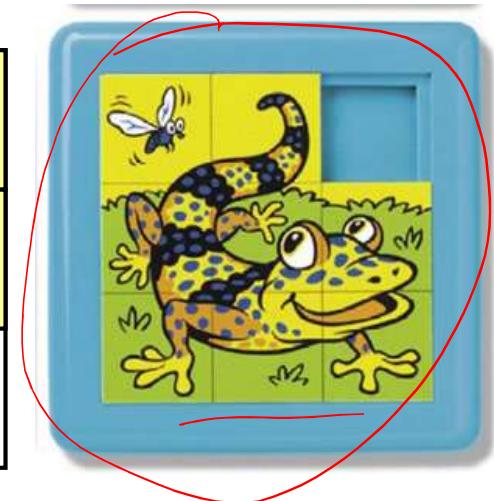


Example: 8-Puzzle

State: Any arrangement of 8 numbered tiles and an empty tile on a 3x3 board

8	2	empty
3	4	7
5	1	6

1	2	3
4	5	6
7	8	



Initial state

Goal state



! there are several standard goals states for the 8-puzzle

1	2	3
4	5	6
7	8	

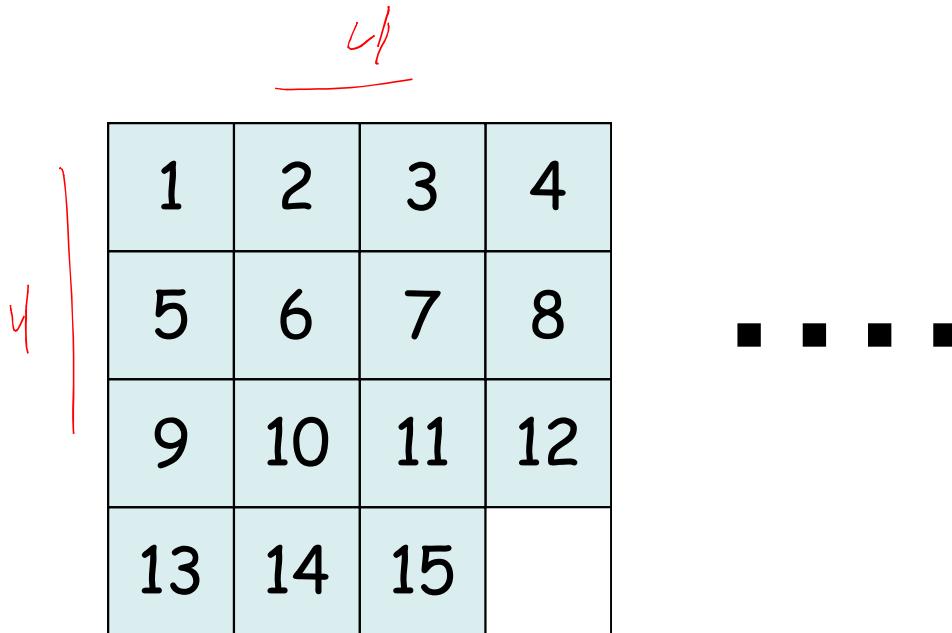
1	2	3
8		4
7	6	5

...

(n^2-1) -puzzle

8	2	
3	4	7
5	1	6

8-puzzle



15-Puzzle

Invented in **1874** by Noyes Palmer Chapman ...
but Sam Loyd claimed he invented it!



Sam Loyd

see Sam Loyd's book of puzzles:
https://archive.org/stream/CyclopediaOfPuzzlesLoyd/Cyclopedia_of_Puzzles_Loyd#mode/2up

State Space

- Problem is represented by:

1. **Initial State**

- starting state
- ex. unsolved puzzle, being at home

2. **Set of operators / moves / actions**

- actions responsible for transition between states

3. **Goal test function**

- Applied to a state to determine if it is a goal state
- ex. solved puzzle, being at Concordia

4. **Path cost function**

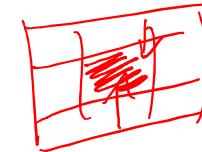
- Assigns a cost to a path to tell if a path is preferable to another

- **State space:**

- the set of all states that can be reached from the initial state by any sequence of action

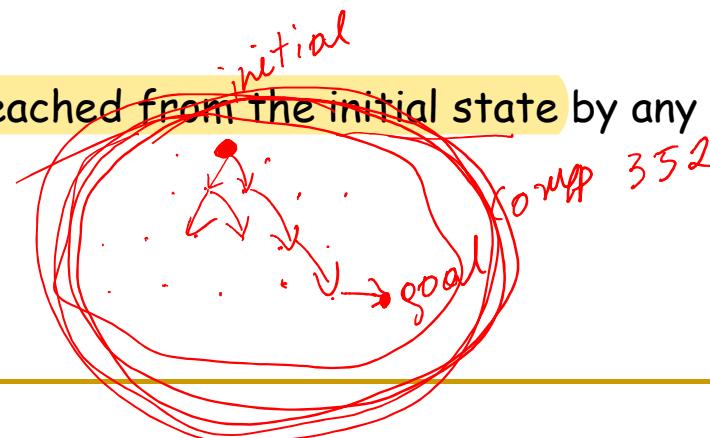
- **Search algorithm:**

- how the search space is visited

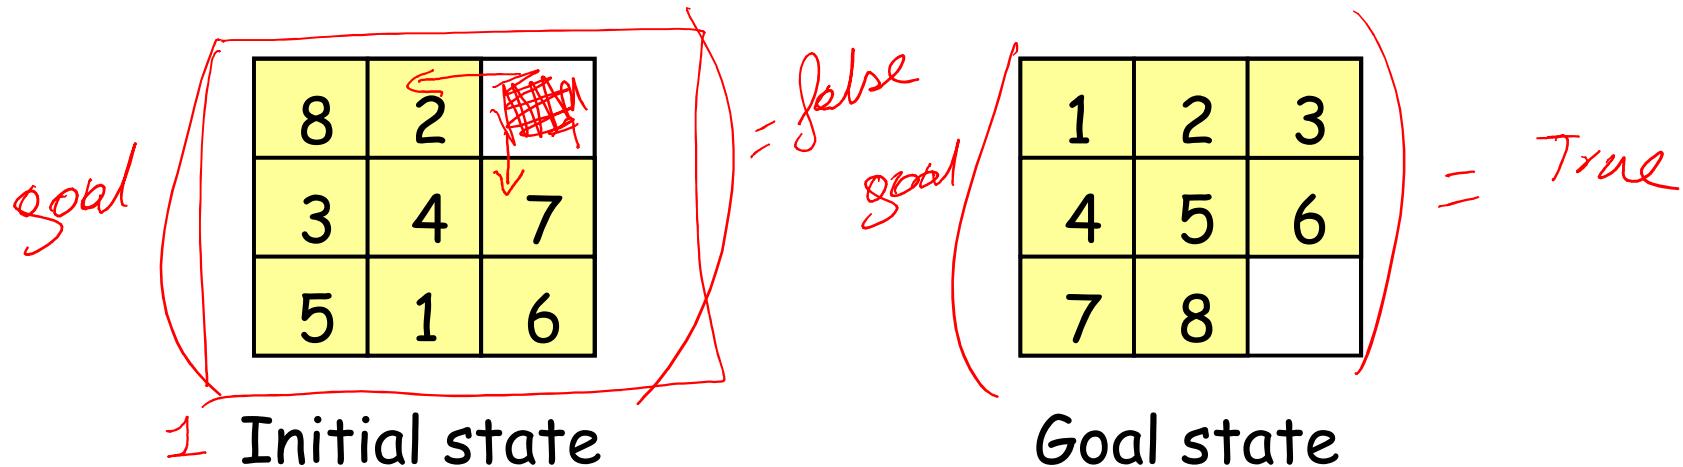


	cost	lost	3
UP	1	1	
down	1		2
left	1		4
right	1	lost	buret
	lost	nb of moves	lost

$$\begin{aligned} \text{goal}(0) &\rightarrow T \\ \text{goal} \left(\begin{array}{|c|c|c|} \hline 1 & 2 & 3 \\ \hline 4 & 5 & 6 \\ \hline 7 & 8 & 9 \\ \hline \end{array} \right) &= T \end{aligned}$$



Example: The 8-puzzle



2 - Set of operators:

blank moves up, blank moves down, blank moves left, blank moves right

3 - Goal test function:

state matches the goal state

4 - Path cost function:

each movement costs 1

so the path cost is the length of the path (the number of moves)

8-Puzzle: Successor Function

{children}

successor(n)
} n_1, n_2, n_3

node n

8	2	7
3	4	7
5	1	6

n_1

8	2	
3	4	7
5	1	6

n_2

8	2	7
3	4	6
5	1	

n_3

8	2	7
3		4
5	1	6

up

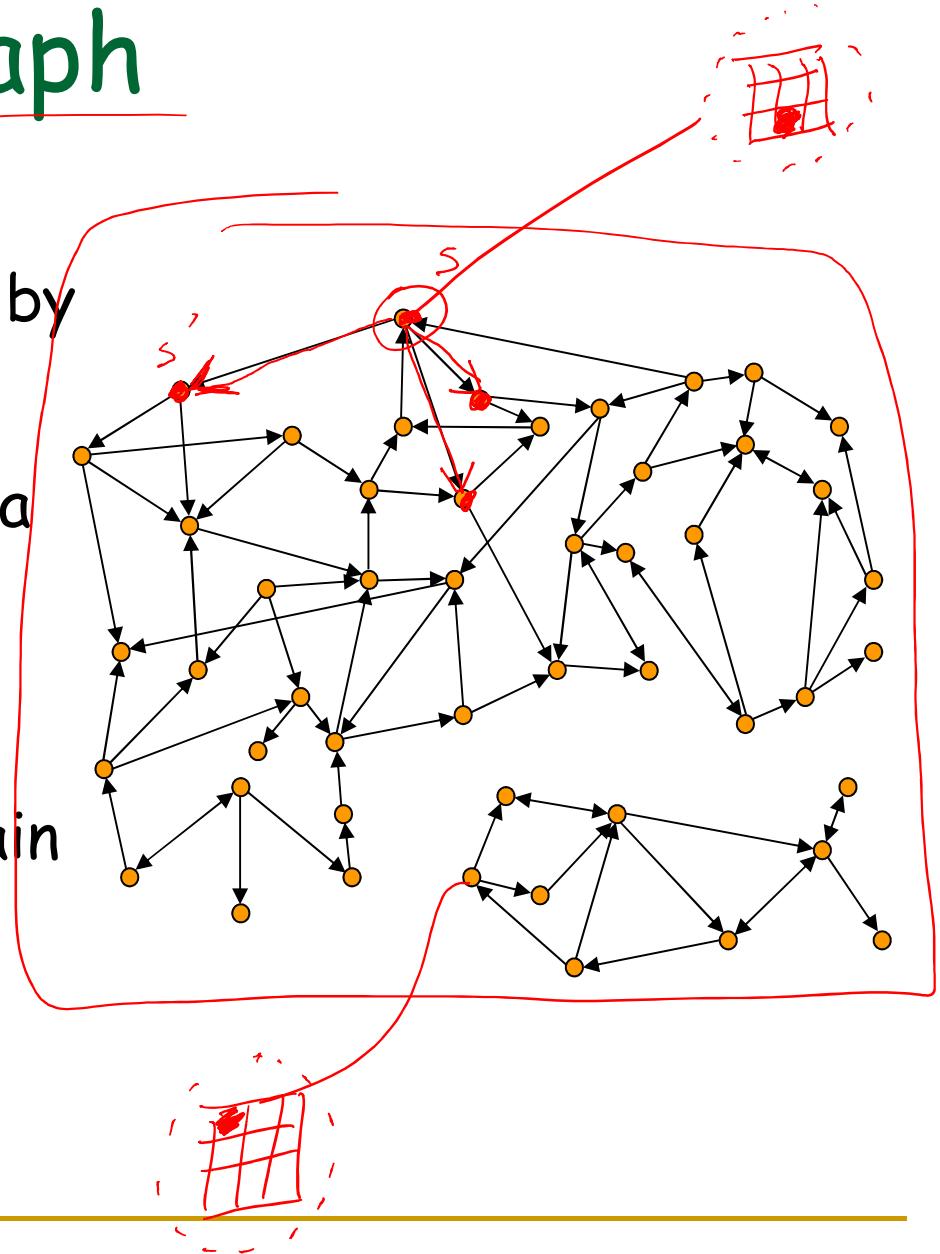
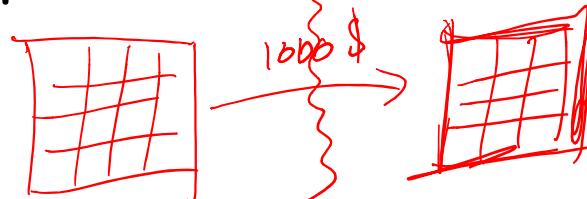
down

left

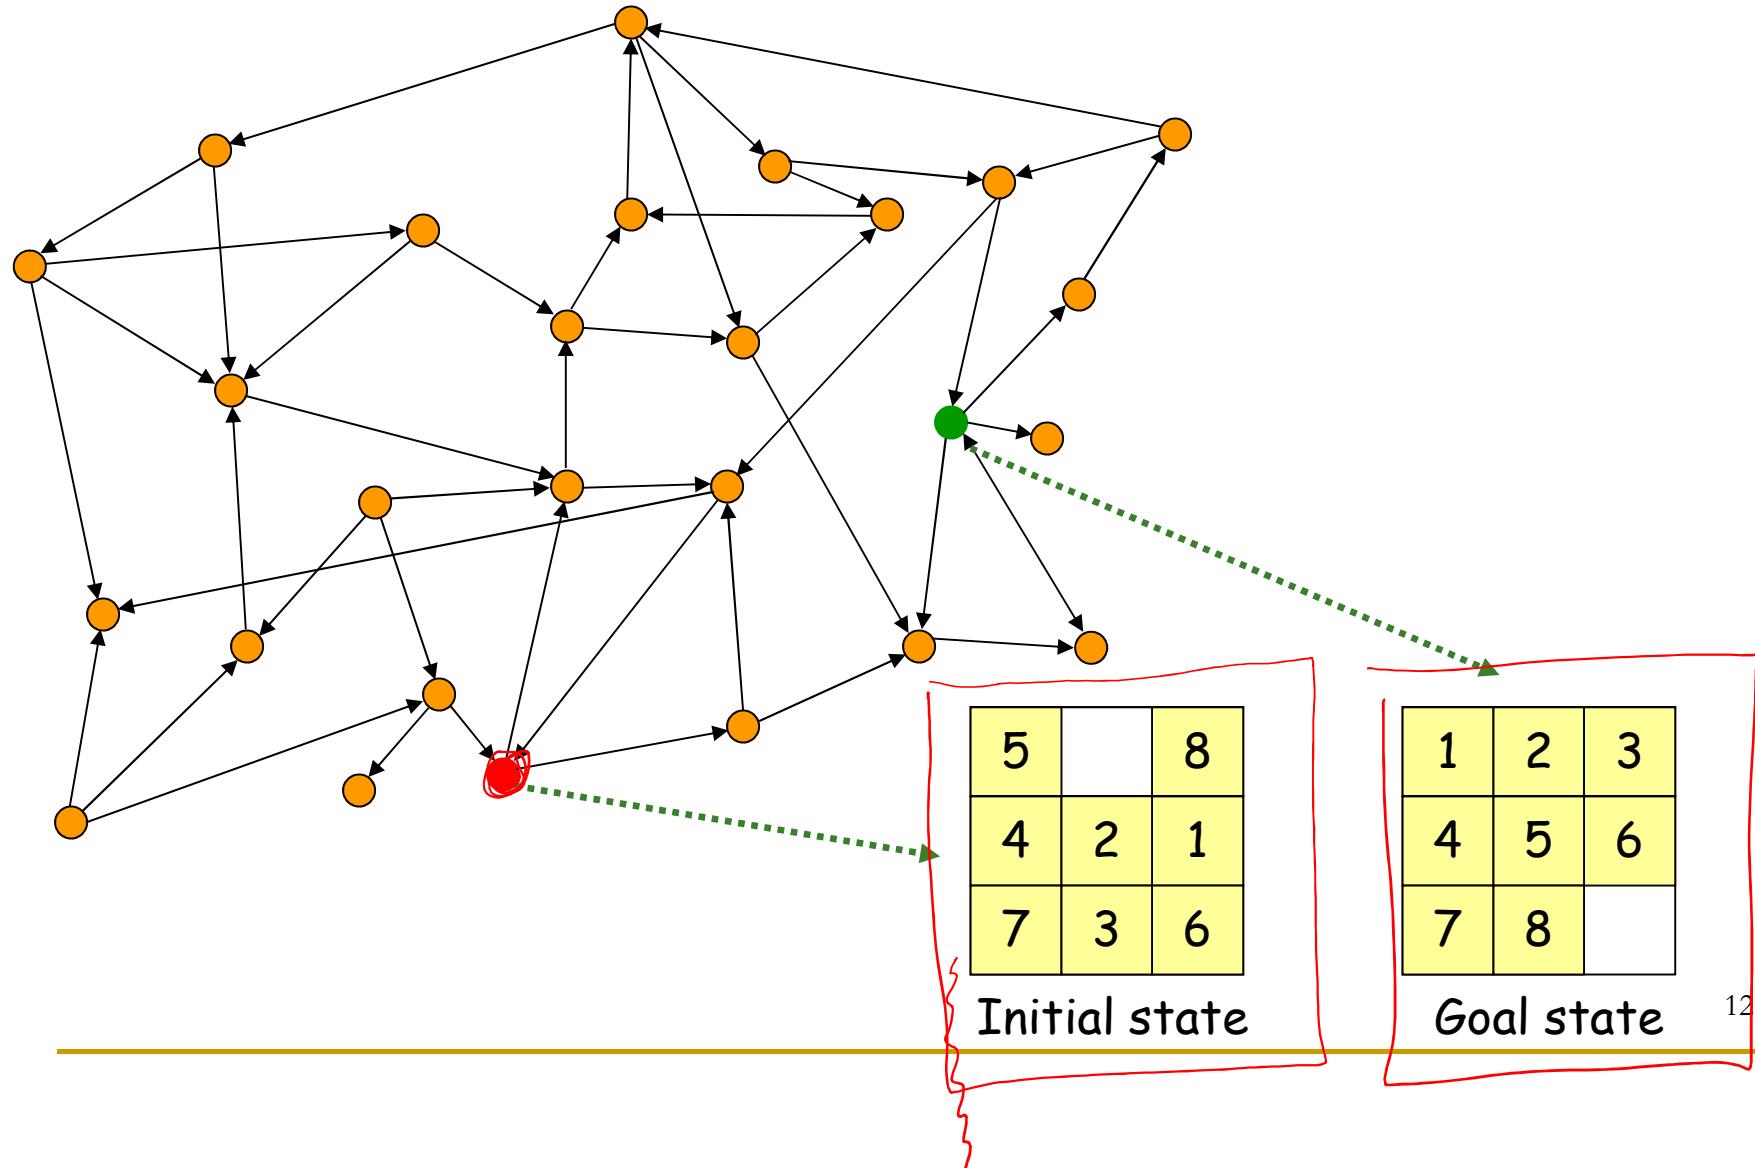
Search is about the exploration of alternatives

State Space Graph

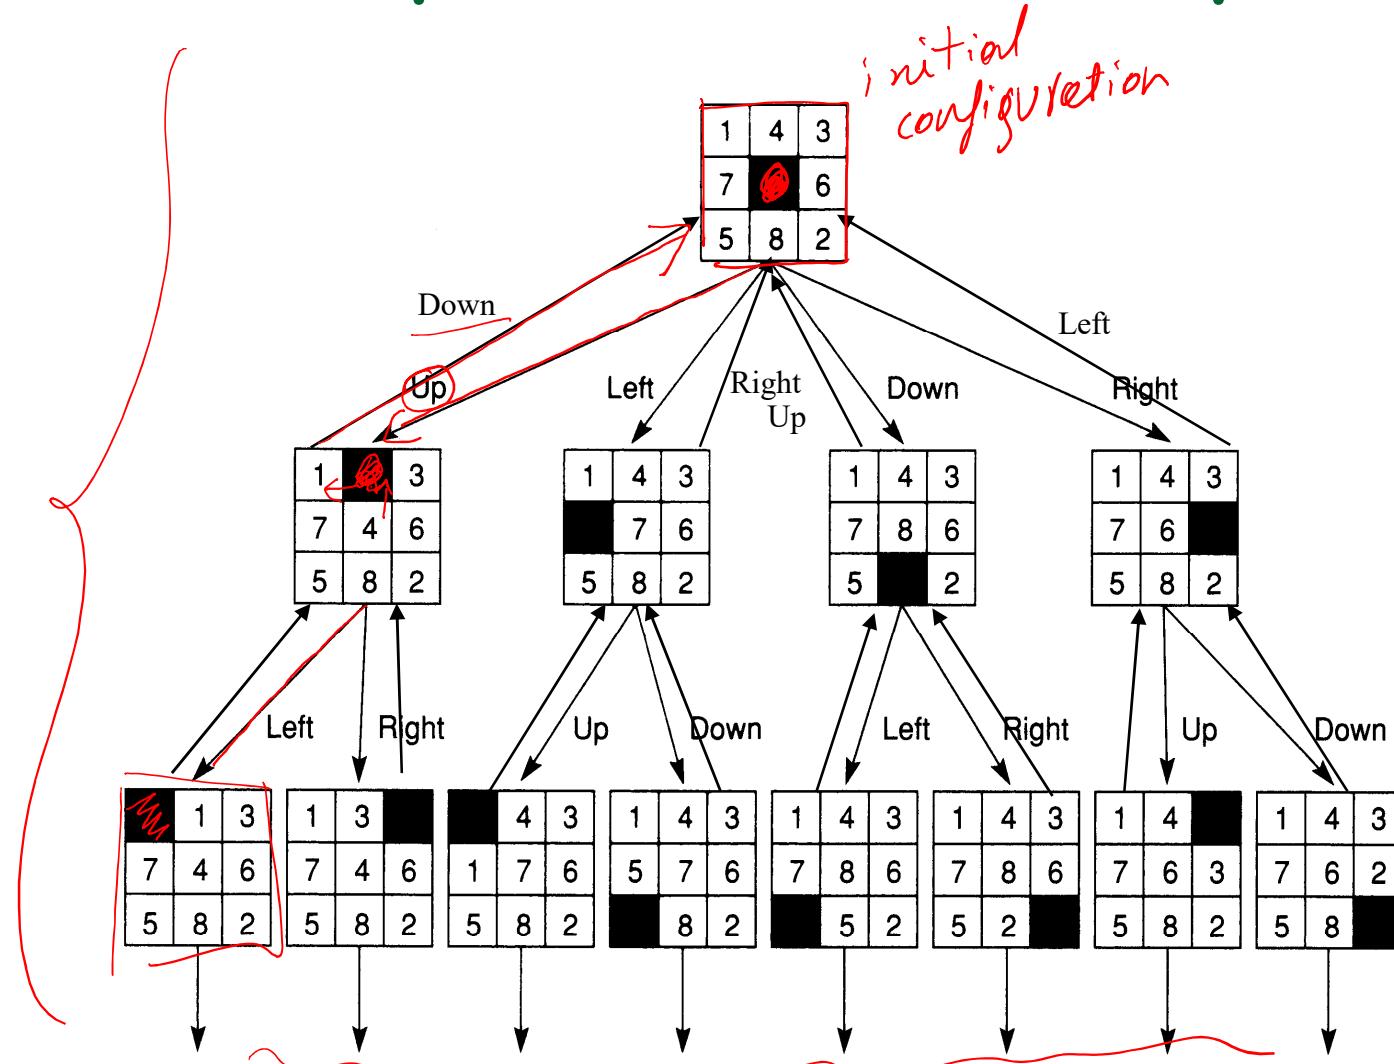
- Each state is represented by a distinct node
- An arc (or edge) connects a node s to a node s' if $s' \in \text{SUCCESSOR}(s)$
- The state graph may contain more than one connected component



Just to make sure we're clear...



State Space for the 8-puzzle

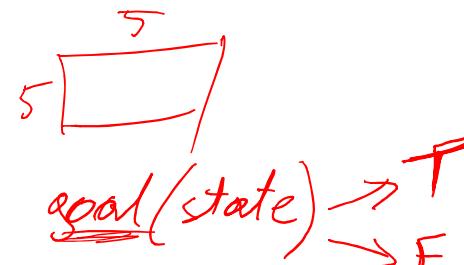


source: G. Luger (2005)

Size of state spaces

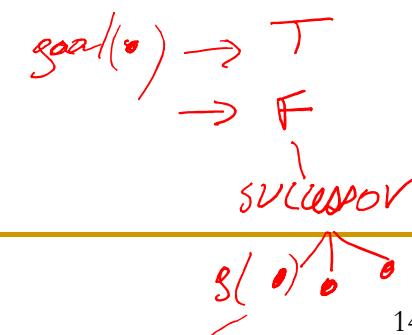
- For the (n-1)-puzzle:

- Nb of states:
 - 8-puzzle $\rightarrow 9! = 362,880$ states
 - 15-puzzle $\rightarrow 16! \sim 2.09 \times 10^{13}$ states
 - 24-puzzle $\rightarrow 25! \sim 10^{25}$ states
- At 100 millions states/sec:
 - 8-puzzle $\rightarrow 0.036$ sec
 - 15-puzzle $\rightarrow \sim 55$ hours
 - 24-puzzle $\rightarrow > 10^9$ years

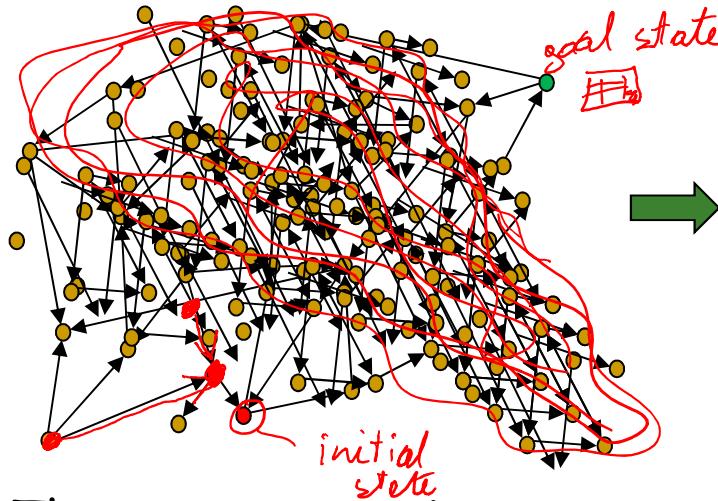


- For real problems:

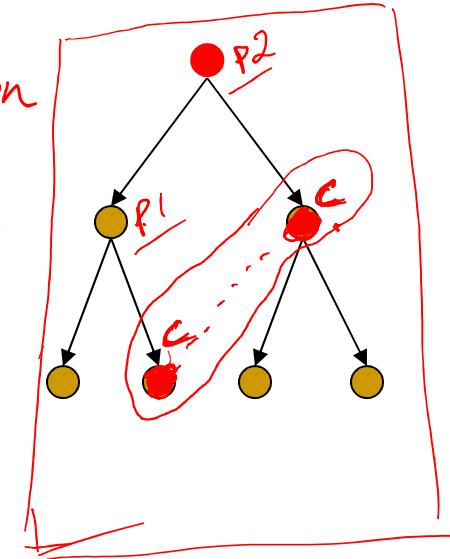
- state spaces are way too large to be generated in advance and searched after
- so it is generated dynamically while searching.



State Space Graph as a Search Tree



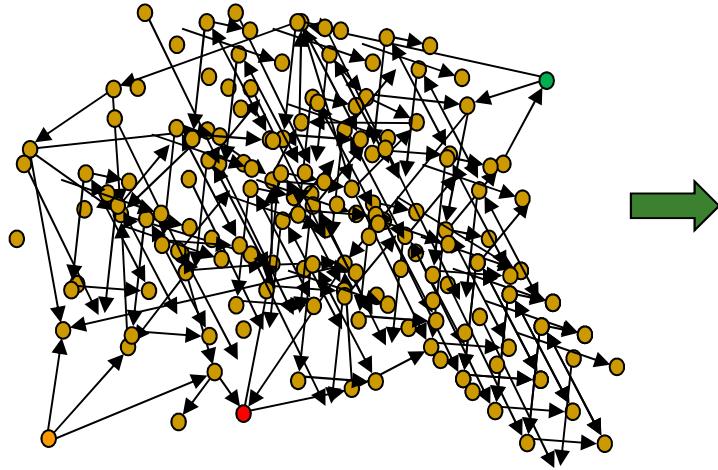
1) $\text{goal}(\bullet) = F$
2) successor function
children
 $\text{goal}(\bullet) = F$



- The state space is generated dynamically while searching.
 - we explore a node:
 - if it the goal, stop and trace back the path from the initial node
 - if it is not the goal, then generating its successors/children and explore these recursively
 - to avoid cycles, the search algorithm will check for duplicate nodes.

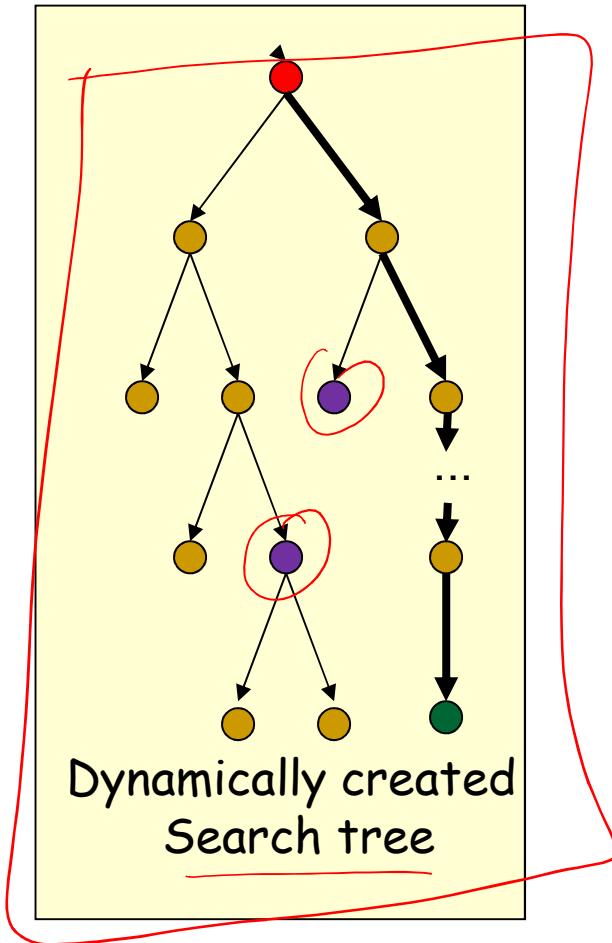
Search tree

State Space Graph as a Search Tree



Theoretical State Space Graph

So now, we just need an efficient search algorithm



Today

1. State Space Representation



2. State Space Search

a) Overview

b) Uninformed search

1. Breadth-first and Depth-first
2. Depth-limited Search
3. Iterative Deepening
4. Uniform Cost

c) Informed search

1. Intro to Heuristics
2. Hill climbing
3. Greedy Best-First Search
4. Algorithms A & A*
5. More on Heuristics

d) Summary

Up Next

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 - 1. Breadth-first and Depth-first
 - 2. Depth-limited Search
 - 3. Iterative Deepening
 - 4. Uniform Cost
 - c) Informed search
 - 1. Intro to Heuristics
 - 2. Hill climbing
 - 3. Greedy Best-First Search
 - 4. Algorithms A & A*
 - 5. More on Heuristics
 - d) Summary

video #2

COMP 352

novelty

COMP 472 Artificial Intelligence

State Space Search *part #3*

Uninformed Search *video #2*

- Russell & Norvig - Section 3.4
- see also: <https://www.javatpoint.com/ai-uninformed-search-algorithms>

Today

1. State Space Representation

2. State Space Search

a) Overview



b) Uninformed search

1. Breadth-first and Depth-first
2. Depth-limited Search
3. Iterative Deepening
4. Uniform Cost

c) Informed search

1. Intro to Heuristics
2. Hill climbing
3. Best-First
4. Algorithms A & A*
5. More on Heuristics

d) Summary

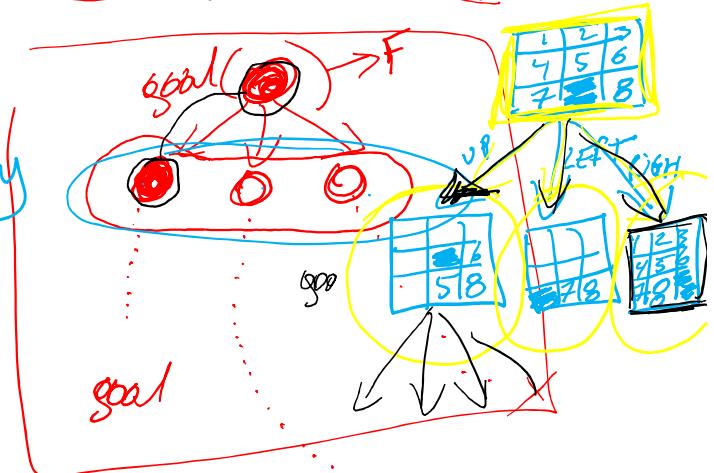
Uninformed VS Informed Search

■ Uninformed search

- ❑ all nodes are equally promising, so we explore them systematically
- ❑ aka: systematic/blind/brute force search
- ❑ many algorithms:

1. Breadth-first search
2. Depth-first search
3. Uniform-cost search
4. Depth-limited search
5. Iterative deepening search
6. ...

BJ2 video today

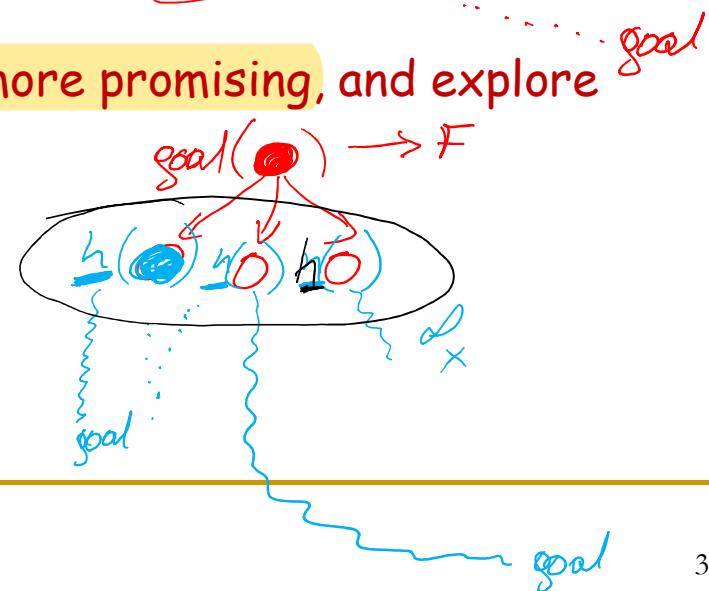


■ Informed search (heuristic search)

- ❑ we try to identify which nodes seem more promising, and explore these first

- ❑ many algorithms:

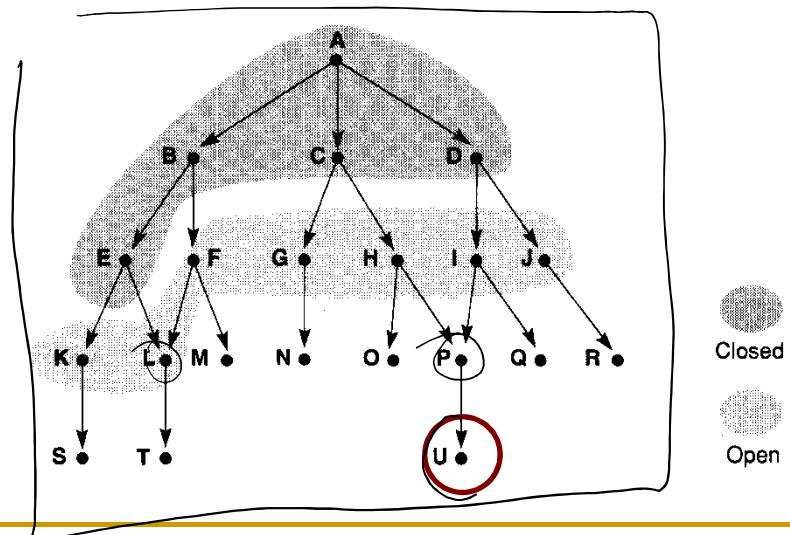
1. Hill climbing
2. Greedy Best-First search
3. Algorithms A and A*
4. ...



Data Structures

- Most search strategies require:
 - open list (aka the frontier) *To-DO*
 - lists generated nodes not yet expanded
 - order of nodes controls order of search
 - closed list (aka the explored set)
 - stores all the nodes that have already been visited (to avoid cycles).
- ex:

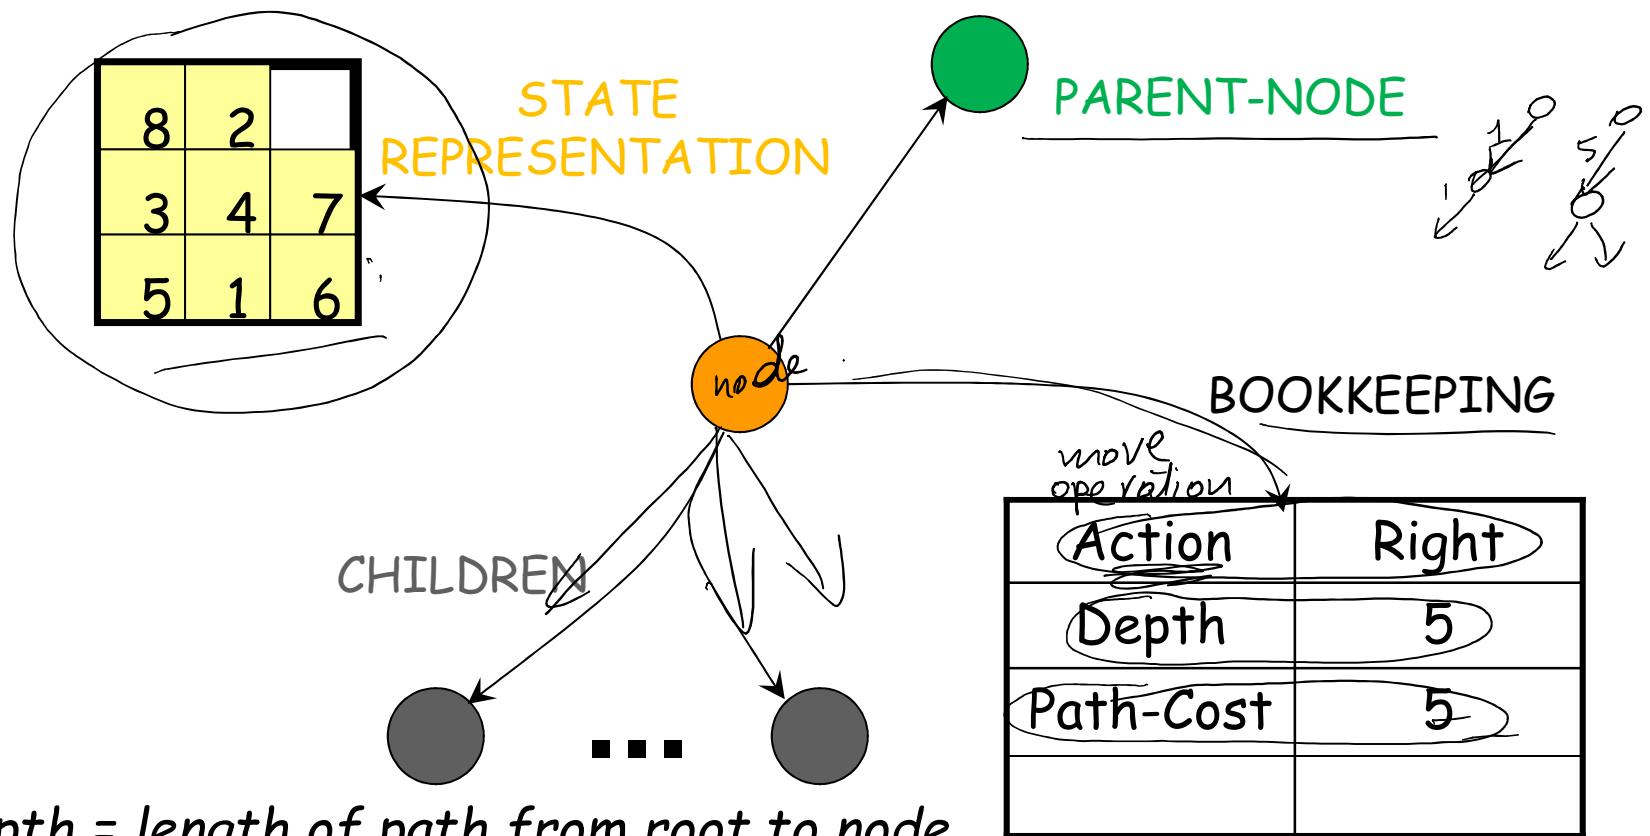
Closed = [A, B, C, D, E]
Open = [F, G, H, I, J, K, L]



source: G. Luger (2005)

Data Structures

- state space representation: To trace back the solution path after the search, each node in the lists contain:



Depth = length of path from root to node

robotics.stanford.edu/~latombe/cs121/2003/home.htm

Generic Search Algorithm

1. Initialize the $\{ \text{open list} \}$ with the initial node s_0 (top node)
2. Initialize the closed list to empty
3. Repeat $\{ \text{done / visited} \}$
 - a) If the open list is empty, then exit with failure.
 - b) Else, take the first node s from the open list.
 - c) If s is a goal state, exit with success. Extract the solution path from s to s_0 .
 - d) Else, insert s in the closed list (s has been visited / expanded)
 - e) Insert the successors of s in the open list in a certain order if they are not already in the closed and/or open lists (to avoid cycles)

Notes:

- The order of the nodes in the open list depends on the search strategy

Today

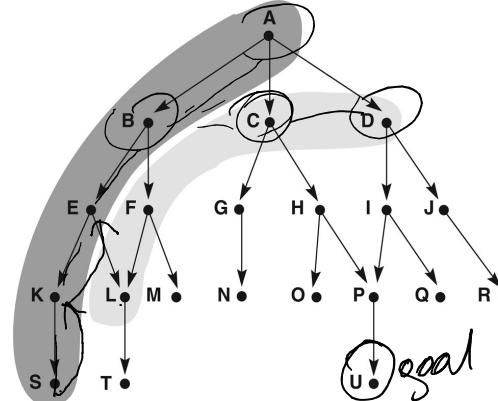
1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first Search and Depth-first Search
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Best-First
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary



Depth-first vs Breadth-first Search

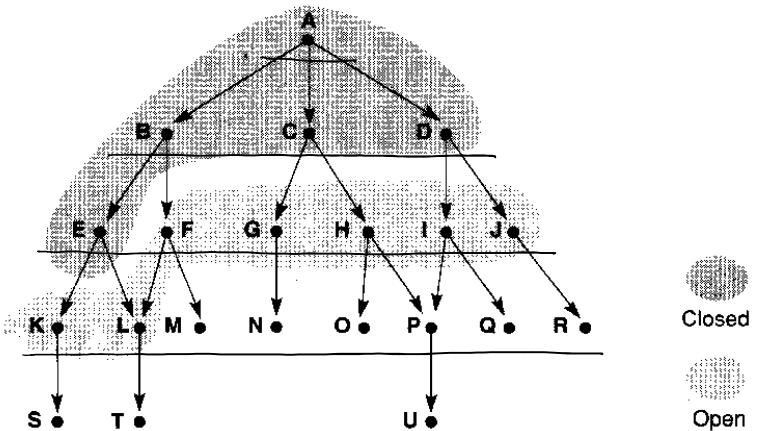
- Depth-first (DFS):

- visit successors before siblings
 - Open list is a stack



- Breadth-first (BFS):

- visit siblings before successors
 - ie. visit level-by-level
 - open list is a queue



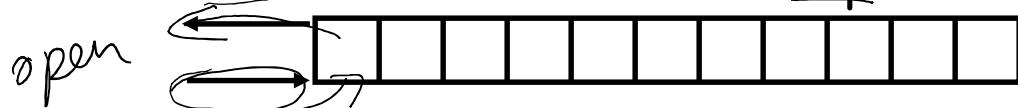
source: G. Luger (2005)

DFS and BFS

- ~~DFS and BFS~~ differ only in the way they order nodes in the open list:

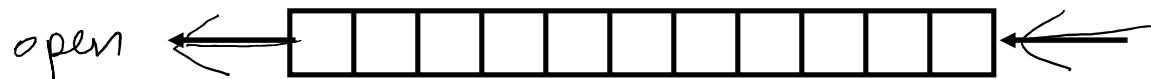
- DFS uses a stack:

- nodes are added on the top of the ^{open} list.



- BFS uses a queue:

- nodes are added at the end of the list.



Breadth-First Search

```
begin
    open := [Start];                                % initialize
    closed := [ ];
    while open ≠ [ ] do                            % states remain
        begin
            remove leftmost state from open, call it X;
            if X is a goal then return SUCCESS          % goal found
            else begin
                generate children of X;
                put X on closed;
                discard children of X if already on open or closed;
                put remaining children on right end of open
            end
        end
    return FAIL
end.
```

Breadth-First Search Example

BFS: (open is a queue)

Assume U is goal state

1. open = [A-null] closed = []
2. open = [B-A C-A D-A] closed [A]
3. open = [C-A D-A E-B F-B] closed = [B A]
4. open = [D-A E-B F-B G-C H-C] closed = [C B A]
5. open = [E-B F-B G-C H-C I-D J-D] closed = [D C B A]
6. open = [F-B G-C H-C I-D J-D K-E L-E] closed = [E D C B A]
7. open = [G-C H-C I-D J-D K-E L-E M-F] as L is already in open closed = [F E D C B A]
8. and so on until either U is found or open = []

→ { search path = A B C D E F G H I J K L M N ... U } // closed list
 solution path = A C H P U with a cost of 4 // extract by following the parent pointer

source: G. Luger (2005)

Depth-First Search

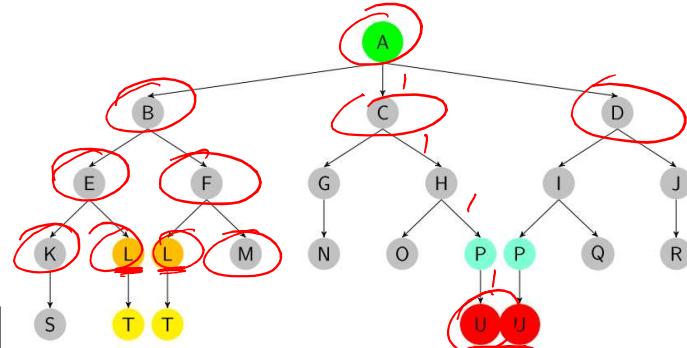
```
begin
  open := [Start];                                % initialize
  closed := [ ];
  while open ≠ [ ] do                            % states remain
    begin
      remove leftmost state from open, call it X;
      if X is a goal then return SUCCESS          % goal found
      else begin
        generate children of X;
        put X on closed;
        discard children of X if already on open or closed; % loop check
        put remaining children on left end of open
      end
    end;
    return FAIL
end.
```

Depth-First Search Example

- DFS: (open is a stack)

Assume U is goal state

1. open = [A-null] closed = []
2. open = [B-A C-A D-A] closed = [A]
3. open = [E-B F-B C-A D-A] closed = [B A]
4. open = [K-E L-E F-B C-A D-A] closed = [E B A]
5. open = [S-K L-E F-B C-A D-A] closed = [K E B A]
6. open = [L-E F-B C-A D-A] closed = [S K E B A]
7. open = [T-L F-B C-A D-A] closed = [L S K E B A]
8. open = [F-B C-A D-A] closed = [T L S K E B A]
9. open = [M-F C-A D-A] as L is already on closed closed = [F T L S K E B A]
10. open = [C-A D-A] closed = [M F T L S K E B A]
11. open = [G-c H-c D-A] closed = [C M F T L S K E B A]



search path = A B E K S L U // closed list
solution path = A C H P U with a cost of 4
MPU

Depth-first vs. Breadth-first

Breadth-first:

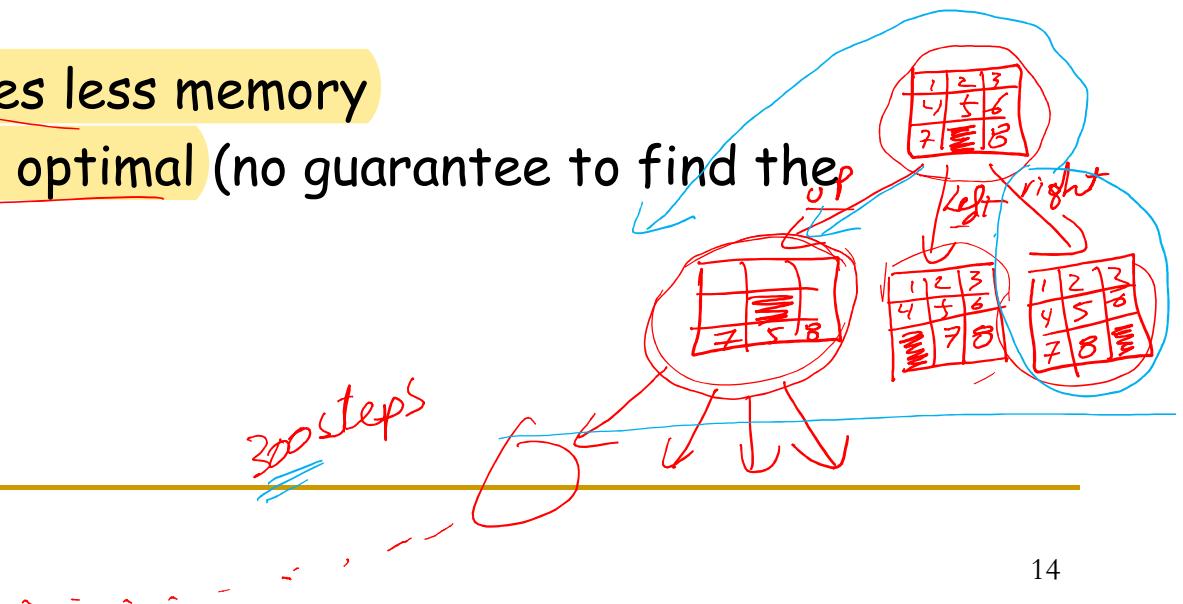
- advantage: optimal, i.e. will always find shortest path
- disadvantage:
 - high memory requirement as we need to keep all states of a level before expanding to the next level
 - exponential space for states required B^n // B =branching factor, n =level

solution

Depth-first:

- advantage: Requires less memory
- disadvantage: Not optimal (no guarantee to find the shortest path)

average #
of successors



Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first Search and Depth-first Search
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Best-First
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Depth-Limited Search

- Compromise for DFS :
 - Do depth-first but
 - with depth cutoff k (depth at which nodes are not expanded)
- Three possible outcomes:
 - Solution *within your k limit* :
 - Failure (no solution)
 - Cutoff (no solution found within cutoff) \times
- advantage: memory efficient - it's a DFS
- disadvantage: may not find a solution if it is below the cutoff

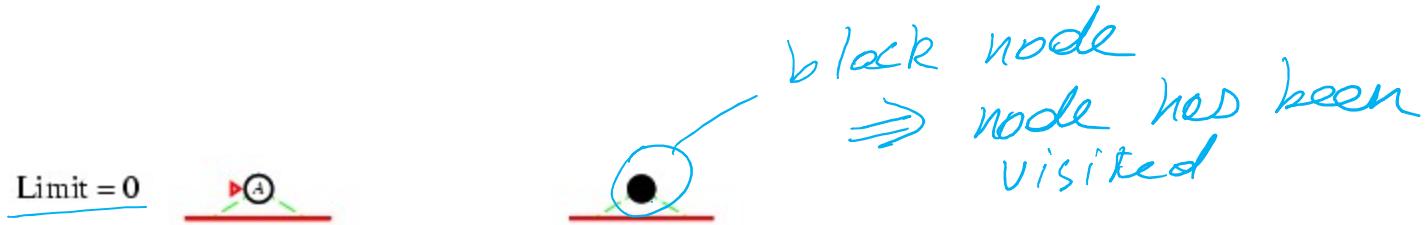
Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first search and Depth-first search
 2. Depth-limited Search
 3. Iterative Deepening 
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Best-First
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Iterative Deepening

- Combination of BFS and DFS:
 - do depth-first search, but
 - with a maximum depth before going to next level
 - i.e. Repeats depth first search with gradually increasing depth limits
- advantage:
 - Requires little memory (fundamentally, it's a depth first)
 - optimal: will find the shortest path (limited depth)
- disadvantage:
 - repeated traversal of the tree top

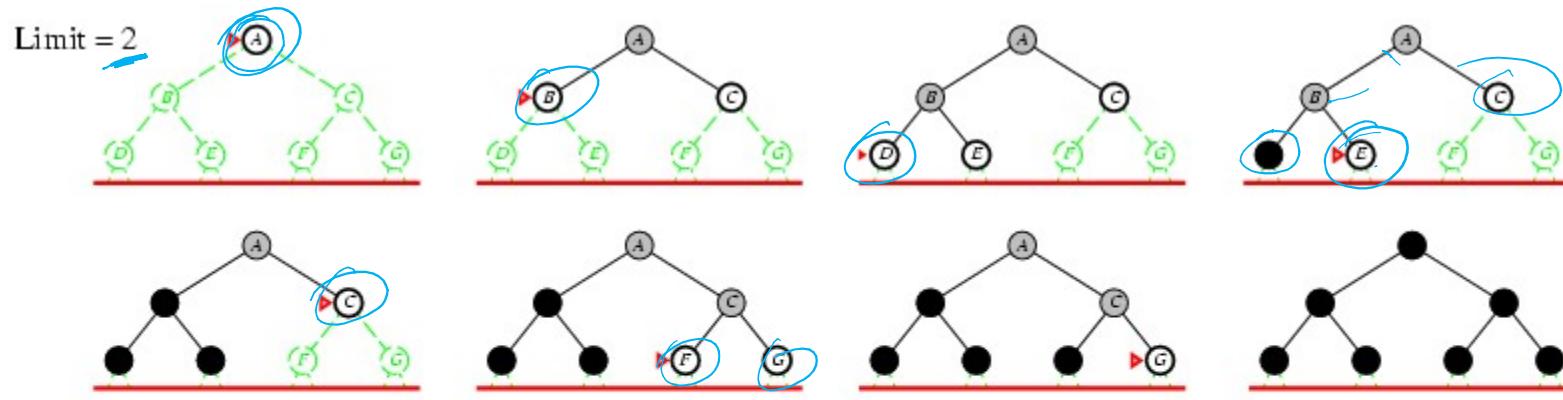
Iterative Deepening: Example



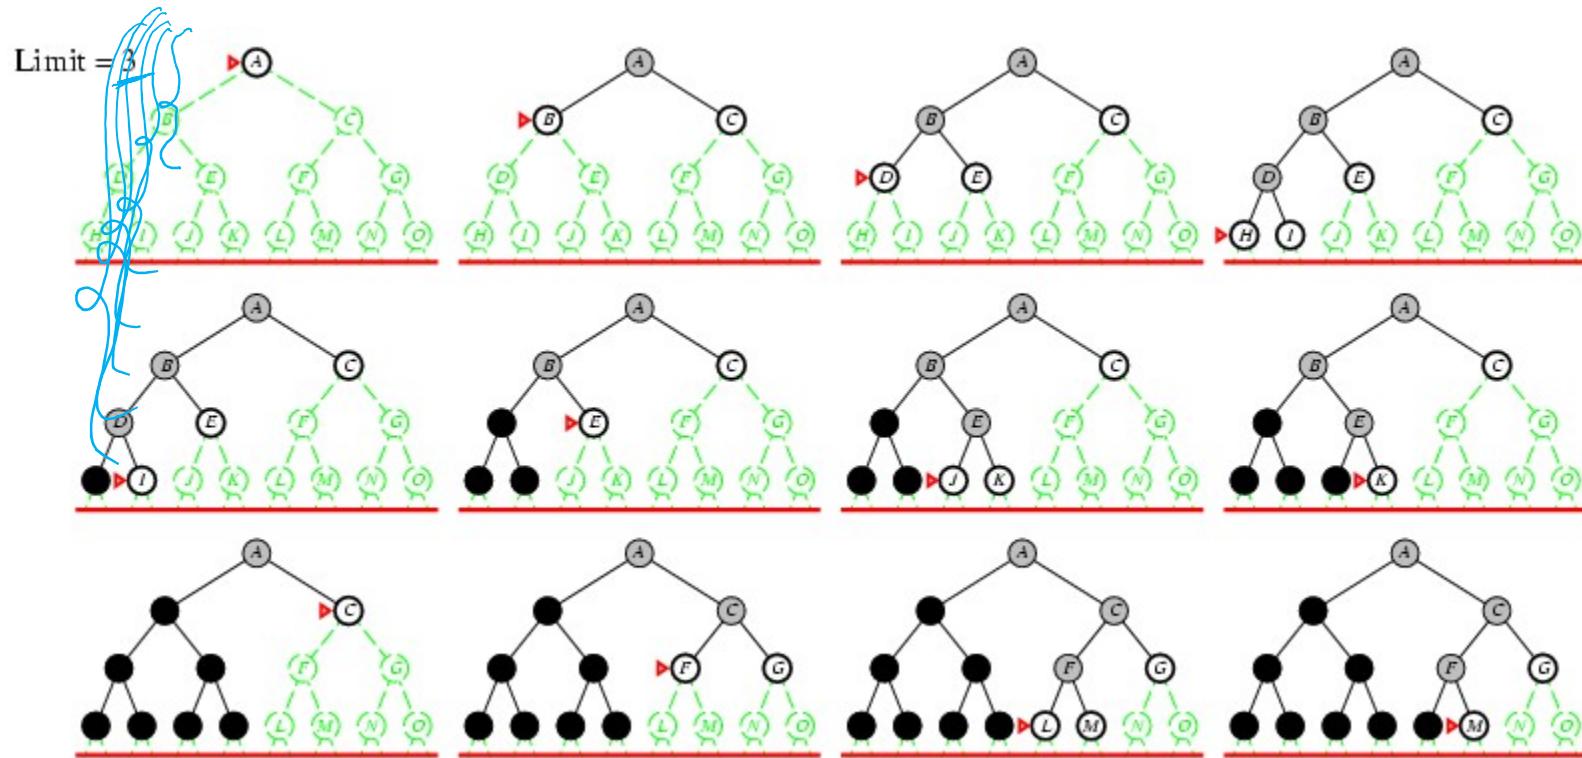
Iterative Deepening: Example



Iterative Deepening: Example



Iterative Deepening: Example



source: Russel & Norvig (2003)

Today

1. State Space Representation

2. State Space Search

a) Overview

b) Uninformed search

1. Breadth-first and Depth-first

2. Depth-limited Search

3. Iterative Deepening

4. Uniform Cost 

c) Informed search

1. Intro to Heuristics

2. Hill climbing

3. Best-First

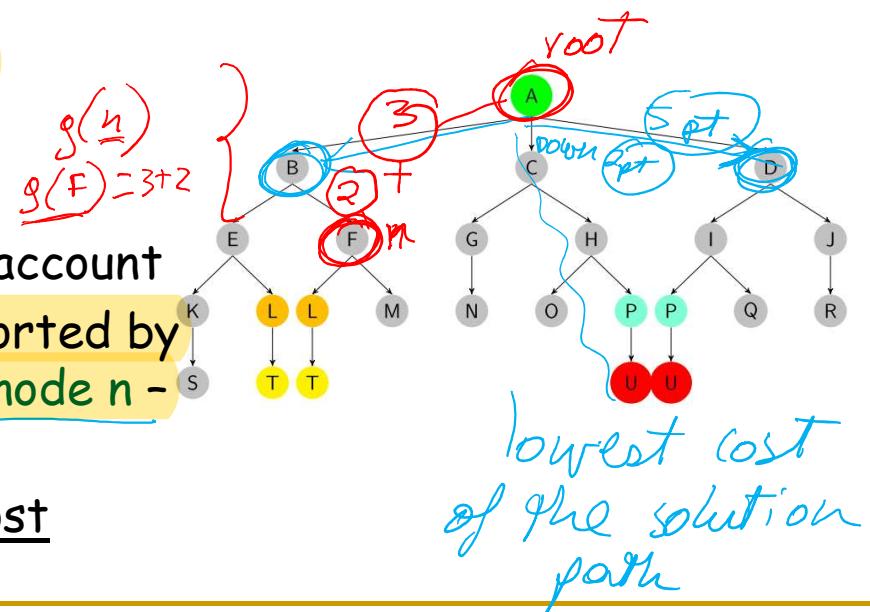
4. Algorithms A & A*

5. More on Heuristics

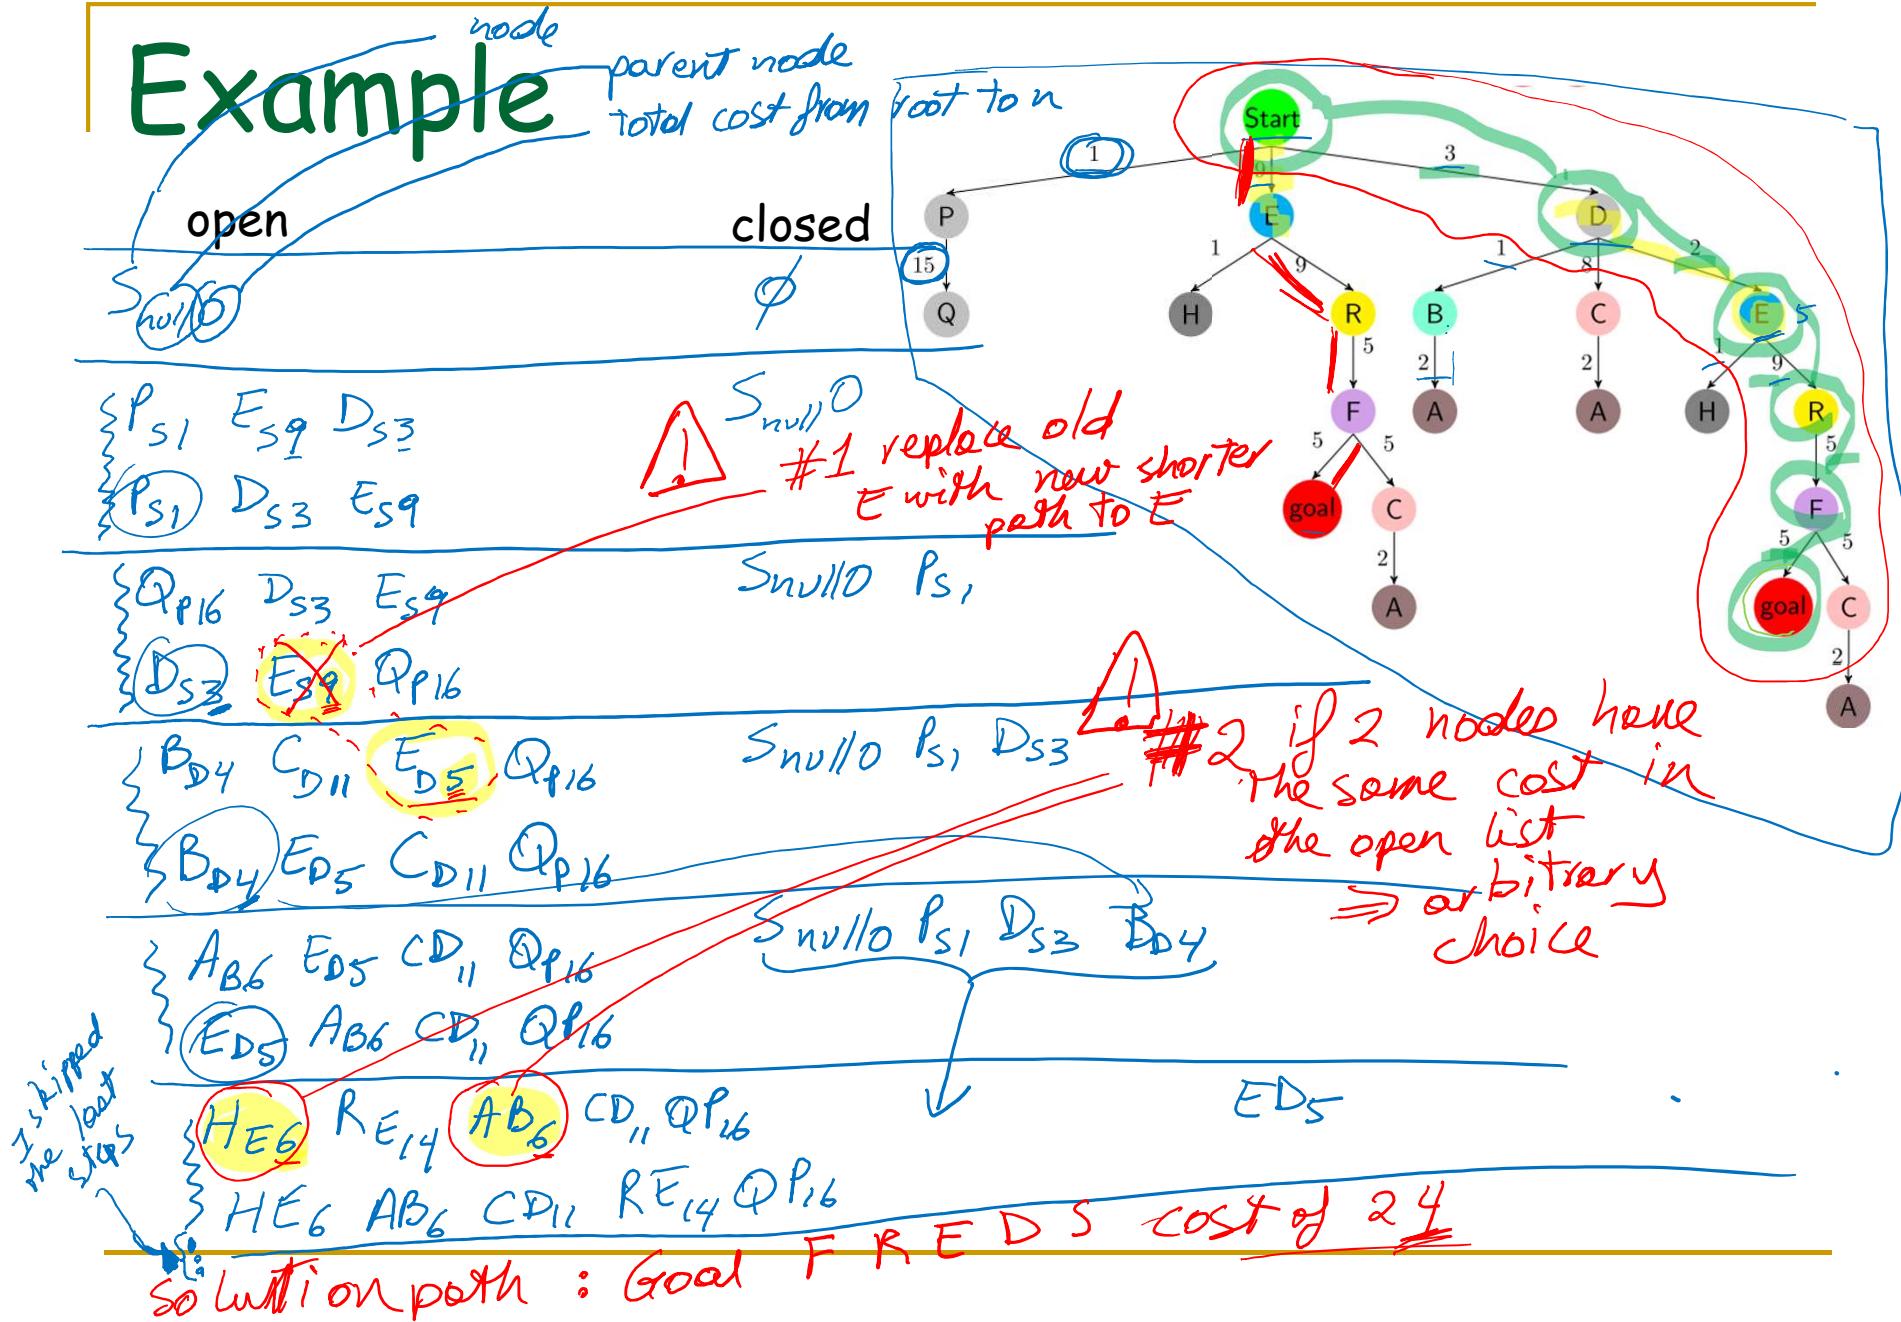
d) Summary

Uniform Cost Search

- all algorithms so far assume that all edges have the same cost
- but what if they have different costs?
 - eg: move UP → 2pts but move DOWN → 1 pt
 - eg: cost(residential road) > cost(commercial road)
- Breadth First Search
 - uses OPEN as a priority queue sorted by the depth of nodes
 - guarantees to find the shortest solution path
- Uniform Cost Search
 - takes the cost of the edge into account
 - uses OPEN as a priority queue sorted by the total cost from the root to node n - later we will call this $g(n)$
 - guarantees to find the lowest cost solution path



Example



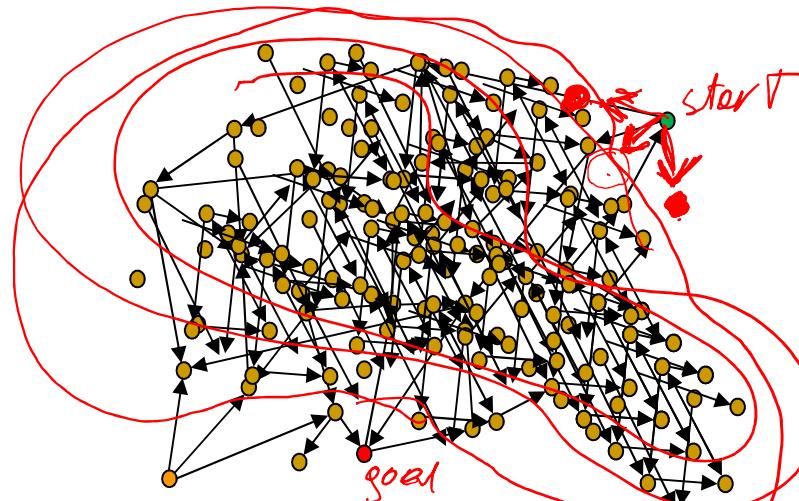
Today

1. State Space Representation 
2. State Space Search 
 - a) Overview 
 - b) Uninformed search 
 - 1. Breadth-first and Depth-first 
 - 2. Depth-limited Search 
 - 3. Iterative Deepening 
 - 4. Uniform Cost 
 - c) Informed search
 - 1. Intro to Heuristics
 - 2. Hill climbing
 - 3. Best-First
 - 4. Algorithms A & A*
 - 5. More on Heuristics
 - d) Summary

352

Problem with Uninformed Search

- inefficient for most AI problems, the state space is too large!
 - e.g. state space of all possible moves in chess = 10^{120}
 - 10^{75} = nb of molecules in the universe
 - 10^{26} = nb of nanoseconds since the "big bang"



- we need a way to try the most promising nodes first

Up Next

1. State Space Representation

2. State Space Search

a) Overview

b) Uninformed search

1. Breadth-first and Depth-first

2. Depth-limited Search

3. Iterative Deepening

4. Uniform Cost

c) Informed search

1. Intro to Heuristics $\hat{h}(n)$

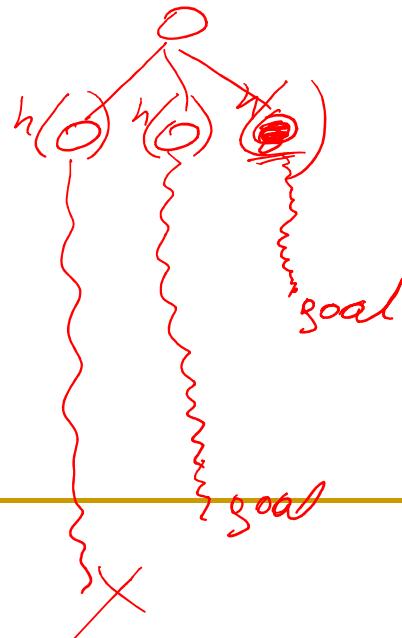
2. Hill climbing

3. Best-First

4. Algorithms A & A*

5. More on Heuristics

d) Summary



COMP 472 Artificial Intelligence

State Space Search *part #3*

Intro to Heuristics *video #3*

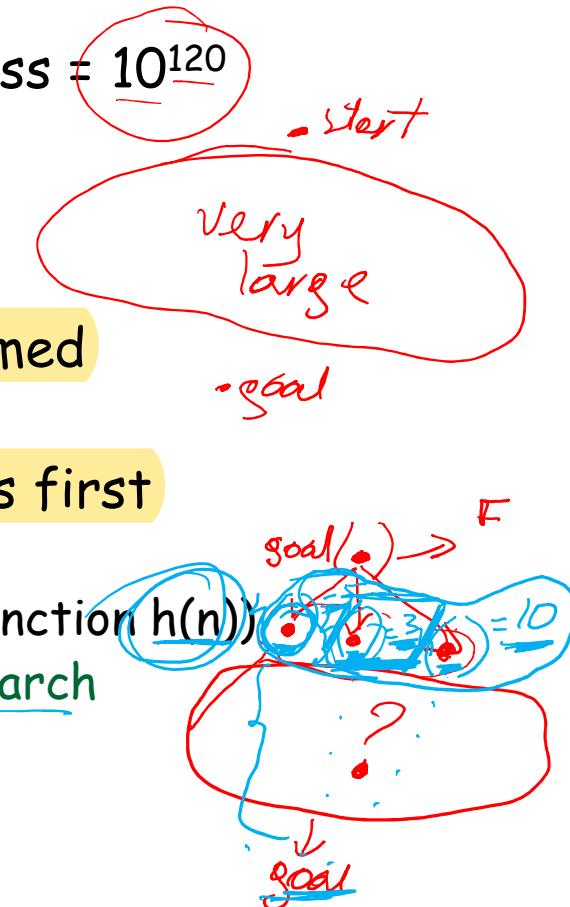
- Russell & Norvig - Sections 3.5.1, 3.5.2, 4.1.1

Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) ~~Uninformed search~~
 1. Breadth-first Search and Depth-first Search
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics $h(n)$
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Informed Search (aka heuristic search)

- Most of the time, it is not feasible to do a systematic search,
the search space is too large
 - e.g. state space of all possible moves in chess = $\underline{10^{120}}$
 - $\underline{10^{75}} = \text{nb of molecules in the universe}$
 - $\underline{10^{26}} = \text{nb of nanoseconds since the "big bang"}$
- so far, all search algorithms have been uninformed
- ie. all nodes are equally promising
- we need a way to visit the most promising nodes first
 - most-promising = close to the goal state
 - so we need a estimate function (i.e. a heuristic function $h(n)$)
 - so the search is now called informed/heuristic search



Heuristic - Heureka!



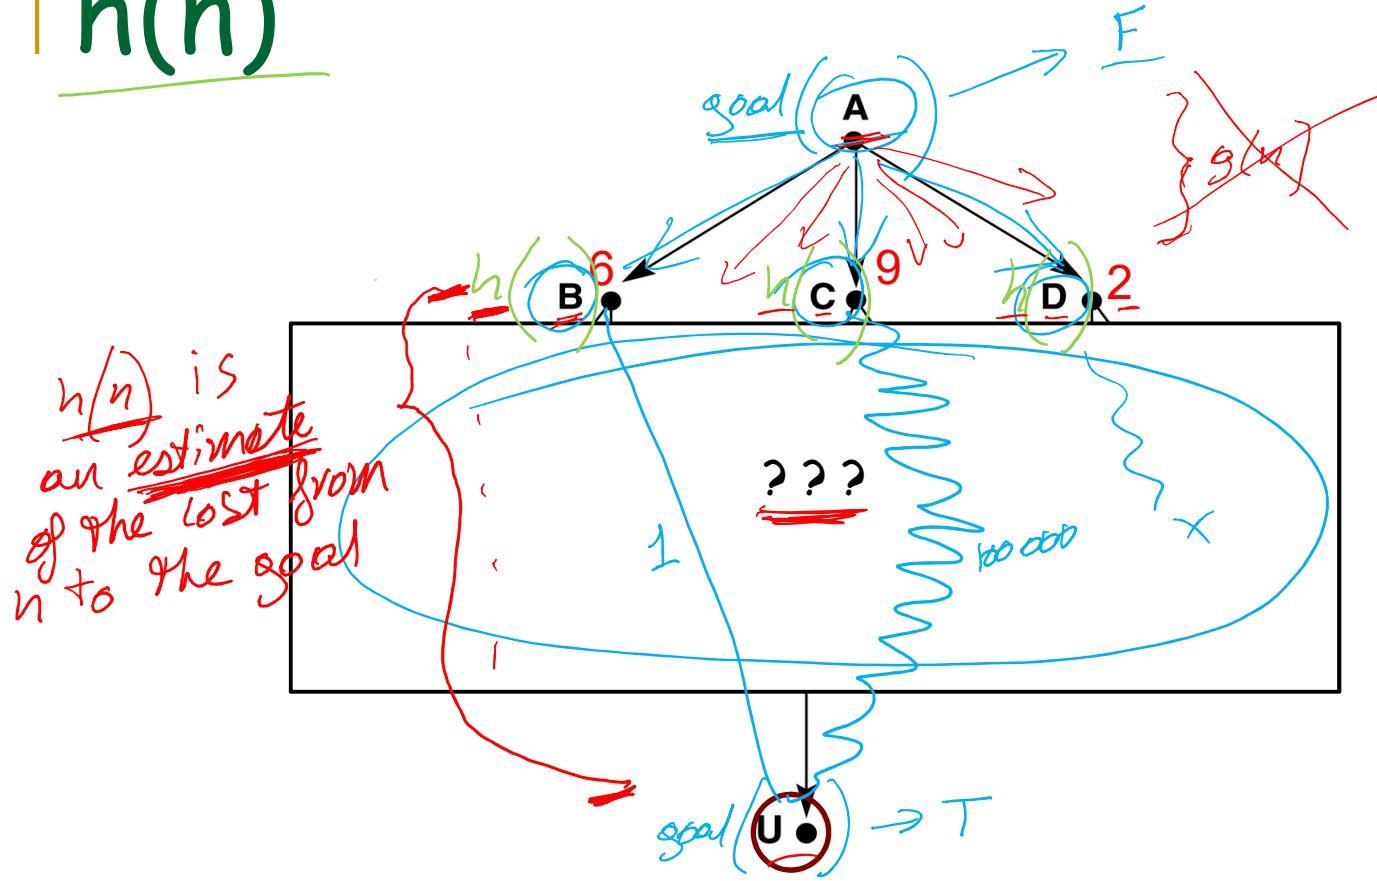
■ Heuristic search:

- A technique that improves the efficiency of search
- Focus on nodes that seem most promising according to some function $h(n)$
- Need an evaluation function (heuristic function) to estimate how close a node is to the goal
of the cost

■ Heuristic function $h(n)$:

- a rule of thumb, a good bet, an estimate
- but has no guarantee to be correct
- an approximation of the lowest cost from node n to the goal

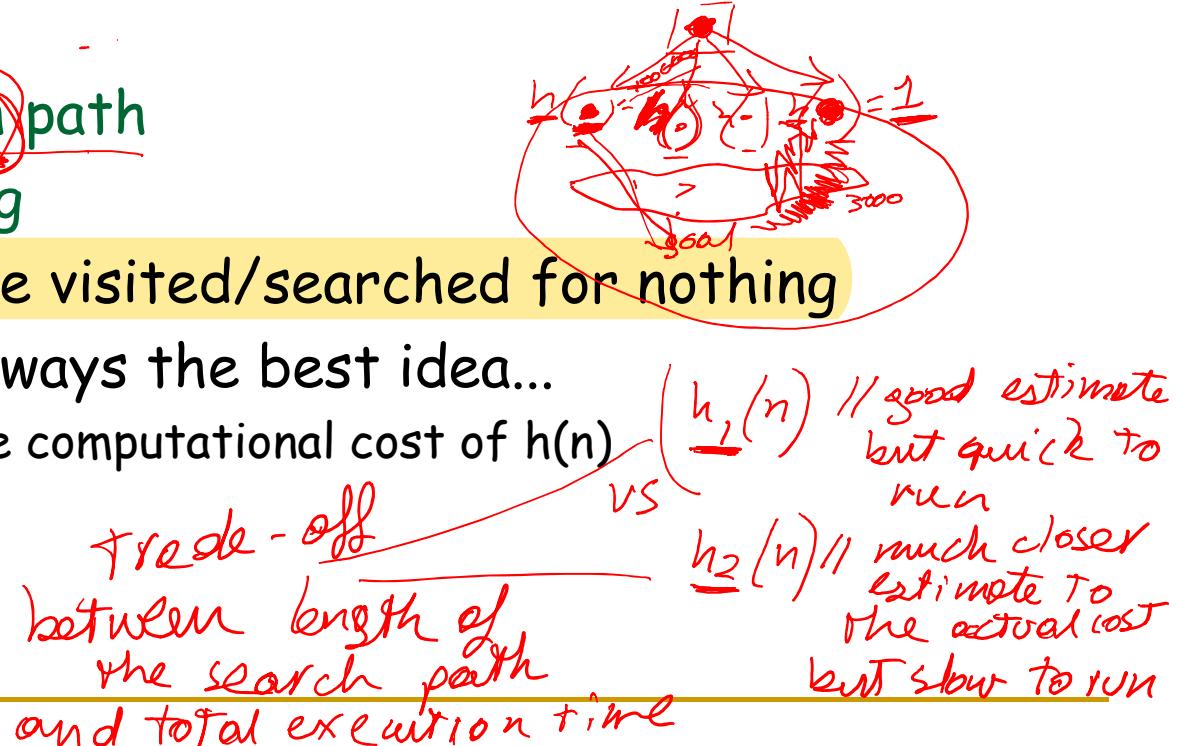
$h(n)$



- $h(n) = \text{estimate of the lowest cost from } n \text{ to goal}$

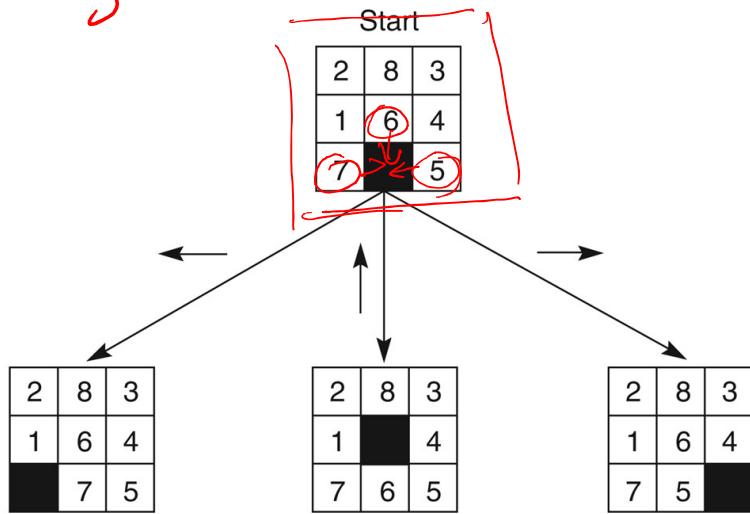
Designing Heuristics

- $h(n)$ are highly dependent on the search domain
- A $h(n)$ whose value is closer to the actual cost to the goal will lead to:
 - a shorter search path
 - less backtracking
 - i.e. less nodes are visited/searched for nothing
 - but this is not always the best idea...
 - it depends on the computational cost of $h(n)$

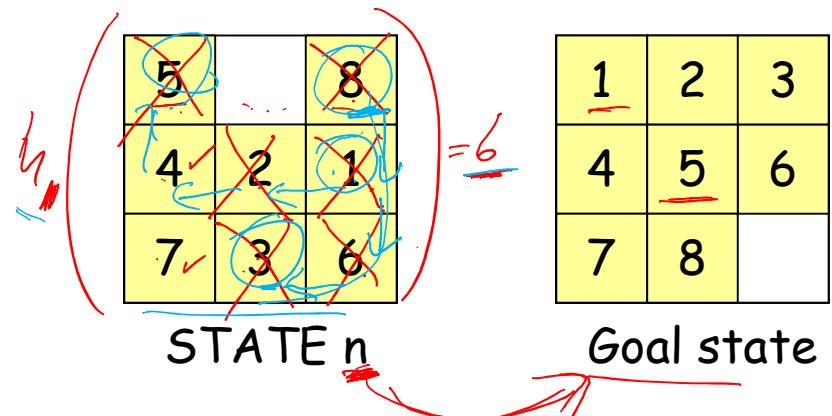


Example: 8-Puzzle - Heuristic 1

key: relax constraints of the problem to simplify it

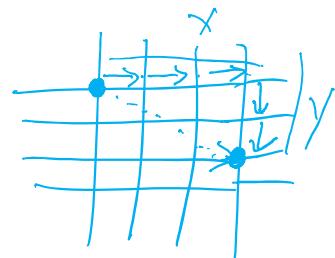
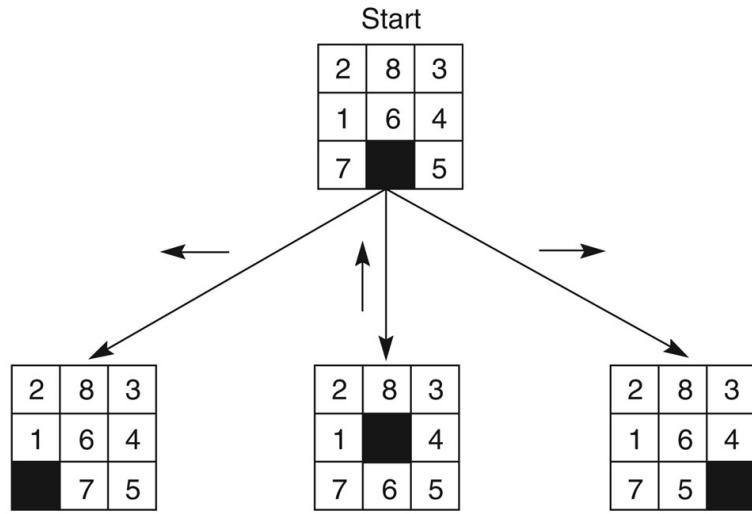


- h_1 : Simplest heuristic *numbered*
- Hamming distance : count number of tiles out of place when compared with goal

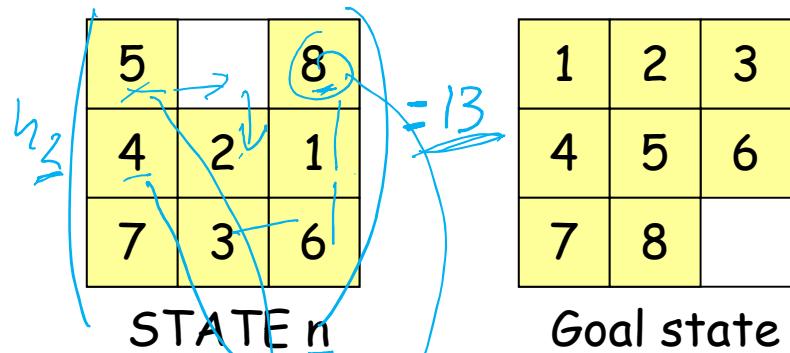


- $h_1(n) = 6$
 - does not consider the distance tiles have to be moved

Example: 8-Puzzle - Heuristic 2

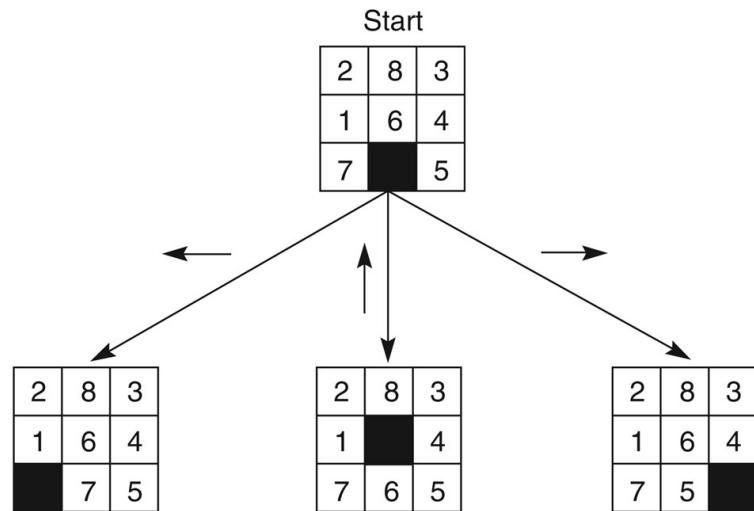


- h_2 : Better heuristic
 - Manhattan distance: sum up all the distances by which tiles are out of place



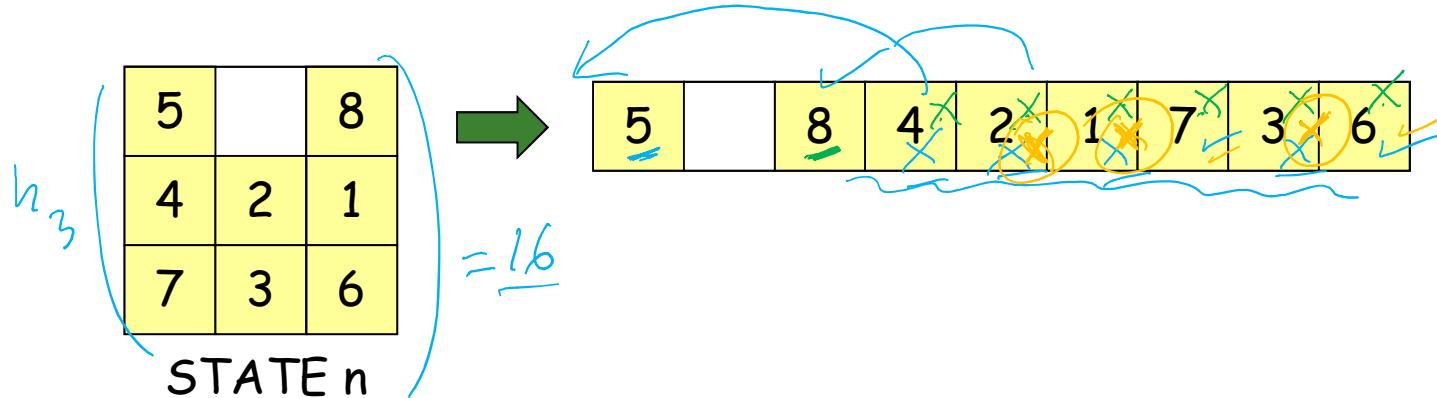
- $$h_2(n) = 2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = 13$$

Example: 8-Puzzle - Heuristic 3



- h_3 : Even Better
 - sum of permutation inversions
 - See next slide...

$h_3(N)$ = sum of permutation inversions



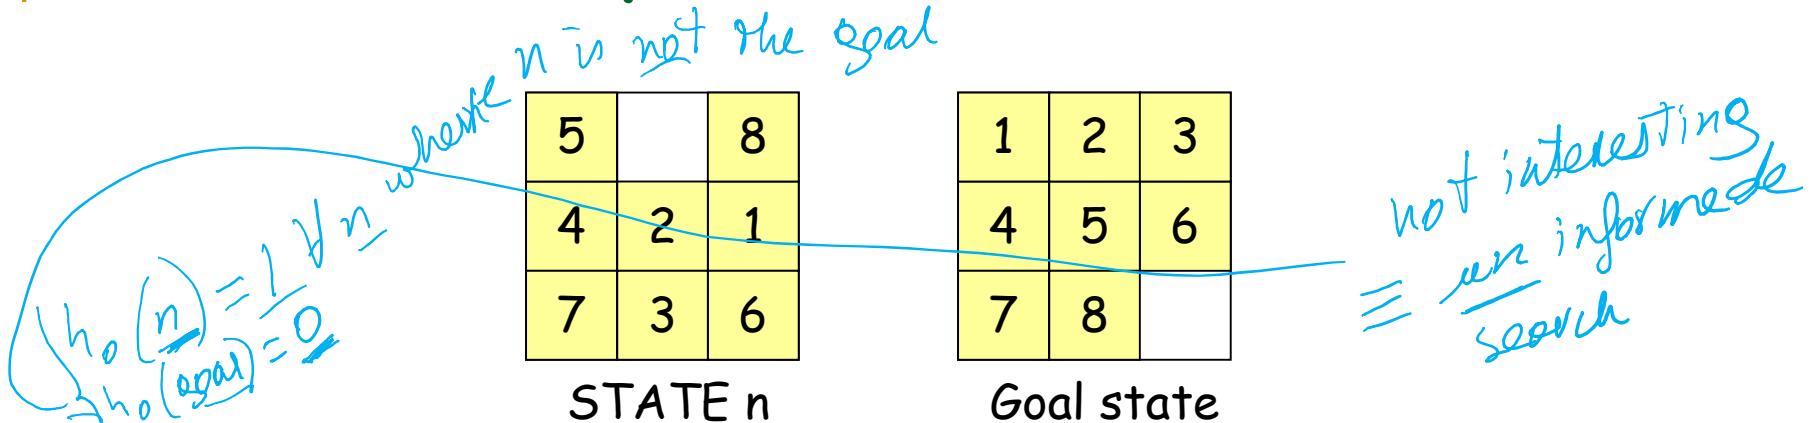
- For each numbered tile, count how many tiles on its right should be on its left in the goal state.

$$\begin{aligned} h_3(n) &= n_5 + n_8 + n_4 + n_2 + n_1 + n_7 + n_3 + n_6 \\ &= 4 + 6 + 3 + 1 + 0 + 2 + 0 + 0 \\ &= 16 \end{aligned}$$

1	2	3
4	5	6
7	8	

Goal state

Heuristics for the 8-Puzzle



- $\underline{h_1(n)} = \text{misplaced numbered tiles}$
~~= 6~~
- $\underline{h_2(n)} = \text{Manhattan distance}$
 $= 2 + 3 + 0 + 1 + 3 + 0 + 3 + 1 = \underline{13}$
- $\underline{h_3(n)} = \text{sum of permutation inversions}$
 $= n_5 + n_8 + n_4 + n_2 + n_1 + n_7 + n_3 + n_6$
 $= 4 + 6 + 3 + 1 + 0 + 2 + 0 + 0 = \underline{16}$

is $h_3(n)$ better?

- $h_3(n)$ may return a better estimate ≈ closer to the actual cost
 - → less backtracking / shorter search path
- BUT, $h_3(n)$ may be longer to compute // longer execution time
 - maybe overall longer to get the solution path
 - solution path may NOT be necessarily of lower cost (more on this later)

Today

1. State Space Representation 
2. State Space Search 
 - a) Overview 
 - b) Uninformed search 
 1. Breadth-first and Depth-first 
 2. Depth-limited Search 
 3. Iterative Deepening 
 4. Uniform Cost 
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Up Next

1. State Space Representation

2. State Space Search

a) Overview

b) Uninformed search

1. Breadth-first and Depth-first

2. Depth-limited Search

3. Iterative Deepening

4. Uniform Cost

c) Informed search

1. Intro to Heuristics

2. Hill climbing

3. Greedy Best-First Search

4. Algorithms A & A*

5. More on Heuristics

d) Summary

$h(n)$

Artificial Intelligence: State Space Search *part 3* Informed Search Hill Climbing *video #4*

- Russell & Norvig - Sections 3.5.1, 3.5.2, 4.1.1

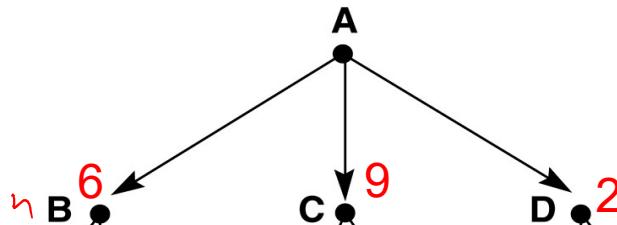
Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first Search and Depth-first Search
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill Climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

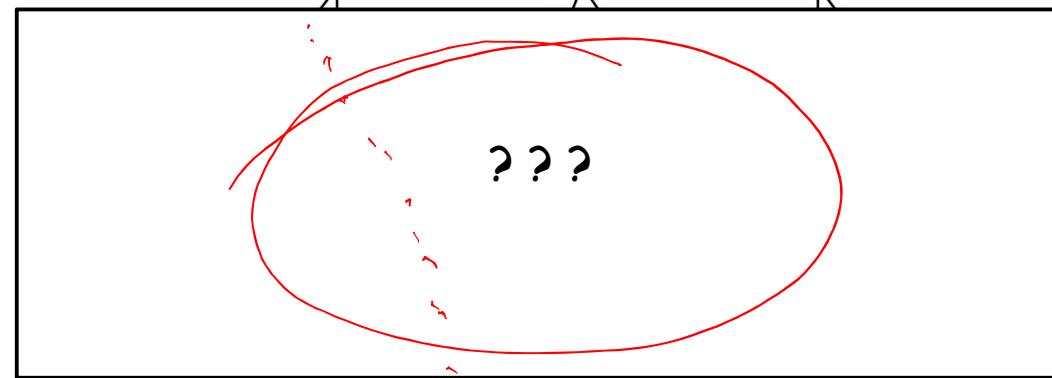


$h(n)$

~~A~~



$h(n)$

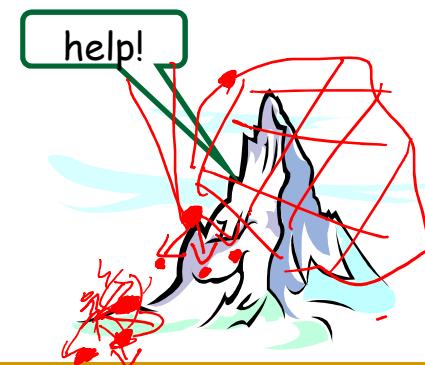


- $h(n)$ = estimate of the lowest cost from n to goal

Hill Climbing

- General idea:

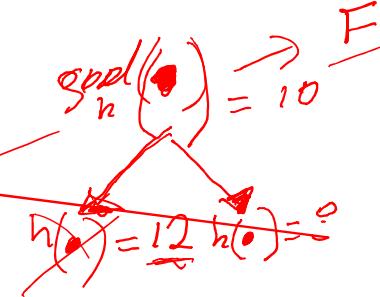
- Similar to climbing a mountain in the fog with amnesia ...
 - in the fog
 - --> only 1-step view of what is to come, so
 - if next step seems higher than where you are now -> go
 - otherwise, you assume you are at the top of the mountain -> stop
 - with amnesia -->
 - if you ever want to try other path, you can't because you did not keep track of where you came from



Vanilla HC vs Steepest Ascent HC

General Hill Climbing

- uses $h(n)$
- does not use an open list (amnesia)



1. Vanilla Hill Climbing

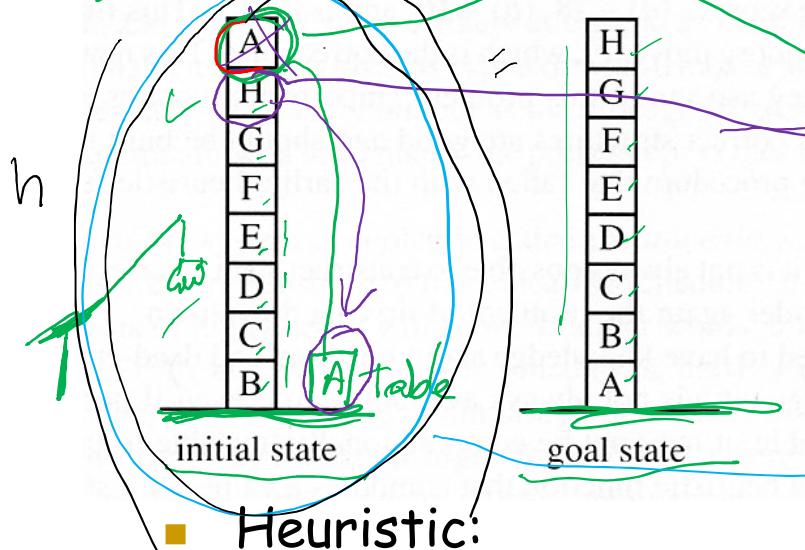
- take 1st successor s with better $h()$ than current state n
- i.e. if lower $h(n)$ is better, chose 1st s with $h(s) < h(n)$ // deep diving

2. Steepest ascent hill climbing:

- generate all successor states S
- run $h()$ on all $s \in S$
- among all successors s with better $h()$ than current state n , take the successor s with the best $h(n)$

Example: Hill Climbing

Block World



Operators:

- **pickup&putOnTable (Block)**
- **pickup&stack (Block1,Block2)**

pickup from top
of a stop & place
on the table

from the top of a stack
& stack on top of a stack

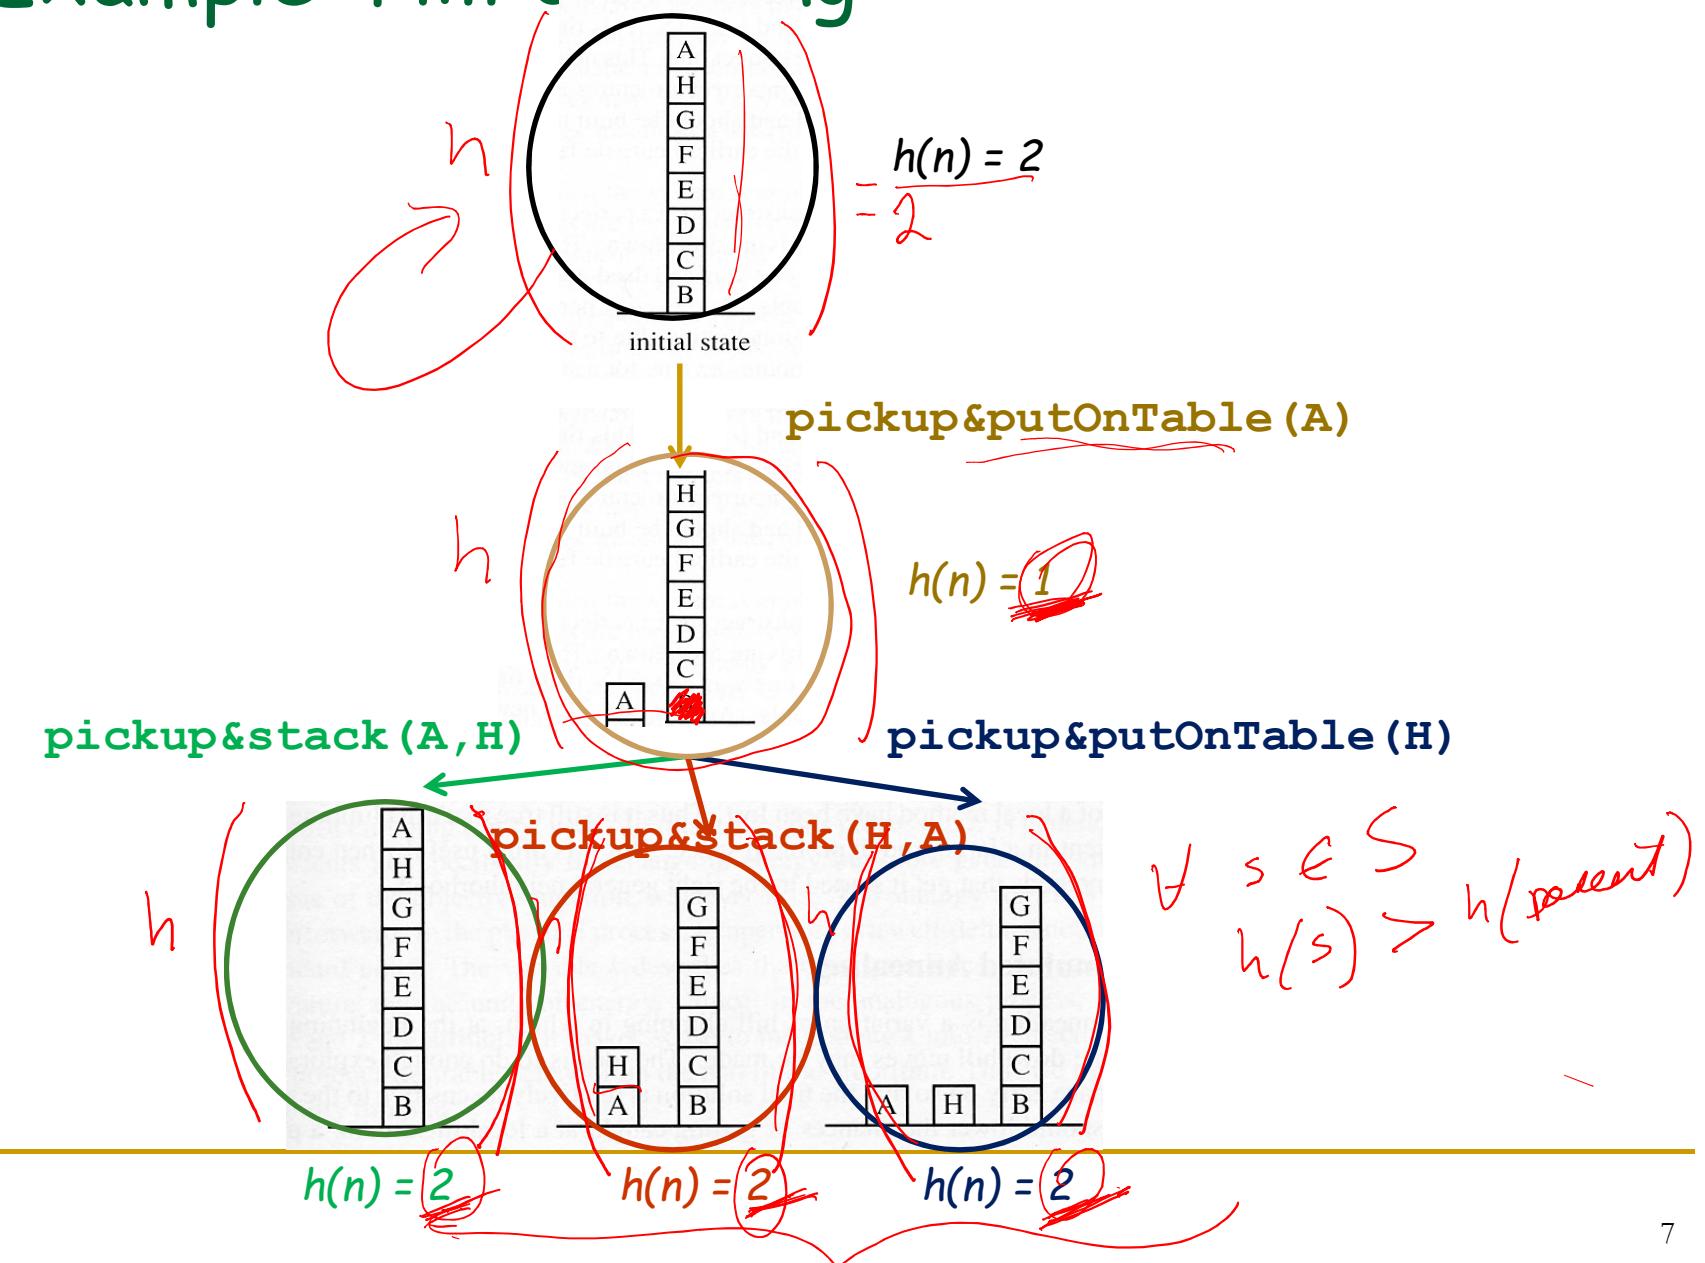
Heuristic:

- Opt if a block is sitting where it is supposed to sit
- +1pt if a block is NOT sitting where it is supposed to sit
- so lower $h(n)$ is better
 - $h(\text{initial}) = 2$
 - $h(\text{goal}) = 0$

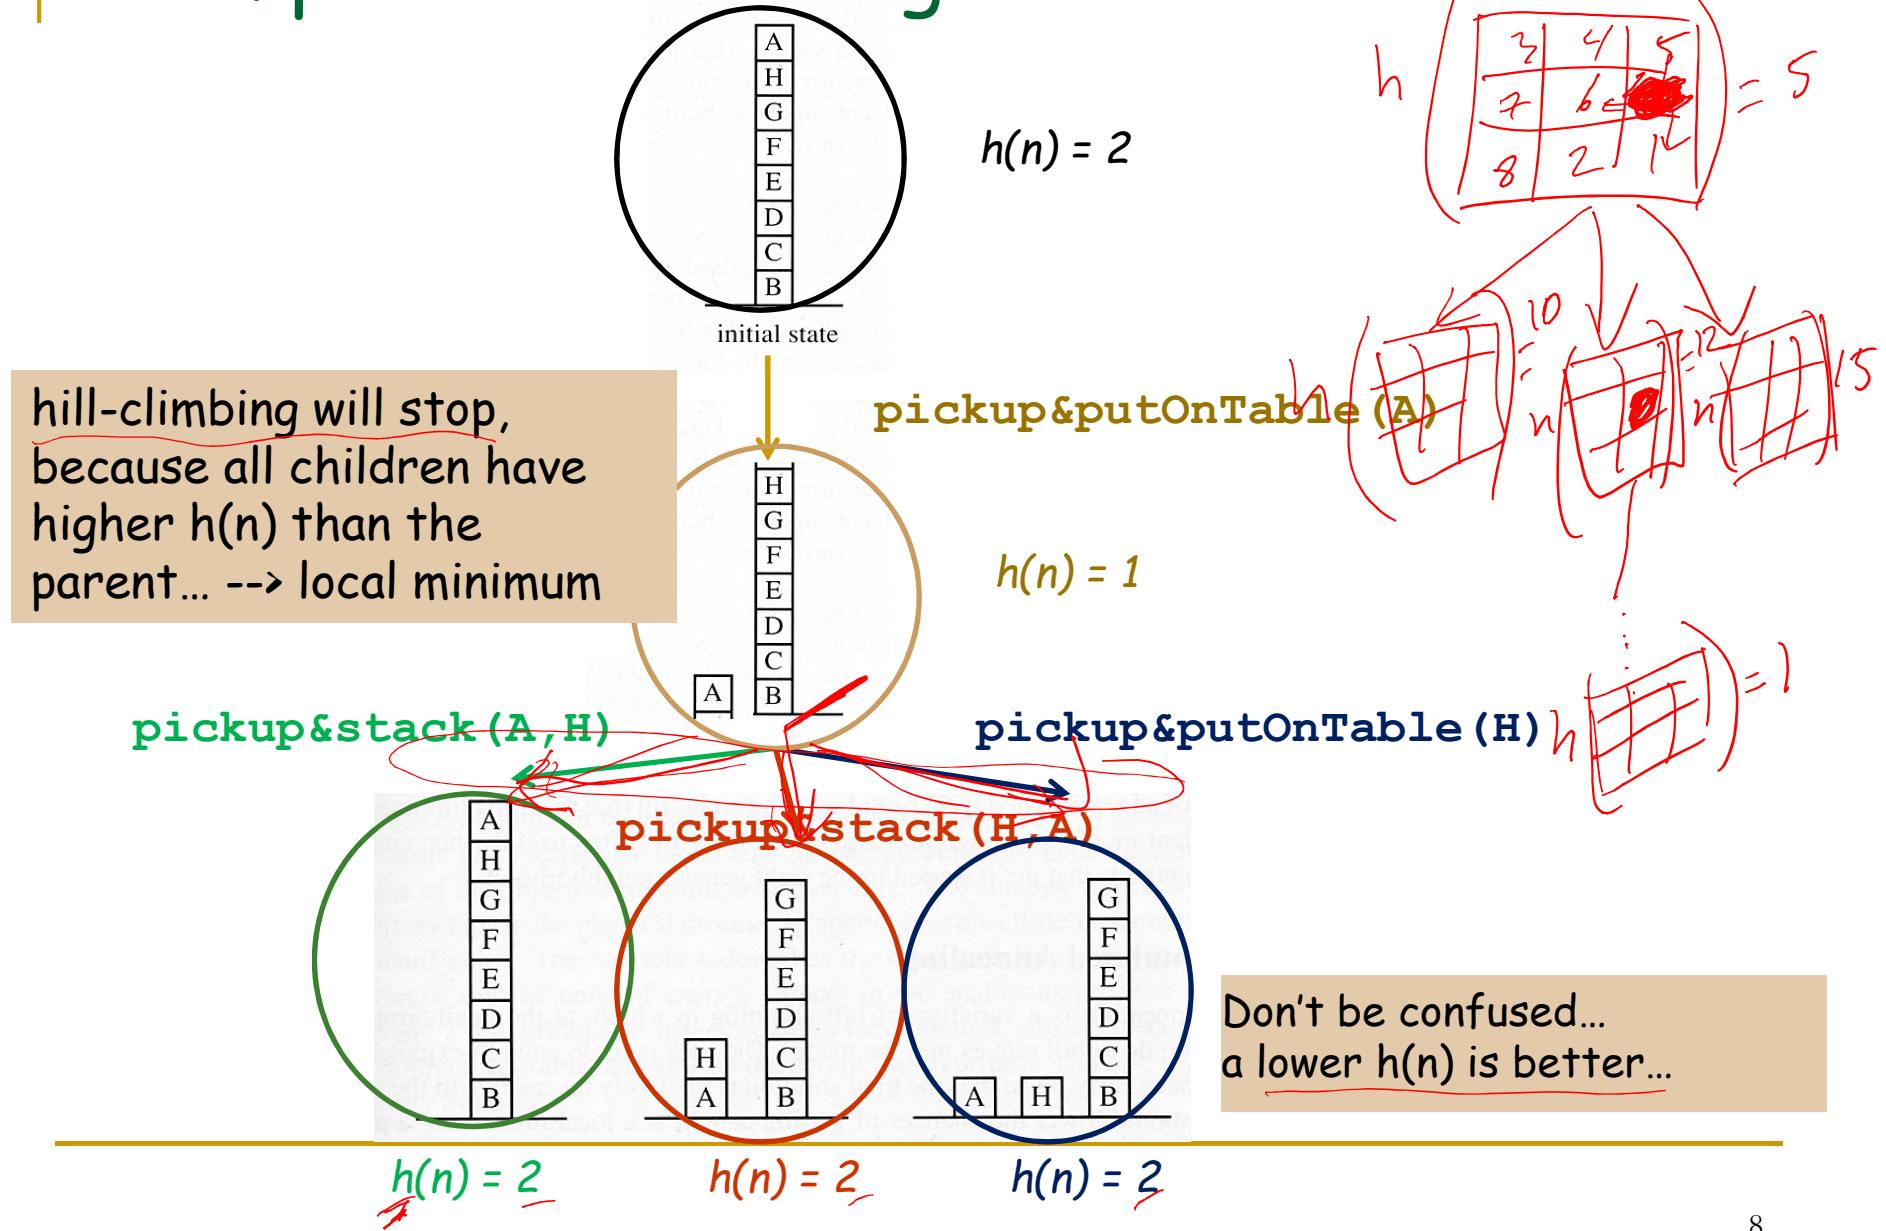


source: Rich & Knight, Artificial Intelligence, McGraw-Hill College 1991.

Example: Hill Climbing



Example: Hill Climbing



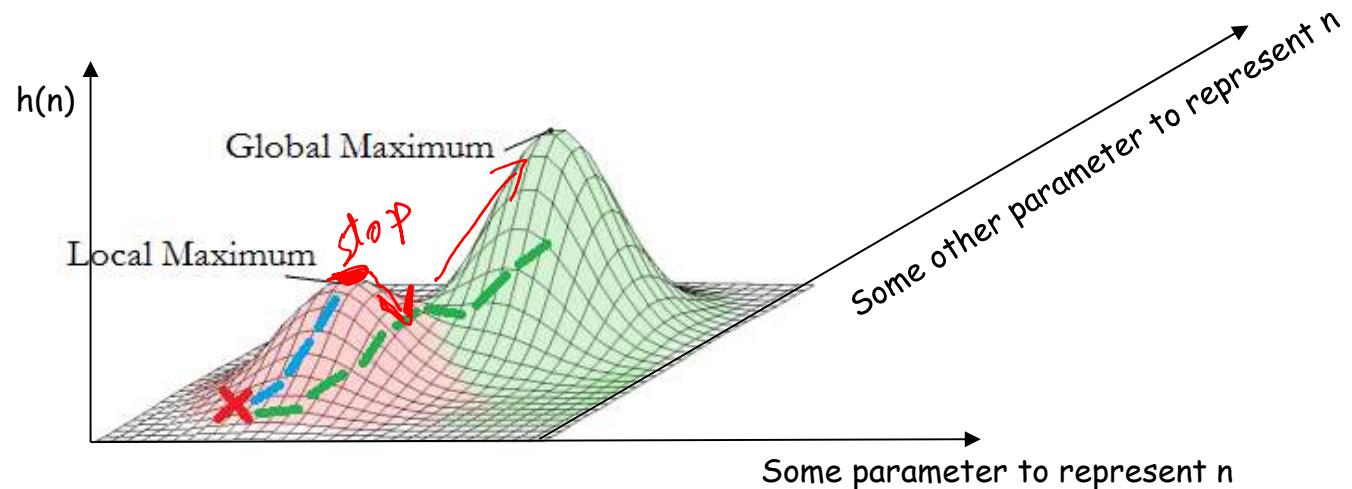
Steepest Ascent Hill Climbing

```
currentNode = startNode;  
loop do  
    L = CHILDREN(currentNode);  
    nextEval = +INFINITY;  
    nextNode = NULL;  
  
    for all c in L  
        if (HEURISTIC-VALUE(c) < nextEval) // lower h is better  
            nextNode = c;  
            nextEval = HEURISTIC-VALUE(c);  
  
    if nextEval >= HEURISTIC-VALUE(currentNode)  
        // Return current node since no better child state exist  
        return currentNode;  
  
    currentNode = nextNode;
```

Problems with Hill Climbing

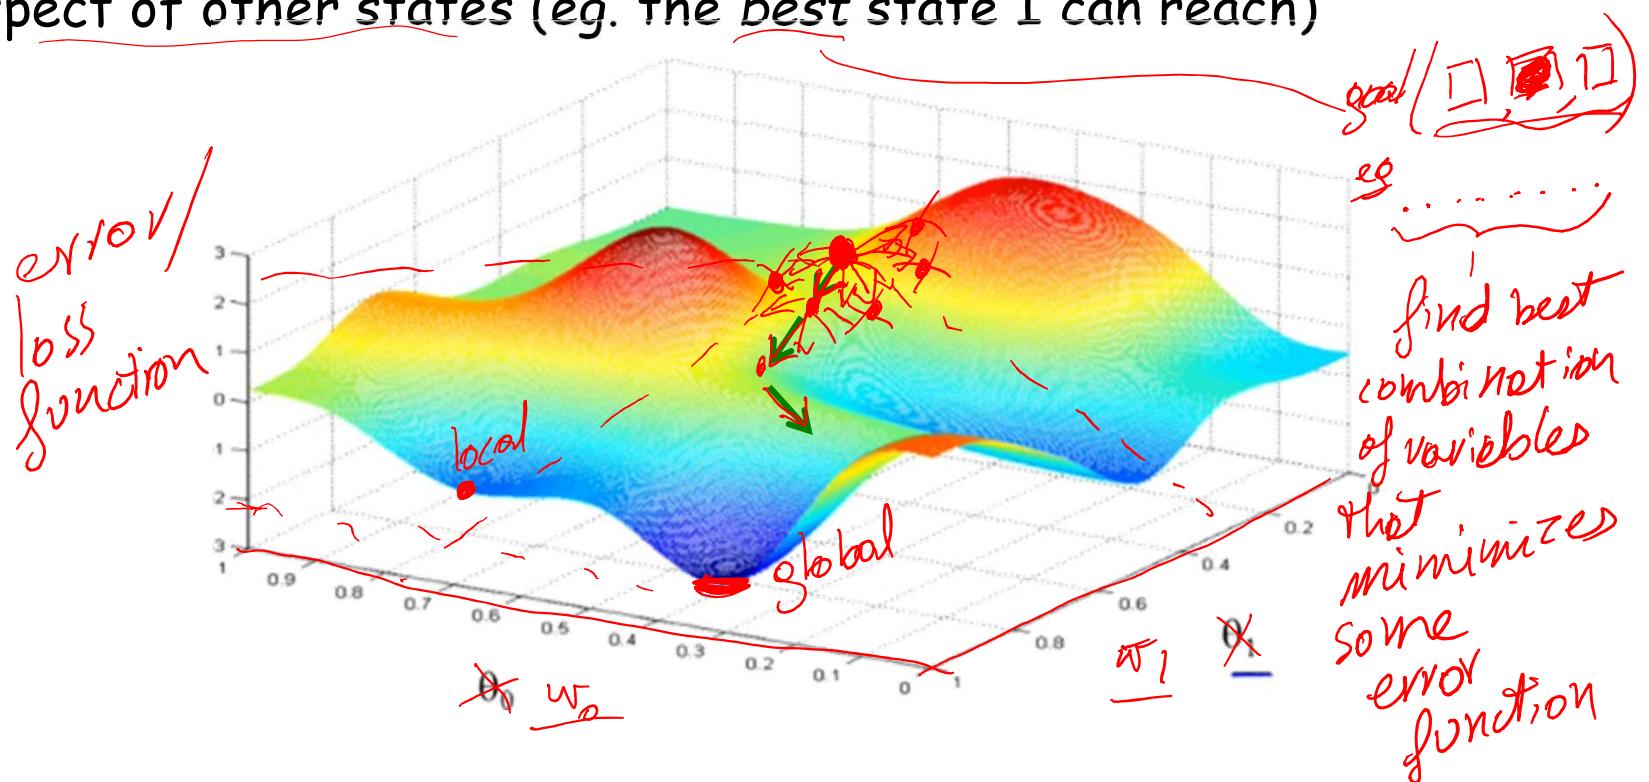
■ Foothills (or local $\overset{\text{optimum}}{\text{maxima}}$)

- ❑ reached a local maximum, not the global maximum
- ❑ a state that is better than all its neighbors but is not better than some other states farther away.
- ❑ at a local maximum, all moves appear to make things worse.
- ❑ ex: 8-puzzle: we may need to move tiles temporarily out of goal position in order to place another tile in goal position



Use of Hill Climbing

- mostly for optimization problems
- i.e. goal defined not as a function of the state alone, but with respect of other states (eg. the best state I can reach)



Today

1. State Space Representation 
2. State Space Search 
 - a) Overview 
 - b) Uninformed search 
 1. Breadth-first Search and Depth-first Search 
 2. Depth-limited Search 
 3. Iterative Deepening 
 4. Uniform Cost 
 - c) Informed search 
 1. Intro to Heuristics 
 2. Hill climbing  
 3. Greedy Best-First Search $h(n)$
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Problem with Hill-Climbing

- used mostly for optimization problems
 - where the goal state is defined with respect to other states
 - ex. shortest path, longest....
- if goal state is independent of other states
 - we should be able to backtrack, and find another path to the goal
 - i.e. we should use an OPEN list
 - i.e. ~~Greedy~~ Best First Search

Up Next

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first and Depth-first
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

Artificial Intelligence: State Space Search } part 3 Informed Search Greedy Best First Search and } video Algorithms A and A* } #5

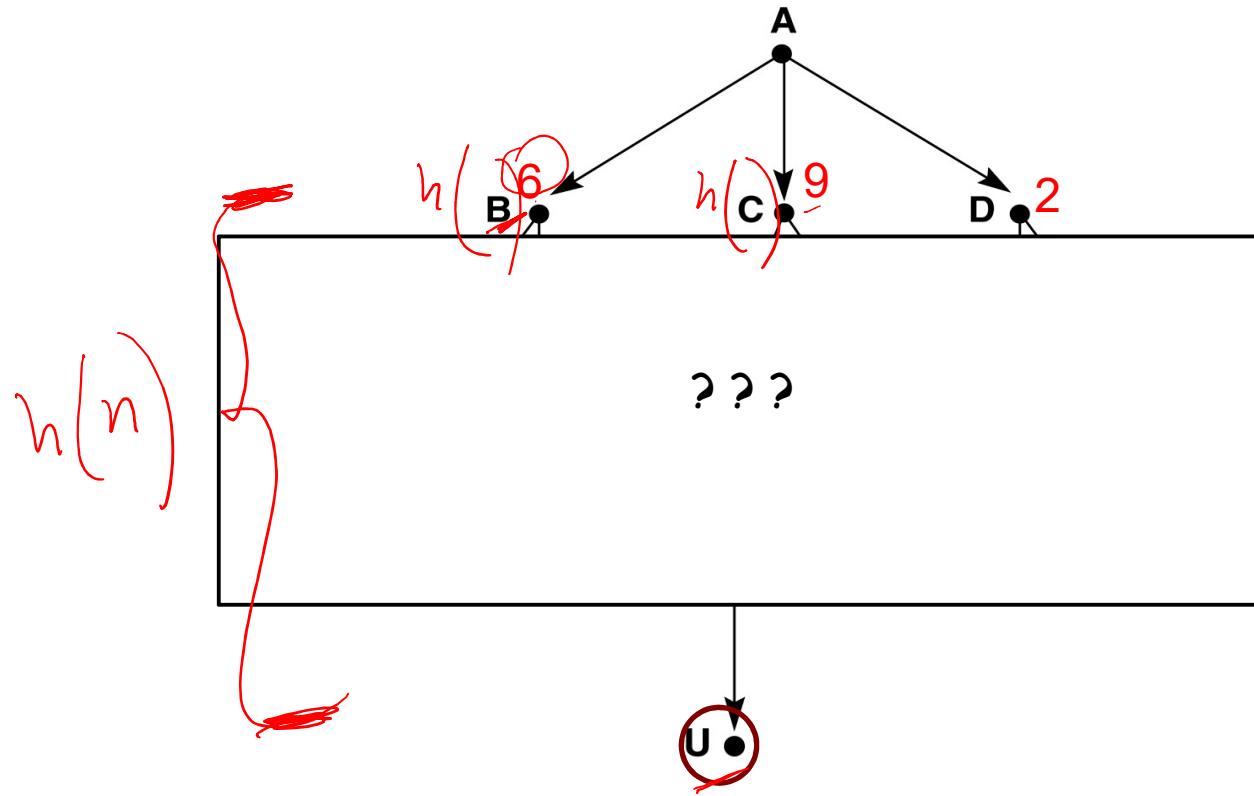
- Russell & Norvig - Sections 3.5.1, 3.5.2, 4.1.1

Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first Search and Depth-first Search
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary



$h(n)$



- $h(n)$ = estimate of the lowest cost from n to goal

Greedy Best-First Search

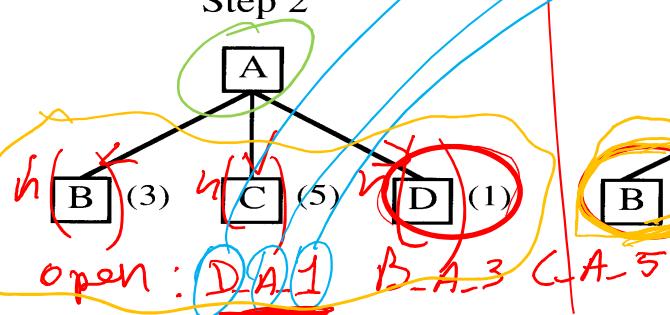
- } problem with hill-climbing:
 - no open list
 - --> can't backtrack
 - one move is selected and all others are forgotten
- solution to hill-climbing:
 - use "open" as a priority queue $h(n)$.
 - this is called best-first search
- Best-first search:
 - Insert nodes in open list so that the nodes are sorted in ascending $h(n)$
 - Always choose the next node to visit to be the one with the best $h(n)$ -- regardless of where it is in the search space lowest

GBF: Example

Step 1

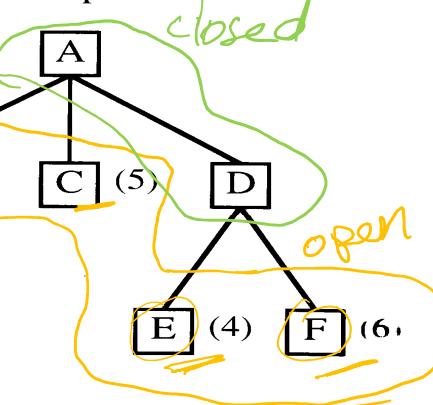


Step 2

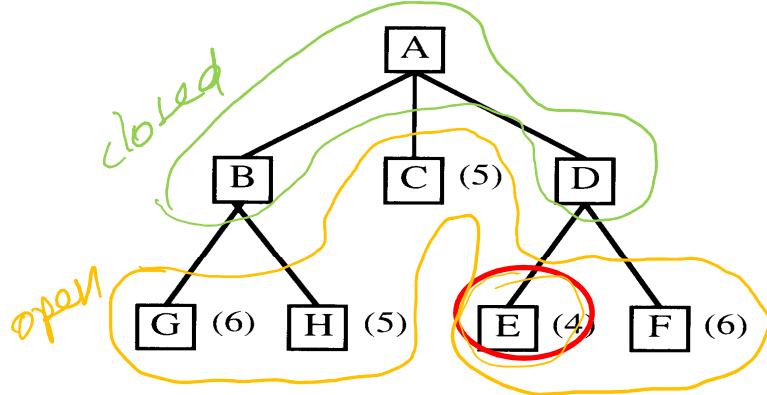


Lower $h(n)$ is better

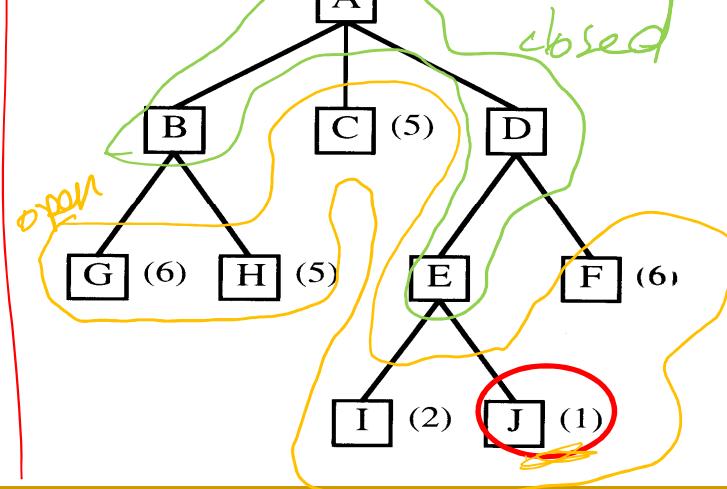
Step 3



Step 4



Step 5



source: Rich & Knight, Artificial Intelligence, McGraw-Hill College 1991.

Notes on GBF

- If you have a good $h(n)$, best-first can find a solution very quickly
- The solution may not be the optimal one (lowest cost) but there is a good chance of finding it quickly

GBF Search: Example



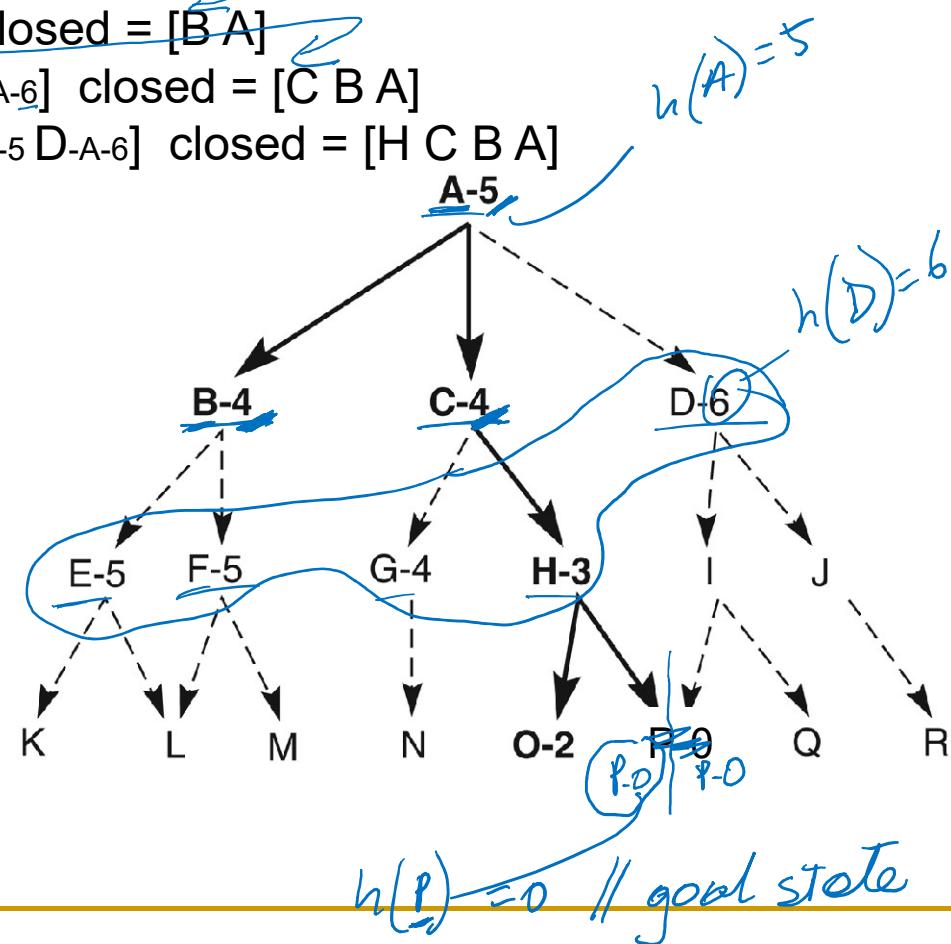
node parent h(node)

1. open = [A-null-5] closed = []
2. open = [B-A-4 C-A-4 D-A-6] (arbitrary choice) closed = [A]
3. open = [C-A-4 E-B-5 F-B-5 D-A-6] closed = [B A]
4. open = [H-C-3 G-C-4 E-B-5 F-B-5 D-A-6] closed = [C B A]
5. open = [P-H-0 O-H-2 G-C-4 E-B-5 F-B-5 D-A-6] closed = [H C B A]
6. goal P found

solution path: A C H P

priority queue
sorted by $h(n)$

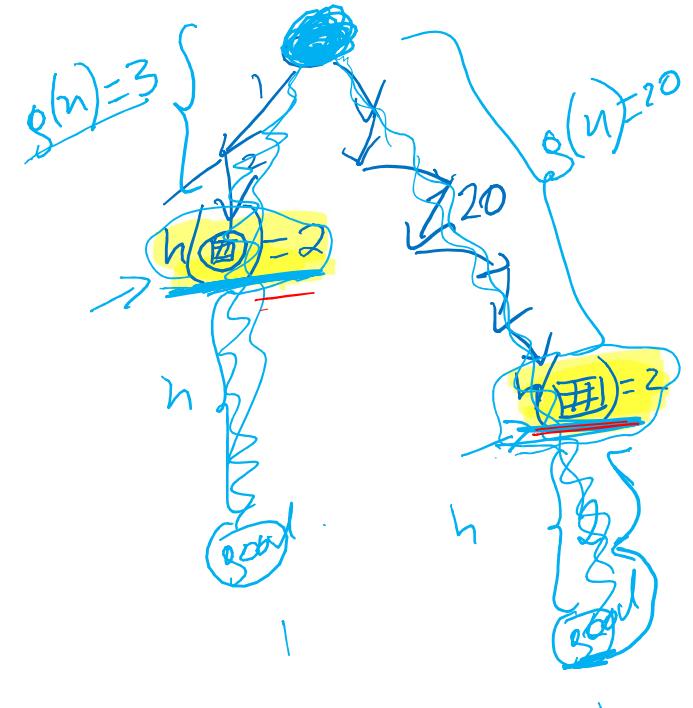
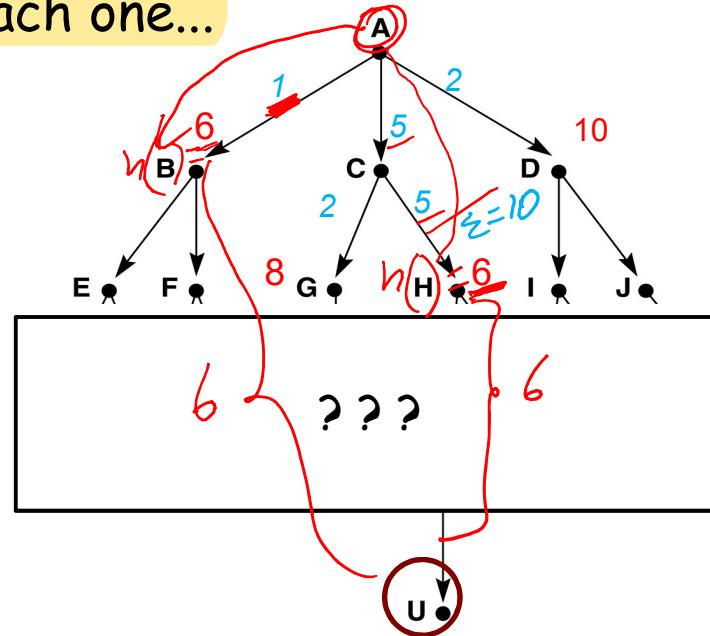
Lower $h(n)$ is better



source: adapted from G. Luger (2005)

Problem with GBF search

- if 2 nodes have the same $h(n)$, no preference to the closest/least costly to reach one...



- Solution:
 - Maintain a cost count - $g(n)$
 - i.e. give preference to nodes with least expensive paths from root to n
 - i.e. combine $h(n)$ and $g(n)$

Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first Search and Depth-first Search
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary



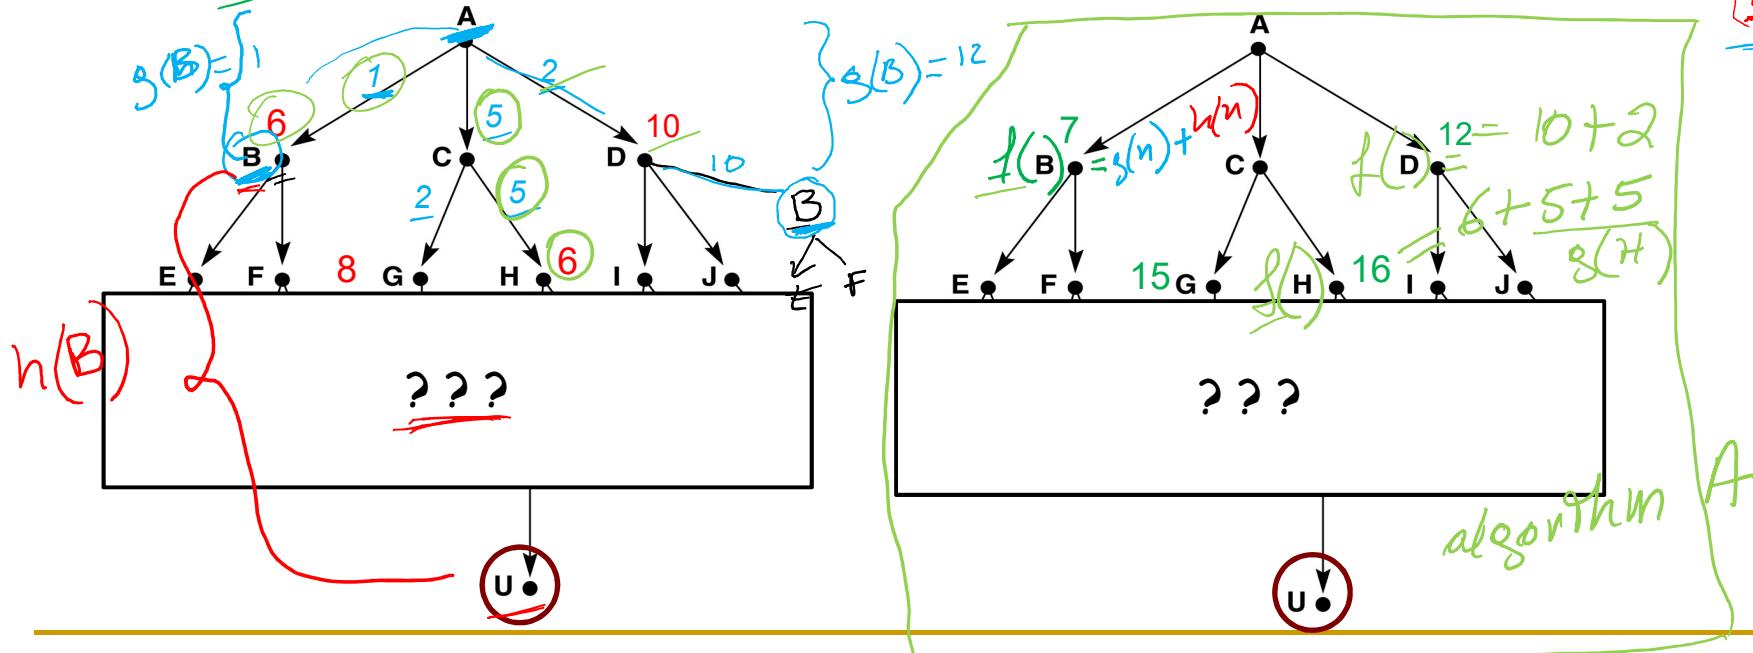
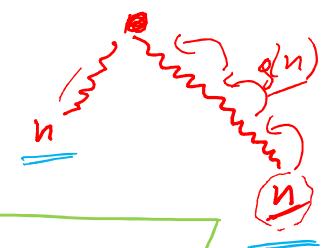
$$f(n) = h(n) + g(n)$$

- Modified evaluation function f :

$$f(n) = \boxed{g(n) + h(n)}$$

- $f(n)$ estimate of total cost along path through n
- $g(n)$ actual cost of path from start to node n
- $h(n)$ estimate of cost to reach goal from node n

from the start + to
the goal



Algorithms A and A*

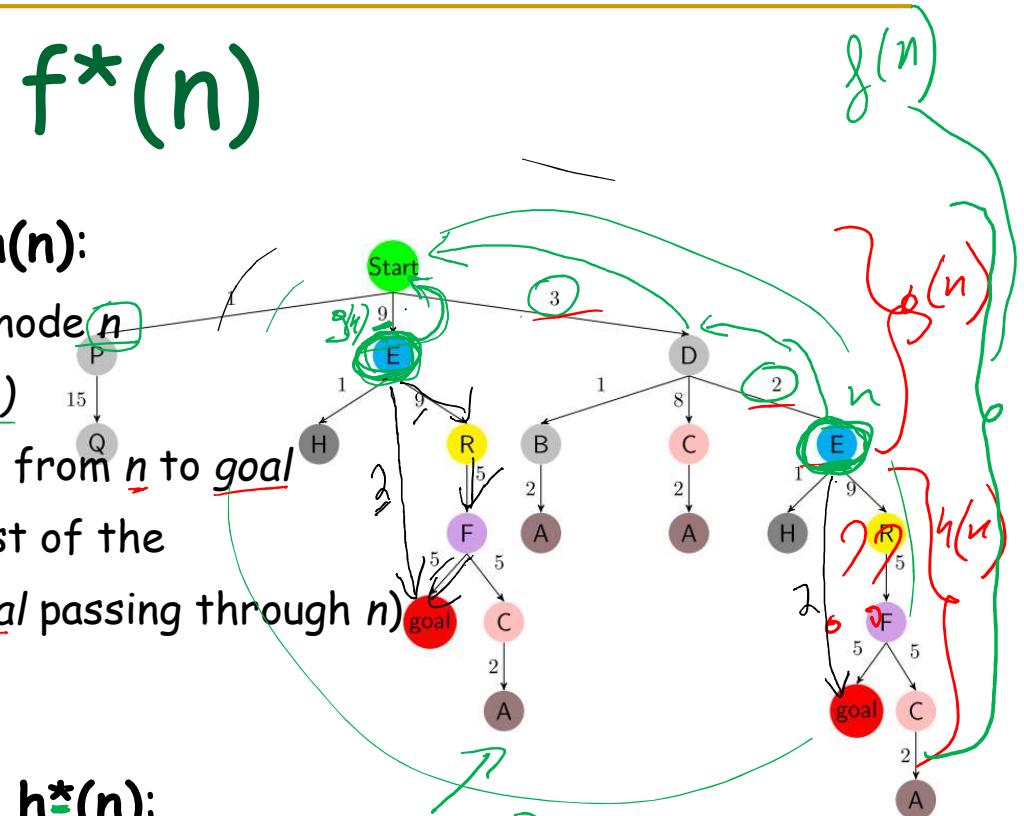
? (see next slides)

- similarly to Greedy Best first search:
 - keep an OPEN list as a priority queue
- But
 - OPEN is sorted by lowest $f(n) = h(n) + g(n)$

$g(n)^*$, $h(n)^*$ and $f^*(n)$

- We know that $f(n) = g(n) + h(n)$:

- $g(n)$ current cost from start to node n
(maybe not be the lowest cost)
- $h(n)$ estimate of the lowest cost from n to goal
---> $f(n)$ estimate of the lowest cost of the solution path (from start to goal passing through n)



- Let us define $f^*(n) = g^*(n) + h^*(n)$:

- $g^*(n)$ cost of lowest cost path from start to node n
- $h^*(n)$ actual lowest cost from n to goal // unknown... what $h(n)$ is trying to estimate
---> $f^*(n)$ actual cost of lowest cost of the solution path (from start to goal passing through n)

for ex in the graph above
 $h^*(E) = 2$

Algorithm A vs Algorithm A*

■ IF

- $\underline{g(n) \geq g^*(n)} \quad \forall n$

// i.e. if the cost from the root to n is considered

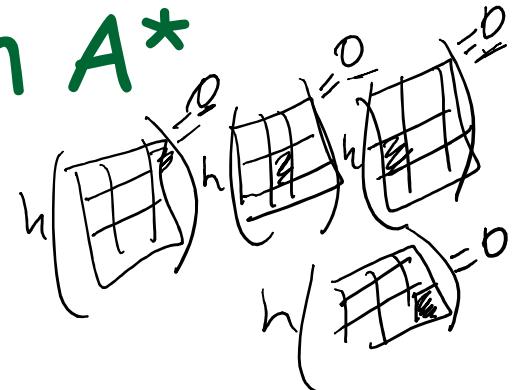
■ AND $\overset{\text{overestimate}}{h(n)}$ $\overset{\text{actual cost}}{h^*(n)}$

- $\underline{h(n) \leq h^*(n)} \text{ for all } n \quad \forall n$

// i.e. $h(n)$ never overestimates the true lowest cost from n to the goal

■ THEN

- algorithm A is called algorithm A*



$h^*(n)$

uniform cost

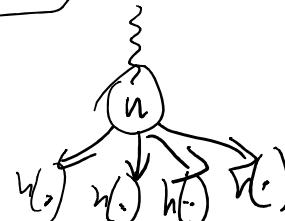
- uses $g(n)$

- uses $h(n)$ where

$h(n) = 0 \quad \forall n$

what's the big deal?

--> algorithm A* is admissible



--> i.e. it guarantees to find the lowest cost solution path from the initial state to the goal

Algorithm A* vs GBF search

- given the same $h(n)$:
 - A* guarantees to find the lowest cost solution path
 - GBF does not
- so is A* always "better" in real life?
 - not necessarily
 - computing $g(n)$ can take time to compute
 - if client is not looking for the optimal (lowest cost) solution
 - a good-enough solution faster (i.e. GBF search) might be preferable

Today

1. State Space Representation 
2. State Space Search 
 - a) Overview 
 - b) Uninformed search 
 1. Breadth-first and Depth-first 
 2. Depth-limited Search 
 3. Iterative Deepening 
 4. Uniform Cost 
 - c) Informed search
 1. Intro to Heuristics 
 2. Hill climbing 
 3. Greedy Best-First Search 
 4. Algorithms A & A* 
 5. More on Heuristics
 - d) Summary

Up Next

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first and Depth-first
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

COMP 472 Artificial Intelligence

State Space Search

Informed Search ^{part 3}

More on Heuristics & Summary ^{video b}

- Russell & Norvig - Section 3.5.2

Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first and Depth-first
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary



Evaluating Heuristics

1.

Admissibility:

- ❑ "optimistic"
- ❑ $h(n)$ never overestimates the actual cost of reaching the goal
- ❑ guarantees to find the lowest cost solution path to the goal (if it exists) A^*

$$\forall n \ h(n) = 0$$

$$\forall n \ h(n) \leq h^*(n)$$

2.

Monotonicity:

- ❑ "local admissibility"
- ❑ guarantees to find the lowest cost path to each state n visited (i.e. popped from OPEN)

not informed

3.

Informedness:

- ❑ measure for the "quality" of a heuristic
- ❑ the more informed, the less backtracking, the shorter the search path

Admissibility

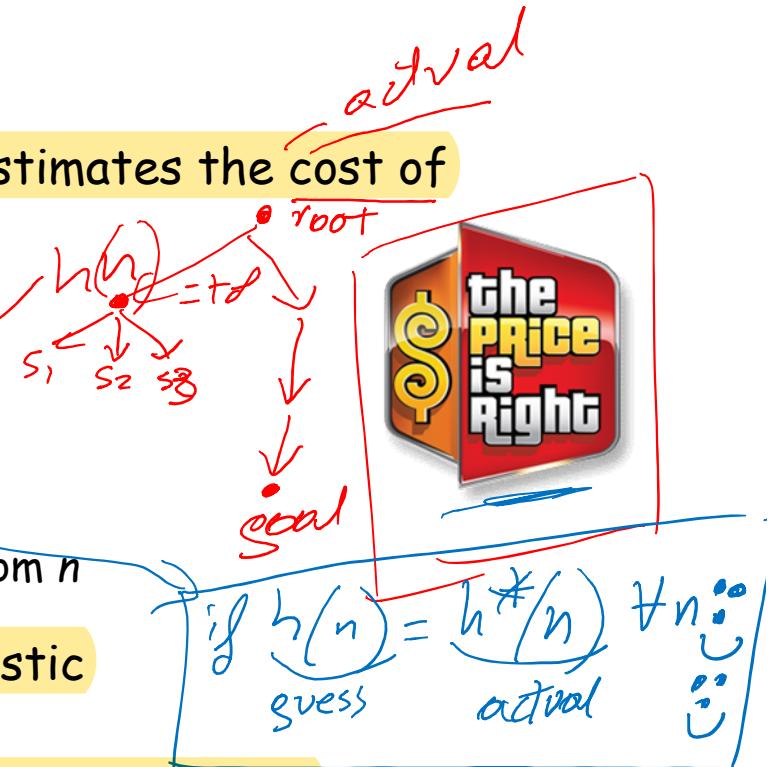
- A heuristic is **admissible** if it never overestimates the **cost** of reaching the goal

- i.e.:
 - $h(n) \leq h^*(n)$ for all n
- hence
 - $h(\text{goal}) = h^*(\text{goal}) = 0$
 - $h(n) = \infty$ if we cannot reach the goal from n

- Algorithm A that uses an admissible heuristic

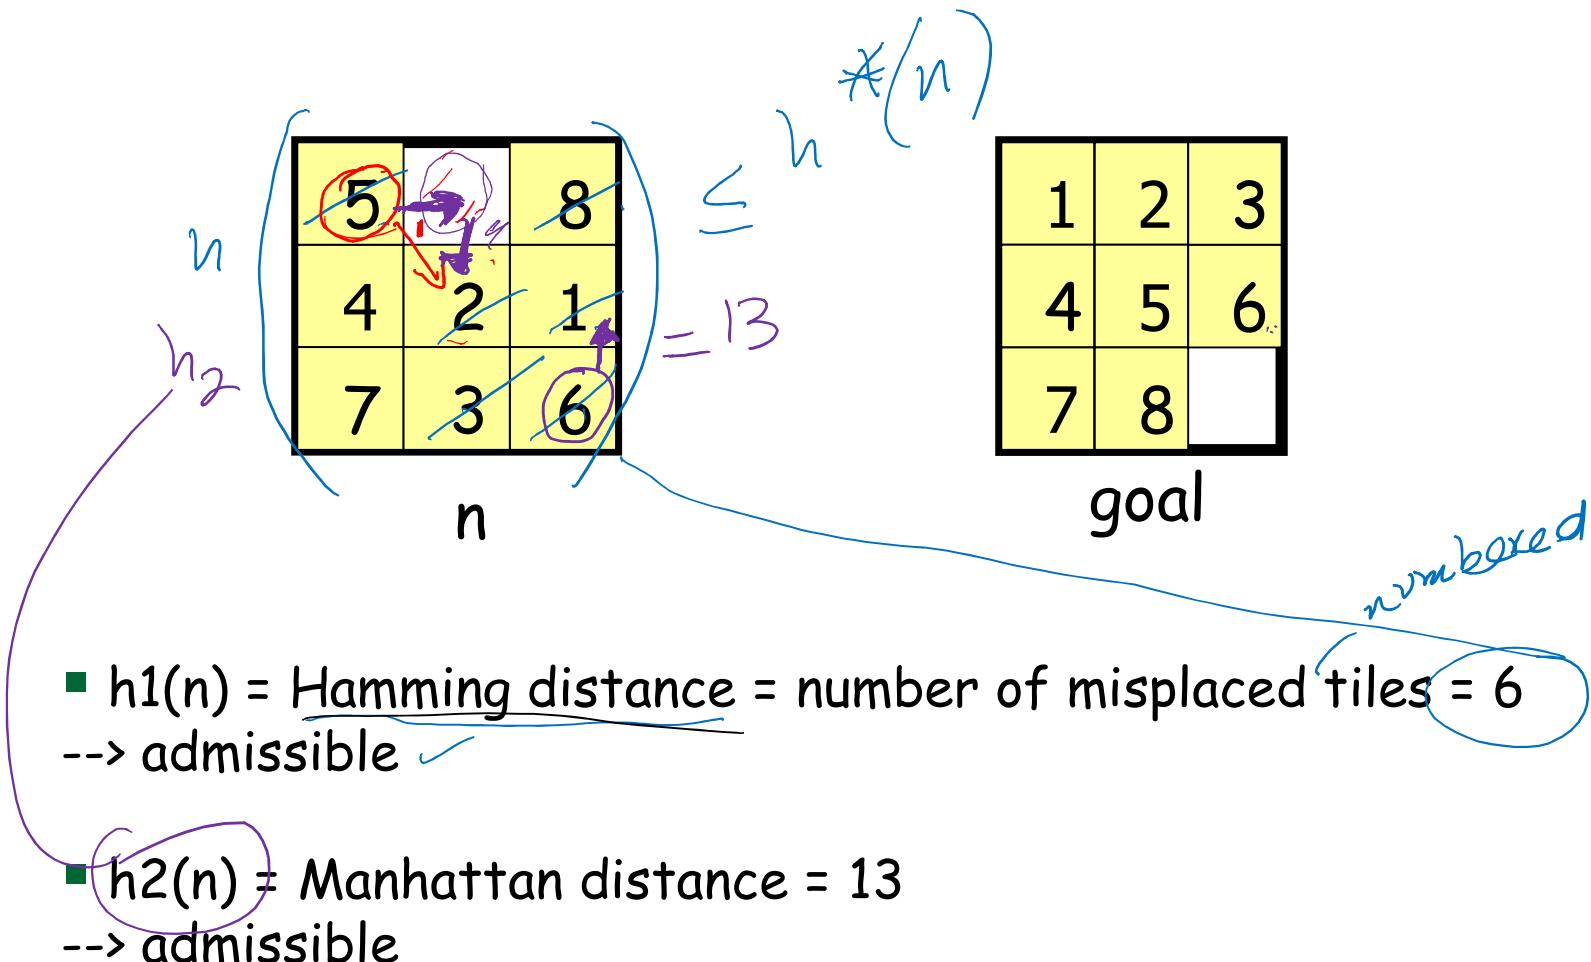
- is called algorithm **A***
- guarantees to find the lowest cost solution path to the goal (if it exists)
- note: does not guarantee to find the lowest cost search path
- e.g.: uniform cost is admissible -- it uses $f(n) = g(n) + 0$

i.e. you can
backtrack



if $h(n) = 0$ then
it is admissible
but uninformative

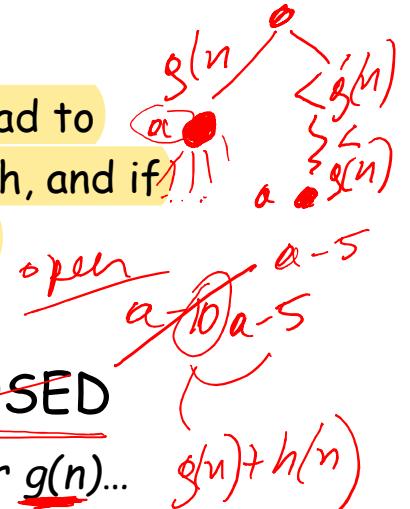
Example: 8-Puzzle



Problem with Admissibility

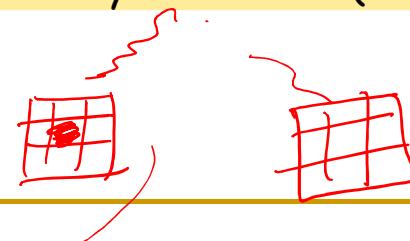
- Admissible heuristics may temporarily reach non-goal states along a suboptimal path

- remember with uniform-cost... when we expanded a node, we had to check if it was already in the OPEN list with a higher path cost, and if so, we would replace it with the current path cost/parent info



- With A^* , if we have a node n in OPEN or even in CLOSED

- We may later find n again, but with a lower $f(n)$ (due to a lower $g(n)$)... the $h(n)$ will, by definition be the same).
 - So to ensure that the solution path has the lowest cost,
 - We may need to update the cost/parent info of node n in OPEN or even put n back in OPEN even if it has already been visited (i.e. in CLOSED)... expensive work...



Admissibility and A* Search

- Admissibility: $\forall n \ h(n) \leq h^*(n)$
 - To guarantee to find the lowest cost solution path, when we generate a successors s: i.e. when node n is popped from OPEN
 1. IF s is already in CLOSED
 - IF s in CLOSED has a higher f-value due to a higher g-value i.e $g(s)$
THEN place s and its new lower f-value in OPEN!
 - // we found a lower cost path to s, but we had already expanded s...
 - // to guarantee the lowest cost solution path, we need to put s back in OPEN and re-visit it again
 - ELSE ignore s
 2. ELSE IF s is already in OPEN
 - IF s in OPEN has a higher f-value
THEN replace the old s in OPEN with the new lower f-value
 - // we found a lower cost path to s, and we had not expanded s yet
 - // to guarantee the lowest cost solution path, we need to replace the old s in OPEN with the new lower-cost s
 - ELSE ignore s
 3. ELSE insert s in OPEN
 - // as usual

Monotonicity (aka consistent)

- Admissibility:
 - does not guarantee that every node n that is expanded (i.e. for which we generate the successors s) will have been found via the lowest cost *The first time we expand it*
- Monotonicity
 - guarantees that!
 - Stronger property than admissibility
- If a heuristic is monotonic
 - We are guaranteed that once a node is popped from the OPEN list, we have found the lowest cost path to it *at root* *on min*
 - i.e. we always find the lowest cost path to each node, the 1st time it is popped from OPEN!
 - So once a node is placed in the CLOSED list, if we encounter it again, we do not need to check that the 2nd encounter has a lower cost. We can just ignore it. (more efficient!)

Monotonicity vs Admissibility

- h is monotonic if for every node n and every successor s of n :
 - $h(n) \leq c(n, s) + h(s) \quad \forall n, s$
 - $h(n) - h(s) \leq c(n, s)$ Estimate of cost from n to s
 - $h(n) - h(s) \leq g(s) - g(n)$ Actual cost from n to s

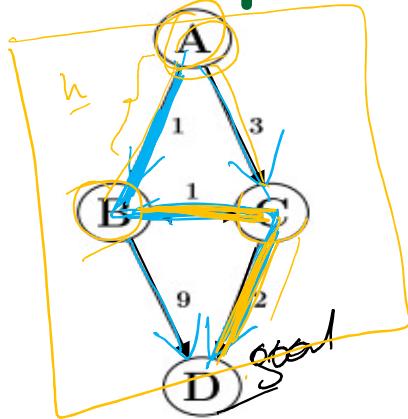
↑ bspw
- \rightarrow monotonic = $h(n)$ is optimistic for all transitions $n \rightarrow s$
- $f(n)$ is non-decreasing along any path
- admissibility = $h(n)$ only needs to be optimistic for $n \rightarrow \text{goal}$
 - $h(n) \leq h^*(n) \quad \forall n$
 - $h^*(n) = g(\text{goal}) - g(n) \rightarrow h(n) \leq g(\text{goal}) - g(n)$
 - $h(\text{goal}) = 0 \rightarrow h(n) - h(\text{goal}) \leq g(\text{goal}) - g(n)$
- Every monotonic $h(n)$ is admissible (but not vice-versa)

Monotonicity and A* Search

■ Monotonicity

- Guarantees to find the lowest cost solution path
- Guarantees to find the lowest cost path to every node, the first time we expand it.
- --> no need to check the CLOSED list again!
- So when we generate a successors s:
 1. ~~IF s is already in CLOSED
IF s in CLOSED has a higher f value
THEN place s and its new lower f value in OPEN!
// we found a lower cost path to s, but we had already expanded s...
// to guarantee the lowest cost solution path, we need to put s back in OPEN and re-visit it again
ELSE ignore s~~
 2. ELSE IF s is already in OPEN
IF s in OPEN has a higher f-value
THEN replace the old s in OPEN with the new lower f-value s
// we found a lower cost path to s, and we had not expanded s yet
// to guarantee the lowest cost solution path, we need to replace the old s in OPEN with the new lower-cost s
ELSE ignore s
 3. ELSE insert s in OPEN
// as usual

Example



ideal $h(n)$
 $h_1(n) = h^*(n) \forall n$

node	h_1	h_2	h^*	Solution paths
A	4	4	4	1. A B D → cost of 10
B	3	3	3	2. A C D → cost of 5
C	2	0	2	3. A B C D → cost of 4
D	0	0	0	4. A C B D → cost of 13

- Admissibility -- $h^*(A)=4$ $h^*(B)=3$ $h^*(C)=2$ $h^*(D)=0$
 - is h_1 admissible? Yes ✓
 - is h_2 admissible? Yes ✓

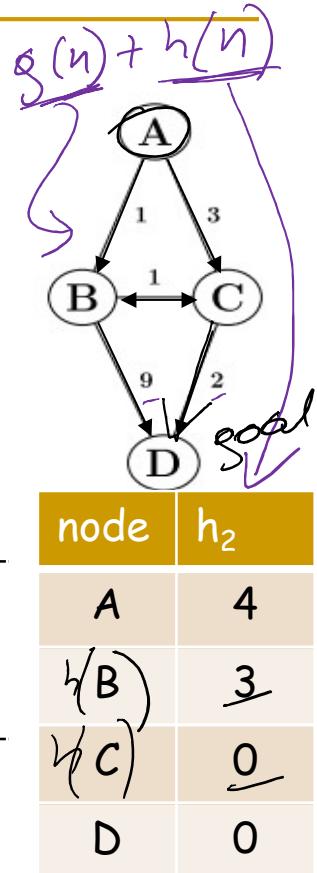
- Monotonic

- is h_1 monotonic? Yes
 - $h_1(A) - h_1(B) \leq g(B) - g(A)$ $4 - 3 \leq 1 - 0$ $1 \leq 1$
 - $h_1(A) - h_1(C) \leq g(C) - g(A)$ $4 - 2 \leq 2 - 0$ $2 \leq 2$
 - $h_1(A) - h_1(D) \leq g(D) - g(A)$ $4 - 0 \leq 4 - 0$ $4 \leq 4$
 - ...

- is h_2 monotonic? No
 - $h_2(A) - h_2(C) \not\leq g(C) - g(A)$ $4 - 0 \not\leq 2 - 0$ $3 \not\leq 2$ $\forall n \text{ s}$
 - $h_2(B) - h_2(C) \not\leq g(C) - g(B)$ $3 - 0 \not\leq 2 - 1$ $3 \not\leq 1$ ✗

Example - h_2

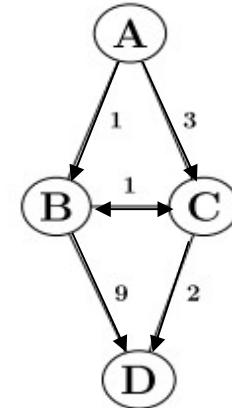
OPEN (unsorted... work in progress)			OPEN	CLOSED
1. $A_{\text{null}} 4_{g=0+h=4}$			1. $A_{\text{null}} 4_{g=0+h=4}$	
2. $B_A 4_{g=1+h=3}$ $C_A 3_{g=3+h=0}$	SORT		1. $C 3_{g=2+h=0} // \text{we will explore } C \text{ directly from } A, \text{ but there is a lower cost path to } C (ABC). h_1 \text{ found it because it is monotonic, but } h_2 \text{ is not monotonic, so it could not guarantee that when we expand a node, we have found the lowest cost path to it...}$ 2. $B_A 4_{g=1+h=3}$	$A_{\text{null}} 4_{g=0+h=4}$
3. $B_C 7_{g=3+1+h=3} // B \text{ is already in OPEN (see below) but with a lower cost path. We do not replace the old, and ignore this version}$ $D_C 5_{g=3+2+h=0}$ $B_A 4_{g=1+h=3}$		1. $B_A 4_{g=1+h=3}$ 2. $D_C 5_{g=3+2+h=0}$		$A_{\text{null}} 4_{g=0+h=4}$ $C_A 3_{g=3+h=0}$
4. $C_B 2_{g=1+1+h=0} // C \text{ was already in CLOSED but with a higher } f\text{-value, we just found a lower cost path to } C... \text{ we need to put this version back in OPEN :-)}$ $D_B 10_{g=1+9+h=0} // D \text{ is already in OPEN (see below) but with a lower cost path. We do not replace the old, and ignore this version}$ $D_C 5_{g=2+2+h=0}$		1. $C_B 2_{g=1+1+h=0}$ 2. $D_C 5_{g=3+2+h=0}$		$A_{\text{null}} 4_{g=0+h=4}$ $C_A 3_{g=3+h=0}$ $B_A 4_{g=1+h=3}$
5. $B_C 6_{g=1+1+1+h=3}$ $D_C 4_{g=1+1+2+h=0} // D \text{ is already in OPEN (see below) but with a higher cost path. We replace the old version with version}$ $D_C 5_{g=1+9+h=0}$		1. $D_C 4_{g=1+1+2+h=0}$ 2. $B_C 6_{g=1+1+1+h=3}$		$A_{\text{null}} 4_{g=0+h=4}$ $B_A 4_{g=1+h=3}$ $C_B 2_{g=1+1+h=0}$
		goal($D_C 4$) = true! Solution path = $D_C C_B B_A A_{\text{null}}$ cost = $2+1+1 = 4$ // lowest cost path found in 5 steps		



lowest cost path found in 5 steps

Example - h_1 admissible + monotonic

	OPEN (unsorted... work in progress)	OPEN	CLOSED
1	• $A_{\text{null}} 4_{g=0+h=4}$	1. $A_{\text{null}} 4_{g=0+h=4}$	
2	• $B_A 4_{g=1+h=3}$ • $C_A 5_{g=3+h=2}$	1. $B_A 4_{g=1+h=3}$ 2. $C_A 5_{g=3+h=2}$	$A_{\text{null}} 4_{g=0+h=4}$
3	• $C_B 4_{g=1+1+h=2}$ // C already in OPEN with a higher f-value, replace old version with this one • $D_B 10_{g=1+9+h=0}$ • $C_A 5_{g=3+h=2}$	1. $C_B 4_{g=1+1+h=2}$ 2. $D_B 10_{g=1+9+h=0}$	$A_{\text{null}} 4_{g=0+h=4}$ $B_A 4_{g=1+h=3}$
4	• $D_C 4_{g=2+2+h=0}$ // D already in OPEN with a higher f-value, replace old version with this one • $B_C 6_{g=2+1+h=3}$ // B already in CLOSED but since h_1 is monotonic, we do not need to check the f-value of the version in CLOSED because we know that the version in CLOSED will have a lower f-value, so can ignore this version • $D_B 10_{g=1+9+h=0}$	1. $D_C 4_{g=2+2+h=0}$ 2. $B_C 6_{g=2+1+h=3}$	$A_{\text{null}} 4_{g=0+h=4}$ $B_A 4_{g=1+h=3}$ $C_B 4_{g=1+h=2}$
		goal($D_C 4$) = true! Solution path = $D_C \rightarrow C_B \rightarrow B_A \rightarrow A_{\text{null}}$ cost = $2+1+1 = 4$ // lowest cost path found in 4 steps	



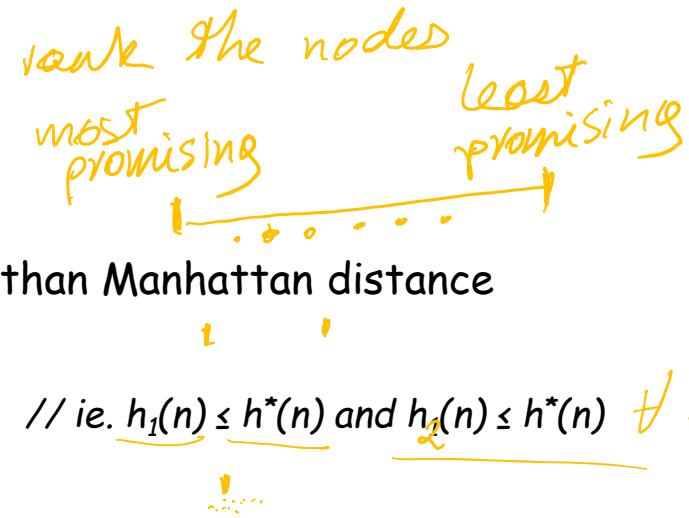
node	h_1
A	4
B	3
C	2
D	0

Admissible +
Monotonic

lowest cost path
found in 4 steps

Informedness

- Intuition:
 - $h(n) = 0$ for all nodes is less informed
 - number of misplaced tiles is less informed than Manhattan distance
- Formally:
 - given 2 admissible heuristics h_1 and h_2 // ie. $h_1(n) \leq h^*(n)$ and $h_2(n) \leq h^*(n)$
 - if $h_1(n) \leq h_2(n)$, for all states n
 - then h_2 is more informed than h_1
 - aka h_2 dominates h_1
- So?
 - a more informed heuristic expands fewer nodes
 - aka the search path is shorter
 - { however, you need to consider the computational cost of evaluating the heuristic... $h(n)$
 - { the time spent computing heuristics must be recovered by a better search



Today

1. State Space Representation
2. State Space Search
 - a) Overview
 - b) Uninformed search
 1. Breadth-first and Depth-first
 2. Depth-limited Search
 3. Iterative Deepening
 4. Uniform Cost
 - c) Informed search
 1. Intro to Heuristics
 2. Hill climbing
 3. Greedy Best-First Search
 4. Algorithms A & A*
 5. More on Heuristics
 - d) Summary

YOU ARE HERE!



Summary

Search	Uses <u>$h(n)$</u> ?	Uses <u>$g(n)$</u> ?	OPEN list
Breadth-first	No ✓	No ✓	Priority queue sorted by level
Depth-first	No ✓	No ✓	Stack
Depth-limited	No ✓	No ✓	Stack
Iterative Deepening	No ✓	No ✓	Stack
Uniform Cost - guarantees to find the <u>lowest cost solution path</u>	No ✓	Yes ✓	Priority queue sorted by <u>$g(n)$</u> When generating successors: - If successor s already in OPEN with higher $g(n)$, replace old version with new s - If successor s already in CLOSED, ignore s ✗
Hill Climbing	Yes ✓	No ✓	N/A
Greedy Best-First - no constraints on $h(n)$ - no guarantee to find lowest cost solution path	Yes ✓	No ✗	Priority queue sorted by <u>$h(n)$</u>
Algorithm A - no constraints on $h(n)$ - no guarantee to find lowest cost solution path	Yes ✓	Yes ✓	Priority queue sorted by <u>$f(n)$</u>
Algorithm A* - $h(n)$ must be admissible - guarantees to find the <u>lowest cost solution path</u>	Yes ✓	Yes ✓	Priority queue sorted by <u>$f(n)$</u> If $h(n)$ is NOT monotonic When generating successors: - If successor s already in OPEN with higher $f(n)$, replace old version with new s - If successor s already in CLOSED with higher $f(n)$, replace old version with new s If $h(n)$ IS monotonic When generating successors: - If successor s already in OPEN with higher $f(n)$, replace old version with new s - If successor s already in CLOSED, ignore it. ✗

Today

1. State Space Representation ✓
2. State Space Search ✓
 - a) Overview ✓
 - b) Uninformed search ✓
 1. Breadth-first and Depth-first ✓
 2. Depth-limited Search ✓
 3. Iterative Deepening ✓
 4. Uniform Cost ✓
 - c) Informed search ✓
 1. Intro to Heuristics ✓
 2. Hill climbing ✓
 3. Greedy Best-First Search ✓
 4. Algorithms A & A* ✓
 5. More on Heuristics ✓
 - d) Summary ✓

Up Next

1. Part 4: Adversarial Search

Artificial Intelligence: Adversarial Search *part 4* Minimax

- Russell & Norvig: ~~Chapter 5~~

Today

■ Adversarial Search

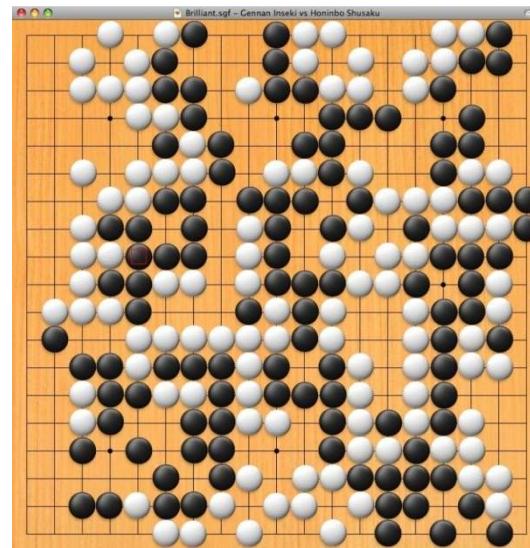
YOU ARE HERE!

Chap 5

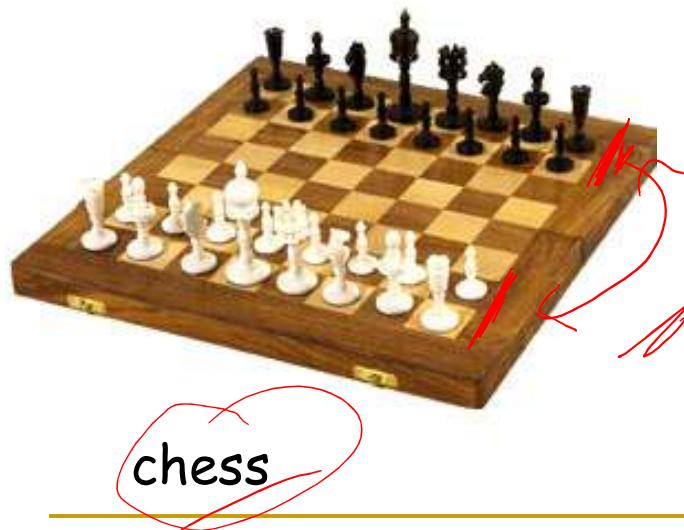
1. Minimax) 4.1
2. Alpha-beta pruning) 4.2
3. Other Adversarial Search
 1. Multiplayer Games
 2. Stochastic Games
 3. Monte Carlo Tree Search

4.3

Motivation



GO



Games

1. Assumptions

1. Zero-sum game:

- no player has a game advantage. If the total gains of one player are added up, and the total losses are subtracted, they will sum to 0.

2. Players play rationally (i.e. to win)

2. Characteristics of Games

Number of players? 1, 2, 3+

Deterministic:

- the outcome of the game is only dependent on the moves of the players

Stochastic:

- chance involved

Perfect information:

- all players know the state of the game and all possible moves a.k.a **fully observable**

Imperfect information:

- eg. a player is hiding their game aka. **partially observable**

	Deterministic	Stochastic
Perfect information	chess, checkers, go	backgammon, monopoly
Imperfect information	battleship	scrabble, poker, bridge

Techniques

Nb of players	Deterministic / Stochastic ?	Perfect / Imperfect information?	Example	Technique
1 -	deterministic	perfect info	8-puzzle	heuristic search per 3 ✓
2	deterministic	perfect info	chess, checkers, go	minimax & alpha-beta, video 4.1 4.2
3+	deterministic	perfect info	Chinese Checkers	max ⁿ
1	stochastic	yes	2048	expectimax
2	stochastic	yes	backgammon	expectiminimax
3+	stochastic	yes		modified expectimax

Techniques

Nb of players	Deterministic / Stochastic ?	Perfect / Imperfect information?	Example	Technique
1	deterministic	perfect info	8-puzzle	heuristic search
2	deterministic	perfect info	chess, checkers, go	minimax & alpha-beta
3+	deterministic	perfect info	Chinese Checkers	max ⁿ
1	stochastic	yes	2048	expectimax
2	stochastic	yes	backgammon	expectiminimax
3+	stochastic	yes		modified expectimax

Today

■ Adversarial Search

1. Minimax 
2. Alpha-beta pruning
3. Other Adversarial Search
 1. Multiplayer Games
 2. Stochastic Games
 3. Monte Carlo Tree Search

Minimax Search

- Game between two opponents, MIN and MAX
 - MAX tries to win, and
 - MIN tries to minimize MAX's score
- Existing heuristic search methods do not work
 - would require a helpful opponent
 - need to incorporate "hostile" moves into search strategy

- 2 flavors:

1. exhaustive Minimax 
2. n-ply Minimax with Heuristic 

Exhaustive Minimax Search

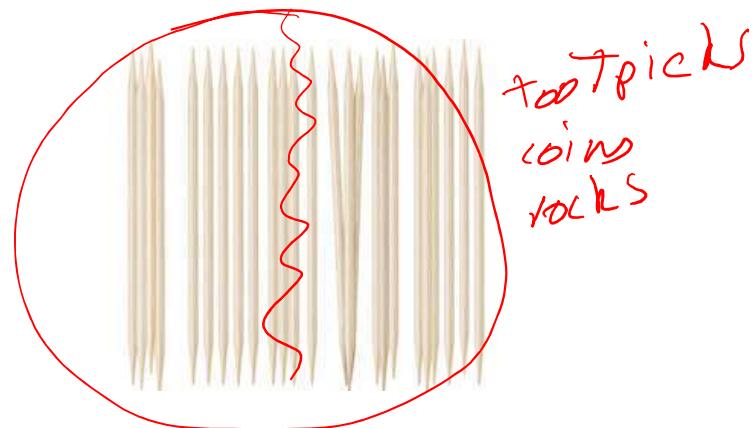


- For small games where exhaustive search is feasible
- Procedure:
 1. build complete game tree
 2. label each level according to player's turn (MAX or MIN)
 3. label leaves with a utility function to determine the outcome of the game
 - e.g., $(0, 1)$ or $(-1, 0, 1)$
 - MAX
turn
 - MAX wins
 - min turn
 - Max wins
 4. propagate this value up:
 - if parent=MAX, give it max value of children
 - if parent=MIN, give it min value of children
 5. Select best next move for player at the root as the move leading to the child with the highest value (for MAX) or lowest values (for MIN)

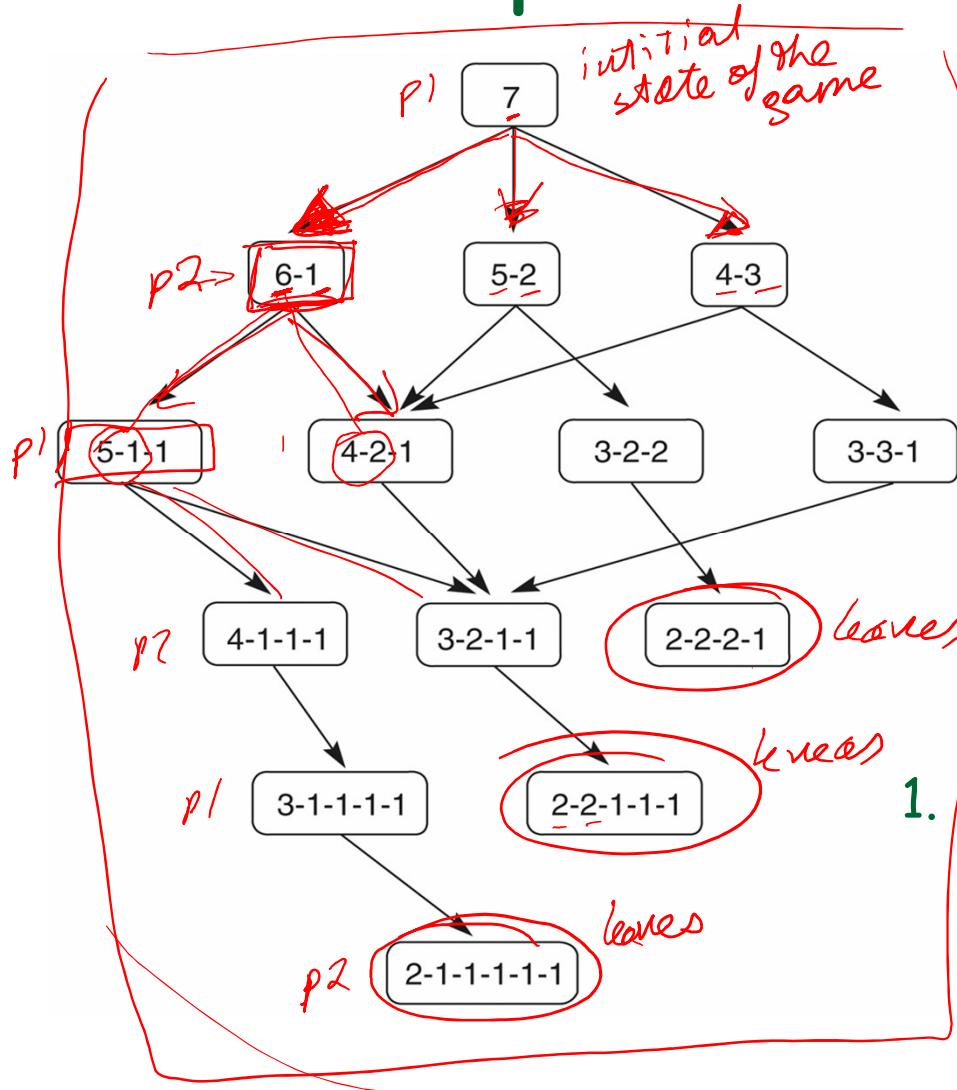
Example: Game of Nim

■ Rules

- 2 players start with a pile of tokens
- move: split (any) existing pile into two non-empty differently-sized piles
- game ends when no pile can be unevenly split
- player who cannot make their move loses



State Space of Game Nim



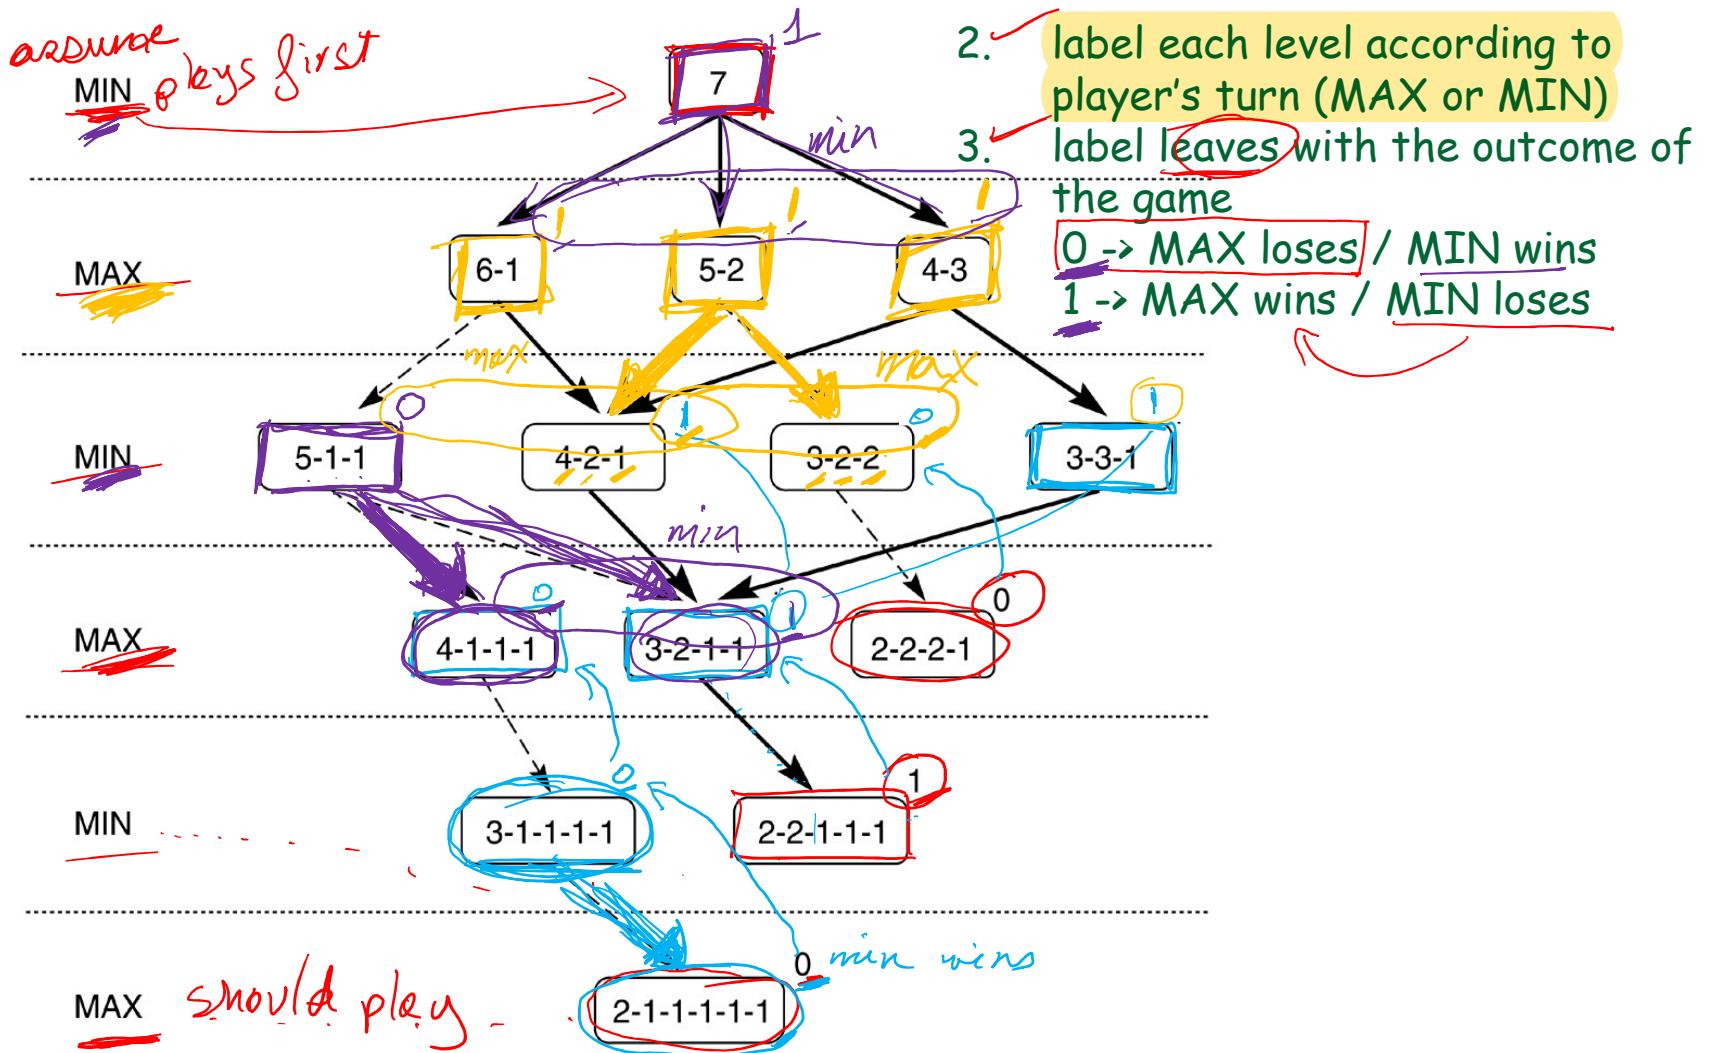
- eg. start with one pile of 7 tokens

- each step has to divide one pile into 2 non-empty piles of different sizes

- player without a move left loses game

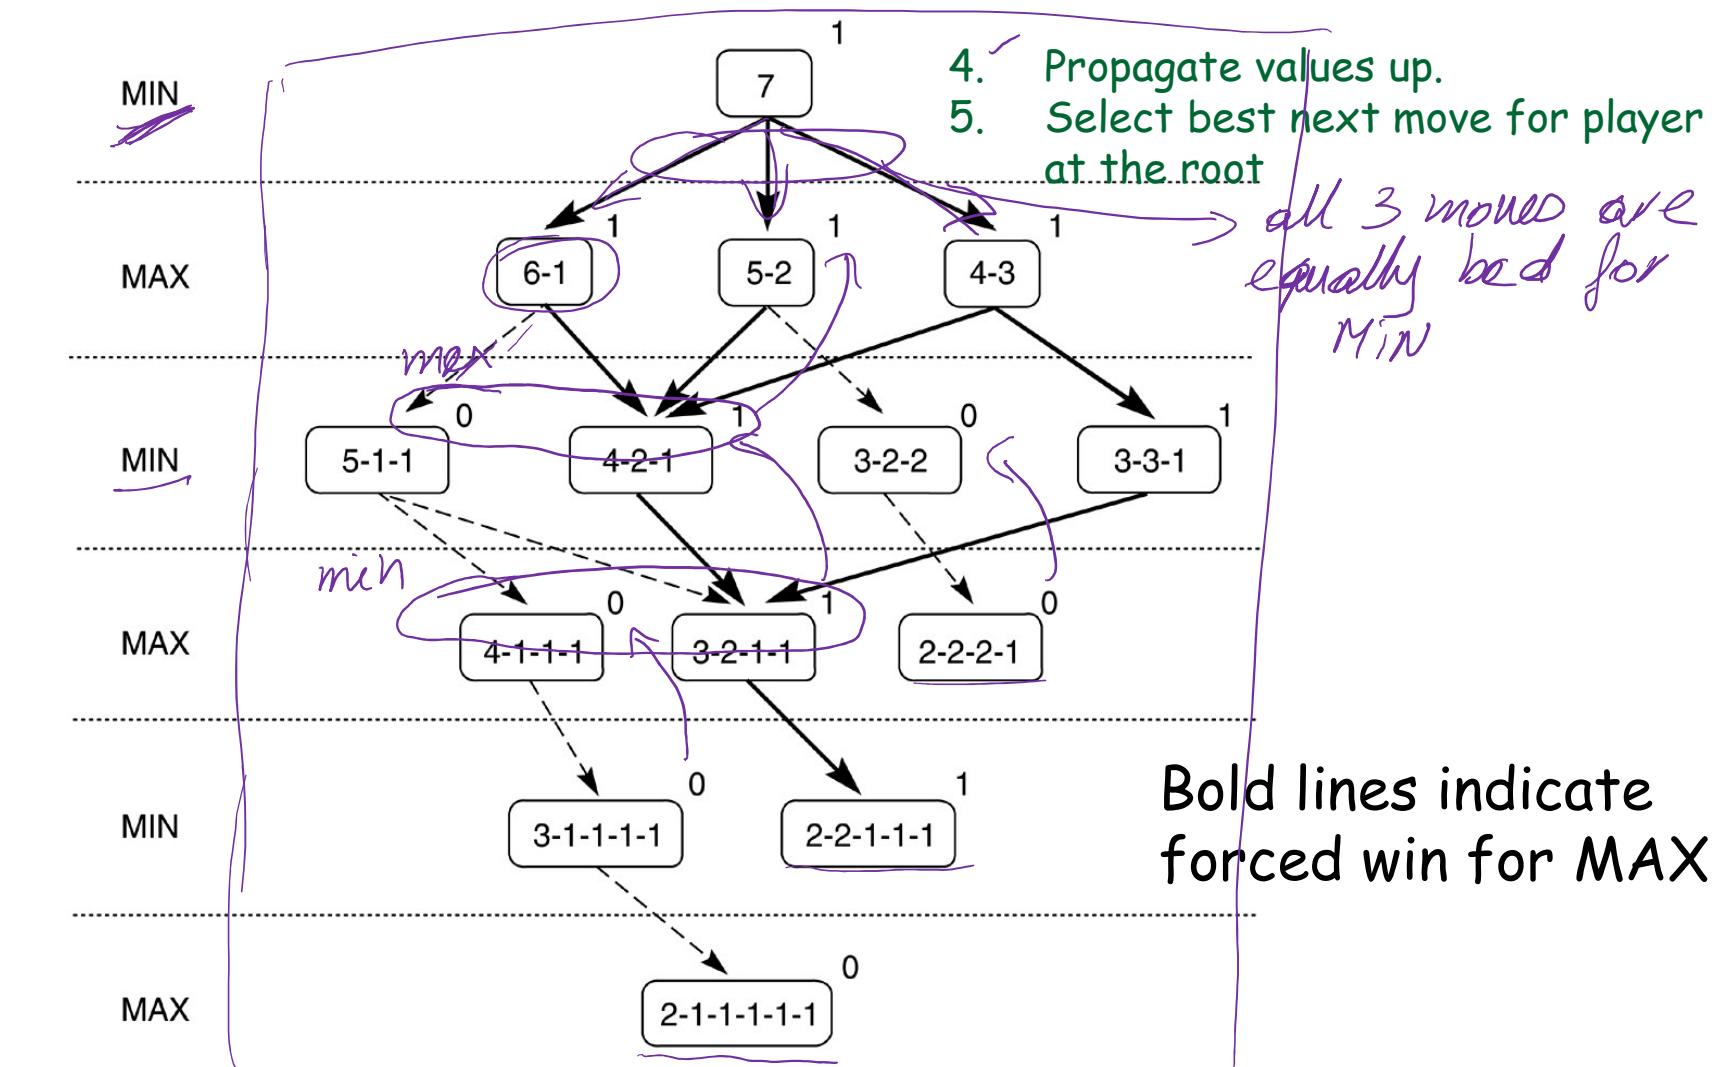
1. build complete game tree ✓

Exhaustive Minimax for Nim



source: G. Luger (2005)

Exhaustive Minimax for Nim



source: G. Luger (2005)

n-ply Minimax with Heuristic

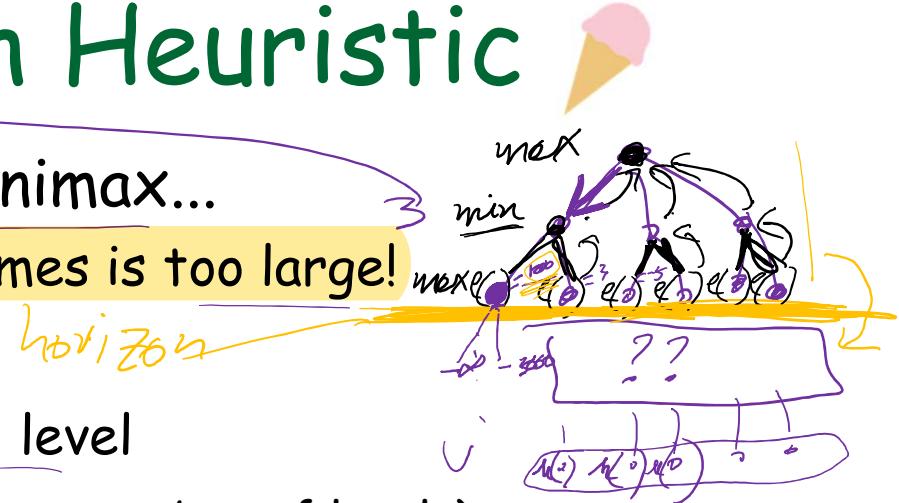
problem with exhaustive Minimax...

- state space for interesting games is too large!

solution:

- search only up to a predefined level
- called n-ply look-ahead ($n = \max$ number of levels)
- not an exhaustive search
- nodes cannot be evaluated with win/loss/tie
- nodes are evaluated with heuristics function $e(n)$
- $e(n)$ indicates how good a state seems to be for MAX compared to MIN
- suffers from the horizon effect

$$e\left(\begin{array}{|c|c|c|} \hline & & \\ \hline & & \\ \hline & & \\ \hline \end{array}\right) = 30$$



<https://www.wallpaperflare.com/beach-shore-sun-sunrise-sea-horizon-ship-water-sky-sunset-wallpaper-sbhsg>

Heuristic Function for 2-player games

- simple strategy:
 - try to maximize difference between MAX's game and MIN's game
- typically called $e(n)$
- $e(n)$ is a heuristic that estimates how favorable a node n is for MAX with respect to MIN
 - $e(n) > 0$ --> n is favorable to MAX
 - $e(n) < 0$ --> n is favorable to MIN
 - $e(n) = 0$ --> n is neutral



Choosing a Heuristic Function $e(n)$

- Usually $e(n)$ is a weighted sum of various features:

$$e(n) = \sum w_i f_i(n)$$

dependent on
actual game

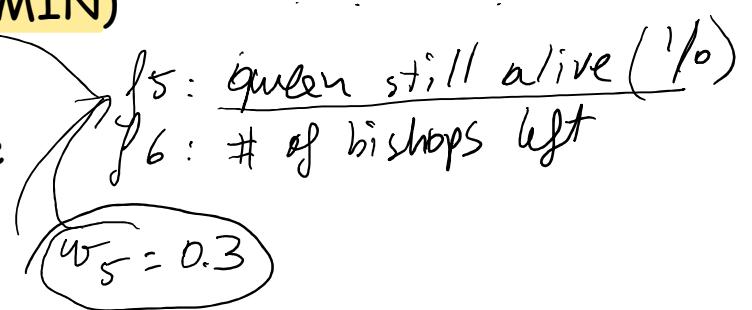
- E.g. of features:

- f_1 = number of pieces left on the game for MAX
- f_2 = number of possible moves left for MAX
- f_3 = -number of pieces left on the game for MIN
- f_4 = -number of possible moves left for MIN

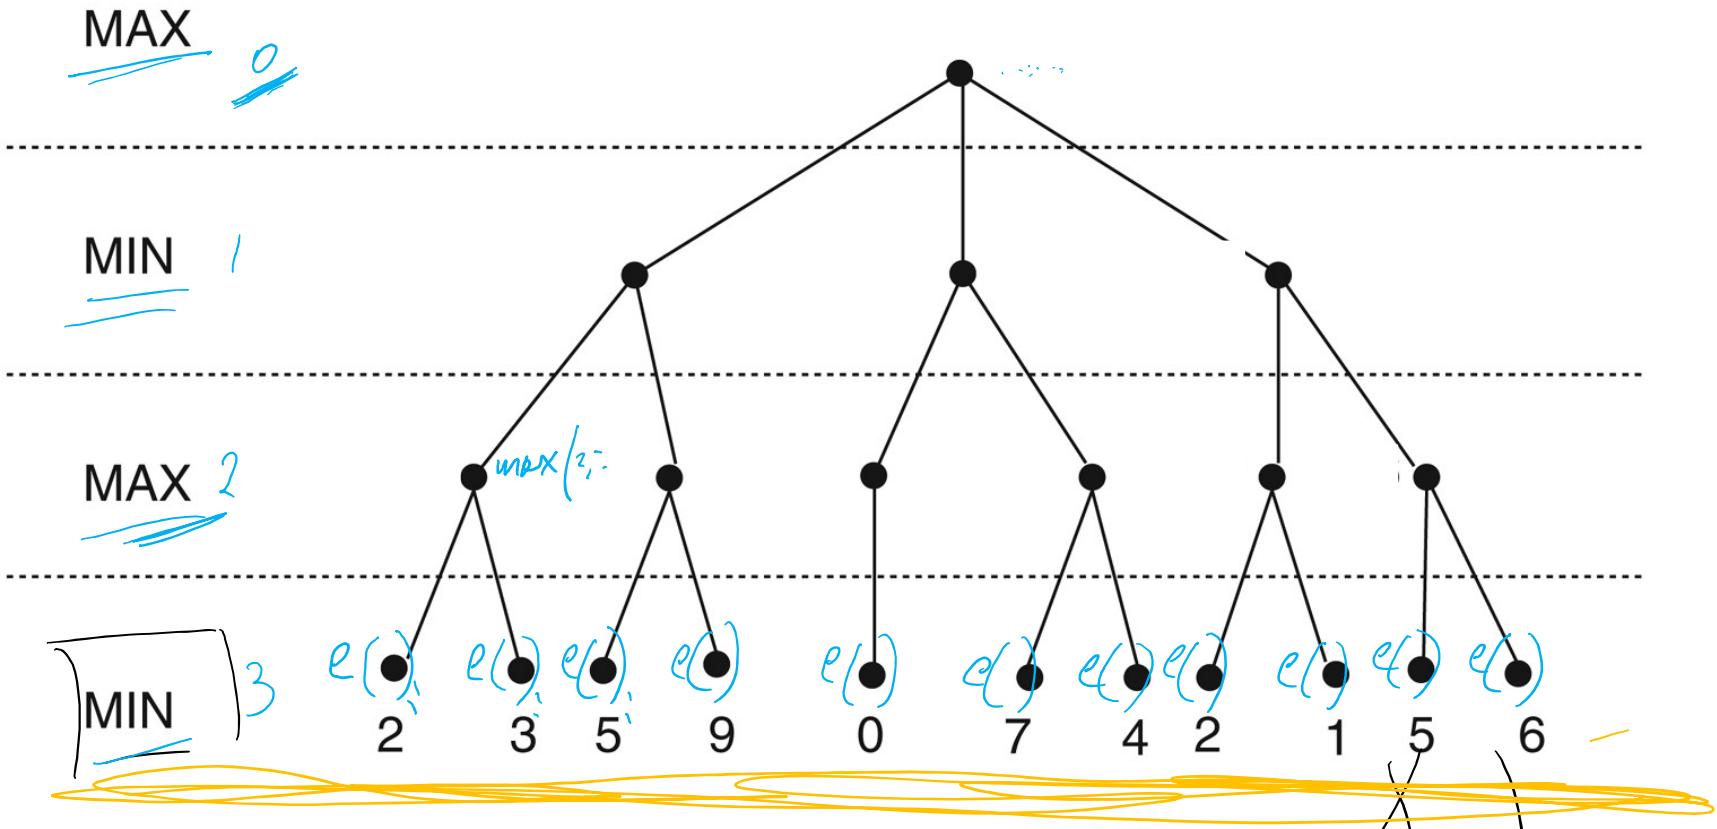


- E.g. of weights:

- $w_1 = 0.5$ // f_1 is a very important feature
- $w_2 = 0.2$ // f_2 is not very important
- $w_3 = 0.2$ // f_3 is not very important
- $w_4 = 0.1$ // f_4 is really not important



Minimax with 3-ply look-ahead



Leaf nodes show the actual heuristic value $e(n)$

source: G. Luger (2005)

end of the game
leaves =

Minimax with 3-ply look-ahead

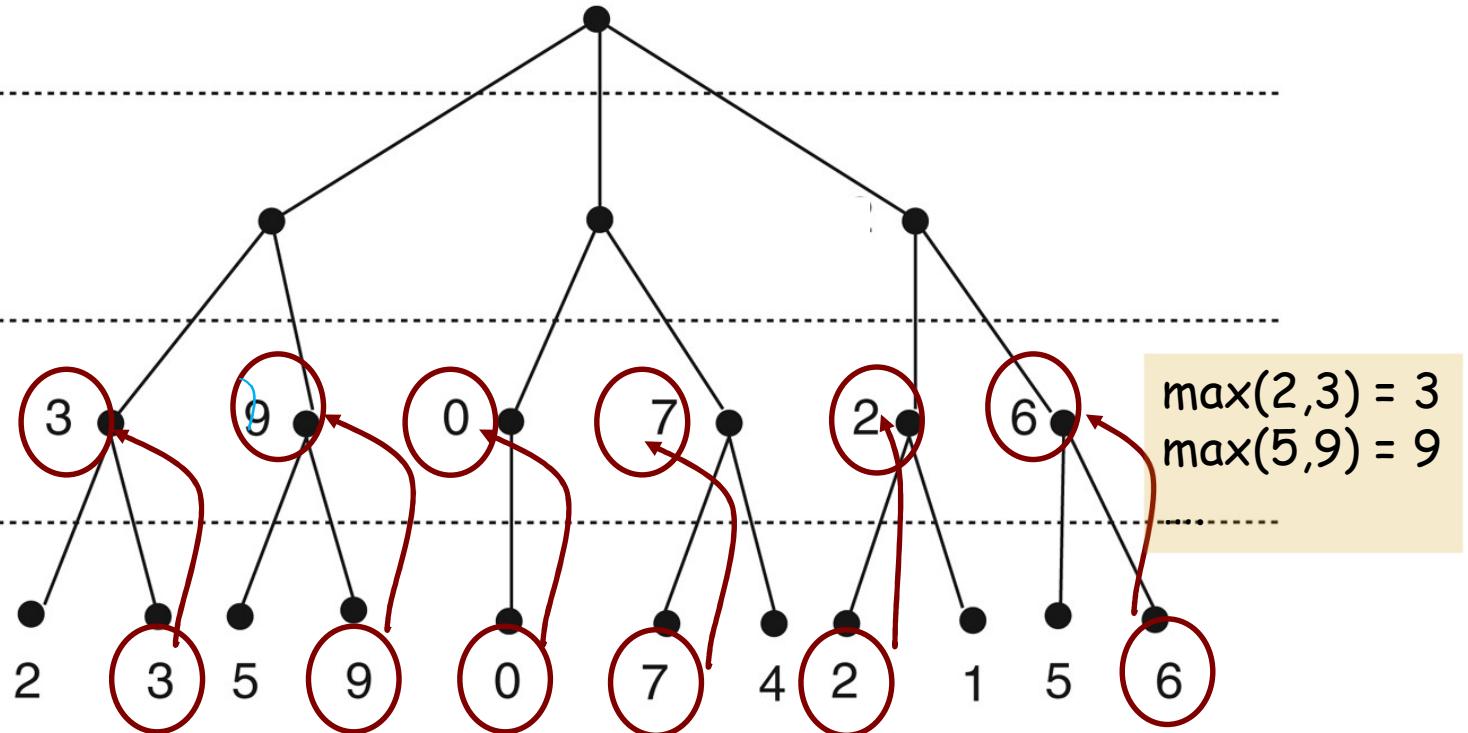


MAX

MIN

MAX

MIN



Leaf nodes show the actual heuristic value $e(n)$
Internal nodes show back-up heuristic value

source: G. Luger (2005)

Minimax with 3-ply look-ahead



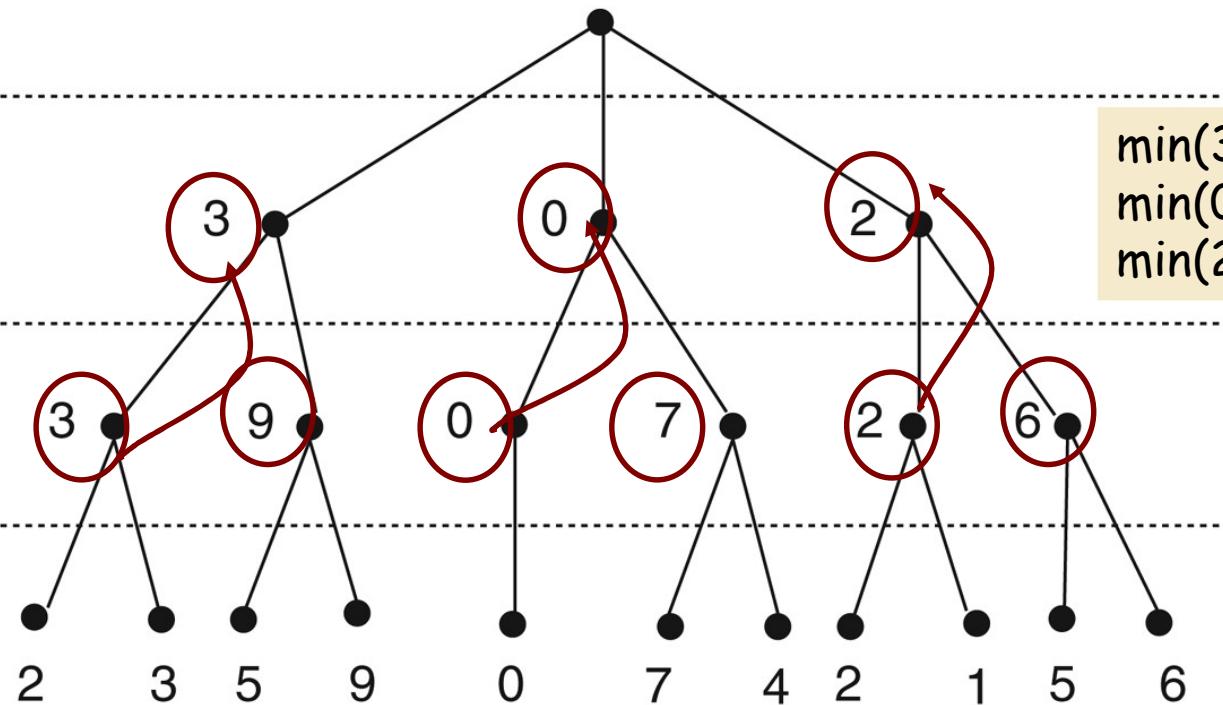
MAX

MIN

MAX

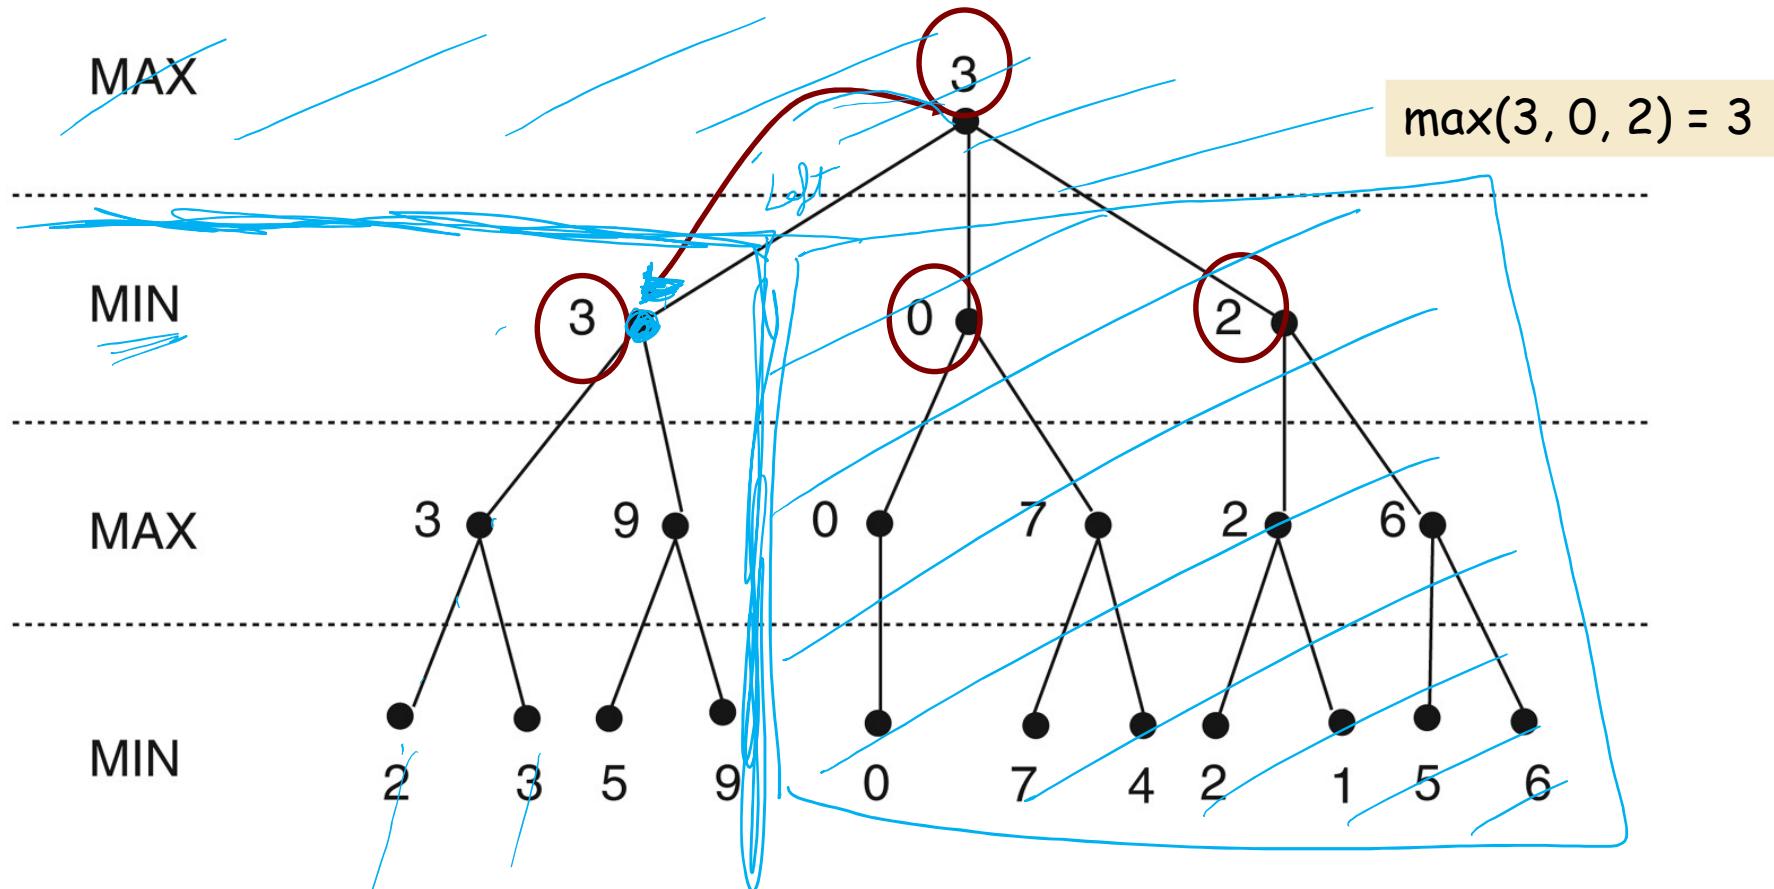
MIN

$$\begin{aligned} \min(3, 9) &= 3 \\ \min(0, 7) &= 0 \\ \min(2, 6) &= 2 \end{aligned}$$



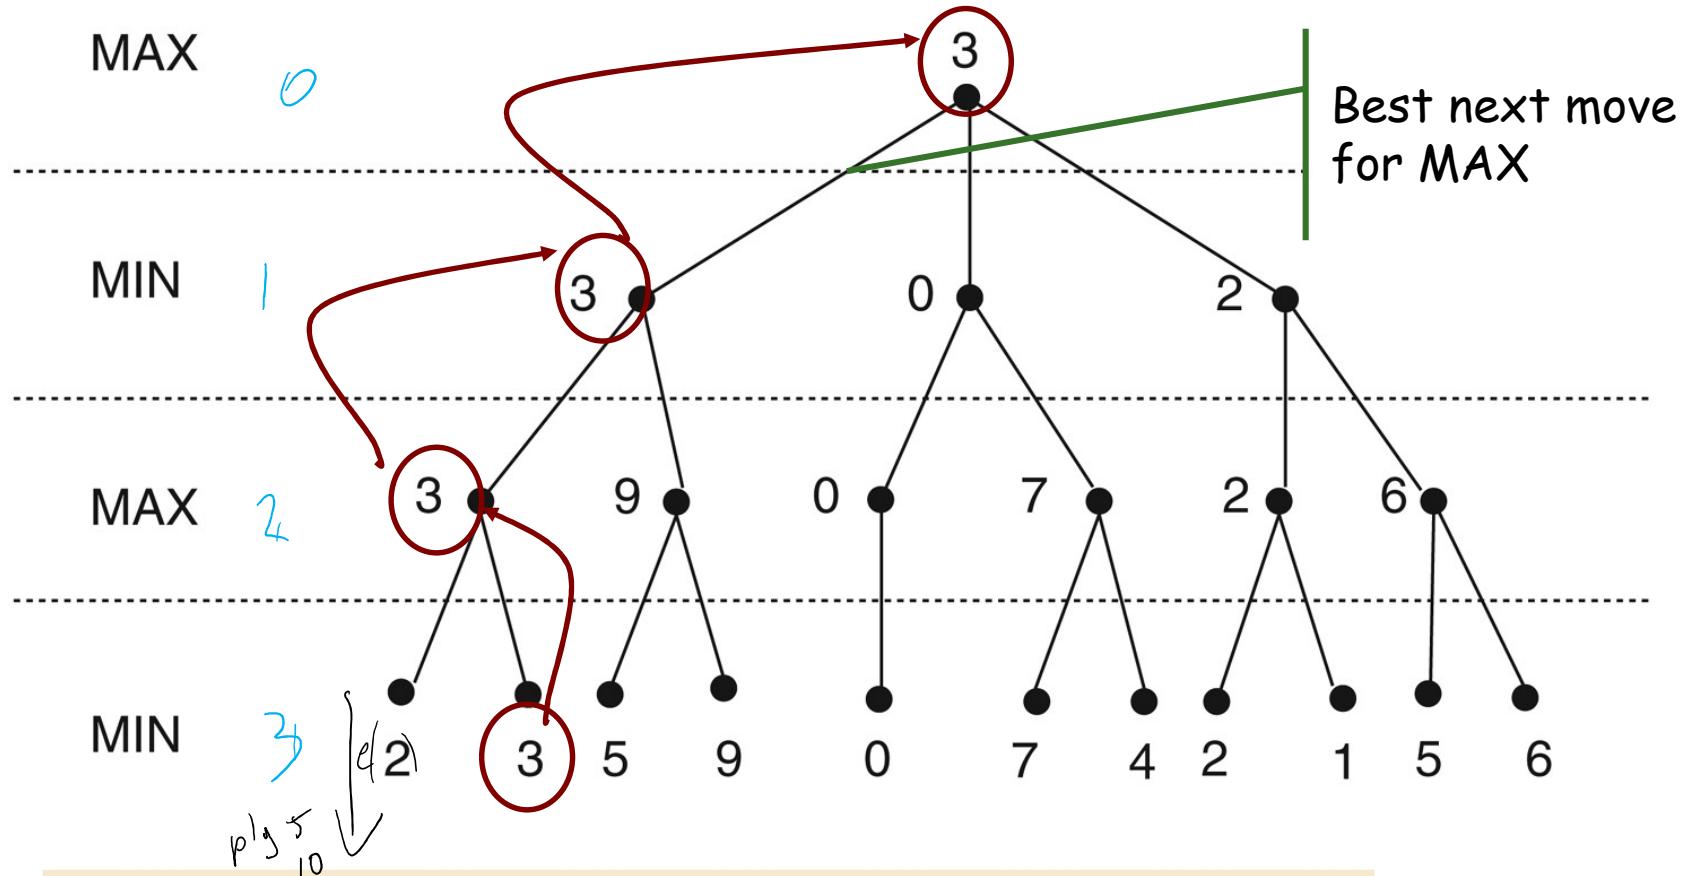
Leaf nodes show the actual heuristic value $e(n)$
Internal nodes show back-up heuristic value

Minimax with 3-ply look-ahead



Leaf nodes show the actual heuristic value $e(n)$
Internal nodes show back-up heuristic value

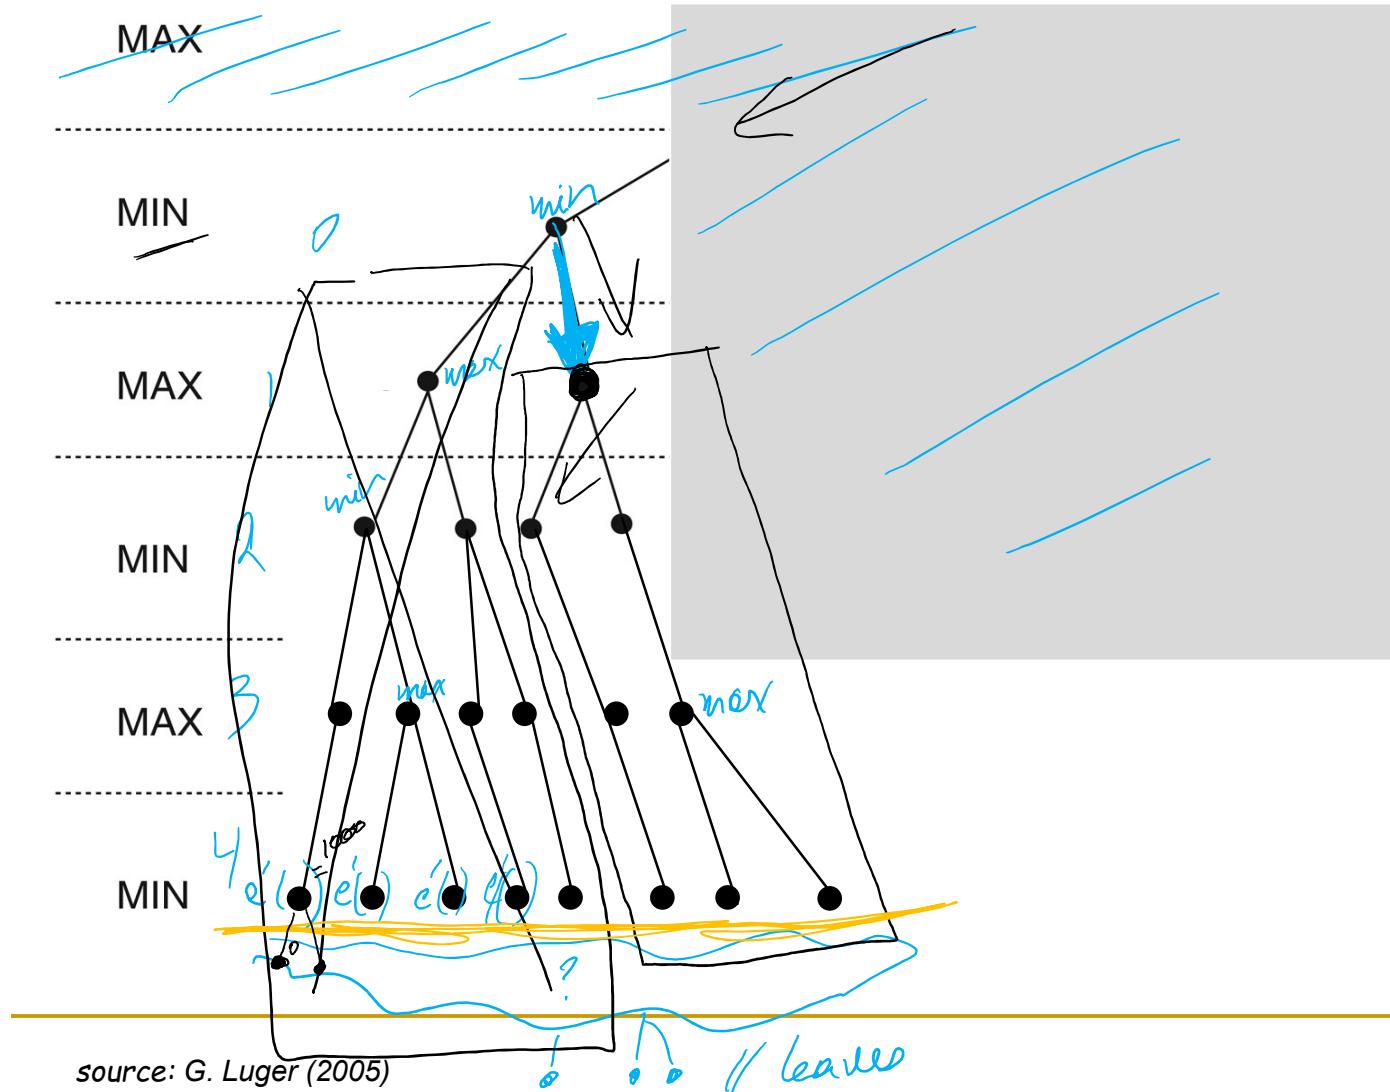
Minimax with 3-ply look-ahead



Leaf nodes show the actual heuristic value $e(n)$
Internal nodes show back-up heuristic value

and then, MIN will play

assume Min can afford to look
4 ply ahead

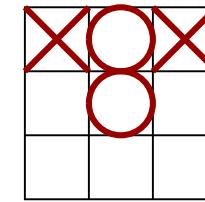
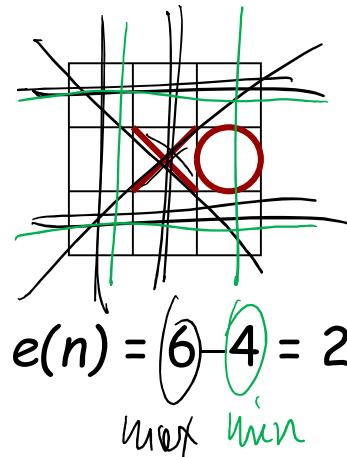
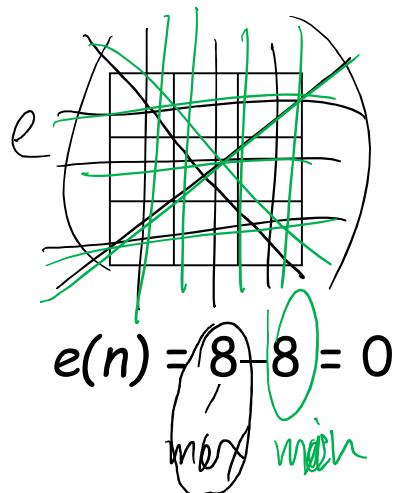


Example: $e(n)$ for Tic-Tac-Toe

- assume MAX plays X

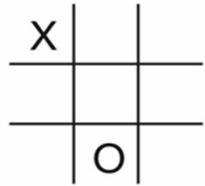
- possible $e(n)$

$$e(n) = \begin{cases} \text{number of rows, columns, and diagonals open for MAX} \\ - \text{number of rows, columns, and diagonals open for MIN} \\ +\infty \text{ if } n \text{ is a forced win for MAX} \\ -\infty \text{ if } n \text{ is a forced win for MIN} \end{cases}$$

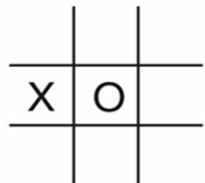
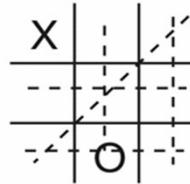
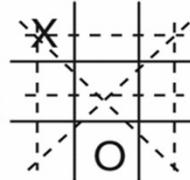


$$\underline{e(n) = 3 - 3 = 0}$$

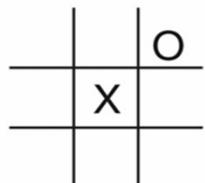
More examples...



X has 6 possible win paths:
O has 5 possible wins:
 $E(n) = 6 - 5 = 1$

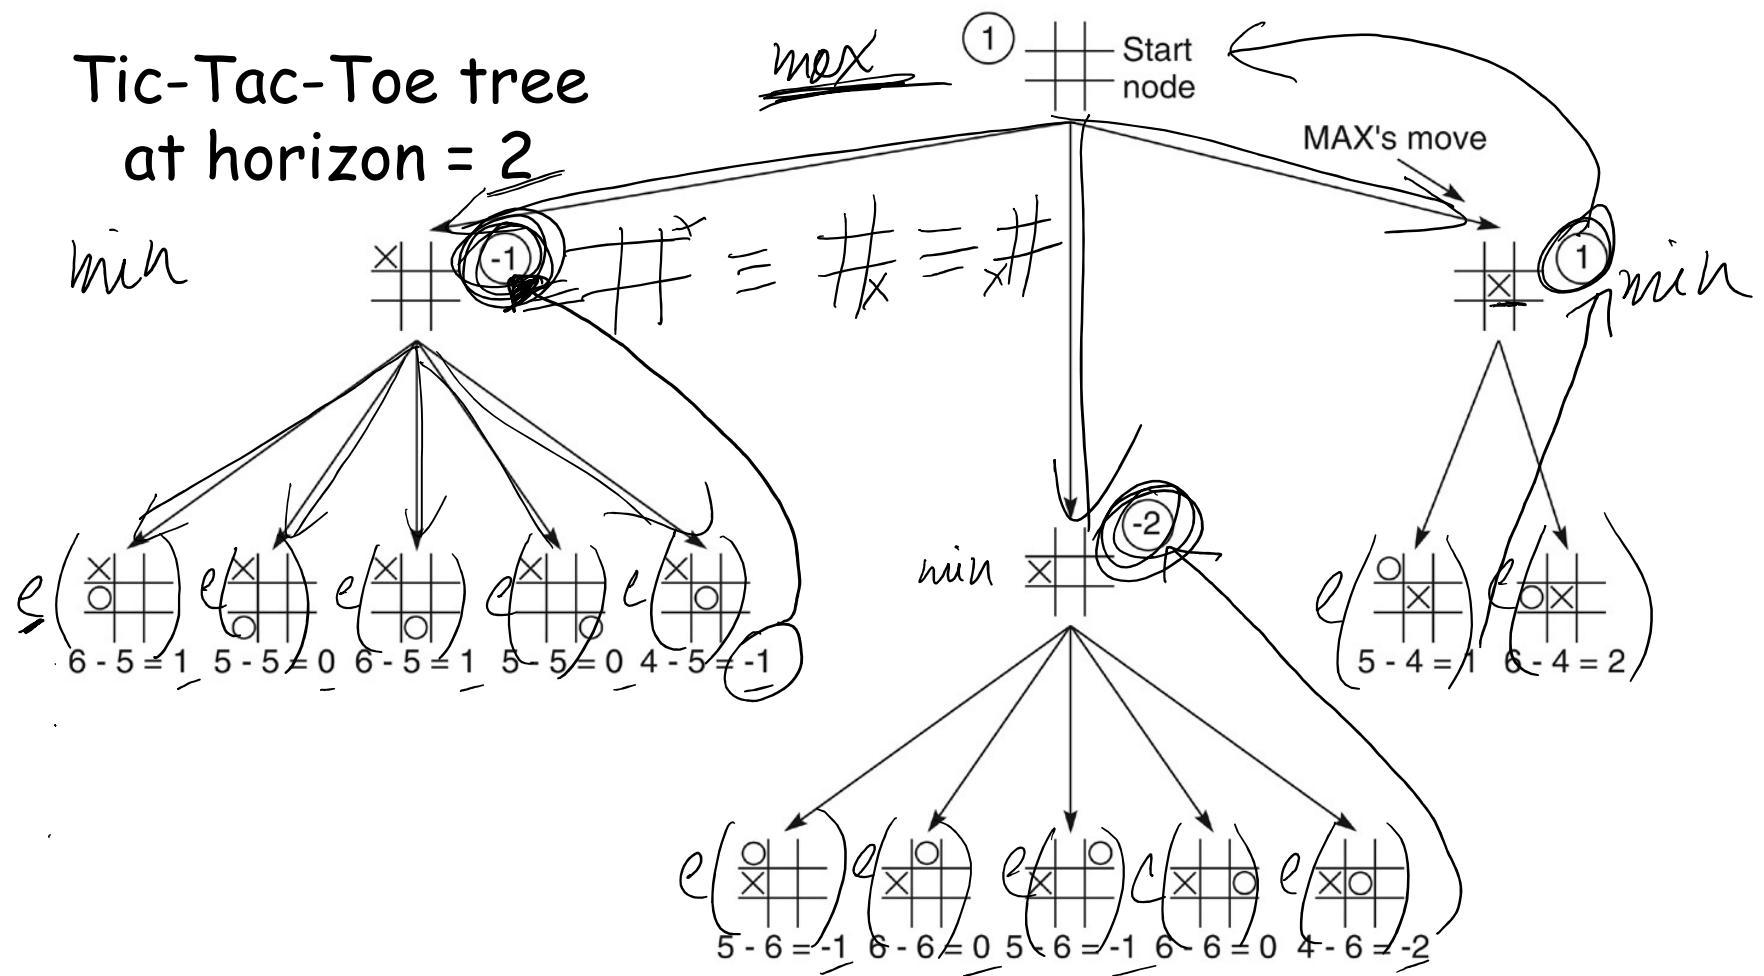


X has 4 possible win paths;
O has 6 possible wins
 $E(n) = 4 - 6 = -2$



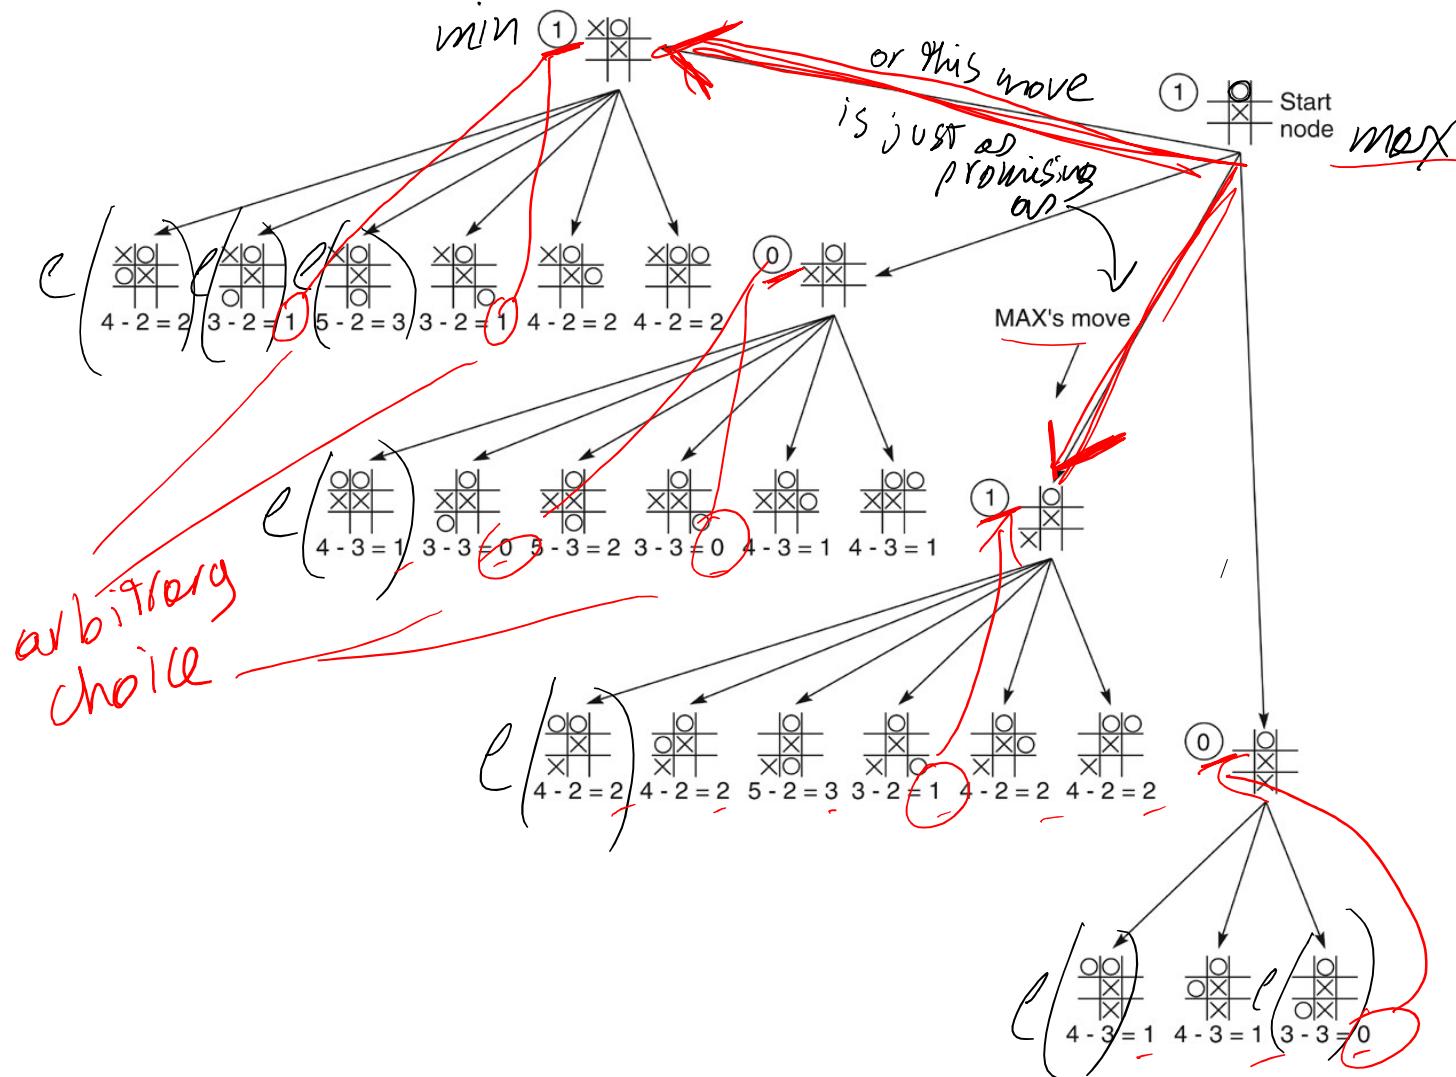
X has 5 possible win paths;
O has 4 possible wins
 $E(n) = 5 - 4 = 1$

2-ply Minimax for Opening Move



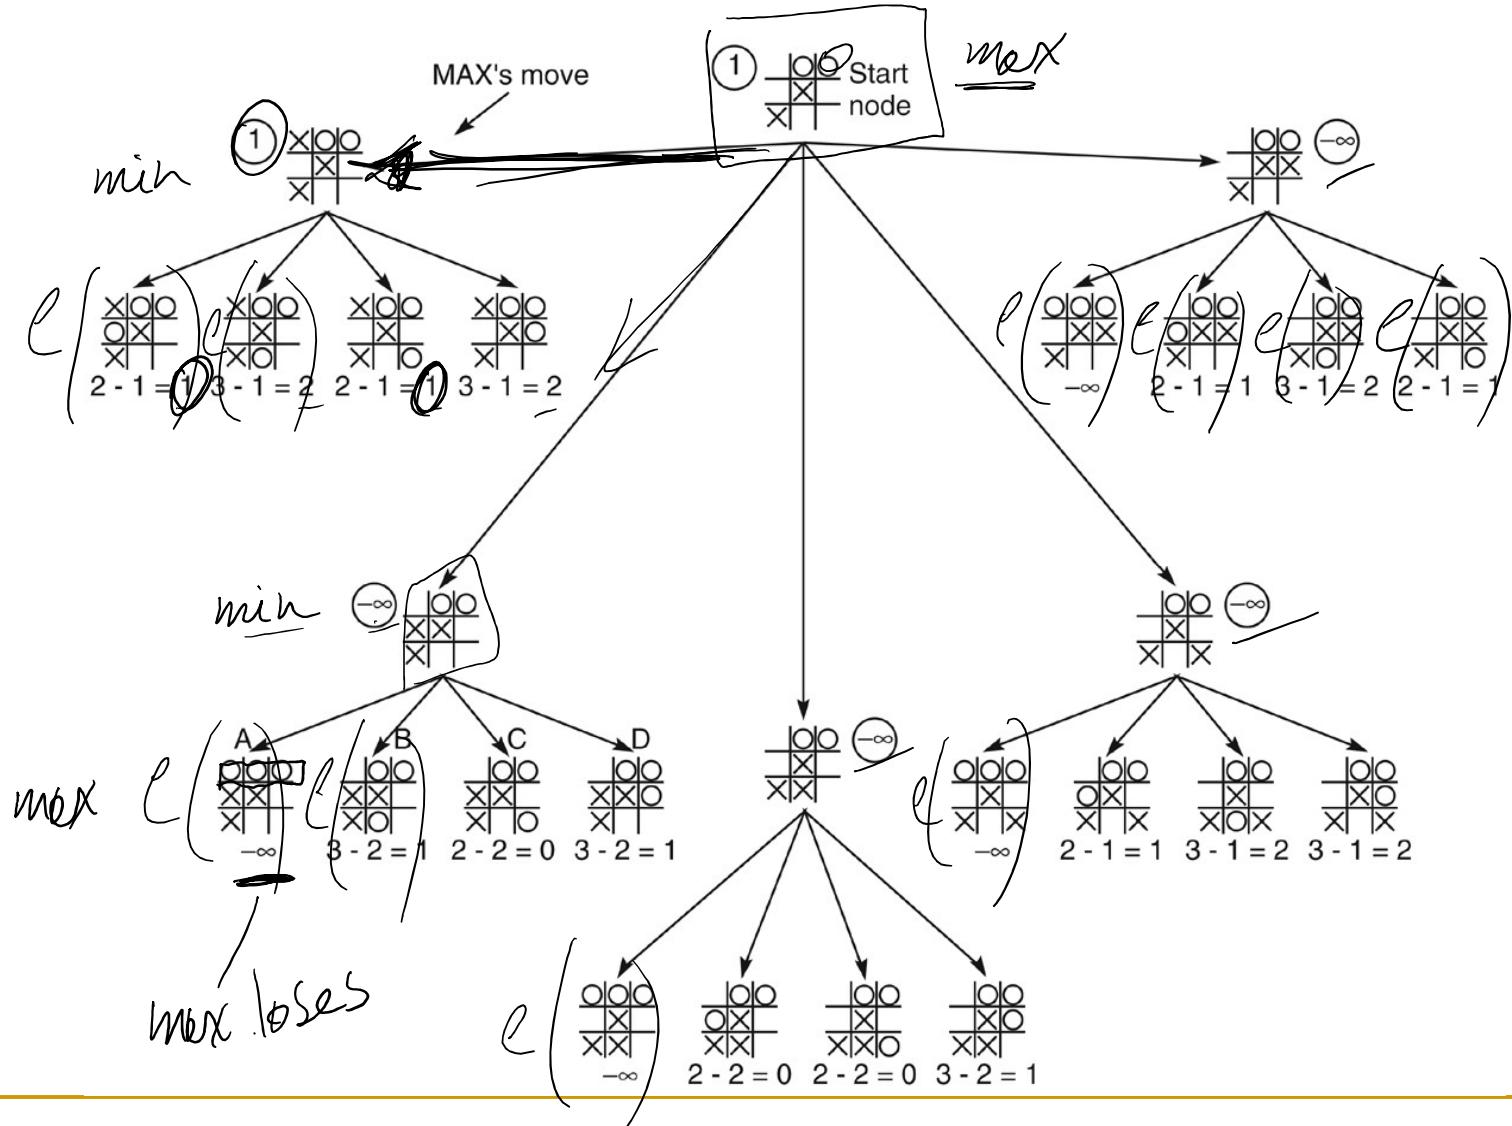
source: G. Luger (2005)

2-ply Minimax: MAX's possible 2nd moves



source: G. Luger (2005)

2-ply Minimax: MAX's move at end



source: G. Luger (2005)

Today

- Adversarial Search ✓
 - 1. Minimax ✓
 - 2. Alpha-Beta Pruning
 - 3. Other Adversarial Search
 - 1. Multiplayer Games
 - 2. Stochastic Games
 - 3. Monte Carlo Tree Search

Up Next

- Adversarial Search
 - 1. Minimax
 - 2. Alpha-Beta Pruning
 - 3. Other Adversarial Search
 - 1. Multiplayer Games
 - 2. Stochastic Games
 - 3. Monte Carlo Tree Search

COMP 472 Artificial Intelligence: Adversarial Search *part 9* Alpha-Beta Pruning *video 2*

- Russell & Norvig: Chapter 5

Today

■ Adversarial Search

1. Minimax
2. Alpha-beta pruning
3. Other Adversarial Searc.
 1. Multiplayer Games
 2. Stochastic Games
 3. Monte Carlo Tree Search



Alpha-Beta Pruning

- Optimization over Minimax, that:

- ignores (cuts off/prunes) branches of the tree that cannot contribute to the solution
 - reduces branching factor ~~effort~~
 - allows deeper search with same effort

↓
↓
↓
↓
↓

Alpha-Beta Pruning: Example 1

yun e(n)

- With Minimax, we look at all nodes at depth n
- With $\alpha\beta$ pruning, we ignore branches that could not possibly contribute to the final decision

1. run $e(n)$ on leaf nodes in a depth first traversal manner on a necessity basis (not on all nodes)

MAX

2. propagate constraints up the tree

MIN

3. to identify leaf nodes (or entire branches) on which running $e(n)$ is pointless

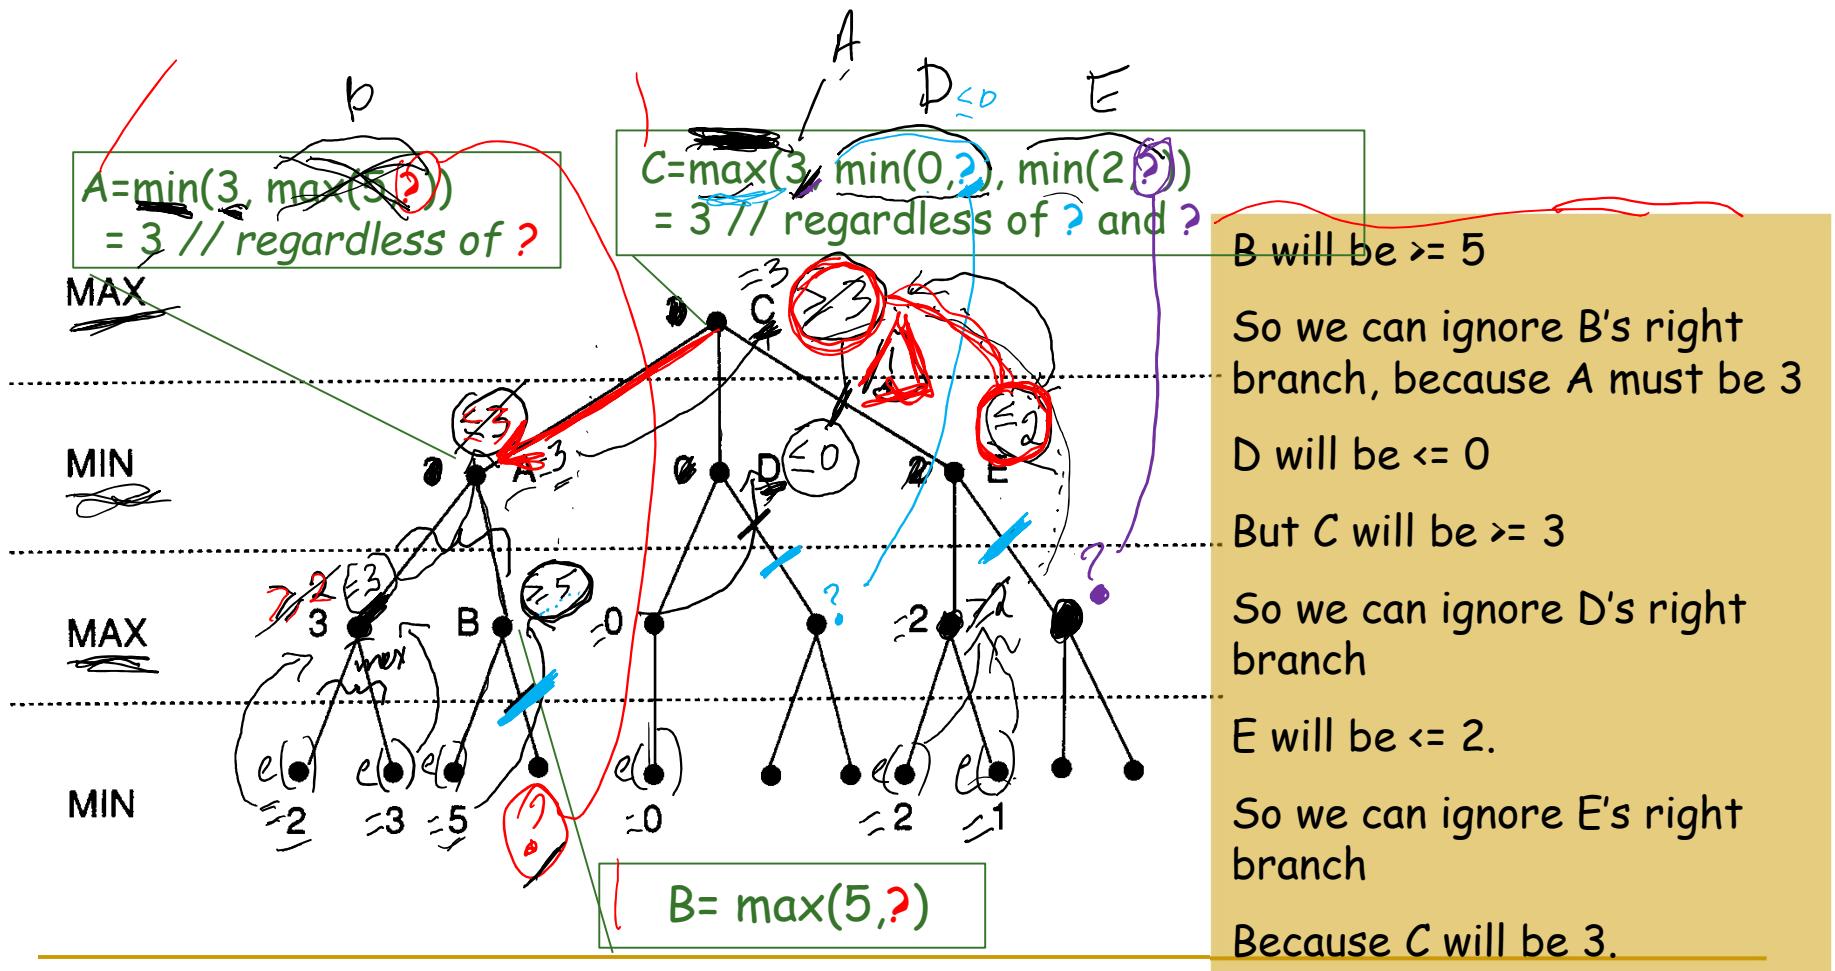
MAX

MIN

11 X e(n)
6 X e(n)

source: G. Luger (2005)

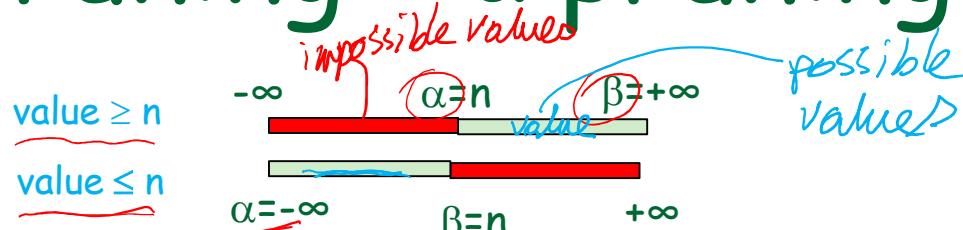
Alpha-Beta Pruning: Example 1



source: G. Luger (2005)

Alpha-Beta Pruning - a pruning

- $\alpha = \text{minimum possible value}$
- $\beta = \text{maximum possible value}$

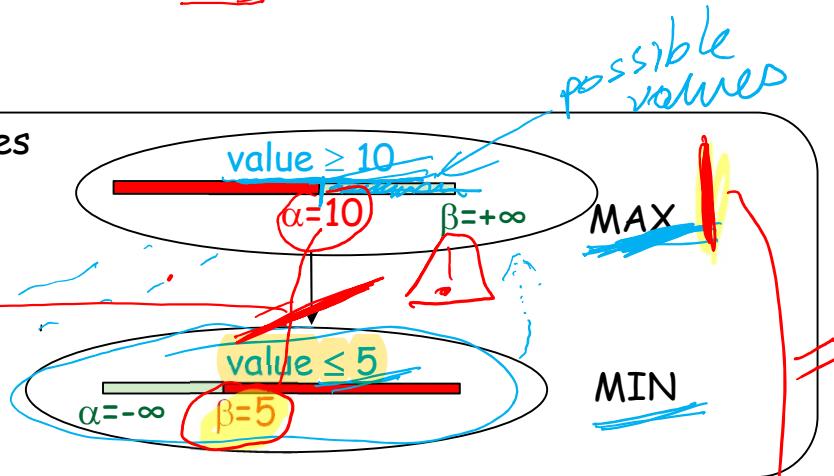


- Alpha pruning - parent is a MAX node

1. if MAX node's $\alpha = n$, then we can prune branches from a MIN descendant that has a $\beta \leq n$.

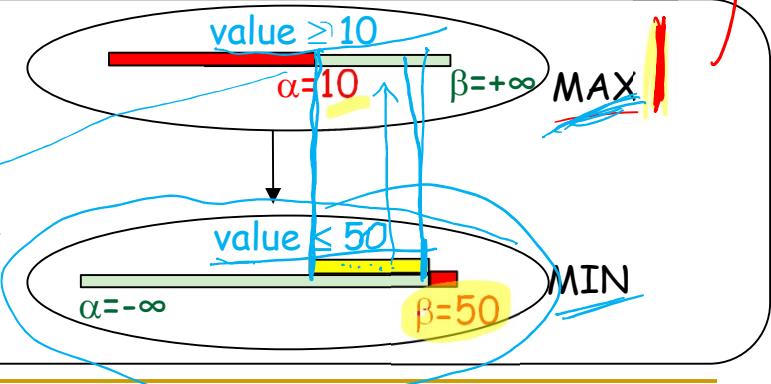
if $(\beta_{\text{child}} \leq \alpha_{\text{ancestor}}) \rightarrow \text{prune}$

incompatible...
so stop searching this branch;
the value cannot come from there!



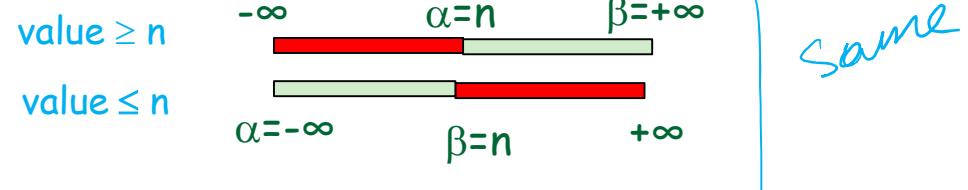
2. if $\beta_{\text{child}} > \alpha_{\text{ancestor}} \rightarrow \text{cannot prune}$

compatible...
we need to search this branch;
the value could come from there!



Alpha-Beta Pruning - β pruning

- α = minimum possible value
- β = maximum possible value

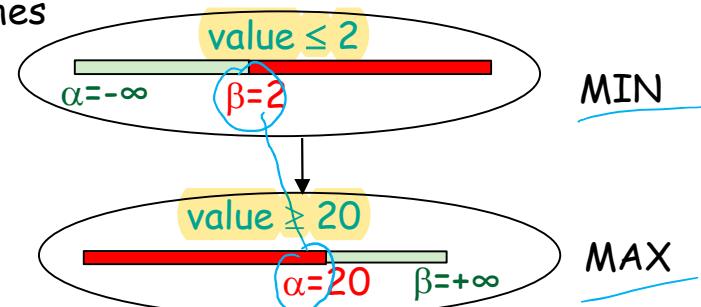


- Beta pruning - parent is a MIN node

1. if a MIN node's $\beta = n$, then we can prune branches from a MAX descendant that has an $\alpha \geq n$.

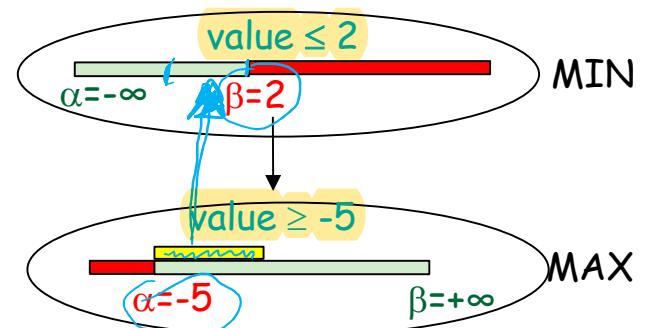
if $(\alpha_{\text{child}} \geq \beta_{\text{ancestor}}) \rightarrow \text{prune}$

incompatible...
so stop searching this branch;
the value cannot come from there!



2. if $(\alpha_{\text{child}} < \beta_{\text{ancestor}}) \rightarrow \text{cannot prune}$

compatible...
we need to search this branch;
the value could come from there!



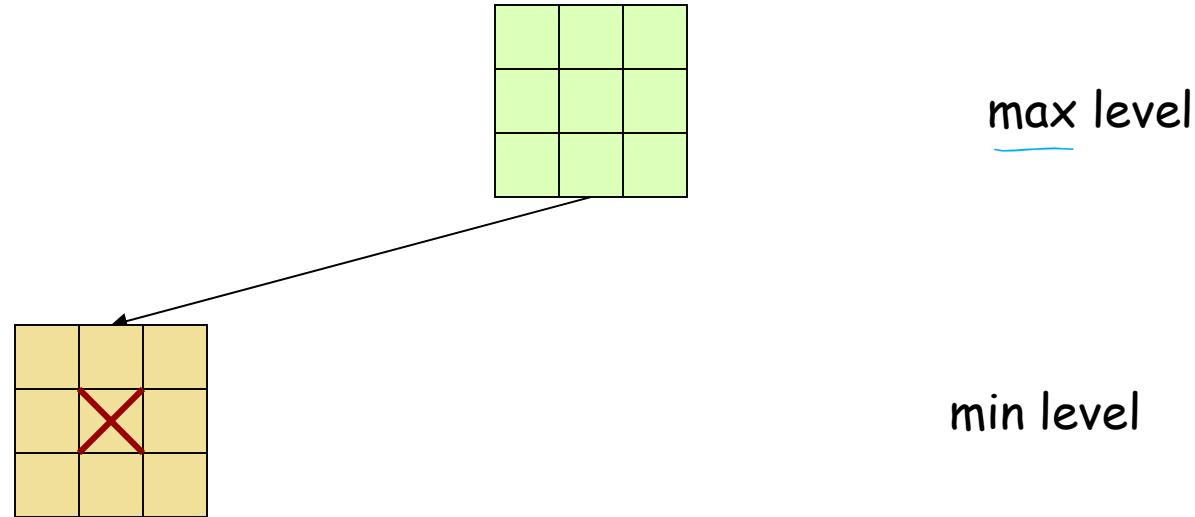
Alpha-Beta Pruning Algorithm

```
01 function alphabeta(node, depth,  $\alpha$ ,  $\beta$ , maximizingPlayer)
02     if depth = 0 or node is a terminal node
03         return the heuristic value of node
04     if maximizingPlayer
05         v := - $\infty$ 
06         for each child of node
07             v := max(v, alphabeta(child, depth - 1,  $\alpha$ ,  $\beta$ , FALSE))
08              $\alpha$  := max( $\alpha$ , v)
09             if  $\beta$   $\leq$   $\alpha$ 
10                 break (*  $\beta$  cut-off *)
11         return v
12     else
13         v :=  $\infty$ 
14         for each child of node
15             v := min(v, alphabeta(child, depth - 1,  $\alpha$ ,  $\beta$ , TRUE))
16              $\beta$  := min( $\beta$ , v)
17             if  $\beta$   $\leq$   $\alpha$ 
18                 break (*  $\alpha$  cut-off *)
19         return v
```

Initial call:

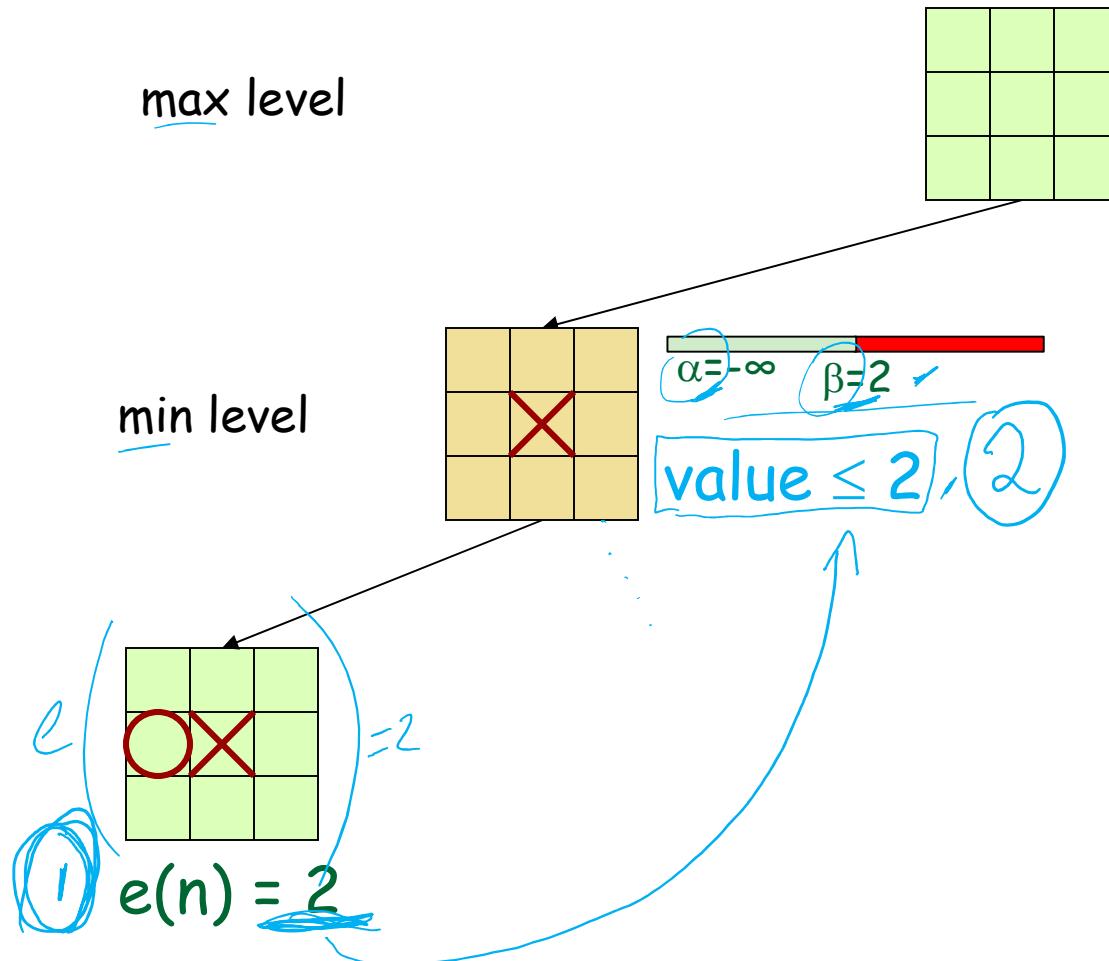
alphabeta(origin, depth, $-\infty$, $+\infty$, TRUE)

Example with tic-tac-toe



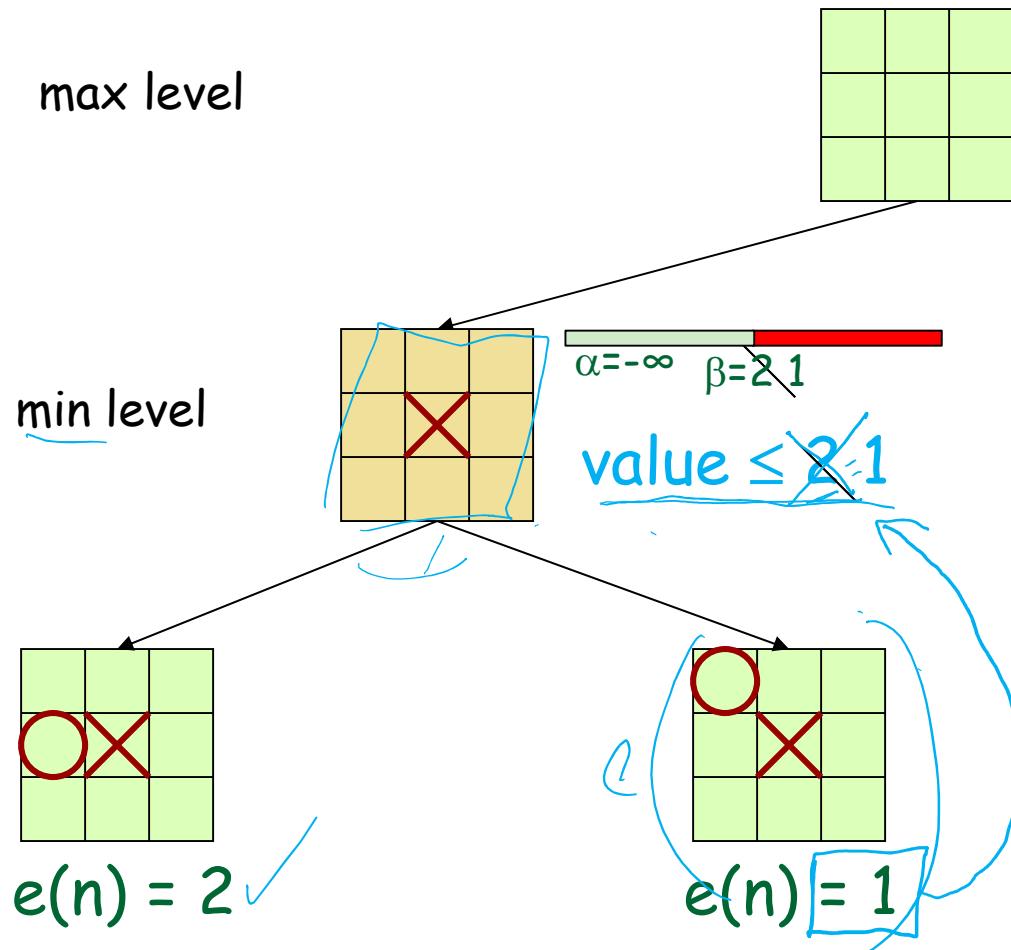
source: robotics.stanford.edu/~latombe/cs121/2003/home.htm

Example with tic-tac-toe



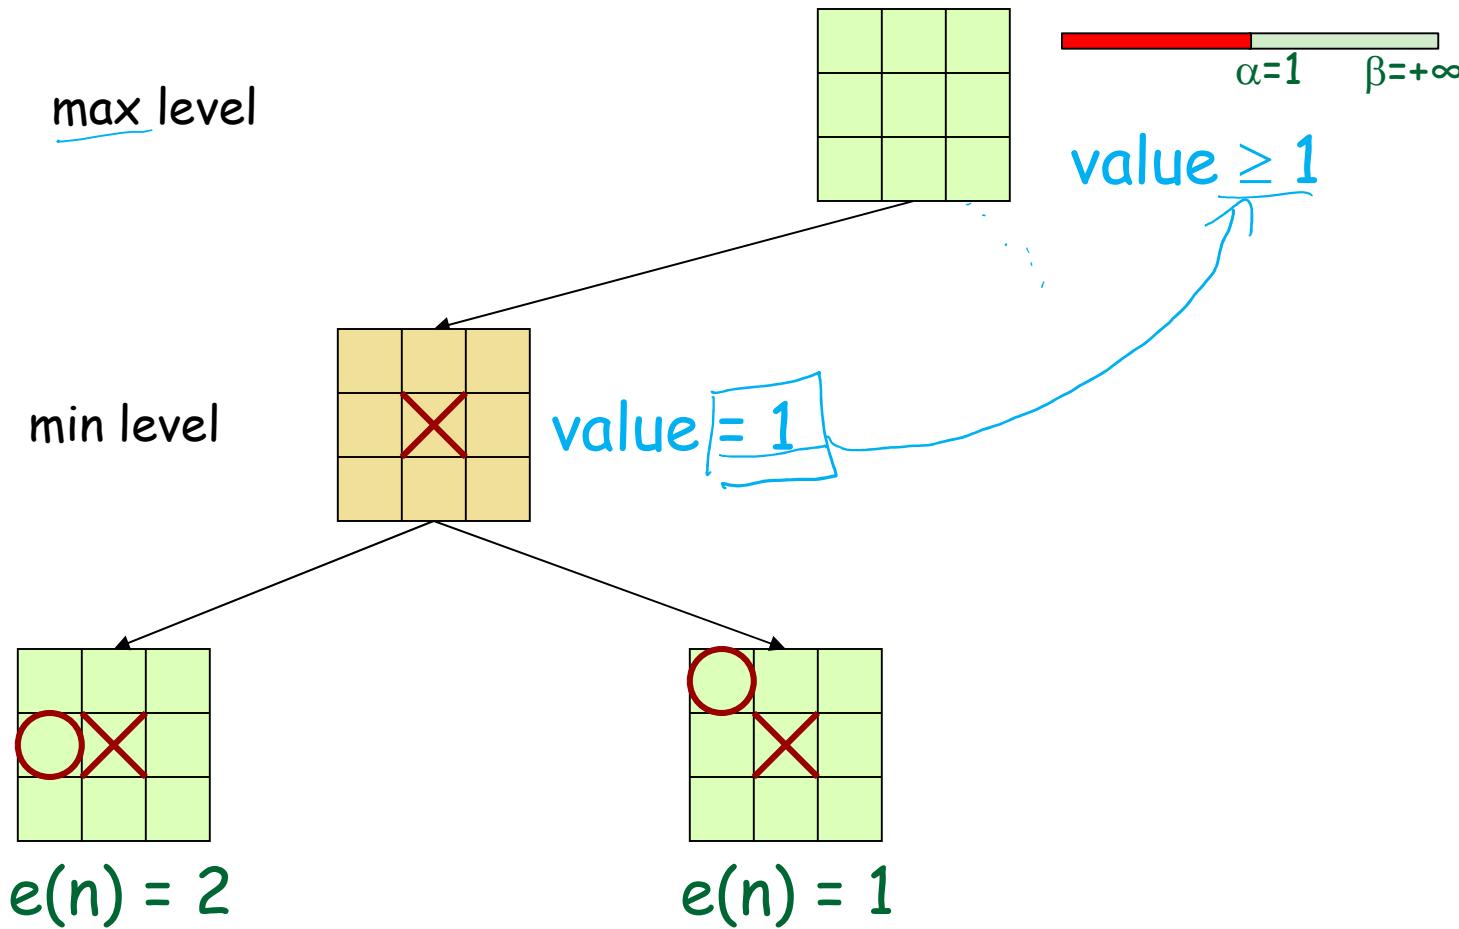
source: robotics.stanford.edu/~latombe/cs121/2003/home.htm

Example with tic-tac-toe

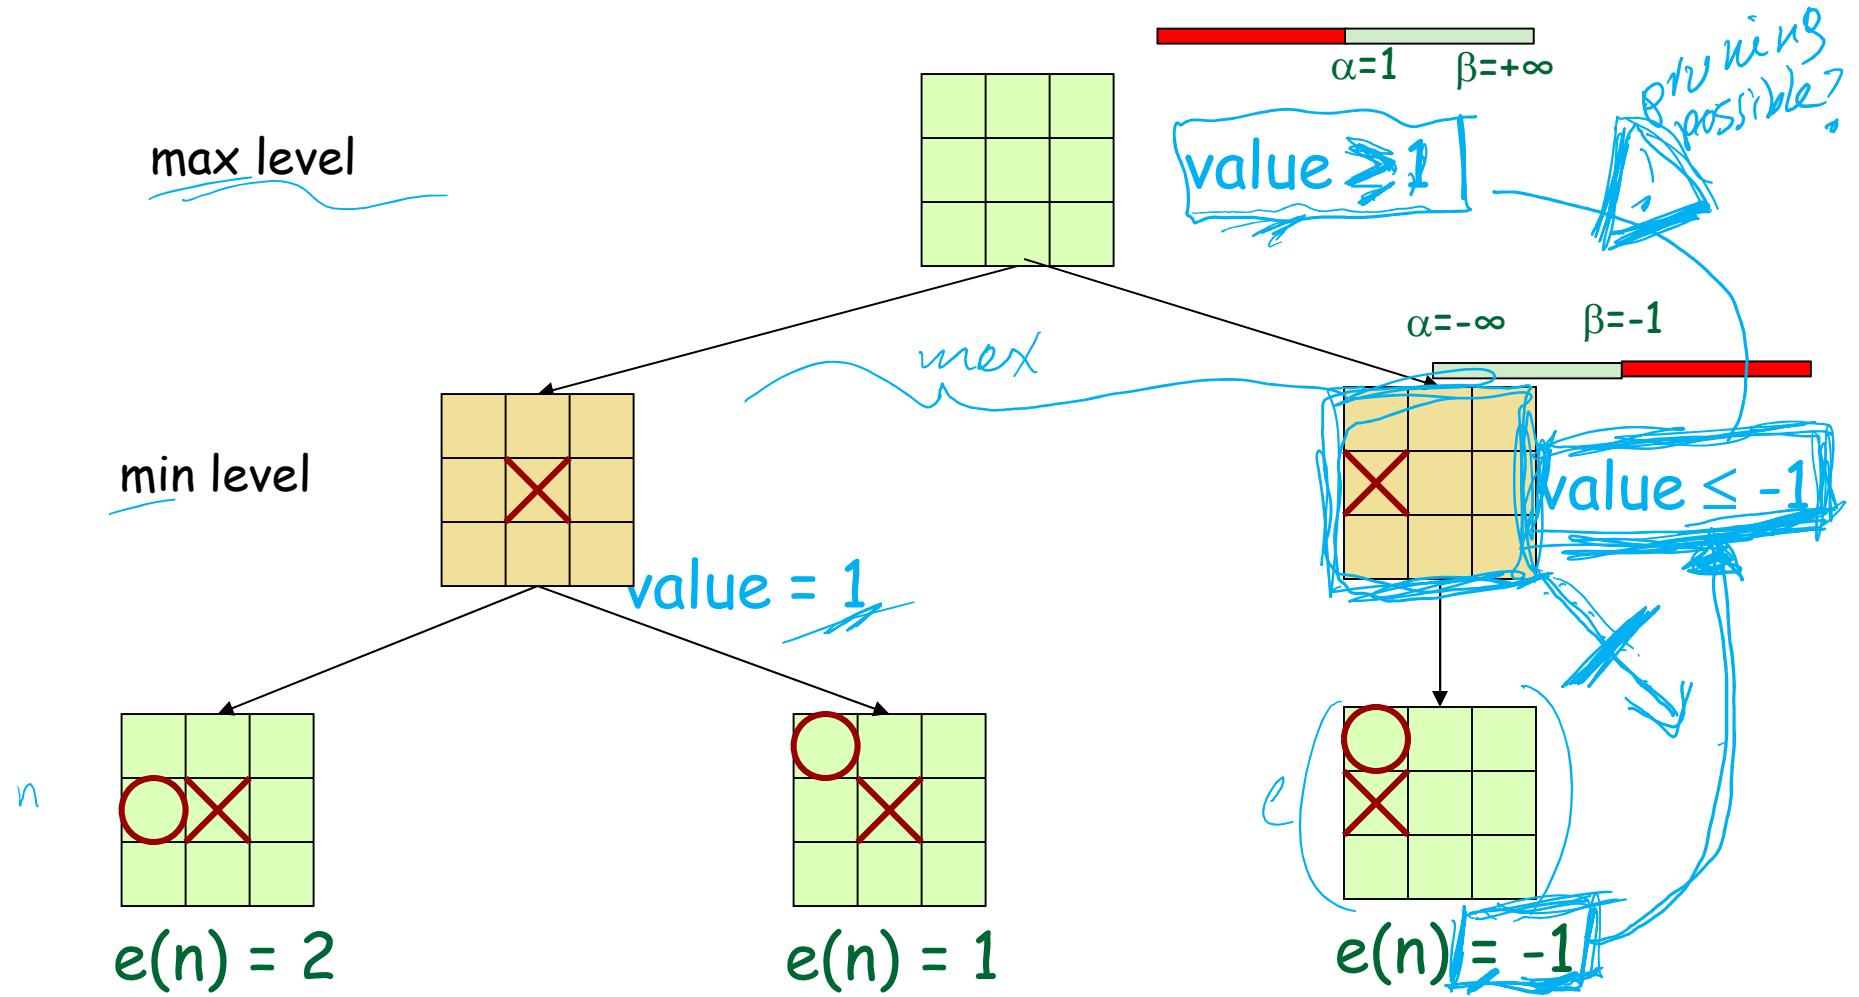


source: robotics.stanford.edu/~latombe/cs121/2003/home.htm

Example with tic-tac-toe



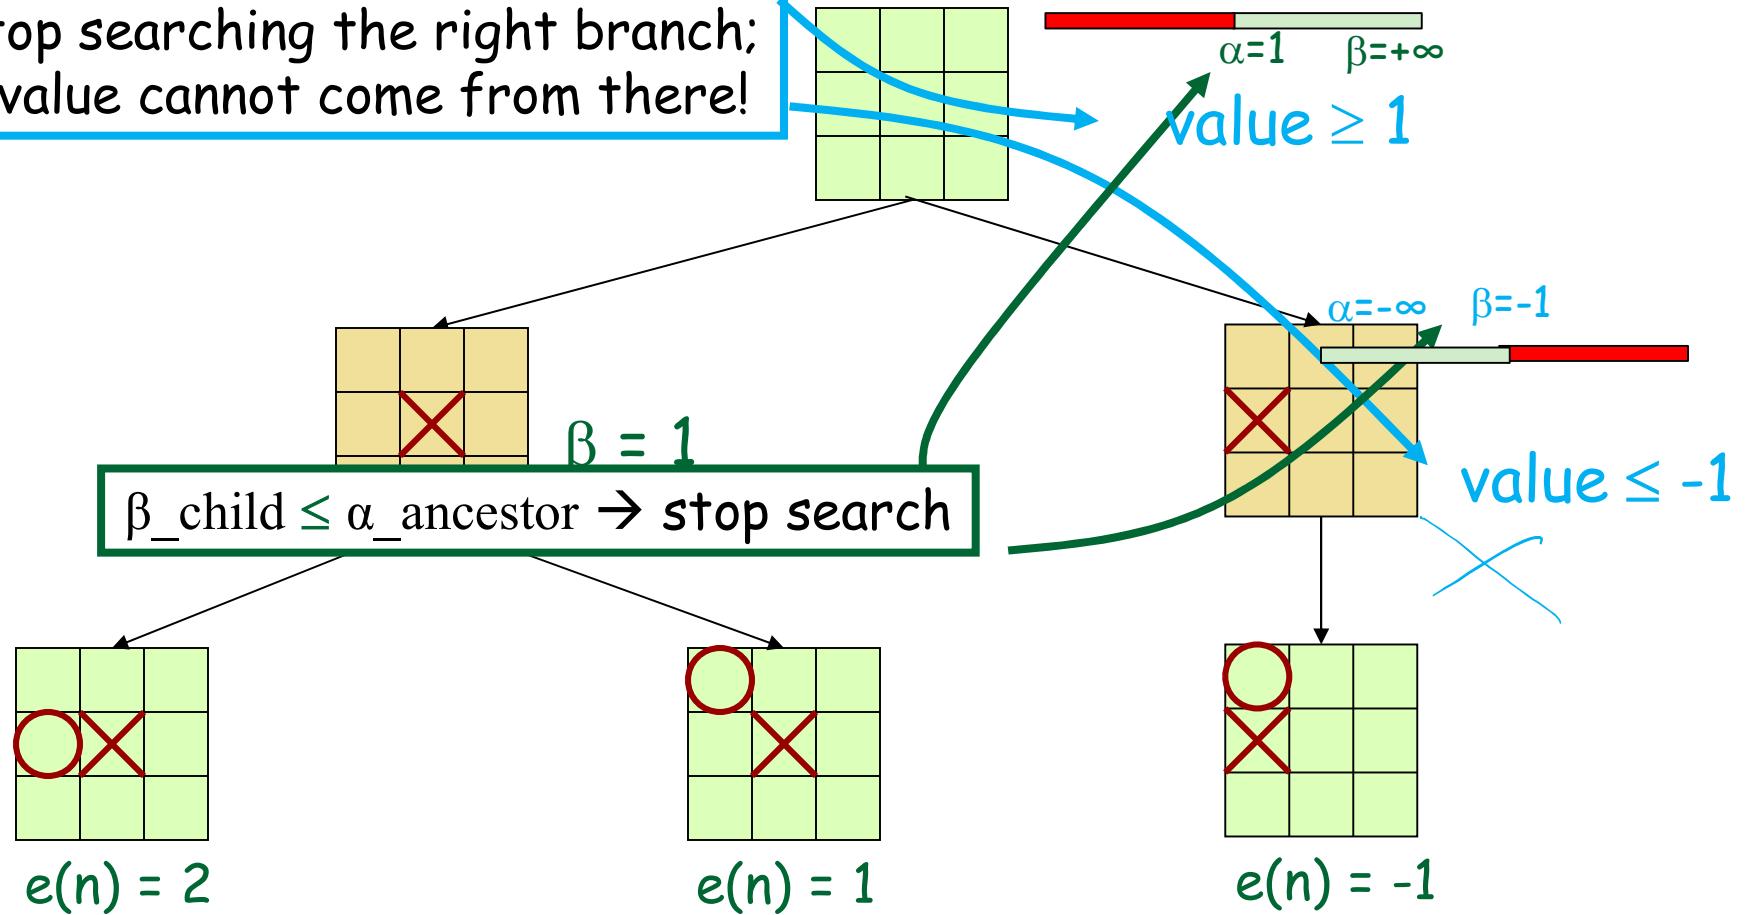
Example with tic-tac-toe



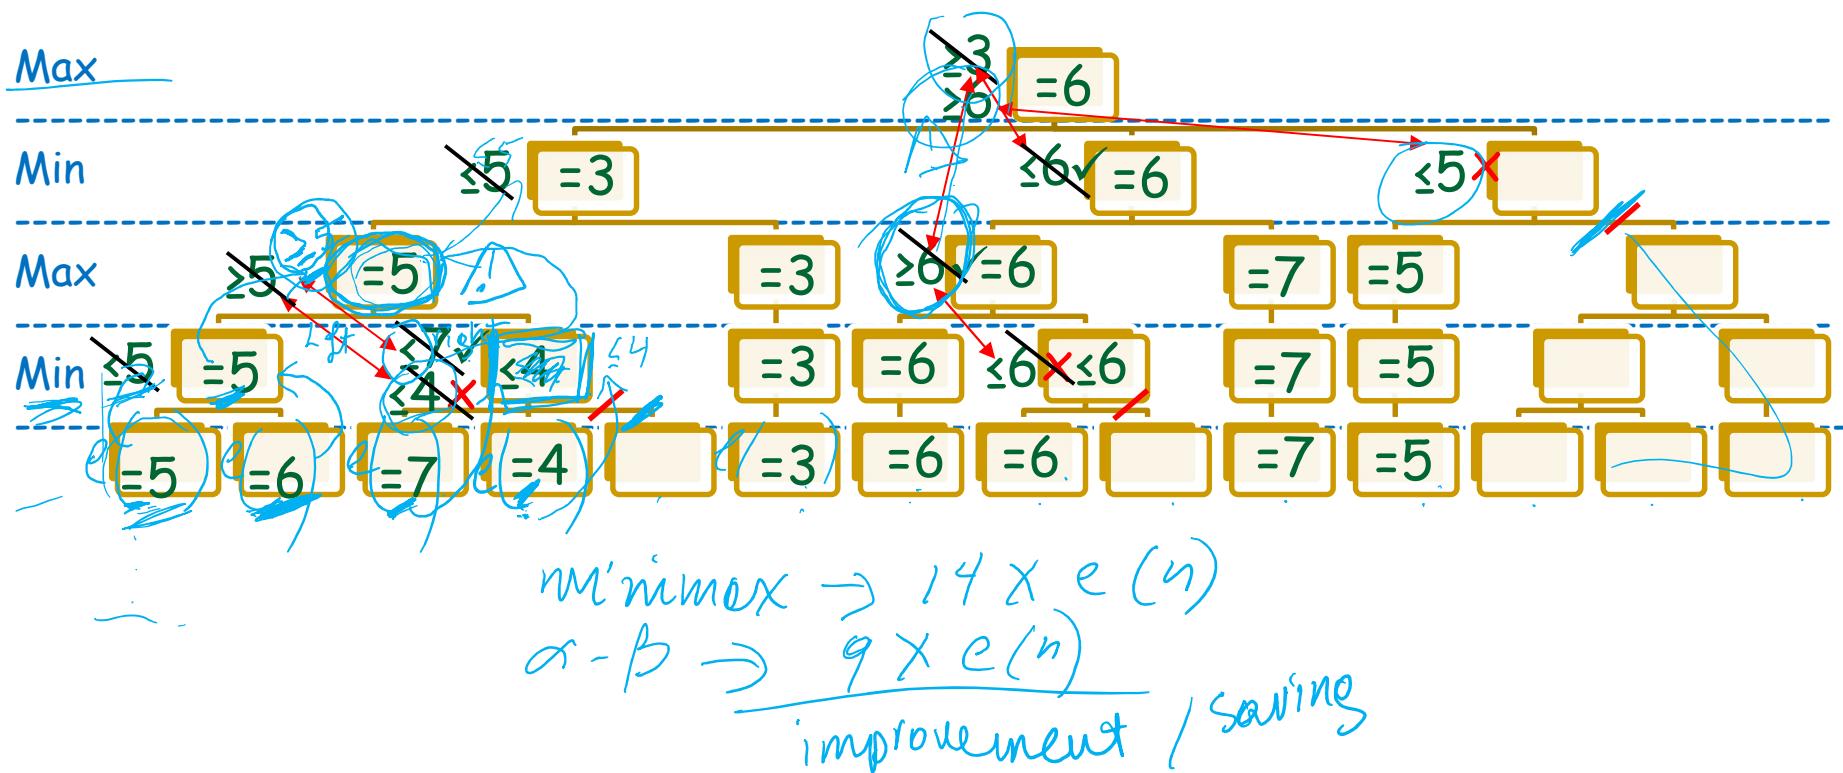
source: robotics.stanford.edu/~latombe/cs121/2003/home.htm

Example with tic-tac-toe

incompatible...
so stop searching the right branch;
the value cannot come from there!

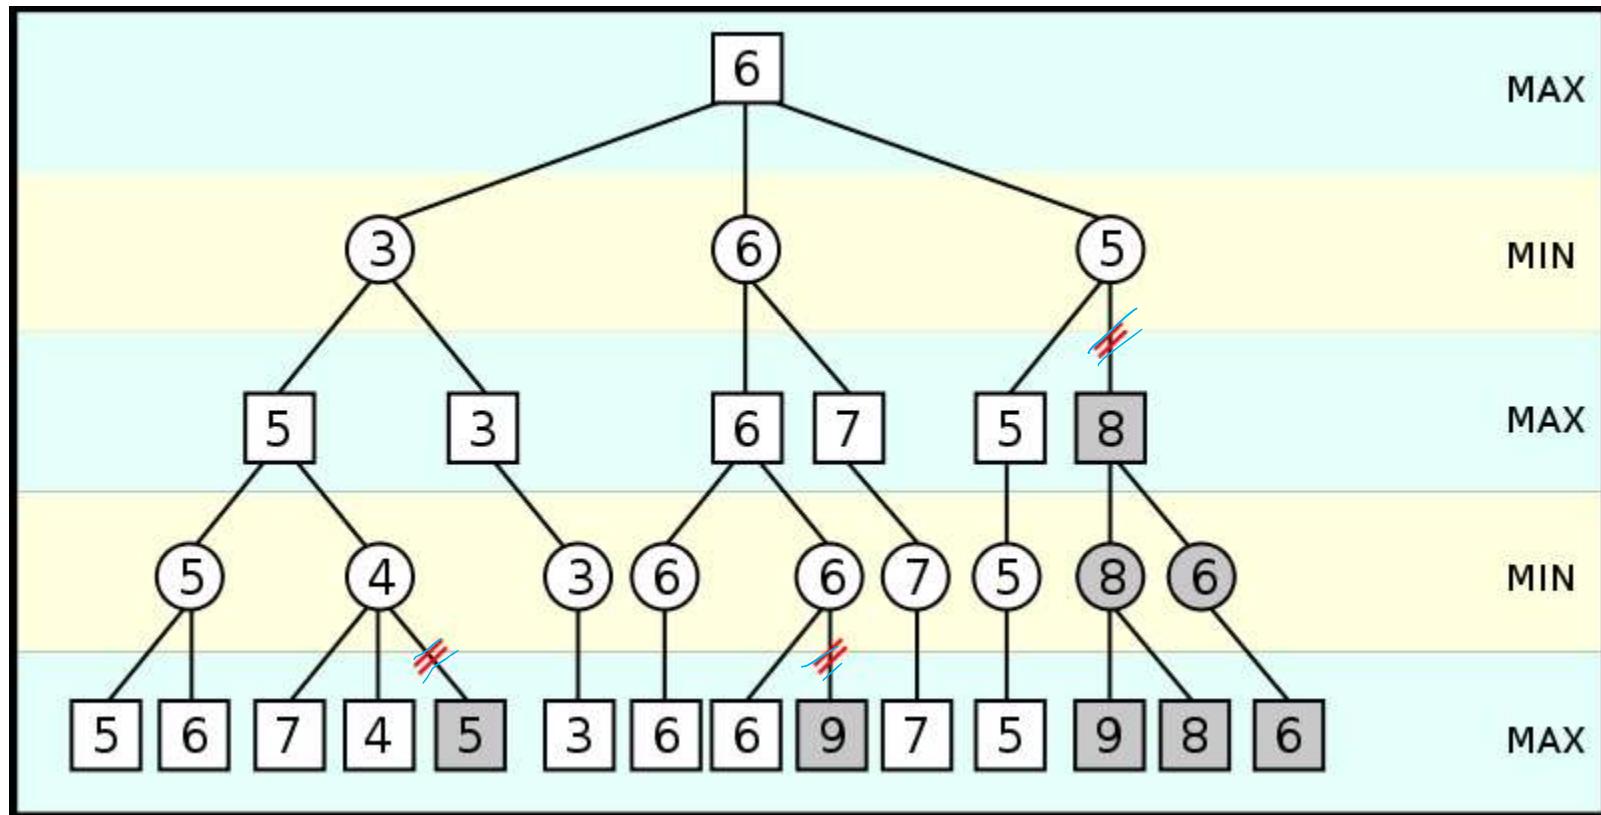


Alpha-Beta Pruning: Example 2



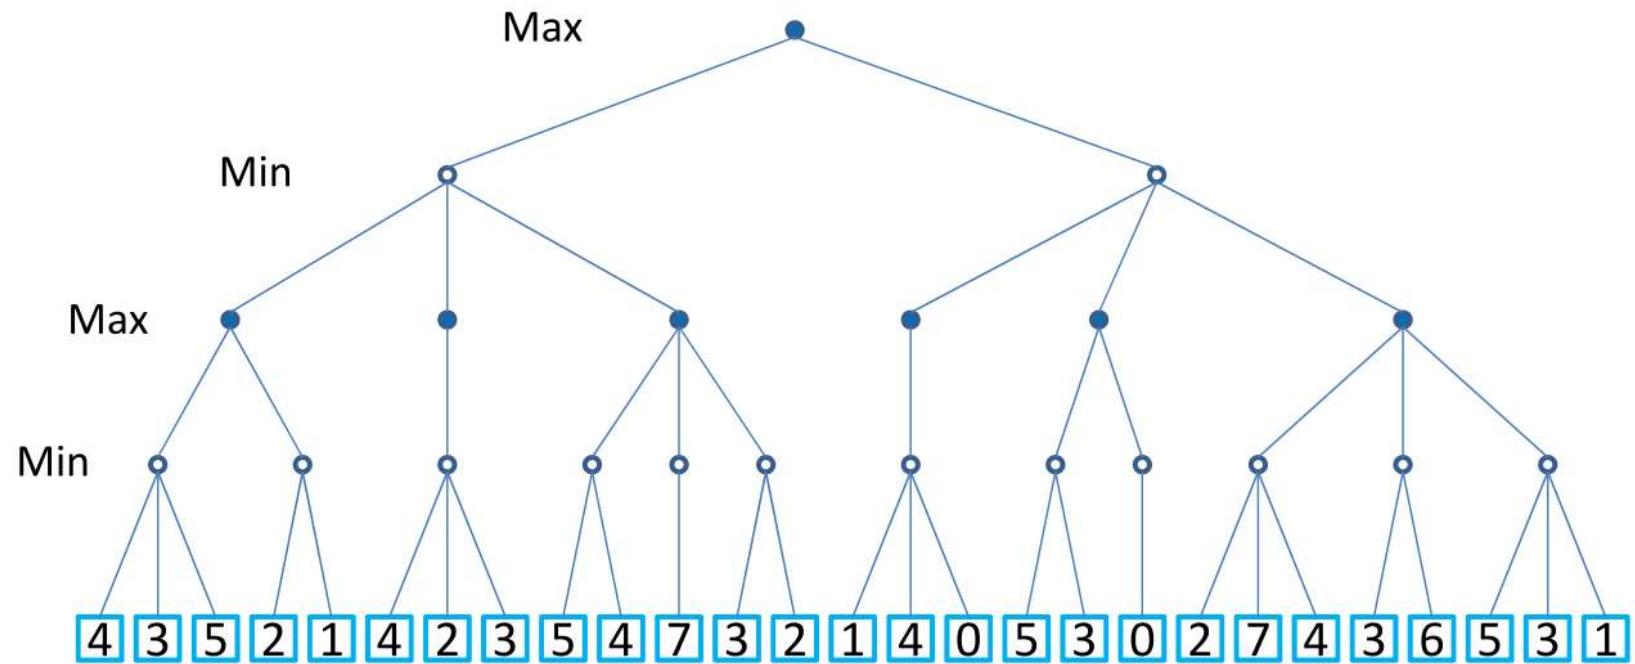
source: http://en.wikipedia.org/wiki/File:AB_pruning.svg

Alpha-Beta Pruning: Example 2

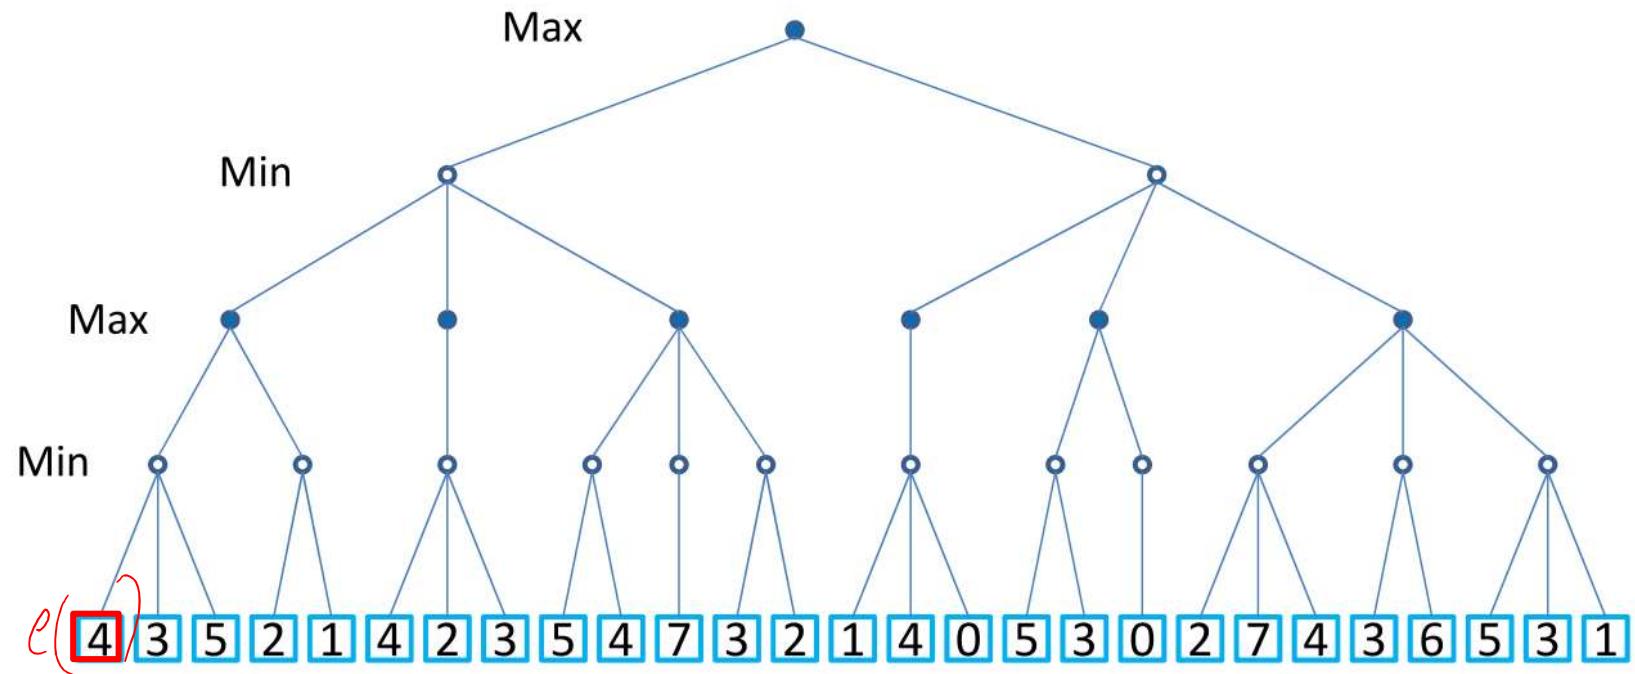


source: http://en.wikipedia.org/wiki/File:AB_pruning.svg

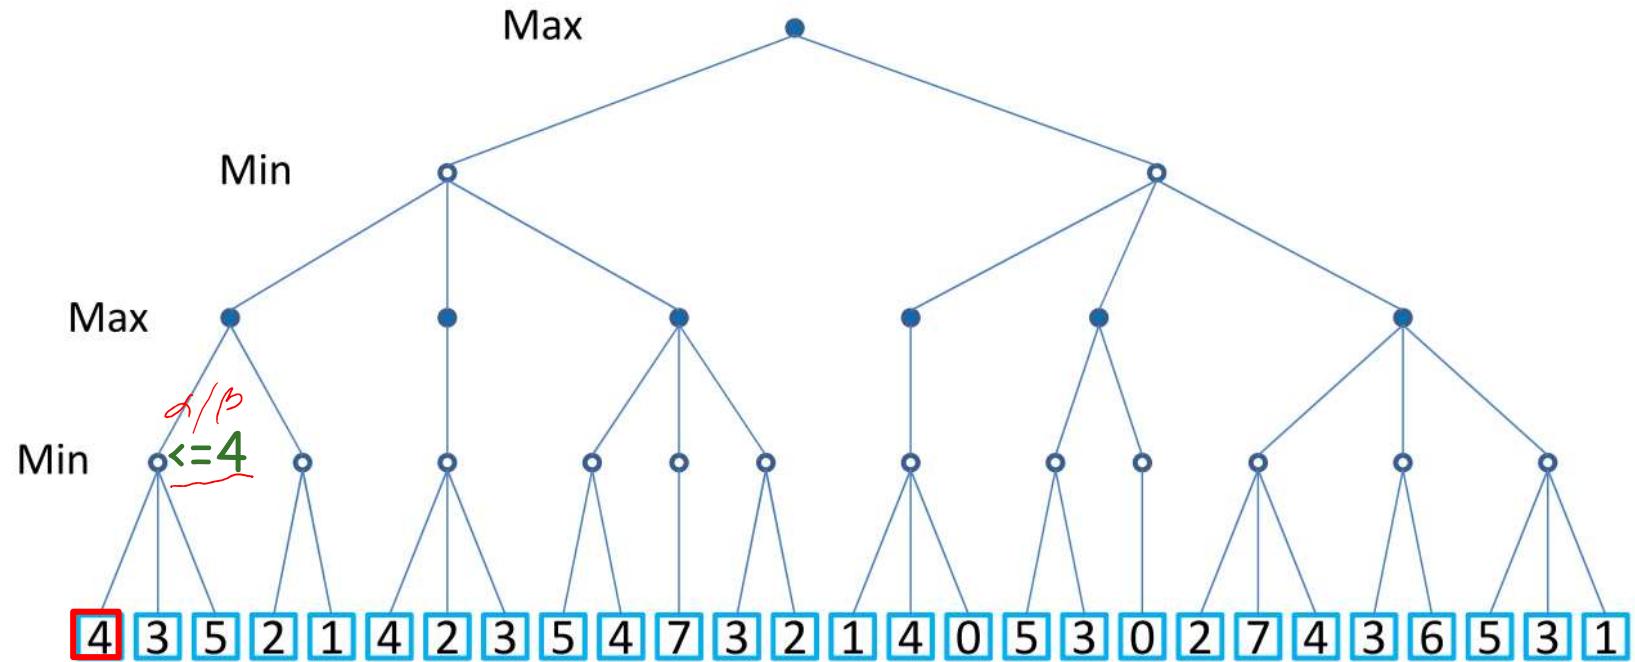
Alpha-Beta Pruning: Example 3



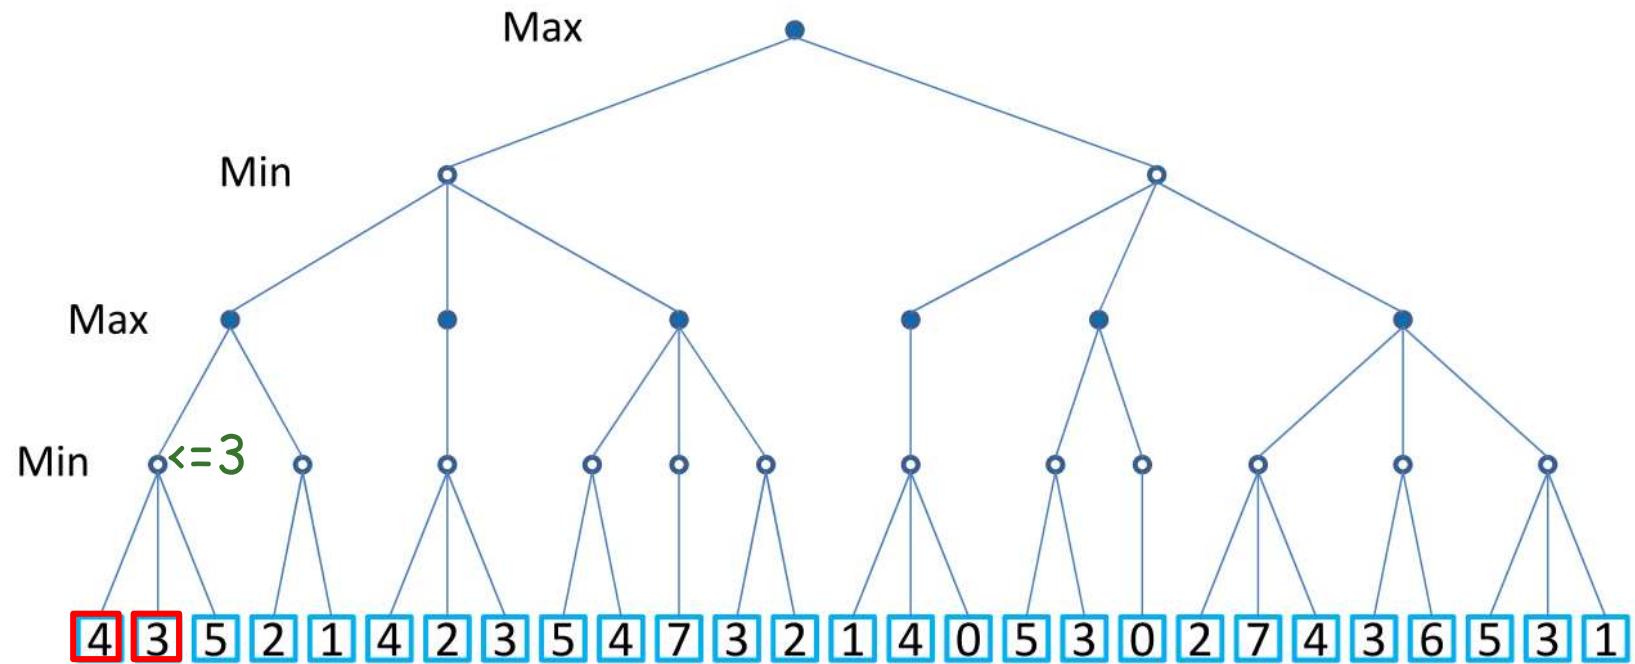
Alpha-Beta Pruning: Example 3



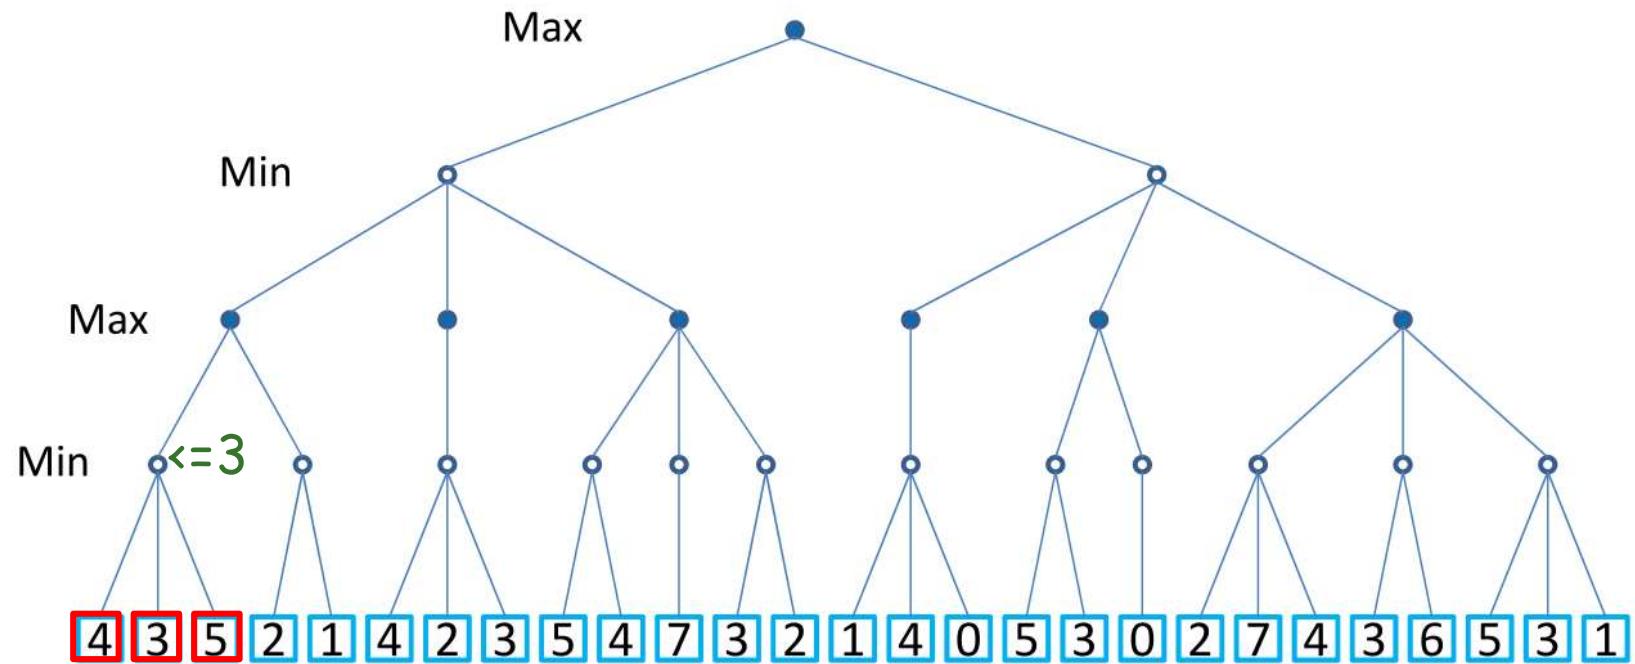
Alpha-Beta Pruning: Example 3



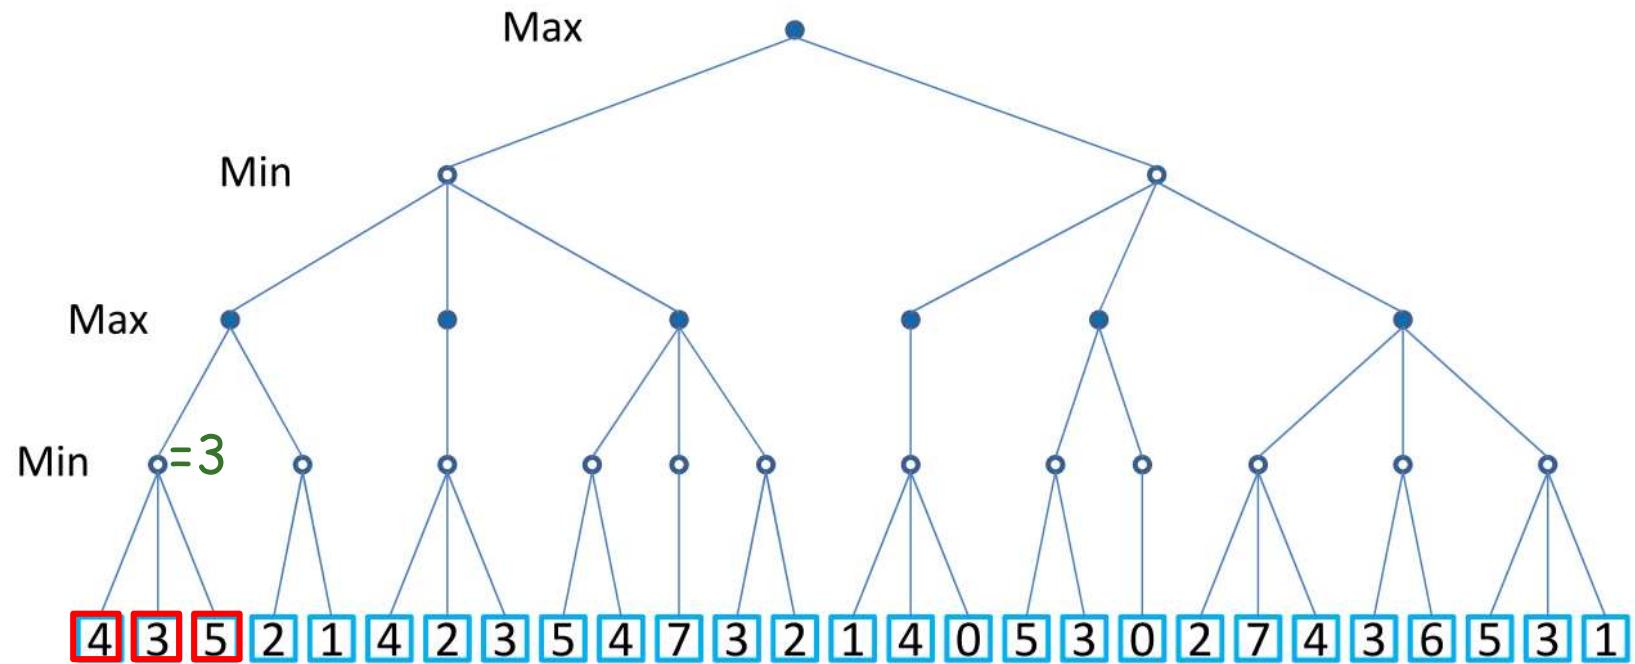
Alpha-Beta Pruning: Example 3



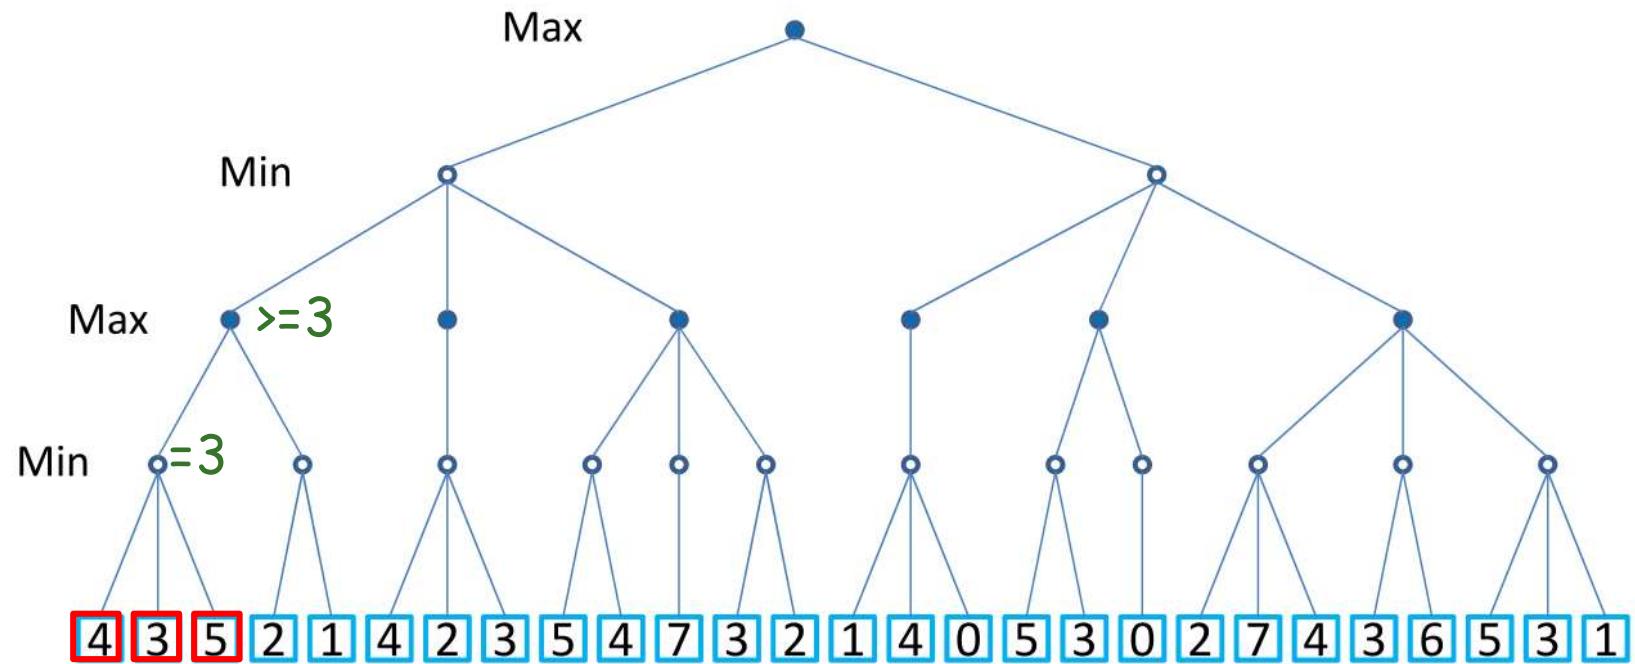
Alpha-Beta Pruning: Example 3



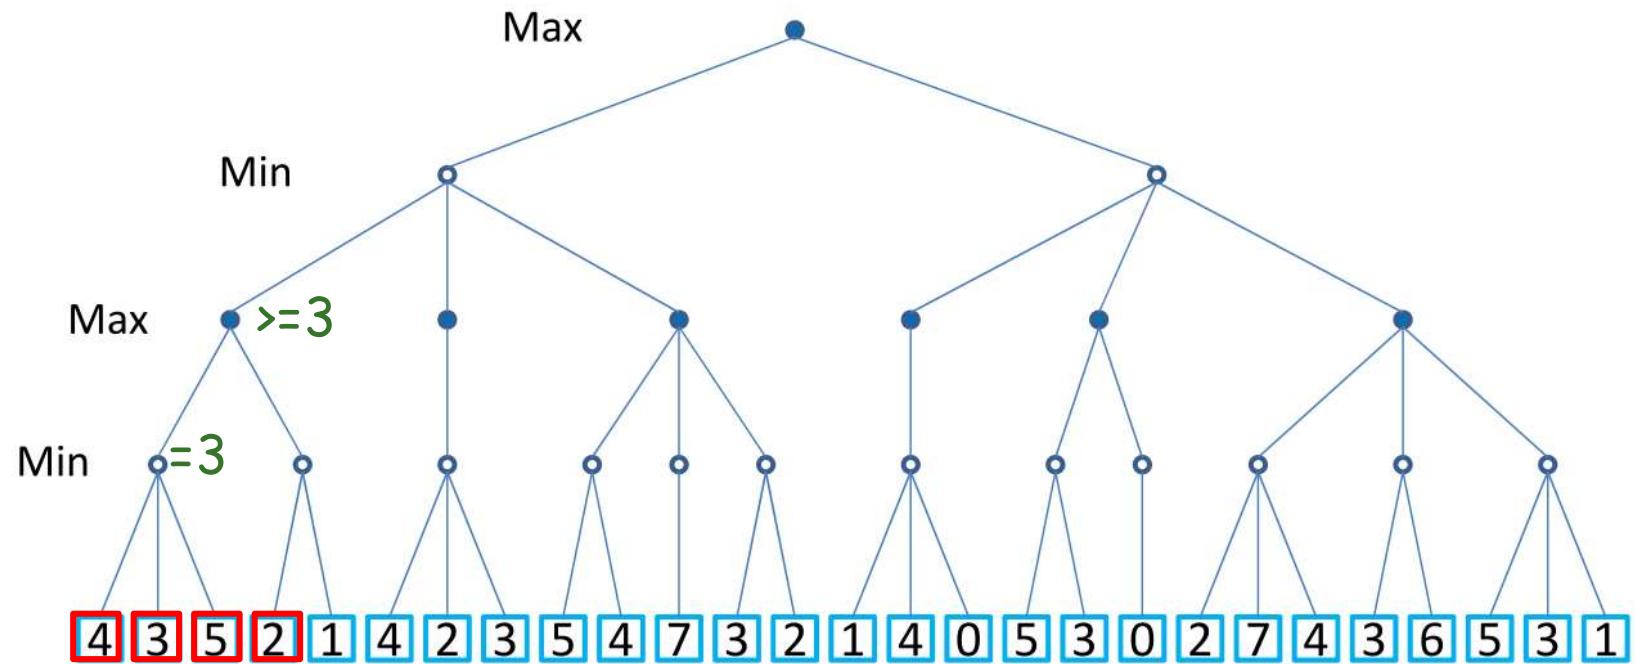
Alpha-Beta Pruning: Example 3



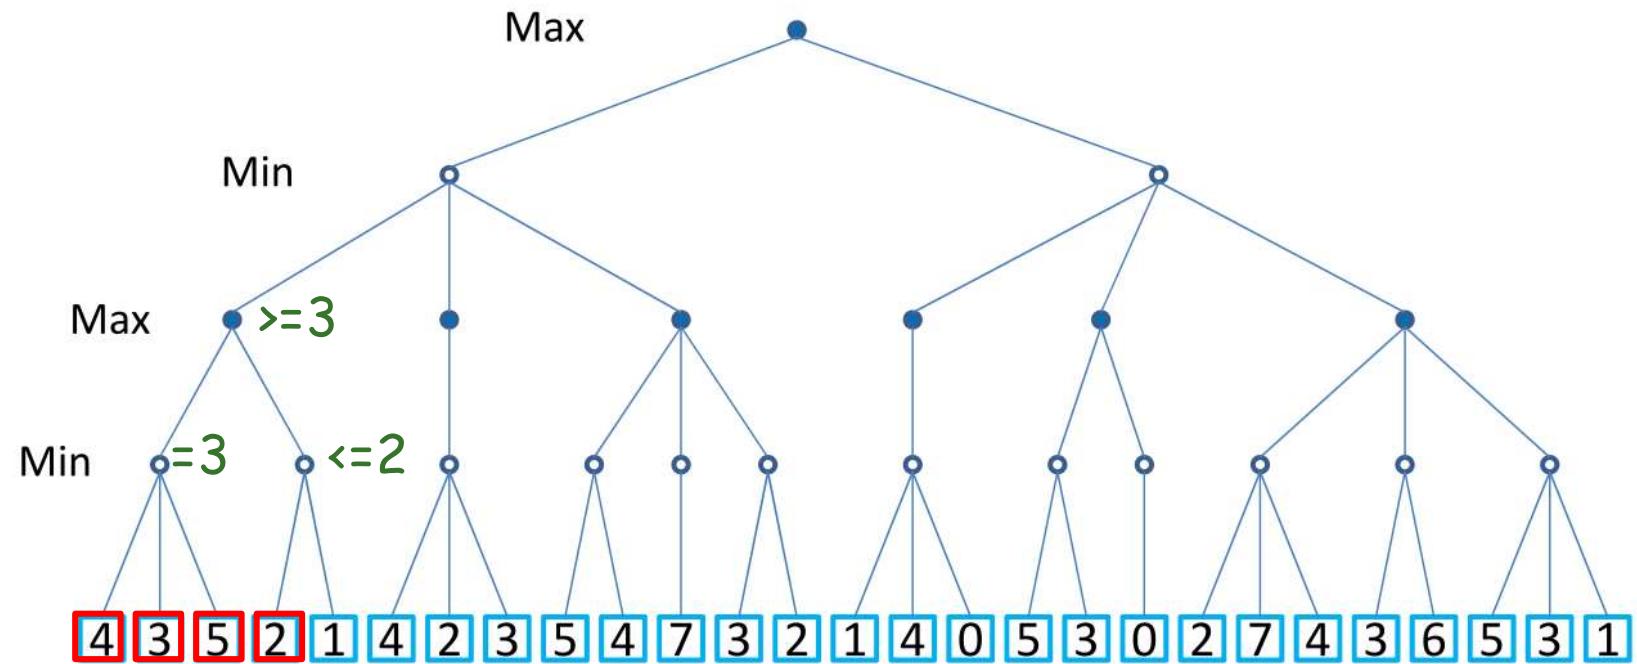
Alpha-Beta Pruning: Example 3



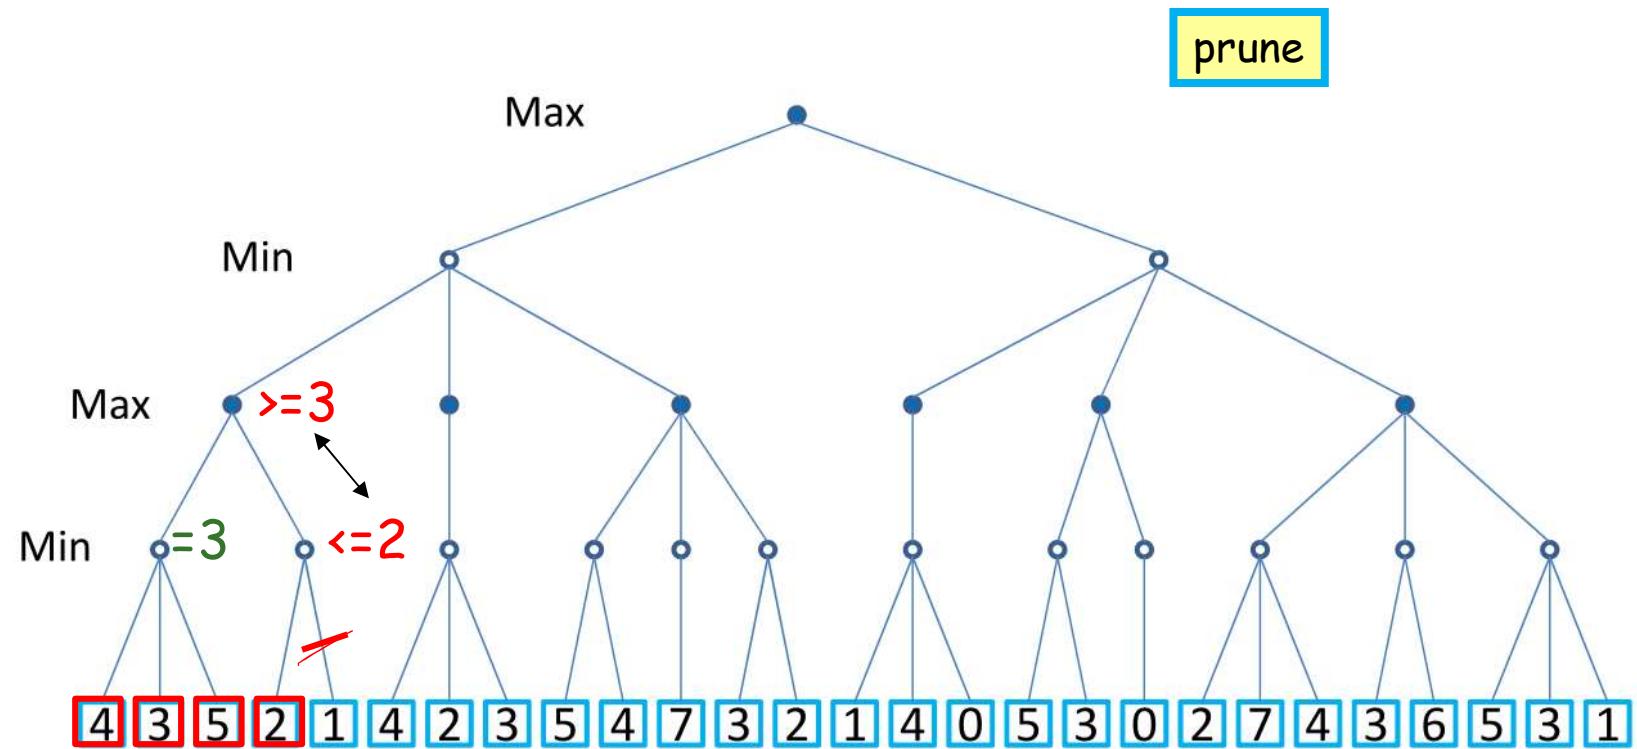
Alpha-Beta Pruning: Example 3



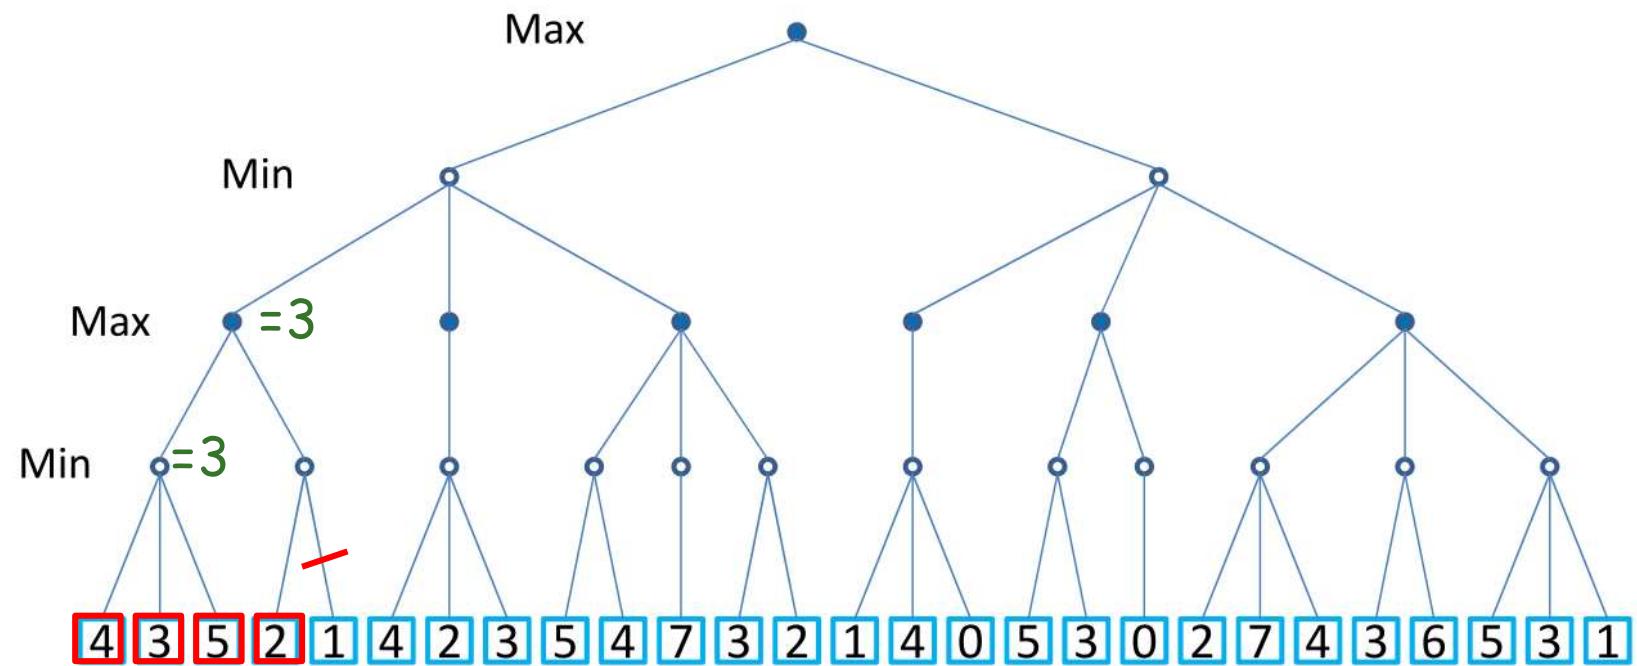
Alpha-Beta Pruning: Example 3



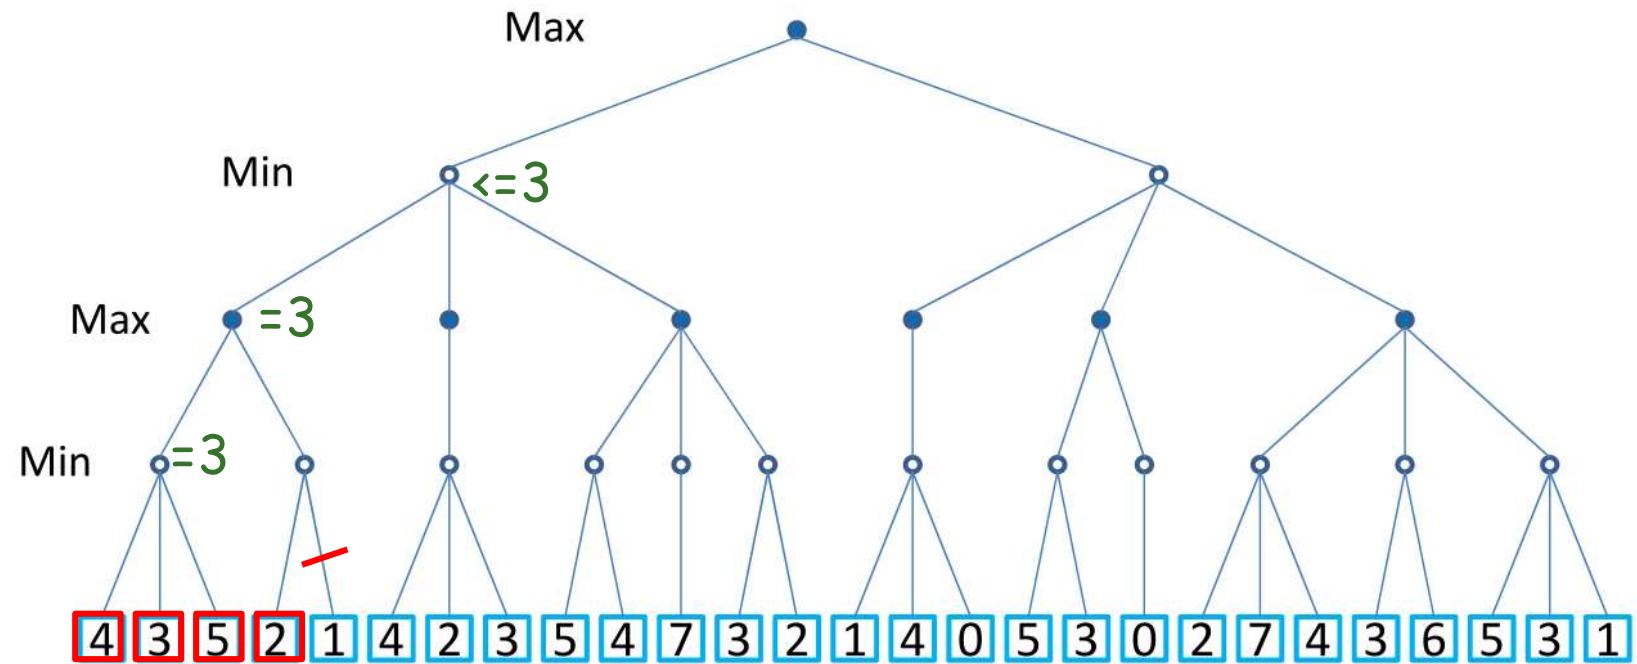
Alpha-Beta Pruning: Example 3



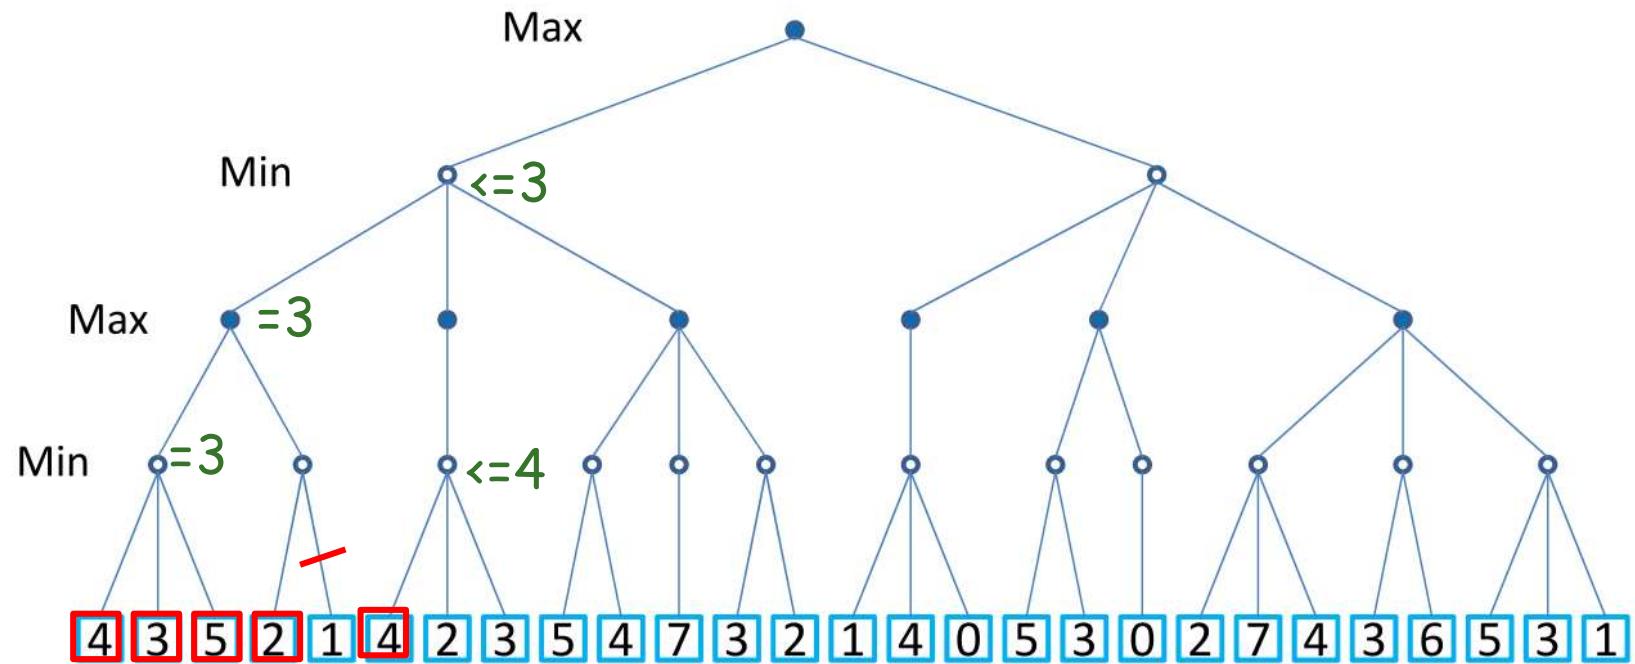
Alpha-Beta Pruning: Example 3



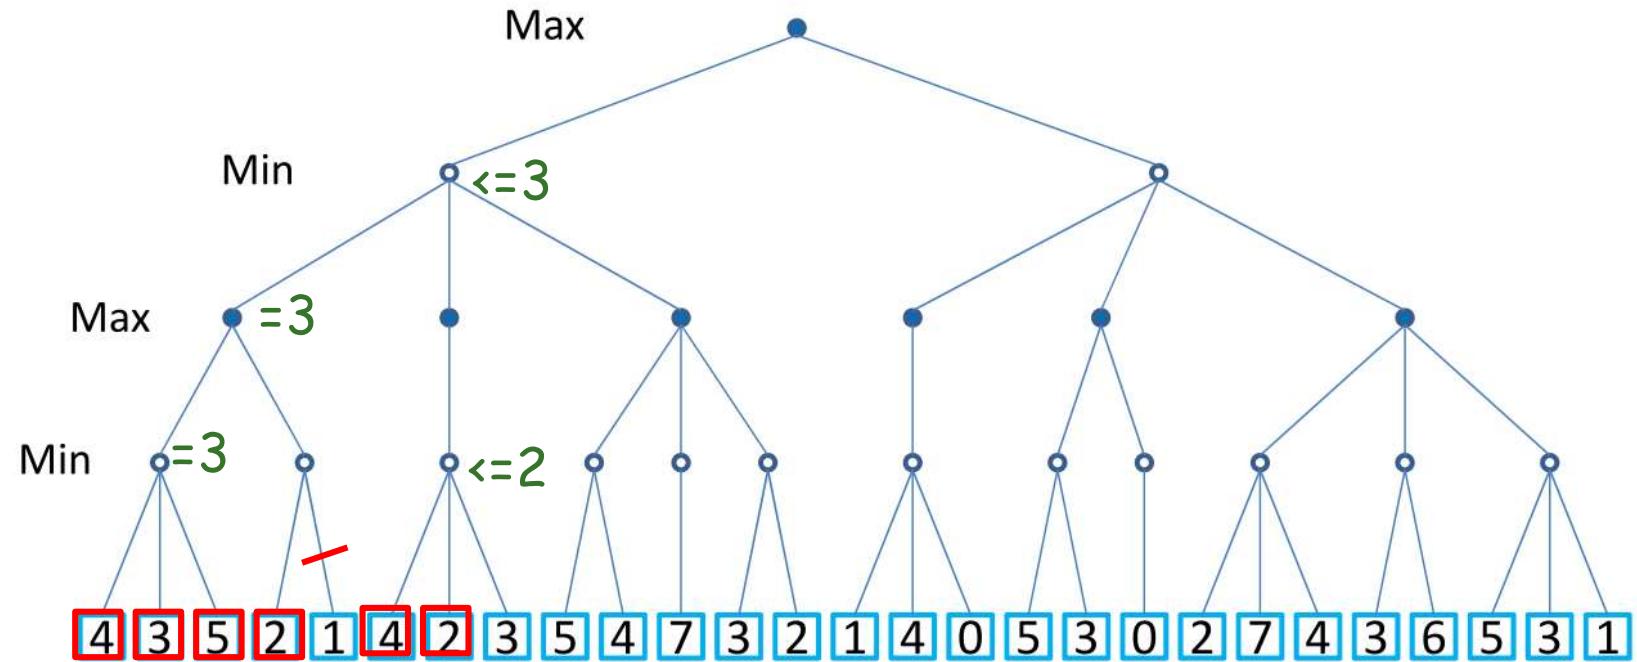
Alpha-Beta Pruning: Example 3



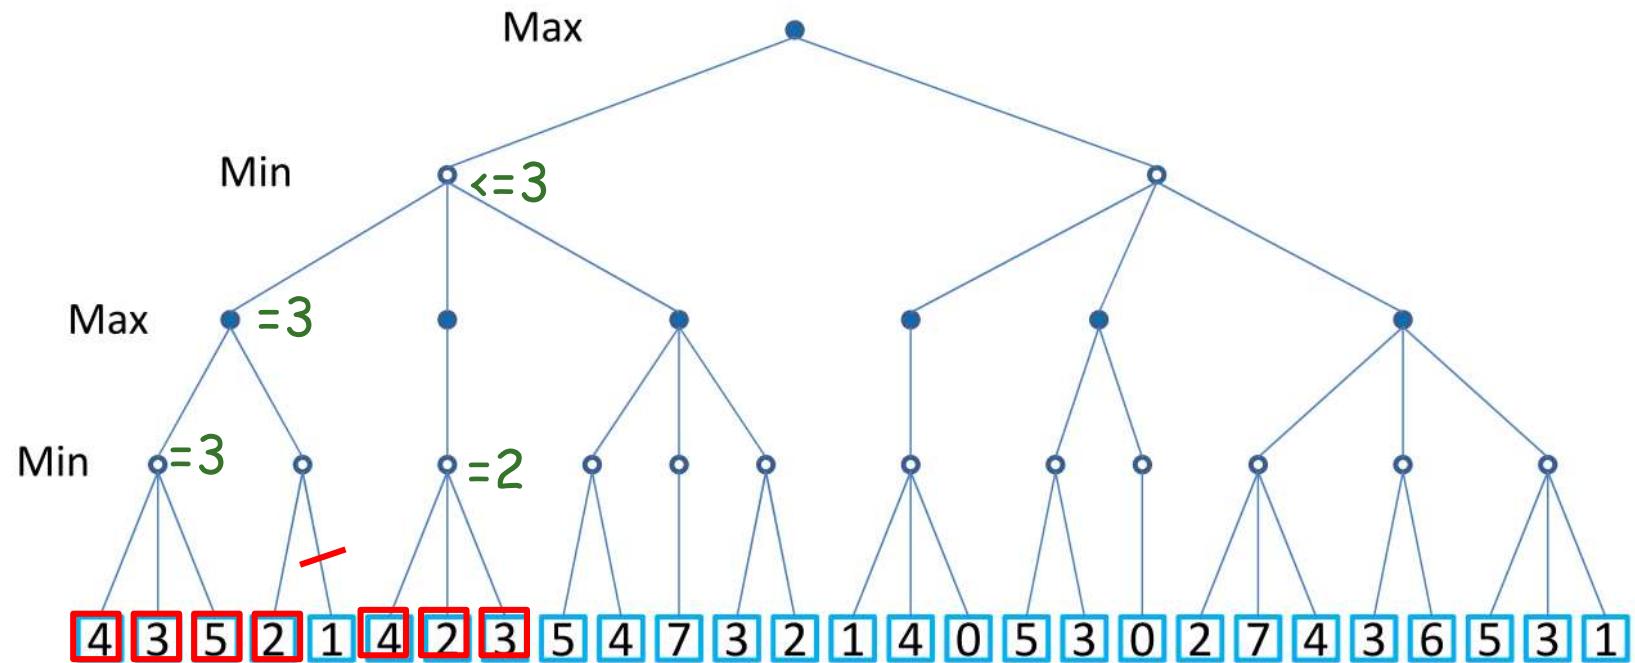
Alpha-Beta Pruning: Example 3



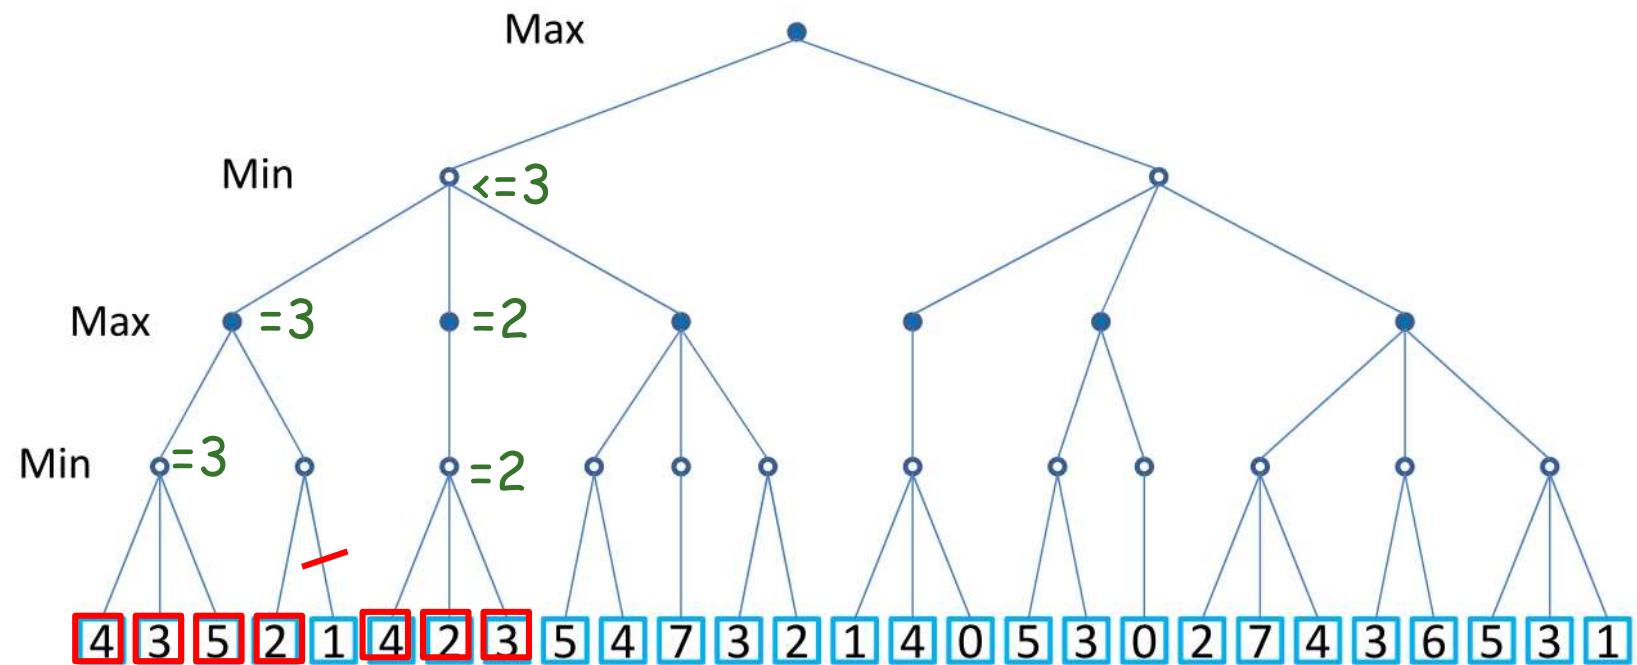
Alpha-Beta Pruning: Example 3



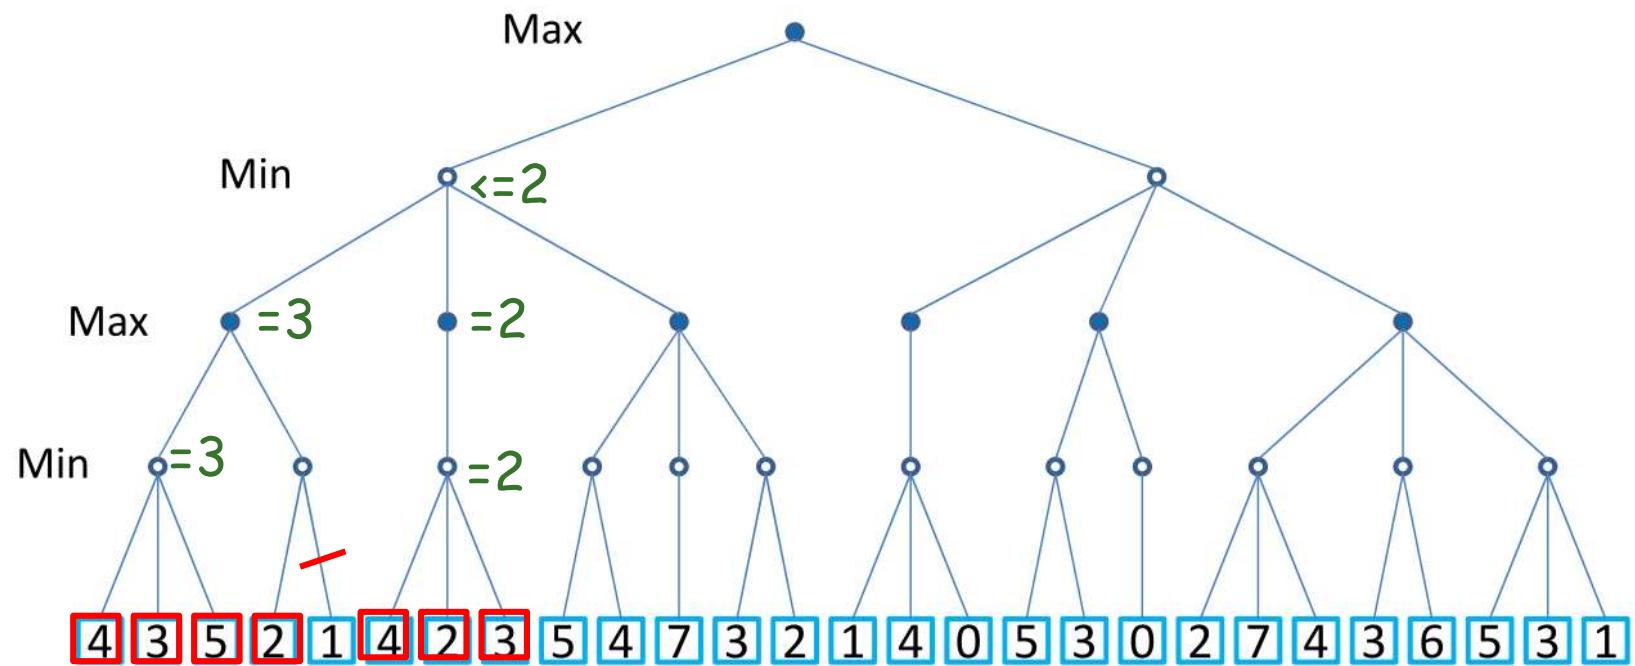
Alpha-Beta Pruning: Example 3



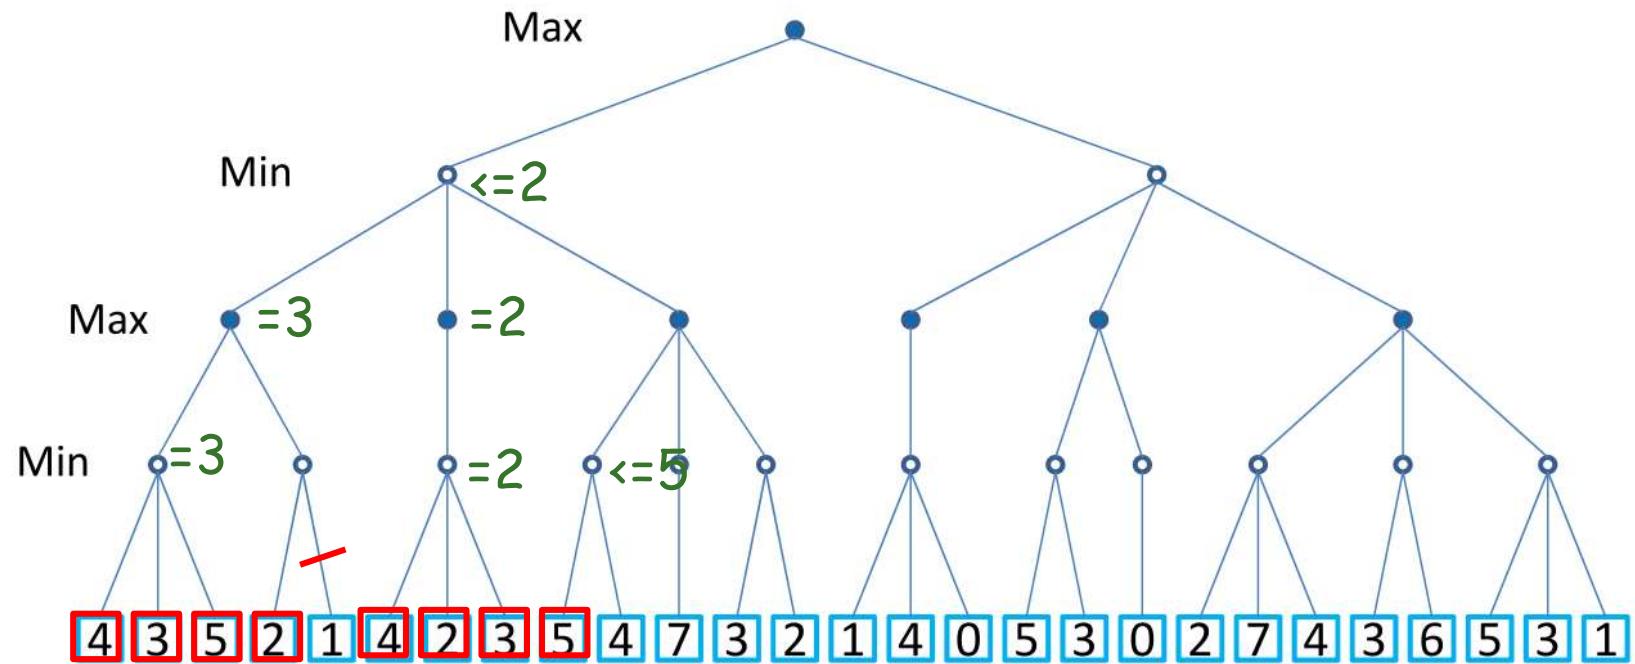
Alpha-Beta Pruning: Example 3



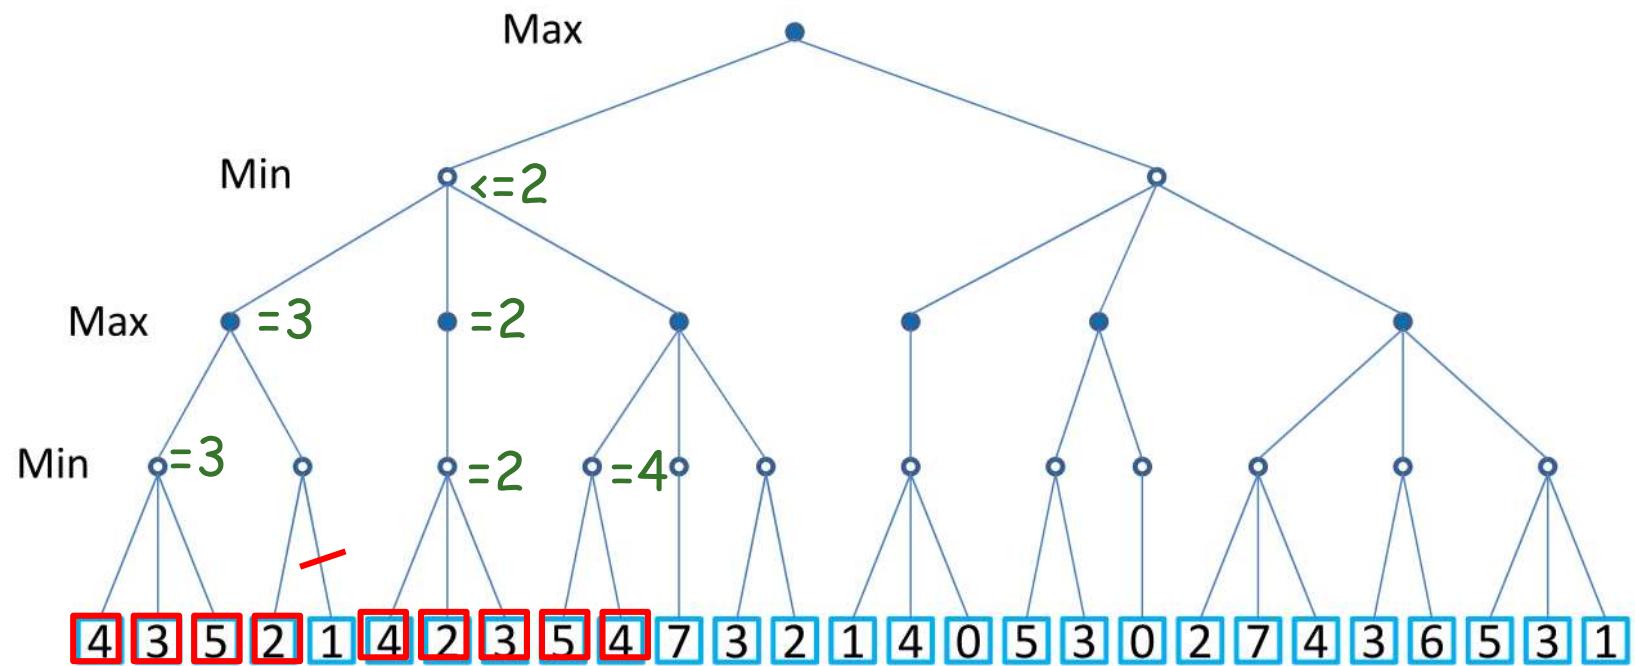
Alpha-Beta Pruning: Example 3



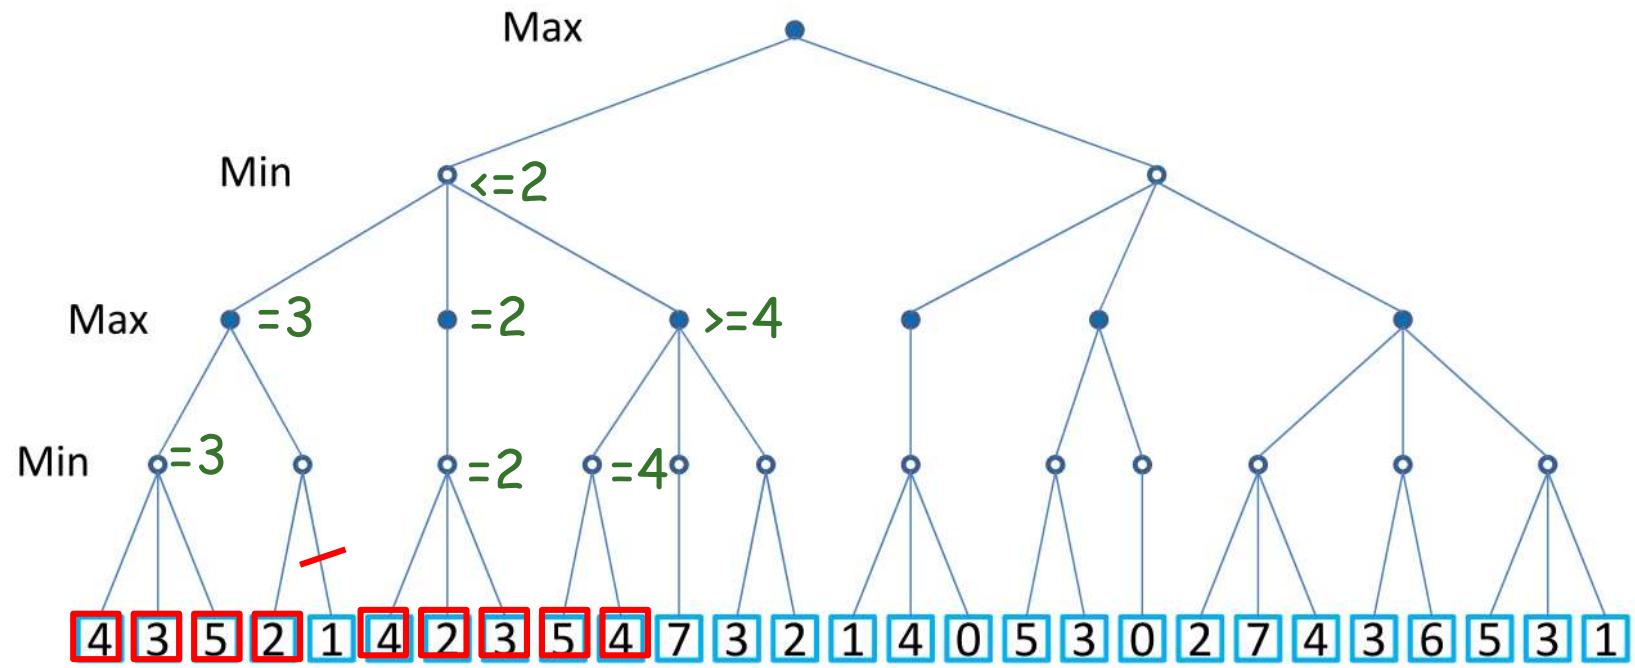
Alpha-Beta Pruning: Example 3



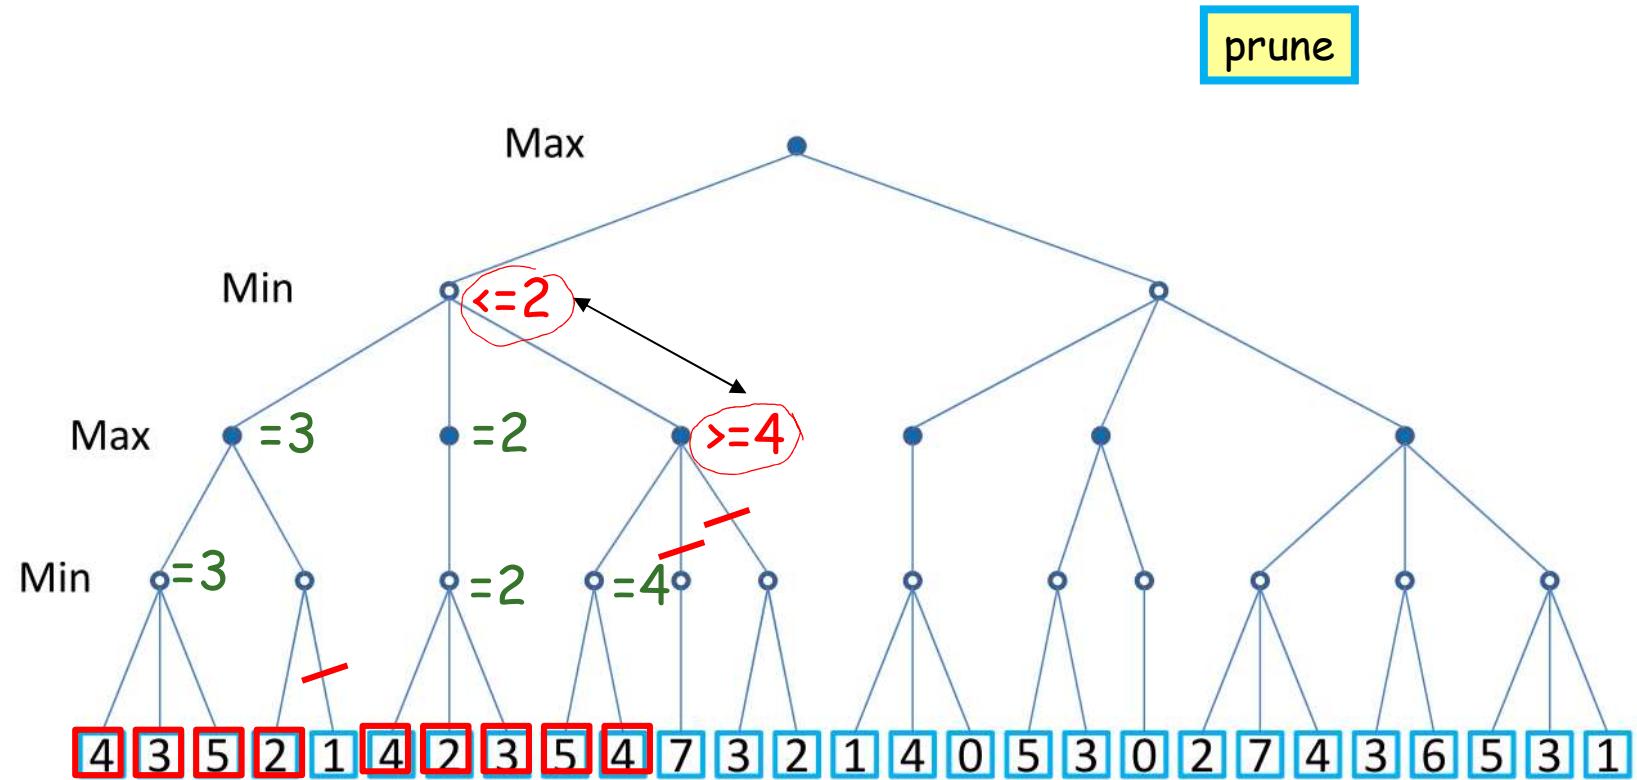
Alpha-Beta Pruning: Example 3



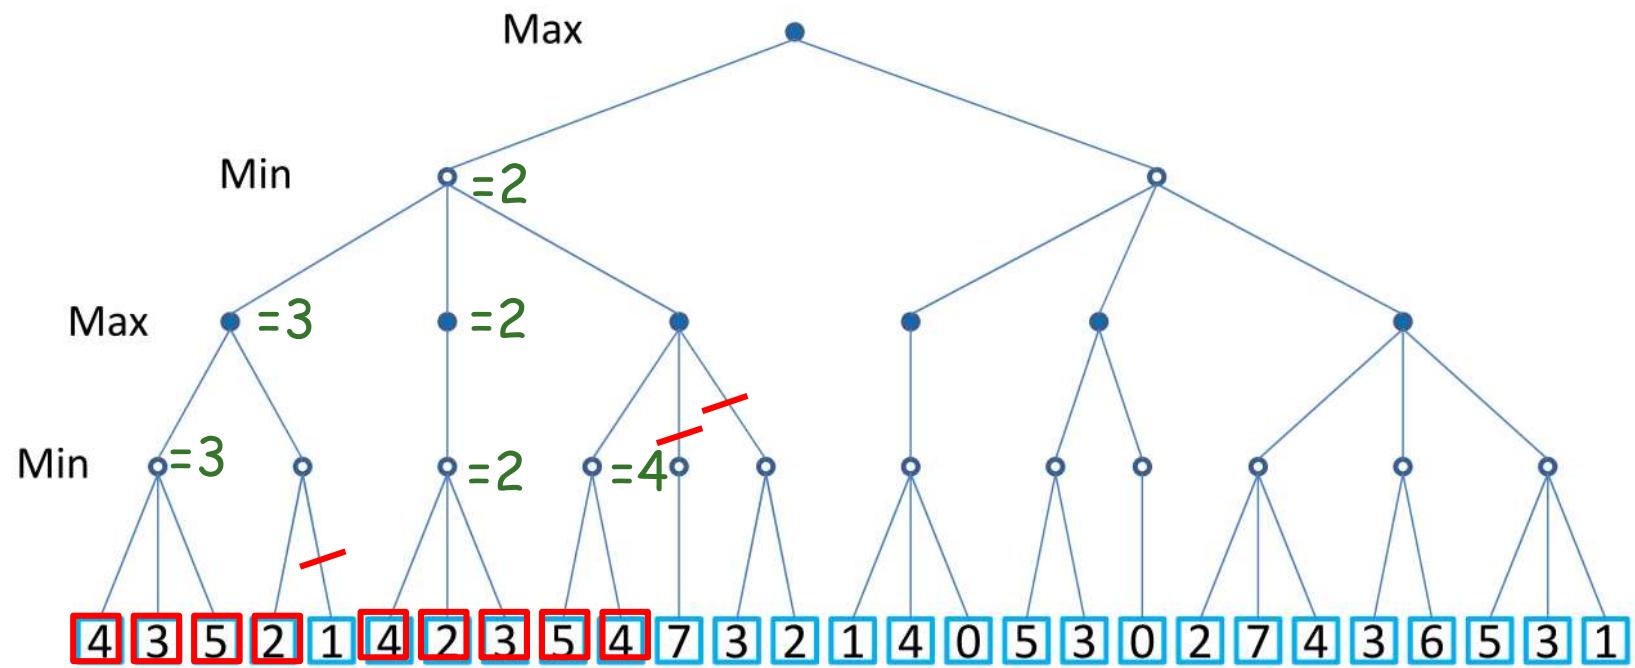
Alpha-Beta Pruning: Example 3



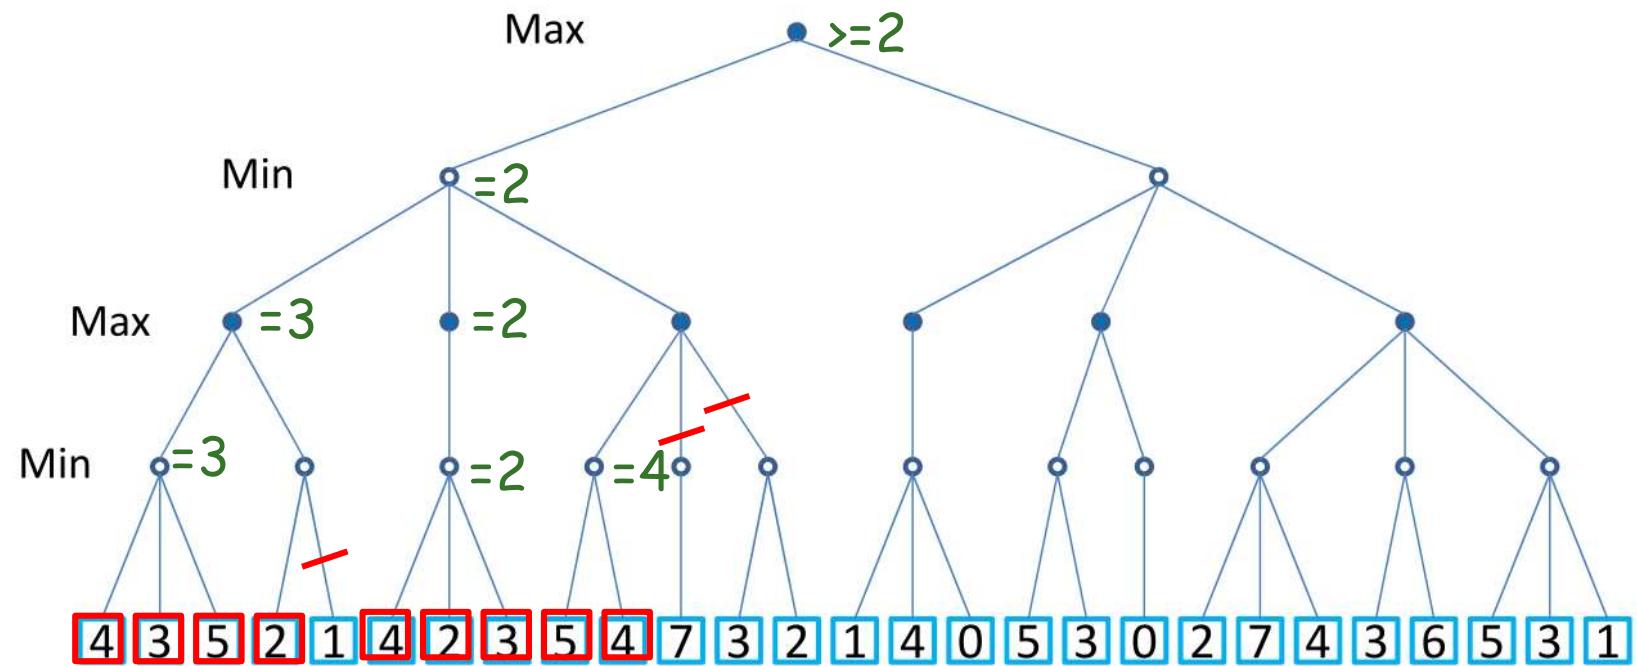
Alpha-Beta Pruning: Example 3



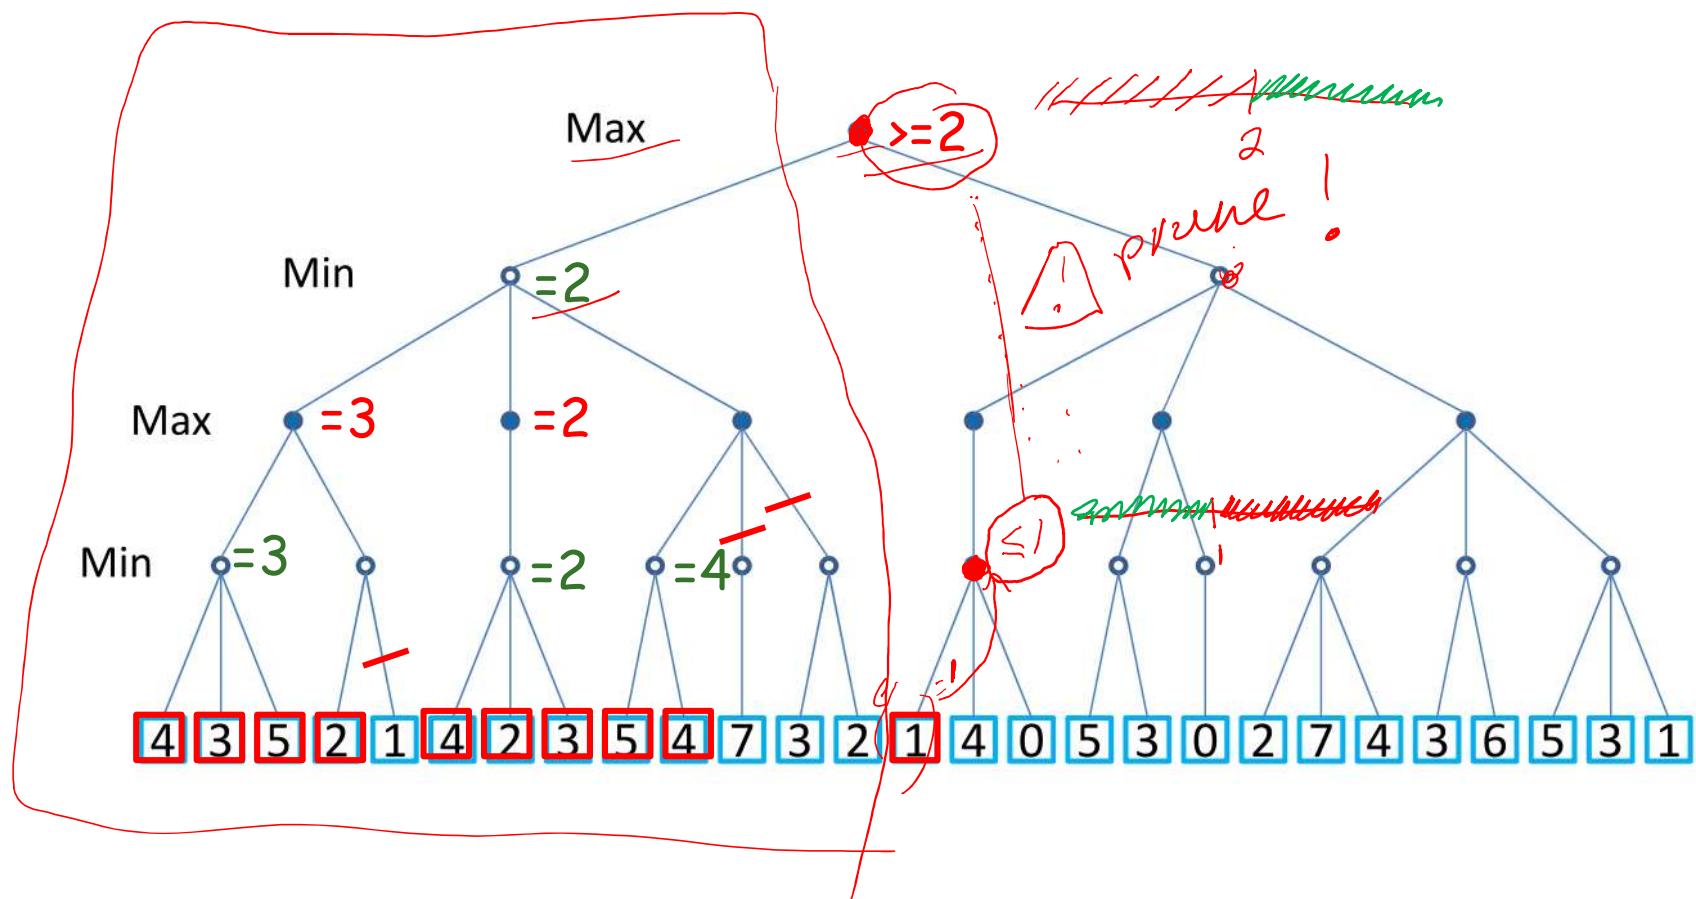
Alpha-Beta Pruning: Example 3



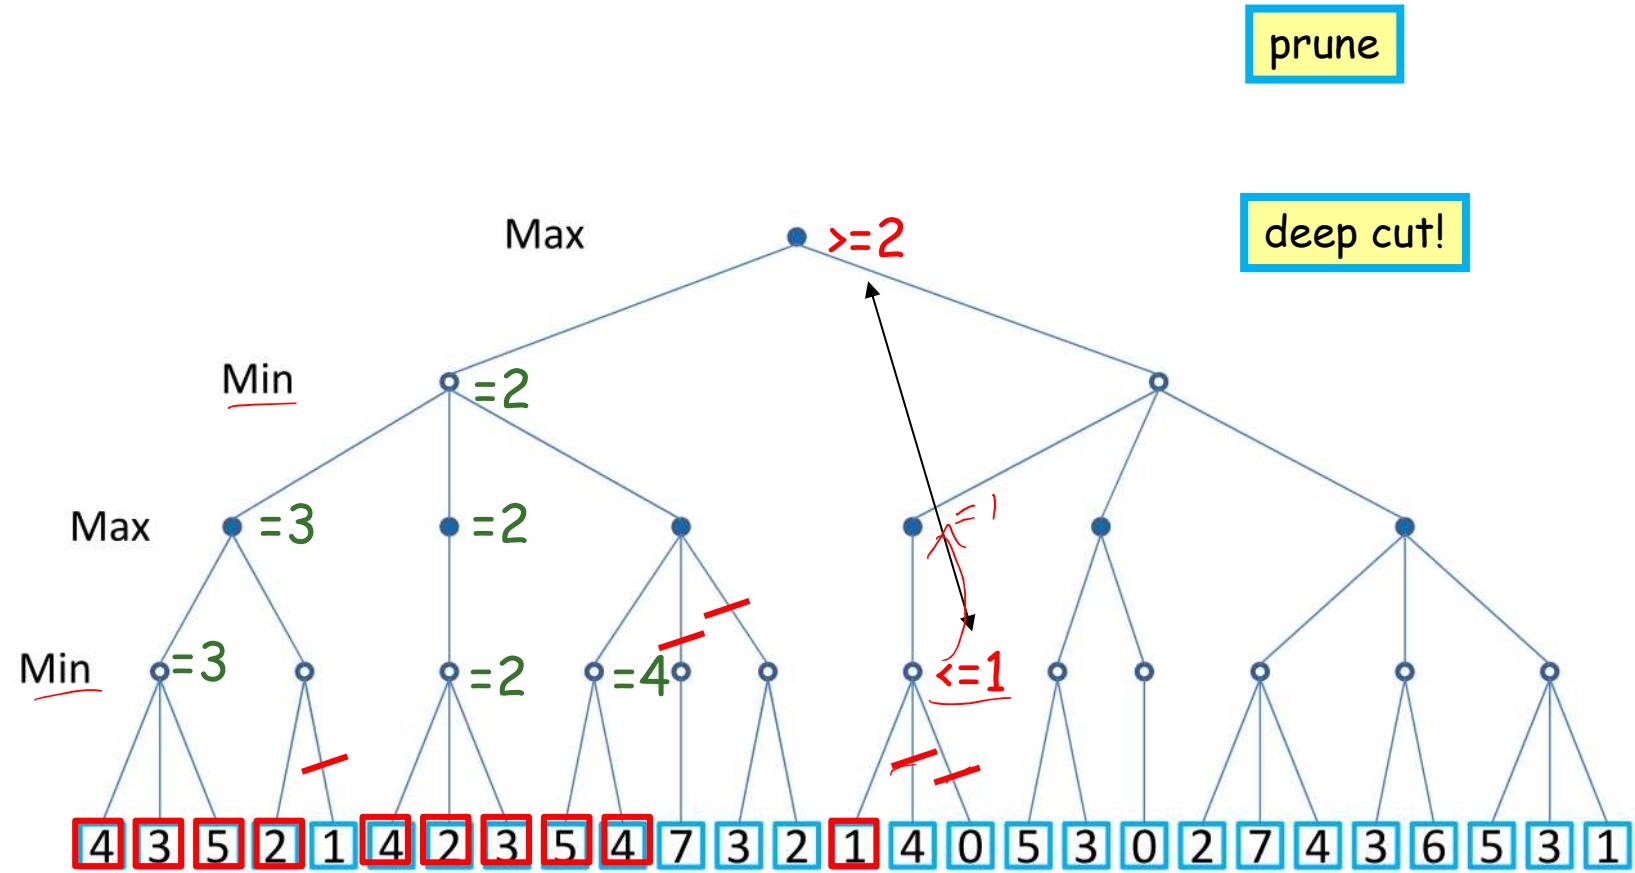
Alpha-Beta Pruning: Example 3



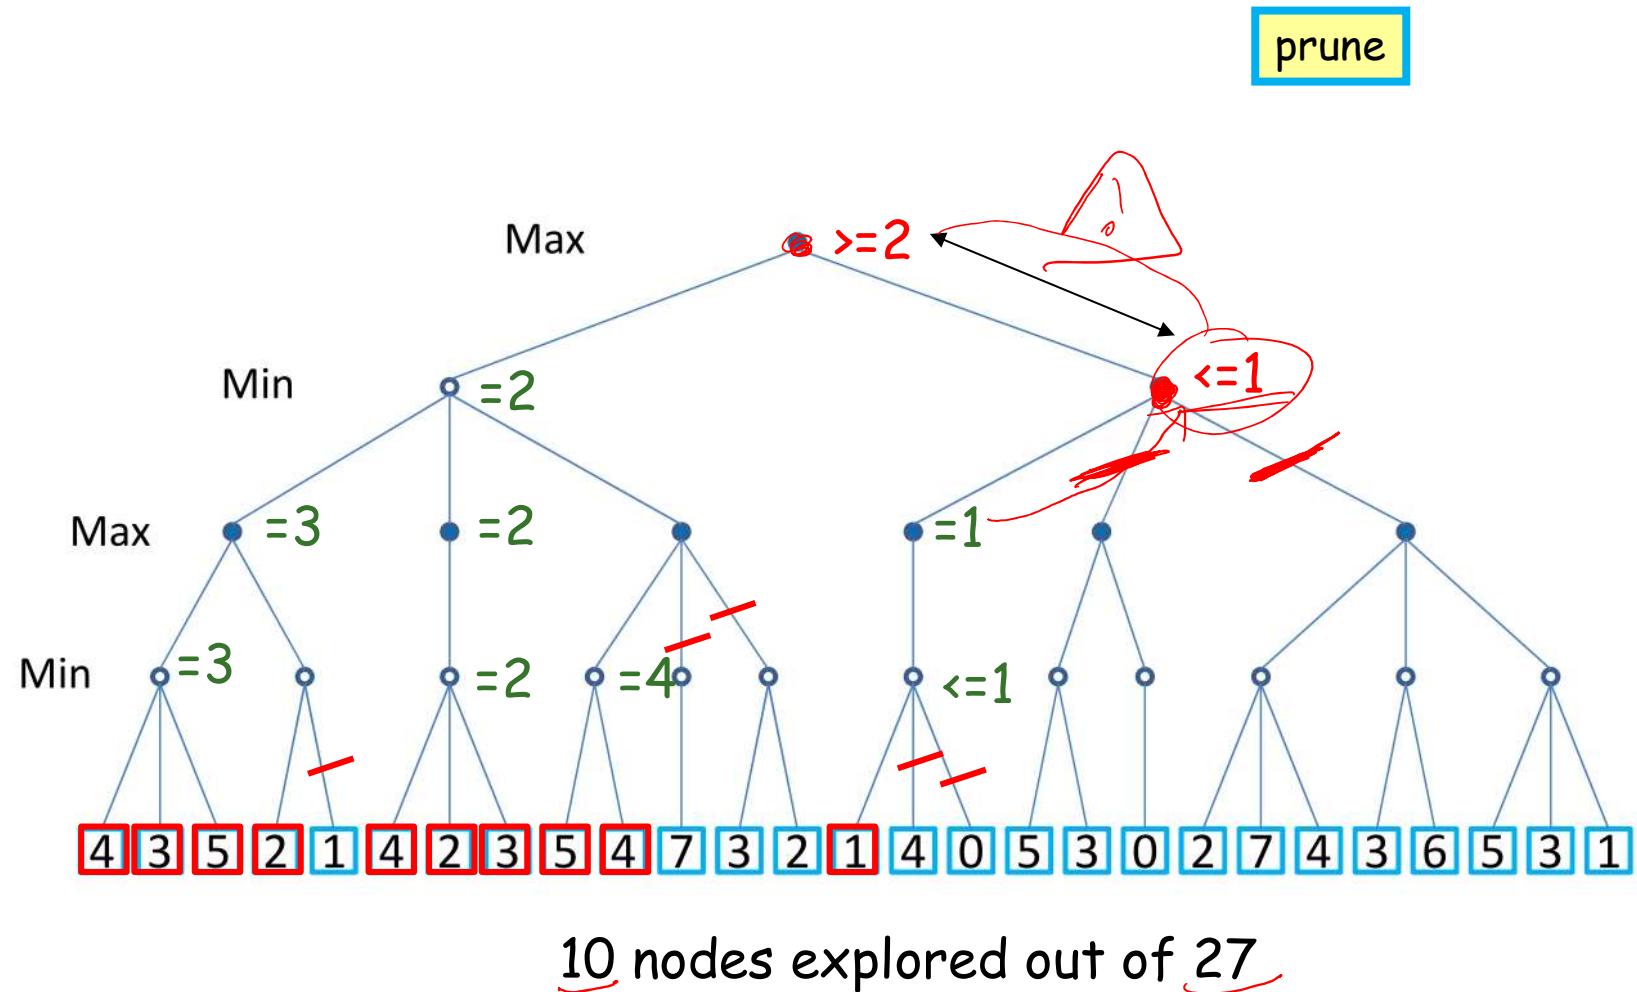
Alpha-Beta Pruning: Example 3



Alpha-Beta Pruning: Example 3

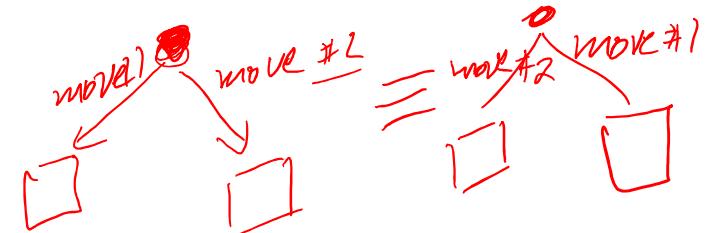


Alpha-Beta Pruning: Example 3



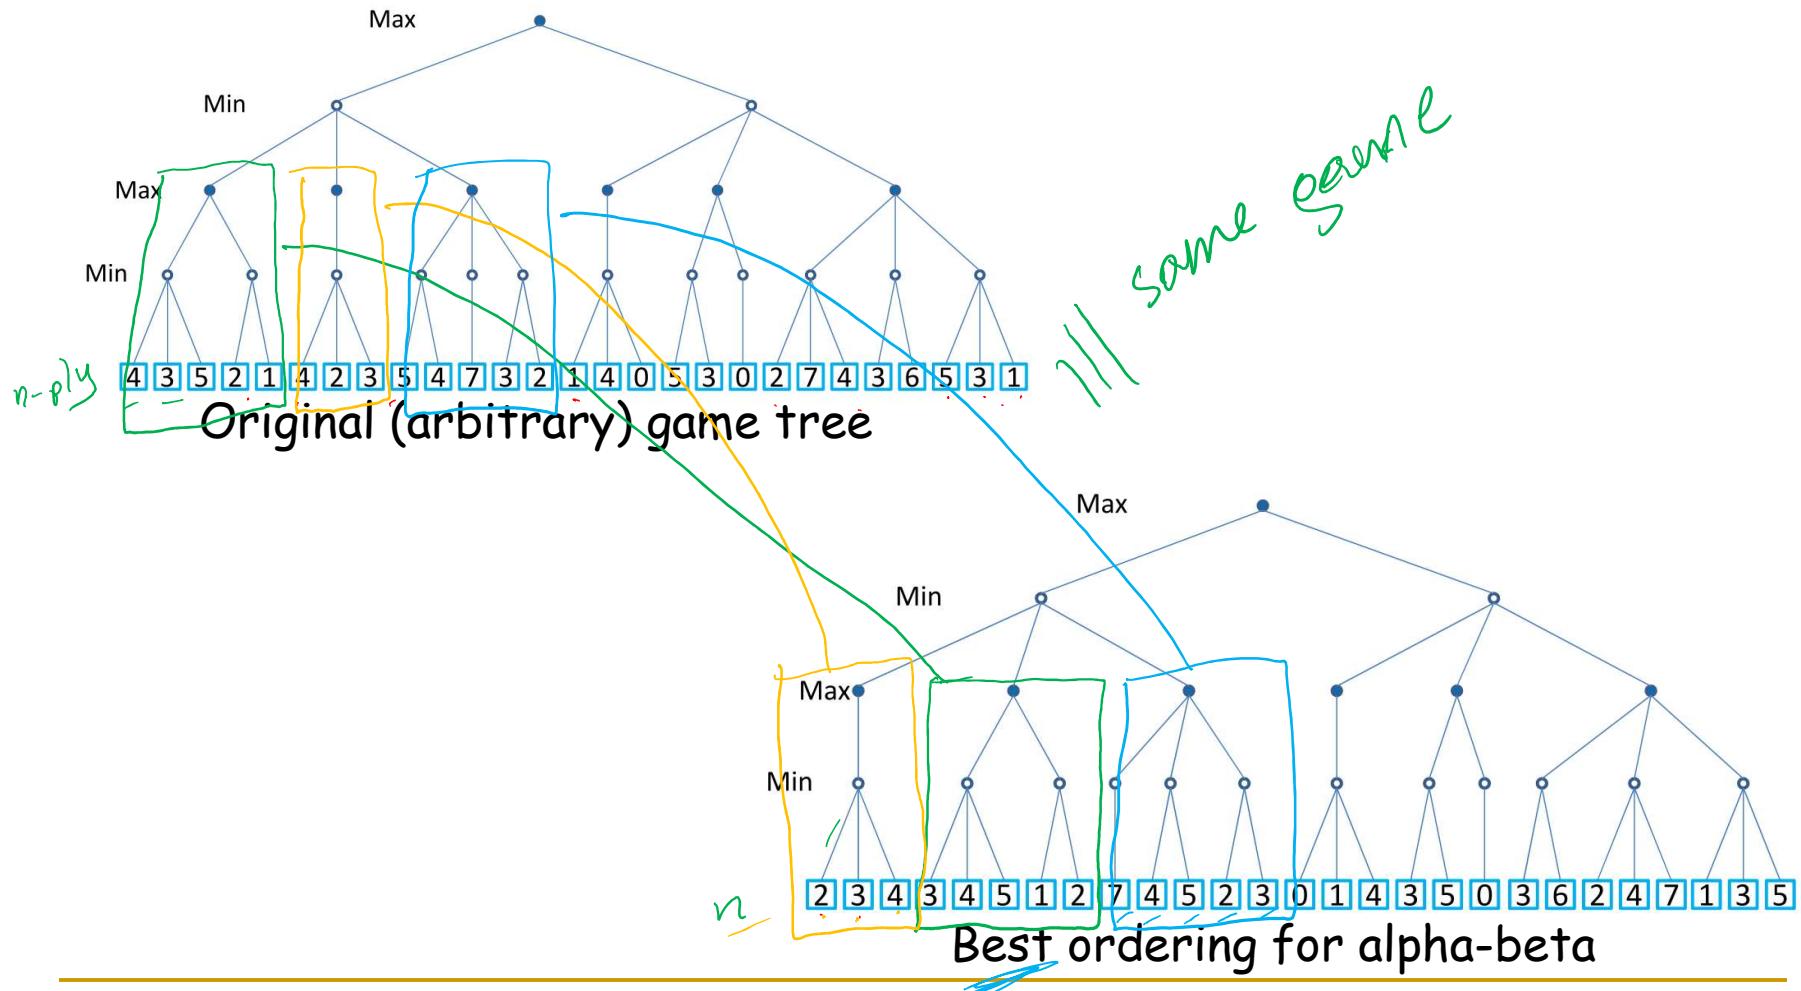
Efficiency of Alpha-Beta Pruning

- Depends on the order the siblings
 - which is an arbitrary choice ;-(
- In worst case:
 - alpha-beta provides no pruning
 - plus extra overhead cost ;-(
- In best case:
 - branching factor is reduced to its square root



$$\Rightarrow \sqrt{b} e(n)$$

Alpha-Beta: Best ordering

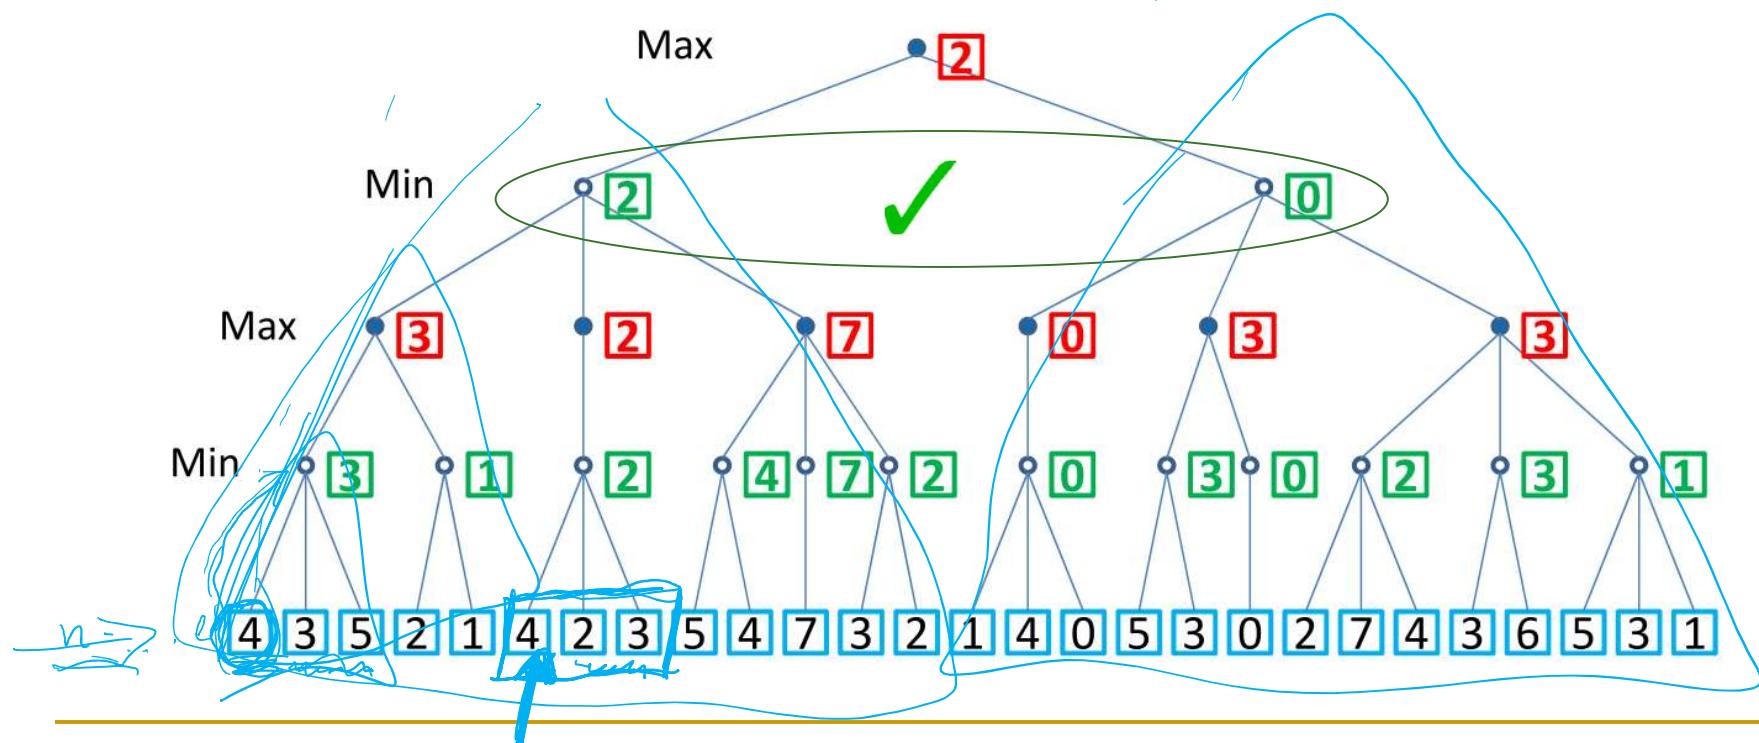


Alpha-Beta: Best ordering

- best ordering:

- strongest constraint placed first, ie:

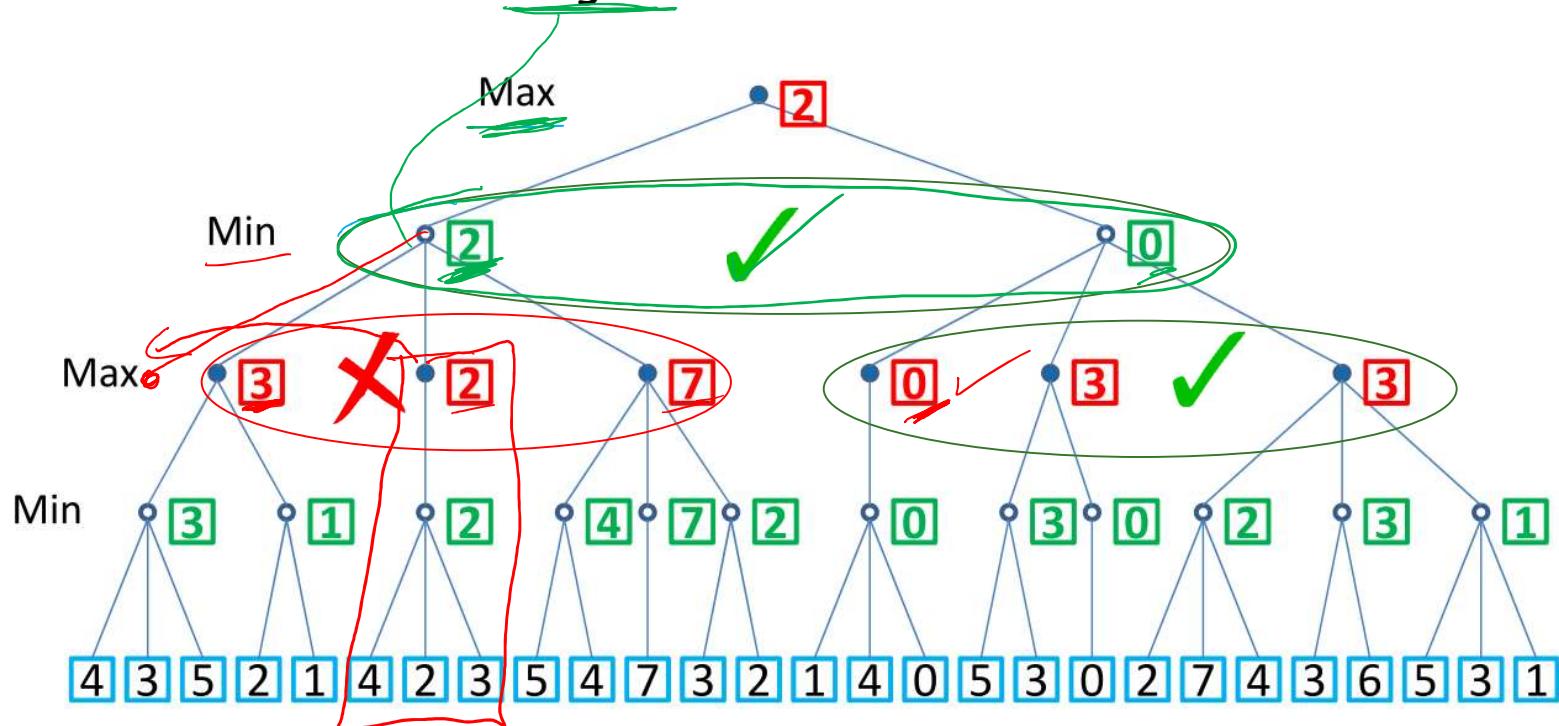
- 1. children of MIN : smallest node first
 - 2. children of MAX: largest node first



Alpha-Beta: Best ordering

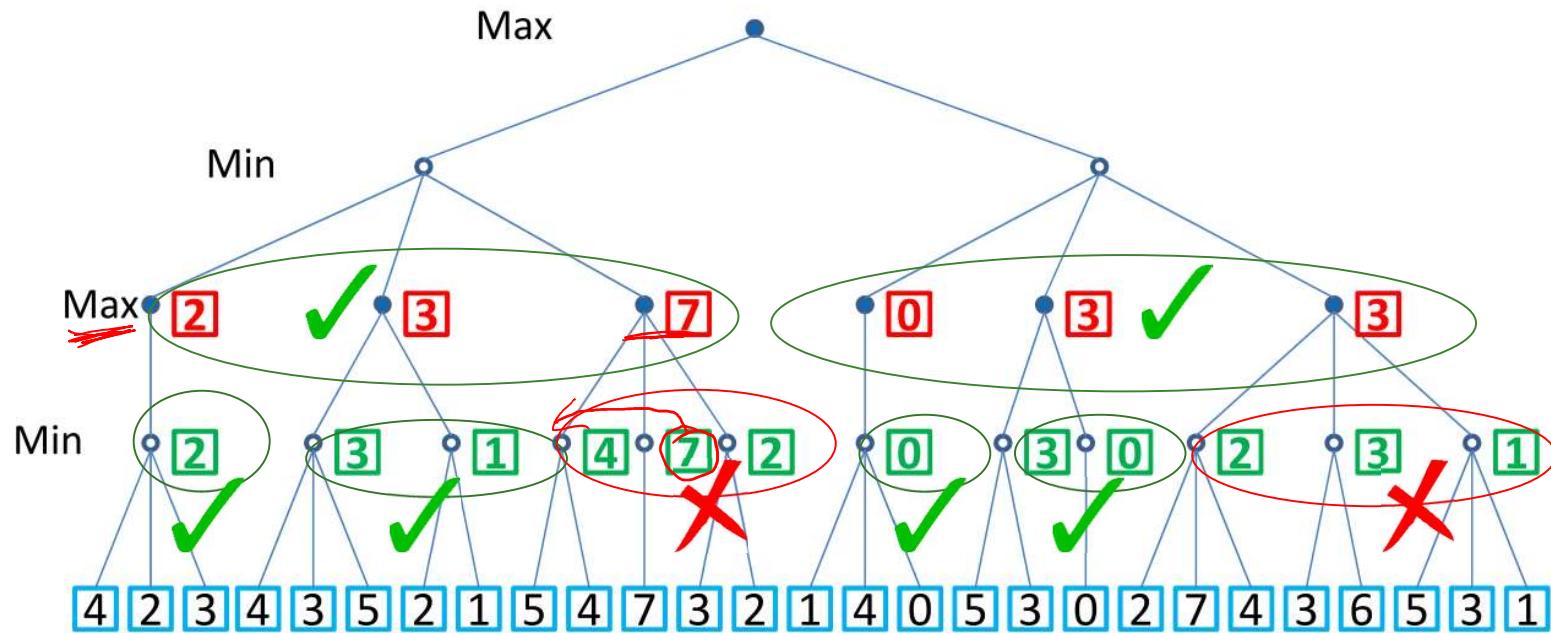
- best ordering:

- children of MIN: smallest node first
- children of MAX: largest node first

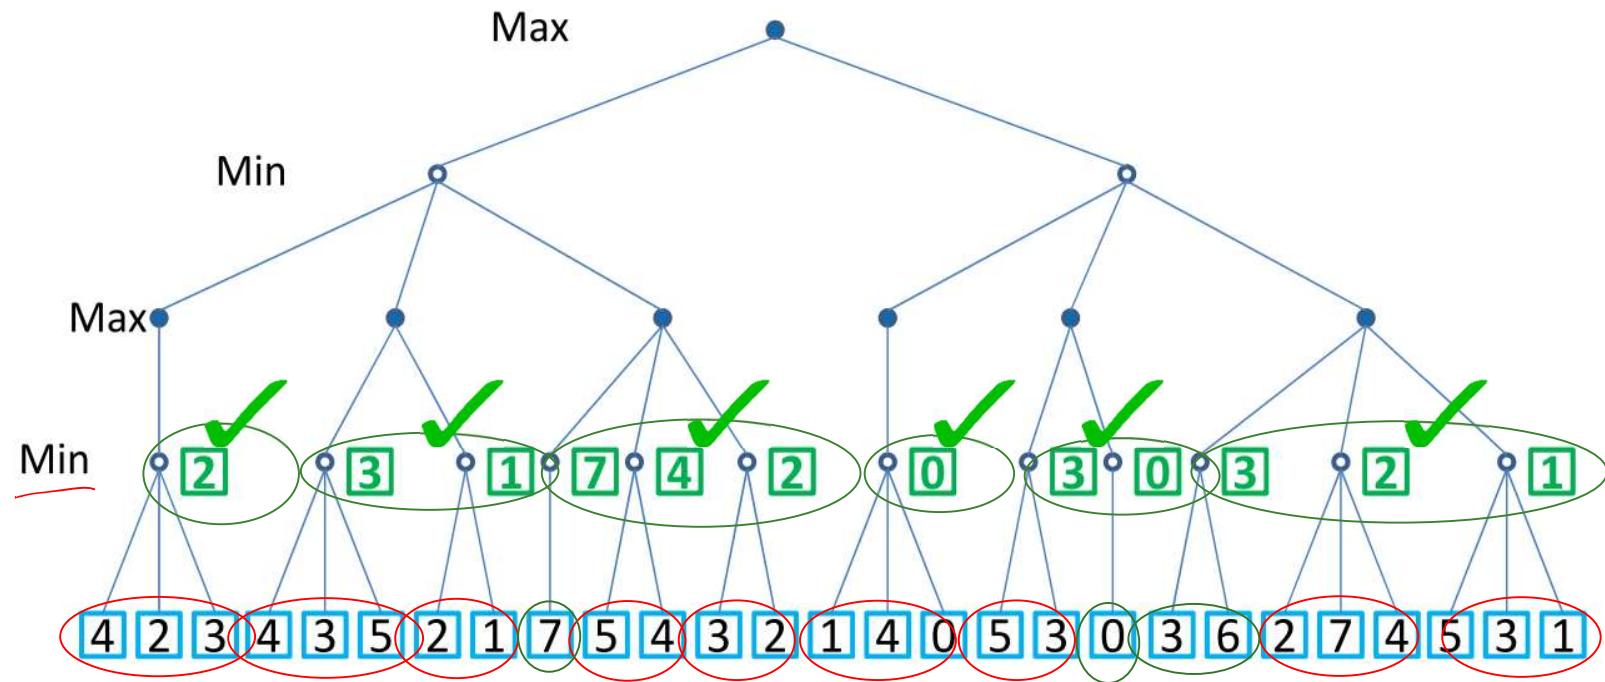


Alpha-Beta: Best ordering

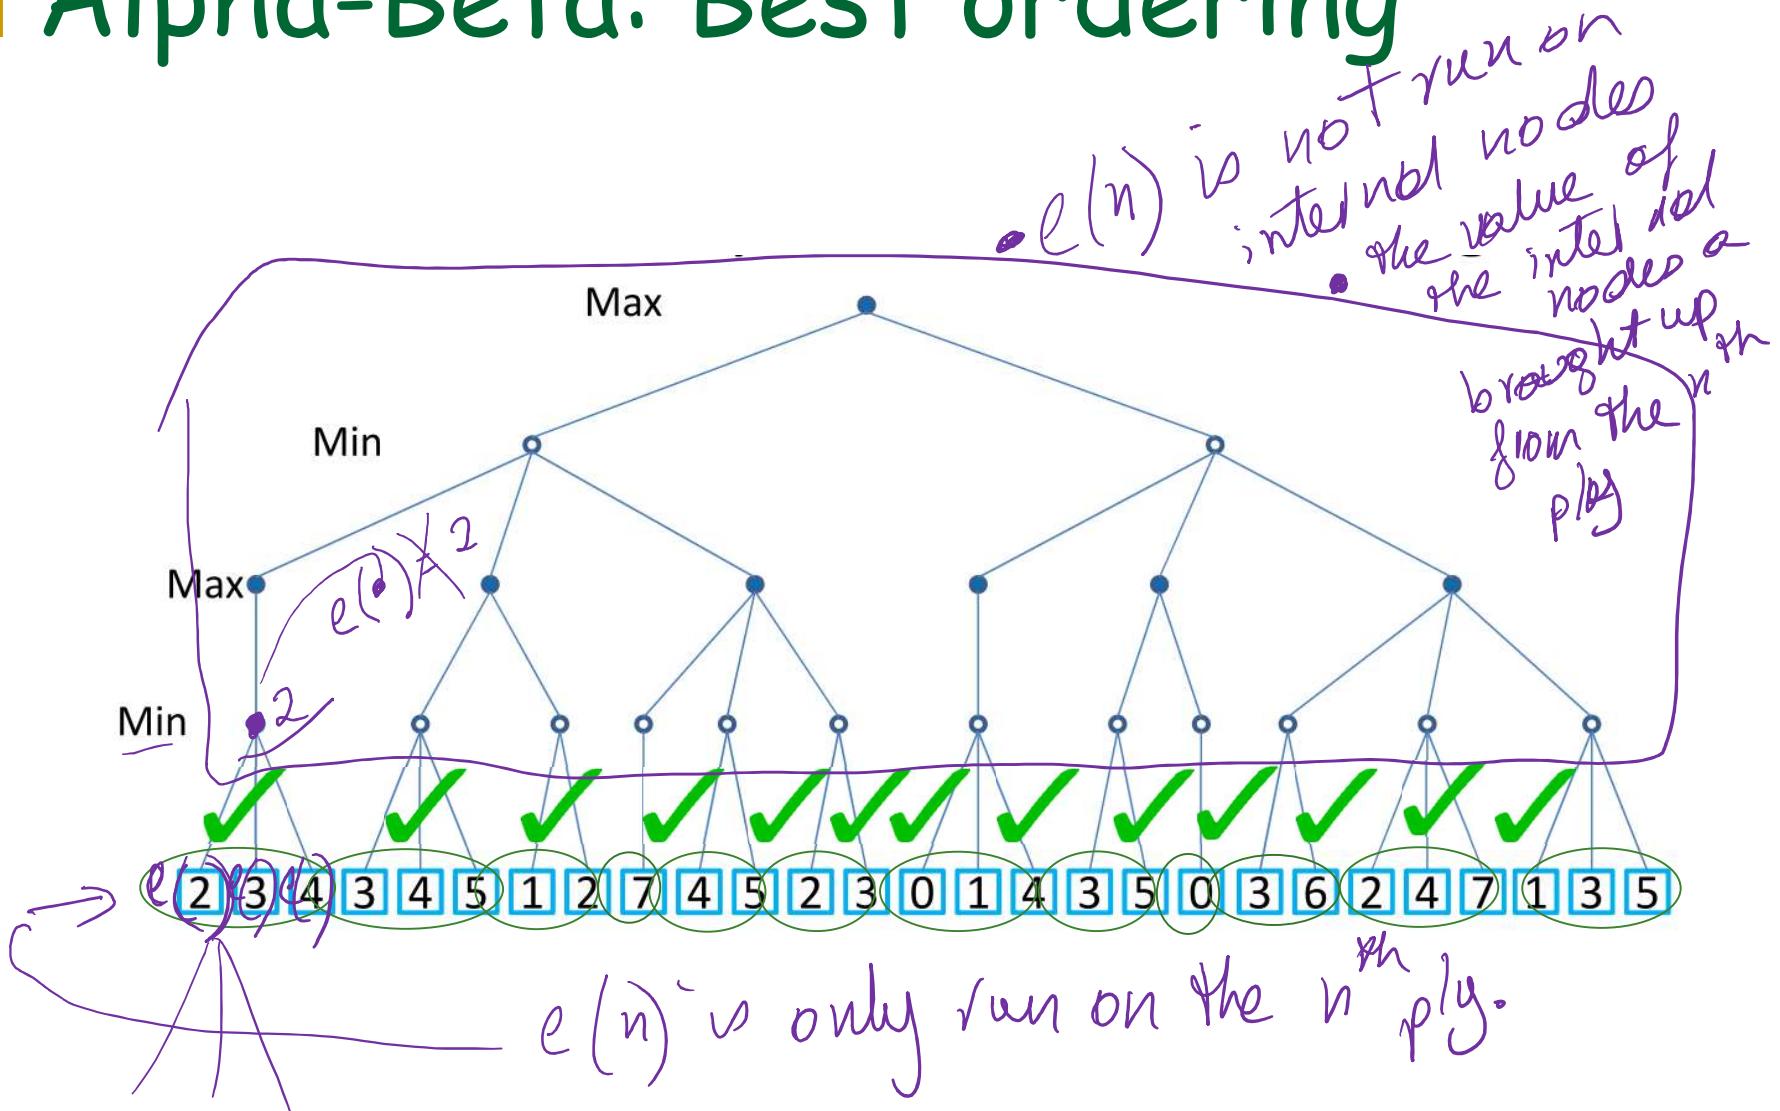
- best ordering:
 - children of MIN: smallest node first
 - children of MAX: largest node first



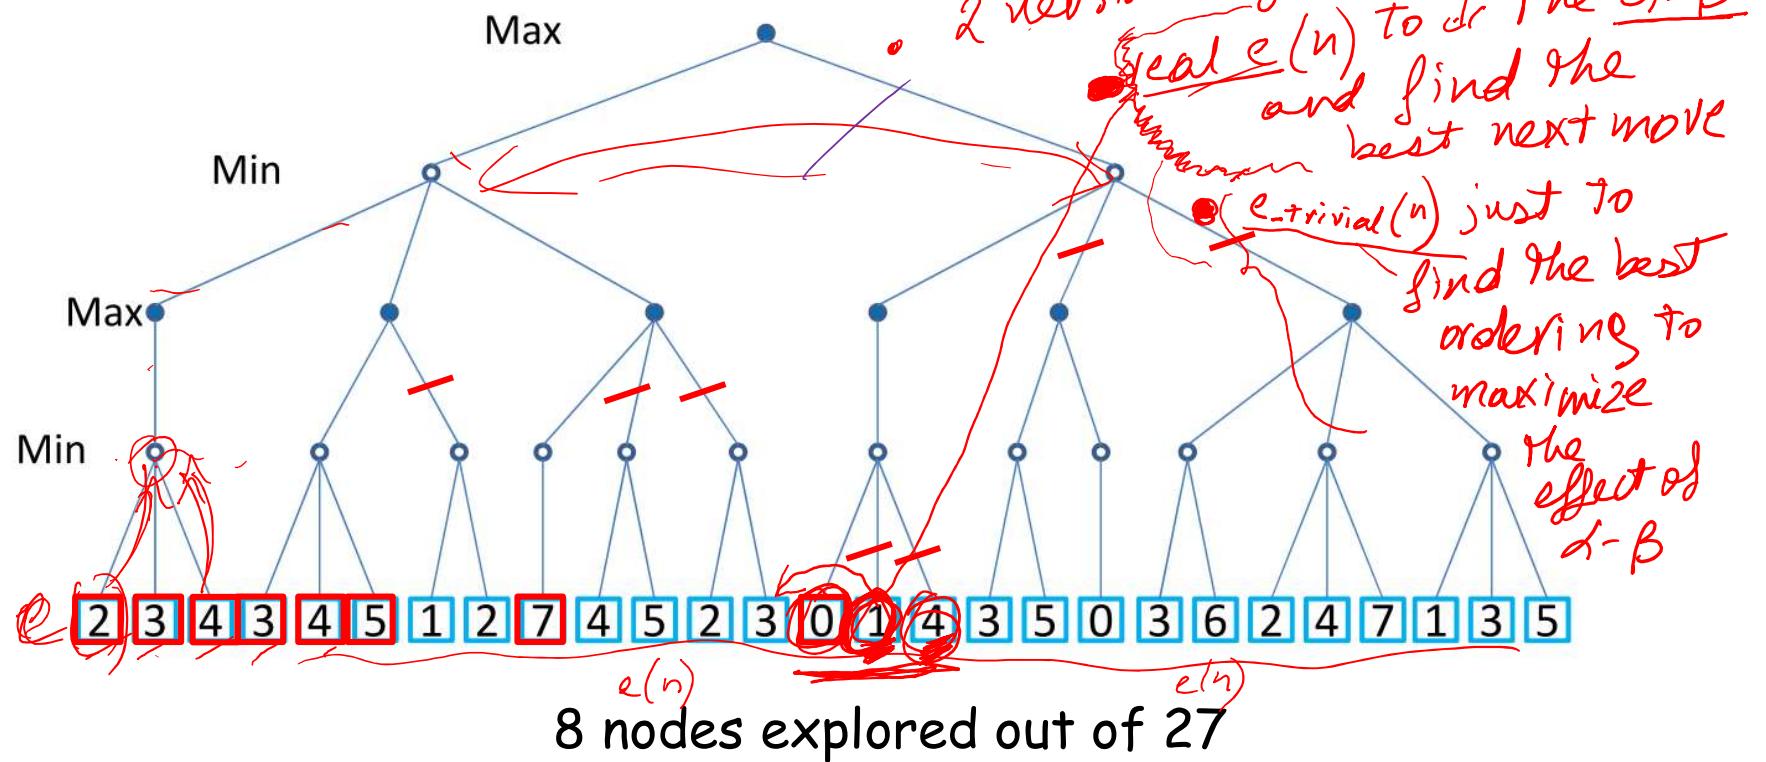
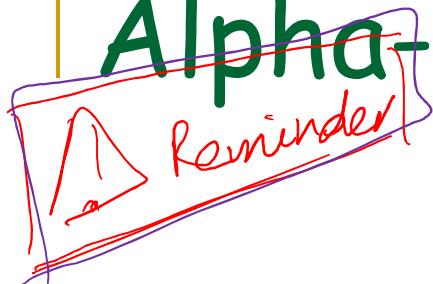
Alpha-Beta: Best ordering



Alpha-Beta: Best ordering



Alpha-Beta: Best ordering



Today

■ Adversarial Search

1. Minimax ✓

2. Alpha-beta pruning ✓

3. Other Adversarial Search

1. Multiplayer Games

2. Stochastic Games

3. Monte Carlo Tree Search

2 players
deterministic
perfect info

4.3

large branching factor
 $S \propto e(n)$

Up Next

- Adversarial Search
 - 1. Minimax
 - 2. Alpha-beta pruning
 - 3. Other Adversarial Search
 - 1. Multiplayer Games
 - 2. Stochastic Games
 - 3. Monte Carlo Tree Search